

BEGINNER EXPLAINER

ChoiceForge

How GPU kernels can accelerate a real travel demand model - from the first idea through a public ActivitySim proof.

Written for readers with no background in transportation modeling, probability, or GPU computing.

2.74x	73.9 s	0
complete-system ratio ordinary CPU / hybrid	launch-to-exit median seconds	disagreements in checked choices

Latest evidence, Phase 63: six balanced old/new pairs on the public 50,000-household model reduce launch-to-exit medians from 82.19 to 73.94 seconds. Every pair improves both clocks. The under-70-second median target is NOT met. Ordinary CPU takes 202.82 seconds. Giving CPU the same new preparation options still leaves a 2.53x hybrid advantage in a separate fresh-worker comparison. All 94 measured models pass their output checks, including changed seeds, coefficients and household counts. Checked decisions, matrices and report values agree under the documented formatting and diagnostic-score contracts. Ten-scenario persistent totals are 590.70 seconds hybrid and 1466.35 seconds CPU. These mixed-size batch totals are not single-run times. Phase 63 safely reuses preparation and fixes retained GPU arrays. Live CPU checks remain. A clean same-machine rebuild downloads public inputs and pins 98 libraries; this is not independent hardware replication. The gains include CPU and file-writing improvements, not only GPU arithmetic.

Who this guide is for

This guide is for a curious high school student. You do not need to know transportation planning, probability, programming, or computer hardware. By the end, you should understand:

- what a travel demand model tries to predict;
- how a computer turns possible choices into probabilities;
- why the same calculation must be repeated millions of times;
- what CPUs, GPUs, and GPU kernels do;
- what ChoiceForge changes;
- how correctness and speed are proven on a public benchmark;
- what the latest live scheduling kernel changes and what the benchmark can and cannot prove;
- why faster modeling could matter to communities.

The one-minute version

Latest result, Phase 63: the complete public 50,000-household model takes 73.94 seconds from launch to exit, compared with 202.82 seconds for newly measured regular CPU ActivitySim using 48-thread capacity. That is 2.74 times faster, or 63.54% less waiting, on this machine and configuration. Against the previous accelerated version in six balanced pairs, time falls from 82.19 to 73.94 seconds, a 10.04% reduction. Every pair improves. All six new runs finish in 73.17-80.00 seconds. **The under-70-second fresh-process target and 65-second stretch target were not met.** Improvement is not the same as meeting every goal. Earlier phases' times came from different sessions.

All 94 measured models pass their output checks, including 80 models with changed seeds, coefficients and household counts. Checked travel decisions, all 115 matrices and values throughout all 24 summary reports agree with independently computed CPU answers. Existing 5 versus 5.0 departure-key formatting allowances and diagnostic-score bounds still apply. The runtime remains a CPU/GPU hybrid; live CPU checks for borderline choices remain.

Phase 63 reuses calculation recipes, not old modeled answers, and fixes a real retained-GPU-memory defect. Giving CPU the same new preparation options still leaves a 2.53-times hybrid advantage: 193.18 versus 76.42 seconds in the separate fresh-worker comparison. Its clock includes extra input and reset checks. Ten-scenario hybrid batches take 590.70 seconds instead of 714.47 with fresh processes, a 17.32% reduction. Those batches mix household counts and are not single-run times. The full-system gains include CPU and file-writing changes, not only GPU arithmetic. Some small calculations still favor CPU.

All 619 main tests pass. A separately rebuilt directory downloads public inputs, installs 98 pinned libraries and passes 618 tests with one expected historical checkpoint skip. This is same-machine reconstruction, not proof on another computer. Sections 338-346 explain the changes, results, memory limits, preparation costs and what remains to do. Earlier sections preserve the history.

A city may want to know what could happen if it adds a bus line, changes a toll, builds housing, or closes a bridge. It cannot test every idea in the real world first, so planners use a **travel demand model**: a computer simulation of how people may decide where, when, why, and how to travel.

Modern models can represent millions of people and many possible choices. For each person, the computer scores the available choices, turns the scores into probabilities, and uses a controlled random number to select one outcome. It repeats this work again and again for trips, tours, destinations, travel modes, and times of day.

A normal processor, called a **CPU**, is very flexible but has a modest number of powerful processing cores. A **GPU** has many more, simpler cores and is good at doing the same kind of arithmetic for a huge number of records at once. A **GPU kernel** is a small program designed to run across those GPU cores.

ChoiceForge is an open-source experiment that moves a particularly repetitive travel-choice calculation to an NVIDIA GPU. Its most important idea is **fusion**: calculate utility scores, apply availability rules, compute a logsum, and make the choice in one GPU operation, without first storing an enormous table of intermediate scores.

The first synthetic test compared the GPU with a readable NumPy reference. Phase 1 added a much stronger 24-core CPU competitor. Phase 2 captured 4,477 real mandatory-tour scheduling decisions from ActivitySim, and every GPU choice matched. It also found a problem: transferring a 151.6-megabyte table made the GPU slower. Phase 3 replaced that table with 22.0 megabytes of compact ingredients and made the captured calculation 3.83 times faster than a purpose-built 48-thread CPU version. Phases 4 and 5 installed and reused the GPU path in real scheduling. Phase 6 handled real destination work. Phase 7 batched repeated setup and added nested logit on the GPU. Phase 8 moved the work to pinned current ActivitySim and a public 50,000-household run. Phase 9 then used the much larger public MTC geography. It confirmed an exact GPU destination result and also caught a scheduling mismatch before that path could be claimed as correct. Phase 11 repeated the full-geography destination experiment three times and established the supported production result: the complete 50,000-household model was 1.064 times faster with byte-for-byte identical outputs. Phase 12 built the foundation for moving the large utility equations themselves onto the GPU. Phase 13 completed the strict CPU answer key. Phase 14 generated GPU code from that same recipe and matched the CPU exactly on 30 real public-model batches. Phase 15 connected it to the real model, removed the giant intermediate table, proved a small destination-component win, and found an honest scale limit. Phase 16 added compact inputs, caches, and a published FP32 policy. Its CPU and GPU matched every checked cell, and three large public pairs proved a repeated 1.025-times trip-destination component speedup with exact modeled decisions. Phase 17 made the compiled GPU work reusable and continued it into trip mode choice. Five large public pairs improved trip destination by 1.040 times with exact modeled decisions; the complete-model median also improved by 1.006 times, although its strict statistical gate still did not pass.

Phase 18 then kept a dependent model-shaped chain on the GPU and processed all 2.875 million public households, but its behavior equations were synthetic. Phase 19 replaced that synthetic boundary with the public calibrated auto-ownership and mandatory-tour-frequency equations. It exactly reproduced 50,000 saved household choices and 78,900 saved person choices. Median GPU modeled work was 17.840 times faster than the independent CPU replay and 12.073 times faster after including one input upload and final download. That is the strong calibrated result, although upstream location/CDAP work remained outside its boundary.

Phase 20 then solved the next structural problem. The GPU turned those 78,900 person choices into 81,983 variable-length tour rows, and every value in all 12 tour columns matched ActivitySim. It fed those exact tour IDs into six real mandatory-scheduling batches containing 15.24 million possible tour-time rows. All 81,983 GPU time choices matched ActivitySim. Tour creation was 11.496 times faster with resident data and 6.272 times faster including transfers. The scheduling kernel was 18.097 times faster with resident compact inputs and 2.935 times faster including compact transfers. ActivitySim still prepares the scheduling logsums, feasible time choices, and timetable facts on the CPU, so this is not yet a whole scheduling-component speed claim.

Phase 21 moved that missing scheduling preparation onto the GPU. It gave every person a small timetable on the graphics card, tested all 190 time choices, kept only the choices that did not collide, rebuilt the seven timetable facts used by ActivitySim, made the choice, and updated the timetable before a later tour was scheduled. Across nine measurements, this complete compact-cache boundary was 8.680 to 10.199 times faster than a compiled parallel CPU version, depending on whether primitive transfers were included. Every one of 15.24 million regenerated rows and all 81,983 time choices matched.

Phase 21 also connected the real raw network skims to the CUDA mode-logsum engine inside ActivitySim. All six live calls used the GPU with no fallback, and the final tour times were unchanged. This required matching Sharrow's fused 32-bit utility arithmetic and ActivitySim's later 64-bit nest arithmetic, because one random draw was close enough to a probability boundary to expose the difference. The 8.680-to-10.199-times result begins after a compact logsum cache exists; the raw-skim integration is a separate correctness proof, not a hidden whole-component speed claim.

Phase 22 joined those two halves. Three paired public runs ended with all 81,983 mandatory schedules unchanged, and the connected GPU path was faster in all three pairs. The middle paired speedup was 1.257 times. The GPU also found 57

choices whose random draws were so close to a probability boundary that CPU and GPU rounding could disagree. The real Sharrow arithmetic resolved only those 57 rows, using 11,400 bytes of transferred logsum data. This makes Phase 22 an exact, mostly-GPU component result - not an absolutely CPU-free model.

Phase 23 then changed the architecture. Instead of repeatedly returning to CPU-owned ActivitySim tables between decisions, it created a versioned model state that stays on the graphics card across calibrated auto ownership, mandatory-tour frequency, tour creation, tour identity linking, and timetable scheduling. A new fused kernel evaluates all 98 mandatory-frequency expressions and the final choice without building a large intermediate feature table. Three independent processes, each with nine measured repetitions, produced a middle resident speedup of **24.405 times**. Even charging one-time setup and final publication to a single run was **1.356 times faster** at the middle result. Every calibrated choice, all 12 tour columns, all 81,983 time choices, and every final timetable value matched. The runtime can also write a self-contained checkpoint, restore it, and continue computing on the GPU.

Phase 24 tackles the largest shared input: network skims. The public file contains 826 compressed matrices, but the reviewed 315-term tour-mode equations need 209 logical skim lookups. Directional views share source data, so 149 physical float32 cubes occupy 6.38 GB and fit under an 8 GiB budget. Across three independent processes, the GPU read and checked 251.76 million real-workload skim values per run with zero mismatches. The middle resident cache-layer speedup was **193.114 times** and the slowest process was still **82.323 times** faster. Even paying the entire upload once gave a middle **1.813-times** win. This is a raw-data-layer proof, not yet the complete utility/logsum stage.

Phase 25 connects that raw-data layer to the real mathematics. Six real 315-term programs now read the resident skims, score 21 travel modes, combine related modes with nested logit, and place the resulting logsums into each tour's 5-by-5 scheduling cache. The cache-placement map is compiled once and kept on the GPU too. Three fresh processes completed 15 measured replays with no changed logsum bit. The middle process time was **0.169 seconds** for all 1.21 million rows. This was **9.655 times faster** than the same process's initial live CUDA setup-and-execution path. It is not a CPU speedup or a whole model speedup. ActivitySim still creates the dense input rows, and the final live scheduler still asks the exact CPU/Sharrow path to settle 57 extremely close probability cases.

Phase 26 connects that producer directly to GPU scheduling and timetable mutation. The GPU now generates the feasible scheduling rows and their compact index, consumes the newly created logsum caches, chooses all 81,983 mandatory tour times, and updates each person's calendar inside one sealed, versioned stage. Three fresh processes completed 15 measured replays at a middle time of **0.201 seconds**, with every logsum bit and final TDD unchanged. The former 57 near-boundary cases no longer travel to a CPU resolver. They use an explicit, public-benchmark Sharrow answer map already held on the GPU; only one needs a correction, and zero boundary bytes are downloaded. This is exact for the qualified frozen benchmark, not a universal arithmetic promise for changed inputs. ActivitySim still creates the dense mode-choice leaves and coordinates before sealing them.

Phase 27 removes those prepared row arrays from the timed graph. It replaces 503.4 MB of repeated dense values and coordinates with 25.0 MB of compact facts: values shared by a whole tour, values shared by a time alternative, and small dictionaries for repeated tour-by-alternative patterns. A CUDA kernel rebuilds the exact arrays in about **0.002915 seconds**. The matched NumPy job takes **0.491203 seconds**, making this reconstruction boundary **168.52 times faster** on the GPU. The entire compact-input-to-calendar graph takes **0.205337 seconds**, only 2.23% longer than Phase 26 even though it now does the missing reconstruction. All 15 full replays keep every logsum bit and all 81,983 final times unchanged. ActivitySim still creates the dense arrays once during qualification so the compact form can be discovered and checked; removing that cold-start dependency is the next phase.

Phase 28 replaces the remaining anonymous response dictionaries with named rules. The GPU now calculates parking cost from a tour rate and duration, and calculates 14 mode-availability fields from raw resident road/transit skims and auto ownership. Three public processes and 15 complete replays finish at a middle **0.211799 seconds**, with every logsum bit and every final tour time unchanged. Compact state falls again to 20.26 MB, a **24.849-times reduction** from the removed rows. Five deliberately changed synthetic populations and skim sets exercise all 15 rules across 8,000 rows with exact, different outputs. The stronger meaning costs time: the full graph is 3.15% slower than Phase 27 because it computes real skim rules instead of looking up remembered patterns. ActivitySim still supplies dense rows before sealing as a qualification answer key, so cold raw-table generation remains unfinished.

Phase 29 takes the larger upstream step. Instead of learning compact facts by looking at the giant prepared answer rows, its compiler starts with one source row per tour, land-use facts, controlled random inputs, possible time slots, and raw network skims. It declares 57 sources for each of six programs and generates all 18 road/transit availability fields on CUDA. Parking rates now come directly from land use plus free-parking status. Across three public processes and 15 replays, the middle complete graph is **0.225311 seconds**; every logsum bit and all 81,983 final times still match. Five changed raw-table populations (10,000 tours) and five changed skim worlds (8,000 rows) also pass. The compact state is 24.85 MB, **20.259 times smaller** than the removed rows. The qualification harness still runs the old dense path separately as an answer key and to expose the compiled utility interface, so the roughly 30.759-second cold ActivitySim run has not disappeared yet.

Phase 30 removes that production dependency. The GPU interface is now built directly from a reviewed equation recipe, declared raw-table types, scalar settings, controlled random draws, and raw skim-cube descriptions. The old dense preprocessor is run only in a separate proof process. Across three fresh public processes, all 1,210,124 dense rows are avoided, all 15 resident replays are exact, and all 81,983 final times match. A byte-level hash test also proves that the native and legacy paths generate identical six logsum arrays. The middle resident time is **0.226712 seconds**. The middle cold run is **30.739 seconds**, essentially unchanged because loading 6.452 GB of network skims is still the largest startup cost. This phase is a major independence and replication result, not a major new cold speedup.

Phase 31 attacks the startup cost that Phase 30 exposed. It converts the needed network measurements into a checked, GPU-ready file once, then skips Sharrow's 6.452 GB dataset in every live run. Three fresh processes keep every logsum bit and all 81,983 schedules exact. Including setup for the small on-device boundary map, the middle cold component time falls from **30.739 seconds to 27.085 seconds: 3.654 seconds saved, 11.887% lower, and 1.135 times faster**. This is the first material cold-start gain since the native runtime work began; it applies to this resumed mandatory-scheduling component, not the whole regional model.

1. What is travel demand modeling?

Imagine a region with homes, schools, offices, stores, roads, sidewalks, rail lines, and buses. Transportation planners ask questions such as:

- How many people might travel during the morning rush?
- Where might they go?
- Will they drive, walk, bike, take transit, or share a ride?
- Which roads or transit services might become crowded?
- Who benefits from a project, and who may be left out?
- How might results change if fuel prices, fares, jobs, or housing change?

A travel demand model does not read minds or predict one person's future with certainty. It creates a plausible regional picture by applying observed patterns and mathematical rules to many simulated people.

People, households, and zones

Models usually begin with a **synthetic population**. These are made-up records that resemble the region's real population without being a list of actual named residents. A record may describe a household's size, income, vehicles, workers, and students.

The region is divided into geographic areas called **zones**. A zone may contain homes, jobs, schools, shops, or transit stations. The model also needs information about travel between zones, such as distance, driving time, transit time, cost, and whether a route is available. These travel measurements are often called **skims**.

Activities, tours, and trips

People travel because they want to do activities. Going from home to school is a trip. Going from home to school, then to a store, then home is often treated as a **tour**: a chain of trips that begins and ends at the same anchor, usually home or work.

An activity-based model may answer a sequence of linked questions:

```
Who is traveling?
-> What activity are they doing?
-> Where will they go?
-> What time will they travel?
-> Which travel mode will they use?
-> Which trips and tours result?
```

These decisions affect one another. A faraway destination may make walking unrealistic. A late return time may make transit unavailable. A household with one vehicle and two workers may face competition for the car.

2. How does a model represent a choice?

Suppose Maya must choose among walking, taking the bus, and riding in a car. The model gives every option a **utility score**. Utility is not electricity here. It is a numerical score for how attractive an option is according to the model.

A simplified utility equation looks like this:

```
utility = constant + (time x time weight) + (cost x cost weight) + other terms
```

A negative time weight means longer trips are less attractive. A negative cost weight means expensive trips are less attractive. The weights are estimated from travel surveys or other observations. Different people can receive different scores because income, age, vehicle access, location, and other facts may matter.

Availability comes first

Not every option is possible. Walking may be unavailable for a very long trip. A transit route may not exist. A person may not have access to a car. The model applies an **availability rule** so an impossible option receives no chance of being chosen.

From scores to probabilities

Raw utility scores are not probabilities. A common model uses this rule:

```
probability of option i = exp(utility i) / sum of exp(utility for every available option)
```

exp is a mathematical function that turns higher scores into much larger positive numbers. Dividing by the total makes all probabilities add to 1, or 100 percent.

Here is a small example:

Option	Utility score	exp(score)	Probability
Walk	0	1.000	9.0%
Bus	1	2.718	24.5%
Car	2	7.389	66.5%

The car is most likely, but it is not guaranteed. This matters because real people with similar situations do not all make the same decision.

Turning probabilities into one choice

The model takes a random draw between 0 and 1. It lays the probabilities end to end:

```
Walk: 0.000 to 0.090
Bus: 0.090 to 0.335
Car: 0.335 to 1.000
```

If the draw is 0.600, the model chooses car. If it is 0.200, it chooses bus. This is called **inverse cumulative distribution sampling**, but the important idea is simple: larger probability ranges are more likely to catch the draw.

ChoiceForge does not invent random numbers inside the GPU. The calling model supplies them. That design lets the CPU and GPU use the exact same draws, which makes fair comparison possible and preserves ActivitySim's rules for repeatable simulations.

What is a logsum?

A **logsum** summarizes how good the full set of choices is:

```
logsum = log(sum of exp(utility for every available option))
```

In the example, the logsum is about 2.407. If a fast, affordable new bus service improves one option, the logsum rises. Models often pass this accessibility-like value into later decisions.

Computers must calculate logsums carefully. Very large utility scores can make `exp(score)` overflow, like a calculator displaying an error for a number that is too large. The standard safe method subtracts the largest utility before exponentiating and adds it back afterward. This is called **stable logsum-exp**. ChoiceForge uses this method.

3. Why does this become a computing problem?

The three-option example is tiny. A regional model may have:

- millions of people, households, tours, or trips;
- dozens of alternatives for mode or time choices;
- hundreds or thousands of possible destinations;
- many variables in every utility equation; and
- dozens of model steps that repeat similar calculations.

If one million choosers each consider 32 alternatives, the model evaluates 32 million chooser-option combinations in just one component. Destination choice can be much larger.

A straightforward program may create a rectangular **utility matrix** with one row per chooser and one column per alternative. For one million choosers and 32 alternatives stored as 32-bit numbers, that table alone uses 128 megabytes. The program may then create more tables for exponentials, probabilities, or cumulative probabilities. Moving and storing these intermediate values can take as much time as the arithmetic.

4. CPU, GPU, and kernel in everyday language

A **CPU** is the general-purpose brain of a computer. Think of it as a small team of highly trained chefs. Each chef can handle complicated instructions, switch tasks quickly, and make decisions.

A **GPU** was designed to calculate many screen pixels at once. Think of it as a very large kitchen crew. Each worker is less independent, but thousands of workers can perform the same short recipe on different ingredients at the same time.

A **GPU kernel** is that short recipe. It says what one group of GPU threads should calculate. Launching a kernel sends many copies of the recipe across the data.

The GPU is not automatically faster. Three costs matter:

- 1 **Transfer:** data may need to cross from ordinary computer memory to GPU memory and back.
- 2 **Launch overhead:** starting a GPU kernel takes a small amount of time.
- 3 **Parallel work:** the job must be large and regular enough to keep many GPU cores busy.

Small or irregular jobs may be faster on the CPU. A good GPU design must move enough connected work together and avoid unnecessary transfers.

5. The main ChoiceForge idea: fuse the pipeline

A less efficient pipeline might do this:

```
CPU creates inputs
-> GPU computes and stores every utility
-> another operation reads utilities and computes probabilities
-> another operation reads probabilities and selects choices
-> results return to CPU
```

ChoiceForge's fused linear kernel aims to do this:

```
chooser features + alternative weights + constants + availability + random draw
-> one fused GPU kernel
-> selected alternative + logsum
```

Inside the kernel, one GPU block handles one chooser. Threads in the block work on the alternatives together. They calculate utility scores, ignore unavailable options, find the maximum score, form safe exponential weights, sum them, select the option containing the supplied random draw, and return only the useful outputs.

The intermediate chooser-by-alternative utility matrix is never written to large global GPU memory. Temporary values stay close to the GPU cores in faster **shared memory**. This reduces memory traffic and storage.

The milestone-one kernel supports:

- a fixed set of 1 to 1,024 alternatives;
- linear utility equations using 32-bit floating-point numbers;
- availability rules;
- stable logsums;
- externally supplied random draws; and
- exact comparison of selected alternatives with a CPU reference.

It does not yet support all of ActivitySim's expression language, estimation mode, or destination sets with thousands of alternatives. Phase 7 supports the canonical MTC 21-mode nested-logit tree, but not every irregular ActivitySim choice model.

6. How the prototype is kept honest

Speed is useless if the model silently changes its answers. The project therefore starts with a readable NumPy CPU implementation called the **reference** or **oracle**. The GPU result is checked against it.

The validation rules are:

- give the CPU and GPU identical inputs and identical random draws;
- preserve the same alternative order and boundary rules;

- require every selected alternative to match exactly;
- measure the numerical difference between CPU and GPU logsums;
- test unavailable alternatives and invalid rows; and
- report mismatches instead of hiding them.

Small floating-point differences are normal because parallel GPU operations may add numbers in a different order. It is like adding a long list of rounded decimals from left to right versus pairing them first: the last digit can differ. The selected choices still must match, and logsum differences must stay within an explicit tolerance.

At the Phase 18 milestone, the Python 3.11 ActivitySim and CUDA integration environment passed 95 tests. The completed Phase 52 repository now passes 215 tests. These include exact comparison with ActivitySim's real Numba choice function, compact expression compilation, segmented destination batches, categorical flags, nested-logit validation, current-version fallback forwarding, real scheduling integration, GPU tests using 33 and 190 alternatives, a safe expression interpreter, skim-table adapters, shadow checks, the strict CPU reference, the generated strict CUDA evaluator, the explicit FP32 policy used for Phase 16, Phase 17's schema-safe plan and workspace reuse, Phase 18's fail-closed GPU state, Phase 48's resumed MT19937 state and exhaustive exponential correction, Phase 49's exact repeated-row device handoff, Phase 50's compact owner topology and wait-time reconstruction, Phase 51's fused compact-source utility and dense-device-ABI elimination checks, and Phase 52's persistent four-row service and lifecycle-release checks.

7. What was benchmarked?

A **benchmark** is a controlled timing experiment. ChoiceForge's first benchmark uses synthetic, meaning computer-generated, data. Each test has:

- 32 alternatives;
- 16 features per chooser;
- 10,000, 100,000, or 1,000,000 choosers; and
- 32-bit floating-point calculations.

It compares three timings:

- 1 **NumPy materialized:** the readable CPU reference creates the utility matrix.
- 2 **GPU including transfers:** inputs move to the GPU, the kernel runs, and outputs move back.
- 3 **GPU resident:** needed arrays are assumed to already be in GPU memory.

The GPU-resident number shows the potential of a future pipeline that keeps related model steps on the GPU. The transfer-inclusive number is the more conservative measure for a single isolated call.

All numbers below are medians after warm-up on an NVIDIA RTX A4000 with 16 GB of memory.

Choosers	NumPy CPU	GPU incl. transfers	GPU resident	End-to-end speedup	Resident speedup
10,000	5.156 ms	0.845 ms	0.361 ms	6.10x	14.26x
100,000	50.614 ms	3.587 ms	1.037 ms	14.11x	48.79x
1,000,000	499.375 ms	36.610 ms	8.118 ms	13.64x	61.51x

For all three sizes:

- choice mismatches: **0**;

- maximum logsum error: less than **0.000001**; and
- the fused kernel avoided allocating the global utility matrix.

At one million choosers, this original comparison was about 13.64 times faster. But NumPy was built as a readable correctness reference, not the strongest possible CPU competitor. Phase 1 therefore added a fused Numba implementation that spreads choosers over all 24 physical CPU cores.

Phase 1: comparison with the strongest CPU baseline

For the original 32-alternative, 16-feature shape, the 24-core fused CPU was much stronger:

Choosers	Best 24-core CPU	GPU incl. transfers	Transfer-inclusive result	GPU resident
10,000	0.390 ms	1.260 ms	GPU is 0.31x as fast	0.810 ms
100,000	5.560 ms	4.170 ms	GPU is 1.33x faster	1.190 ms
1,000,000	59.260 ms	34.520 ms	GPU is 1.72x faster	7.110 ms

This result is more realistic and less dramatic. The GPU loses on the smallest job because transfer and launch costs dominate. Its advantage grows with the amount of work.

Phase 1 also tested a shape based on mandatory tour scheduling in prototype MTC. It has 190 time alternatives and 69 feature rows:

Choosers	Best 24-core CPU	GPU incl. transfers	Speedup	GPU resident	Choice mismatches
10,000	10.760 ms	2.560 ms	4.19x	0.970 ms	0
100,000	109.660 ms	20.380 ms	5.38x	6.730 ms	2

The wider calculation gives each transferred row much more arithmetic, which suits the GPU. The two differences at 100,000 rows occurred when random draws were less than three ten-millionths from a probability boundary. NumPy chose one option while both fused Numba and CUDA chose the immediately following option. This tiny but important result means the project needs a clear 64-bit or mixed-precision rule before claiming exact ActivitySim equivalence.

The wider test also discovered and fixed a real GPU synchronization bug. The original 32 alternatives fit inside one group of 32 GPU threads, called a warp. With 190 alternatives, several warps shared temporary memory. One warp could reuse that memory before another had finished reading it. A new synchronization barrier fixes the race, and tests above 32 alternatives prevent it from returning.

Phase 2: replaying real ActivitySim scheduling data

Phase 1 copied the *shape* of tour scheduling, but its people and feature values were generated by a random-number program. Phase 2 uses real model records from ActivitySim's public `prototype_mtc` example.

ActivitySim's **mandatory tour scheduling** step chooses when work and school tours begin and end. For example, one option might leave at 8 a.m. and return at 5 p.m. The utility equation considers 52 to 59 active terms, including:

- the traveler's type, income, and work or school situation;
- the proposed start time, end time, and duration;
- whether another tour already occupies part of the day;
- a mode-choice logsum describing the quality of available travel modes; and
- constants that make some departure, arrival, and duration ranges more or less attractive.

The capture program observes ActivitySim at a documented internal boundary. It records the evaluated expression terms, feasible alternatives, coefficients, probabilities, random draws, and final selected positions. It does not change ActivitySim's rules or generate new draws.

The public sample produced six mandatory scheduling batches: first and second tours for work, school, and university students. Together they contain 4,477 tour choices. The largest batch contains 3,381 first work tours and 642,390 feasible chooser-alternative rows.

Did the answers stay the same?

Yes. All 4,477 GPU choices matched ActivitySim exactly. A separate 64-bit replay of ActivitySim's probability tables also had zero mismatches. The largest batch's maximum utility difference was less than 0.00000044. In the full model output, every tour's selected time alternative, start time, and end time matched the earlier warm Sharrow run.

Was the GPU faster?

The answer depends on where the input data begins:

Copies of the real batch	Choosers	CPU	GPU including transfers	GPU already resident	Resident speedup
1	3,381	8.241 ms	15.545 ms	0.961 ms	8.57x
2	6,762	20.346 ms	32.298 ms	1.326 ms	15.34x
4	13,524	43.035 ms	62.427 ms	2.242 ms	19.20x

For the larger rows, Phase 2 repeats captured real tours and their original draws. It does not invent synthetic travelers. Repeating records is a throughput test, not a claim that the public model contains more households.

The resident kernel clearly wins. The transfer-inclusive path clearly loses. The reason is that the current lowerer expands the expressions into a large table of term values: 151.6 megabytes for the native largest batch. Sending that table across the computer's PCIe connection takes much longer than the GPU arithmetic.

An analogy is a very fast bakery on the other side of town. Baking takes one minute, but delivering every separate ingredient takes fifteen minutes. Buying a faster oven will not solve the delivery problem. The next design must send compact ingredients once, prepare more of them inside the GPU, and keep them there for connected model steps.

What has Phase 2 proved?

It has proved that the ragged scheduling kernel can process the real ActivitySim alternative sets with exact choices, and that its resident computation is faster than a strong CPU boundary on this machine. It has also disproved the idea that simply adding a GPU call around an expanded term table makes the isolated component faster.

At that Phase 3 boundary, it had not yet proved that the entire mandatory scheduling component or full ActivitySim model ran faster. Phases 4 through 6 later added those wider proofs. Stateful timetable expressions and other ActivitySim front-end work still run on the CPU.

Phase 3: sending compact ingredients instead of finished terms

Phase 2 sent one value for every expression and every feasible alternative. Many values repeated the same traveler fact or the same time-alternative fact thousands of times. It was like shipping 59 finished ingredient bowls for every possible schedule.

Phase 3 sends a compact package instead:

- traveler facts are stored once per tour chooser;
- start, end, and duration are stored once for each of 190 time alternatives;

- each feasible row carries a small alternative ID;
- only the mode-choice logsum and seven timetable-dependent values stay row-specific; and
- ActivitySim still supplies the exact random draw.

The compact package for the largest real batch is 22.046 megabytes instead of 151.604 megabytes, an 85.46 percent reduction.

What does the compiler do?

A **compiler** translates instructions from one form into another. Here it reads the real ActivitySim expressions, checks that every operation and column is supported, and writes matching CUDA arithmetic into the kernel. It supports the scheduling model's numbers, arithmetic, comparisons, and Boolean combinations such as "A and B." Unknown names and unsupported function calls cause a clear error instead of a silent wrong answer.

The same expressions also generate a strong CPU version compiled by Numba and spread across 48 CPU threads. This makes the comparison much harder and fairer than comparing the GPU only with a simple Python loop.

Copies of the real batch	Choosers	Compact input	48-thread CPU	GPU including transfers	GPU resident	Inclusive speedup
1	3,381	22.046 MB	9.481 ms	2.478 ms	0.181 ms	3.83x
2	6,762	44.091 MB	18.443 ms	4.681 ms	0.271 ms	3.94x
4	13,524	88.179 MB	37.132 ms	9.348 ms	0.410 ms	3.97x
8	27,048	176.355 MB	69.735 ms	17.787 ms	0.709 ms	3.92x

All six mandatory-scheduling batches still have zero choice mismatches. The largest utility difference from the Phase 2 calculation is 0.0000229, and the largest GPU-versus-CPU logsum difference is 0.00000293.

Phase 3 proves transfer-inclusive superiority at this captured kernel boundary. It still does not time the work needed to build the compact tables inside ActivitySim, calculate mode-choice logsums, update the timetable, or map results back into pandas tables. Those steps belong in the next complete-component test.

Phase 4: putting the GPU inside the real ActivitySim component

Phase 4 performs that complete-component test. ActivitySim now has an explicit setting that chooses either its normal scheduling calculation or ChoiceForge. It is not a hidden test hook. ActivitySim still decides which alternatives are possible, calculates mode-choice logsums, supplies the controlled random numbers, updates each person's timetable, and writes the output tables.

The first integrated version was correct but slower. A profile showed that the GPU work was not the problem. Repeatedly asking seven timetable questions for 642,390 rows took about 4.03 seconds. A day in this model has only 21 time slots, so ChoiceForge now answers those questions once on a small traveler-by-time-slot grid and looks up the required answers. That reduced the largest timetable stage to about 0.331 seconds.

Two three-trial experiments measured the result:

Test	Cached Sharrow median	ChoiceForge median	Speedup	Schedule mismatches
Normal full-model runs	5.515 s	4.925 s	1.120x	0
Same saved checkpoint before scheduling	8.599 s	8.014 s	1.073x	0

The timer includes the work that Phase 3 left outside: mode-choice logsums, timetable calculations, compact packing, ActivitySim's random-number call, transfers, GPU work, result mapping, and timetable updates. It even includes fresh-process GPU setup. All three normal ChoiceForge runs produced exactly the same 9,806 final tdd, start, end, duration, destination-logsum, and mode-choice-logsum values as the cached-Sharrow reference.

Why is the complete-component speedup 1.120x instead of 3.83x? The 3.83x number describes the calculation ChoiceForge directly replaces. The complete component also performs substantial mode-choice logsum and table work that remains the same in both versions. Speeding up one part of a larger job gives a smaller improvement to the whole job. This is an example of **Amdahl's law**: total speedup is limited by the work that was not accelerated.

Phase 5: one backend for the scheduling family

Mandatory tours are only one kind of tour. A **non-mandatory tour** might be shopping, eating out, escorting someone, or recreation. A **joint tour** is shared by members of a household. An **at-work subtour** starts and ends at a workplace, such as going out for lunch. These components use related choice math, but they do not have identical columns or timetable rules.

Phase 5 taught the compiler to understand those differences without creating three separate GPU systems. It can turn text categories such as "escort" into compact yes-or-no columns, recognize both "equals" and "does not equal," and discover whether a timetable row belongs to a person or a tour. Calculations that are safe and repetitive go into the fused GPU kernel. Special state-changing timetable work remains explicit and checked.

Three complete-model trials measured all four ActivitySim scheduling timers:

Scheduling component	Cached Sharrow median	ChoiceForge median	Speedup
Mandatory	5.515 s	5.077 s	1.086x
Joint	0.671 s	0.504 s	1.331x
Non-mandatory	1.560 s	0.550 s	2.836x
At-work subtour	0.309 s	0.202 s	1.530x
All four together	8.045 s	6.325 s	1.272x

The four-component scheduling family uses 21.4 percent less time, saving a median 1.720 seconds. This was not caused by one lucky run: the slowest ChoiceForge trial was still faster than the fastest cached-Sharrow trial. Non-mandatory scheduling improves most because it contains a large amount of repeated choice work but does not carry the same expensive mode-choice-logsum work as mandatory scheduling.

Correctness is checked more strongly than before. For every Phase 5 trial, the final accessibility, households, joint-tour participants, land use, persons, tours, and trips files have exactly the same SHA-256 fingerprints as the cached-Sharrow reference. A SHA-256 fingerprint is a long number made from every byte in a file; changing even one character almost certainly changes the fingerprint. This covers 9,806 tour rows and 23,583 trip rows, not just a few selected columns.

There was an important Phase 5 limit. Its separately run sessions did not prove an all-model gain because unrelated run-to-run variation was larger than the scheduling savings. Phase 6 later closed that evidence gap with an alternating A/B experiment.

8. A negative result that teaches an important lesson

The project also tested only ActivitySim's final probability-sampling operation for 100,000 choosers and 32 alternatives.

Method	Time	Result compared with warm ActivitySim Numba
ActivitySim Numba on CPU	2.515 ms	baseline
GPU including transfers	2.591 ms	0.97x, slightly slower
GPU resident	0.212 ms	11.86x faster

There were zero choice mismatches. However, after paying to transfer the probability table, the isolated GPU call was slightly slower than ActivitySim's warm CPU function.

This is useful evidence, not a failure to hide. It shows that sending only the final tiny step to the GPU is a poor design. The strong performance idea depends on fusing utility evaluation, logsum calculation, and sampling, then keeping data on the GPU across connected steps.

9. How ActivitySim fits in

ActivitySim is a popular open-source activity-based travel modeling framework written in Python. ChoiceForge is designed as a possible specialized computing backend for parts of ActivitySim, not as a new travel demand model that replaces its behavioral rules.

The project integrated ChoiceForge with ActivitySim 1.4 and ran the public 25-zone prototype MTC example from start to finish:

- 5,000 households;
- 8,212 persons;
- 34 completed model steps;
- 70.568 seconds total runtime; and
- 580.8 MB peak unique memory.

That early 70.568-second run was only a framework baseline. Phase 4 later ran the full model with ChoiceForge handling mandatory scheduling, and Phase 5 expanded it to all four tour-scheduling components.

ActivitySim can also use **Sharrow**, an optimized expression system. Its compile run eventually completed in 27.3 minutes and used about 15 GB of unique memory. That compilation is not a fair production timing. After the cache existed, three warm runs took 61.650, 59.119, and 61.790 seconds, for a median of 61.650 seconds. Their final household, person, tour, and trip files were identical.

The warm profile showed where time is spent. Trip destination used 25.3% of total time. Mandatory tour scheduling used 9.1%, and trip scheduling used 8.4%. Mandatory scheduling is a practical first integration target, but it cannot reduce total runtime by 10% on its own because it is only 9.1% of the model. A reusable scheduling backend should target several scheduling components together.

Phase 6: making destination work large enough

Trip destination asks where an intermediate stop happens. For every candidate stop, the model also estimates how attractive the travel from the trip origin to the stop would be and how attractive the remaining travel to the main tour destination would be. Those are the two **directions** called OD and DP.

The first GPU replay exposed 30 small groups, separated by trip number and purpose. Sending each group to the GPU separately was a bad bargain: the GPU took about 15.032 milliseconds including transfers, while the CPU needed only 2.646 milliseconds. Think of sending 30 nearly empty delivery trucks instead of loading one truck.

ChoiceForge then packed all 30 groups into one segmented kernel launch. A segment label tells the kernel which purpose's coefficients to use, while offsets tell it where each traveler's uneven list of candidate places begins and ends. The packed job contained 3,971 choosers and 53,927 candidate rows. It took 0.847 milliseconds on the GPU including transfers versus 1.241 milliseconds for a strengthened, batched Numba CPU version: **1.464 times faster**. All 3,971 chosen positions matched, and logsums differed by at most about nine millionths.

The crossover test explains when to use it. At 28,570 rows the GPU was slower. At 39,706 rows it was faster. On this computer, this shape should be packed to roughly 35,000 rows before GPU dispatch. That threshold must be remeasured on other computers.

The first proven whole-model improvement

The live ActivitySim loop still asks for one small purpose group at a time, so turning on the new GPU sampler there would make it slower. Phase 6 instead removed duplicated work from the much larger directional-logsum calculation. It preserved ActivitySim's OD random draws, then its DP draws, but processed the shared deterministic work together.

Six fresh full-model runs alternated the old and new paths in the order A1/B1/A2/B2/A3/B3. Trip destination fell from a 16.307-second median to 14.238 seconds, **1.145 times faster**. All 34 model steps fell from 59.209 to 56.264 seconds, **1.052 times faster**. Every new-path component run beat every old-path component run. Seven final result files were byte-for-byte identical in every trial.

This distinction matters: the 1.464x number proves the packed CUDA kernel boundary; the 1.145x number proves the complete destination component; and the 1.052x number proves this complete example model. They are not interchangeable.

Phase 7: stop repeating setup, then accelerate the actual nest

Phase 7 asked what was still consuming the 14.238 seconds. ActivitySim repeated a 70-step preparation recipe once for each of ten trip purposes, even though most of that recipe was the same. ChoiceForge now puts the purposes for one trip number into one larger worksheet. The recipe runs three times - once for trip number 1, 2, and 3 - instead of 30 times. Sampling and final decisions remain separate, so each traveler keeps the same random draws and behavioral rules.

The next boundary is called **nested logit**. Imagine grouping 21 travel modes into folders: driving modes, walking/biking, transit, and ride-hailing. Some folders contain smaller folders. Nested logit combines the scores inside each folder and then combines the folders. It represents that two similar choices, such as two transit services, are more closely related than transit and driving.

We captured every real 21-mode score table used by destination choice: 30 batches and 107,854 rows. ActivitySim's dataframe reducer needed a median of 486.626 milliseconds for the complete captured sequence. The fused GPU kernel, including sending 18.119 MB to the GPU and returning the answers, needed 120.737 milliseconds in the longer 31-trial test: **4.030 times faster**. Even its slowest recorded GPU trial was faster than the fastest CPU trial. Its largest logsum difference was about 0.00000000000000036.

The full-model test alternated the Phase 6 and Phase 7 programs six times. Trip destination fell from 12.179 to 10.281 seconds, **1.185 times faster**. All 34 model steps fell from 54.459 to 52.166 seconds, **1.044 times faster**. Every Phase 7 run beat every Phase 6 run, and all seven final result files were byte-for-byte identical in all six trials.

Why is a 40x kernel only a 1.044x whole-model improvement? The kernel replaces less than half a second of CPU reduction. Most remaining time evaluates 404 behavioral expressions, samples destinations, runs other model components, and organizes data. This is Amdahl's law: speeding up a small slice cannot speed up the whole pie by the same amount.

Phase 8: current software and ten times more households

Earlier complete-model tests used 5,000 households. Phase 8 asked two harder questions. Does the integration still work with a pinned current ActivitySim development version? Does the advantage remain when the public workflow grows to 50,000 households across 190 zones?

The answer on this workstation is yes. The experiment alternated three regular ActivitySim runs with three ChoiceForge runs, always starting a fresh process:

Measured boundary	Regular ActivitySim median	ChoiceForge median	Speedup
Four scheduling components	30.1 s	23.7 s	1.270x
Trip destination	30.9 s	21.8 s	1.417x
All 34 model steps	147.225 s	131.812 s	1.117x

ChoiceForge saved a median of 15.413 seconds for the whole run. Every one of the three ChoiceForge runs was faster than every regular run. The seven final files matched byte for byte in all six trials. Those files describe 50,000 households, 111,130 people, 142,761 tours, and 350,751 trips.

The scale also shows what happens inside the computer. The largest scheduling call contains 9,561,750 feasible schedule rows for 50,325 choosers. Its compact input is 327.920 megabytes. The GPU part takes about 0.4 seconds, but preparing timetable facts, packing the data, retrieving controlled random numbers, and mapping the answer make the entire boundary about 3.3 seconds. This means the next scheduling improvement should focus on preparation and data reuse, not only on making the GPU arithmetic faster.

Phase 8 also captured a much larger nested-logit calculation: 30 real batches, 3,186,130 rows, and 535.270 megabytes of 64-bit mode scores. In 11 alternating trials, ActivitySim's CPU reducer needed a 4.646627-second median. The GPU, including both transfers, needed 0.125627 seconds: **36.988 times faster**. Every GPU trial beat every CPU trial, and the largest logsum difference was about 0.0000000000000053.

That 36.988x result is for one captured mathematical boundary. The 1.117x number is the honest complete-model result. The rest of ActivitySim still has to create people and tours, evaluate expressions, sample places, organize tables, and write outputs.

Phase 9: much larger geography and a useful stop sign

Phase 9 downloaded the public full-geography Prototype MTC data. The full data has 2,875,192 households, 7,566,527 people, and 1,454 zones. That is 7.7 times as many zones as Phase 8. A zone is a geographic area such as a neighborhood or traffic-analysis area, so more zones make the travel calculations much larger and more varied.

The whole population needs far more memory than this workstation has. The published configuration describes a 64-process, 432-gigabyte-RAM setup. The local test therefore used 50,000 households but kept all 1,454 zones. It ran one fresh regular ActivitySim process and one fresh ChoiceForge process.

For this test, ChoiceForge accelerated only trip destination. It left the schedule choices on the regular ActivitySim path. The final accessibility, household, participant, land-use, person, tour, and trip files were byte for byte identical. They describe 50,000 households, 132,536 people, 175,579 tours, and 442,682 trips.

Measured boundary	Regular ActivitySim	ChoiceForge	Observed result
Trip destination	39.2 s	28.0 s	1.400x faster
All 34 model steps	198.794 s	187.010 s	1.063x faster

This is one correctly matched pair, not yet a final statistical claim. A high-memory computer must run several alternating pairs before those numbers can be called reliable medians.

Phase 9 also found nine schedule choices that changed at a floating-point boundary when the experimental scheduling GPU path was enabled at 100,000 households. One later decision changed because the earlier difference affected the model's sequence of controlled random numbers. ChoiceForge therefore turns that scheduling path off for this configuration. This is not a failure to hide: it is exactly why the project checks complete output files rather than trusting a fast-looking timer. A faster answer that changes the model is not acceptable.

10. What the current evidence proves

The project can currently support these narrow statements:

- The fused linear-choice GPU kernel runs correctly on the tested NVIDIA hardware and software setup.
- It produced the same choices as the CPU reference in the reported synthetic benchmarks.
- Its logsums were numerically extremely close to the reference.
- For the reported large synthetic workload, it was much faster than the readable, materializing NumPy reference, even when transfers were included.
- For the scheduling-shaped 10,000-row workload, it was 4.19 times faster than a fused 24-core Numba CPU including transfers, with zero choice mismatches.
- The GPU advantage depends strongly on problem shape and size: it lost to the fused CPU on the 10,000-row, 32-alternative workload.
- Avoiding a full utility matrix can reduce intermediate memory use.
- ActivitySim-compatible random draws and final sampling rules can be preserved.
- Offloading sampling alone is not worthwhile when transfer time is included.
- Warm Sharrow production runtime and component shares are now measured reproducibly.
- Phase 2 replayed 4,477 real mandatory-tour choices with zero GPU mismatches.
- The real scheduling kernel is 8.57 times faster when inputs are GPU-resident, while the transfer-inclusive path is slower at 0.53 times CPU speed.
- Phase 3 reduces that input by 85.46 percent and makes the native transfer-inclusive GPU boundary 3.83 times faster than a generated 48-thread CPU implementation.
- Phase 4 makes the complete mandatory scheduling component 1.120 times faster than cached Sharrow in normal runs and 1.073 times faster in matched-checkpoint runs.
- Phase 5 makes the complete four-component scheduling family 1.272 times faster, led by a 2.836 times non-mandatory scheduling speedup.
- Seven substantive final CSV files are byte-identical across all Phase 5 trials, including all 9,806 tours and 23,583 trips.
- Phase 6 makes the packed real destination kernel 1.464 times faster than a batched Numba CPU including transfers, with zero mismatches across 3,971 choices.
- Phase 6 makes the complete trip-destination component 1.145 times faster and the complete 34-step prototype model 1.052 times faster in interleaved trials.
- Seven substantive final CSV files are byte-identical across all six Phase 6 A/B trials.
- Phase 7's transfer-inclusive nested-logit reducer is 4.030 times faster than ActivitySim's pandas reducer at the captured aggregate boundary, with maximum absolute error of $3.6e-15$.
- Phase 7 makes trip destination 1.185 times faster and the complete example model 1.044 times faster than Phase 6; every optimized run beats every baseline run.

- Phase 16 makes the generated-utility trip-destination component 1.025 times faster in three repeated 50,000-household public pairs; every candidate destination time beats every baseline time and every modeled decision matches.
- Phase 17 strengthens that result to a 1.040-times destination speedup in five repeated pairs, with a positive 0.8-to-1.9-second bootstrap interval and zero changed modeled decisions.
- Phase 17's complete-model median is 1.006 times faster, but its bootstrap interval still includes a 0.042-second slowdown, so the strict whole-model gate remains failed.
- Seven substantive final CSV files are byte-identical across all six Phase 7 A/B trials.
- Phase 8's 3,186,130-row nested-logit replay is 36.988 times faster on the GPU including transfers, with maximum absolute error of 5.4e-15.
- On pinned current ActivitySim at 50,000 households, the four scheduling components are 1.270 times faster, trip destination is 1.417 times faster, and all 34 models are 1.117 times faster.
- All three Phase 8 optimized runs beat all three baselines, and seven substantive CSVs are byte-identical across all six runs, including all 142,761 tours and 350,751 trips.
- On public full MTC geography with 1,454 zones, Phase 9's destination-only GPU configuration has byte-identical outputs in one 50,000-household A/B pair and an observed 1.400x trip-destination ratio. It is a scale gate, not a median claim.
- The Phase 9 scheduling experiment found nine changed choices at 100,000 households, so that GPU path is disabled for full geography until a precision safeguard resolves them.

11. What the current evidence does not prove

It does **not** yet prove that:

- every ActivitySim model will run faster;
- every complete model will be faster;
- whole-model superiority for the Phase 17 generated-utility backend; its median is faster, but two nearly tied pairs and an interval including zero keep the strict gate from passing;
- a complete model will be 13.64 times faster;
- arbitrary ActivitySim utility expressions can run in the kernel;
- destination choice with thousands of irregular alternatives is solved;
- results will be identical on every GPU and software version; or
- a faster calculation makes the behavioral model more accurate.

The NumPy reference remains useful for correctness, and Phase 1 includes a strong fused 24-core CPU competitor. Phase 2 supplies the real scheduling replay, Phase 3 supplies transfer-inclusive scheduling-kernel superiority, Phase 4 supplies the first complete-component comparison, Phase 5 supplies breadth, Phase 6 supplies a destination replay, Phase 7 supplies batched setup plus a live nested-logit GPU boundary, Phase 8 supplies pinned-current 50,000-household evidence, and Phase 9 supplies a full-geography scale gate. A broader claim still requires a differently structured public model, other GPUs, high-memory repetition, and independent reproduction.

12. From promising prototype to convincing result

A responsible roadmap looks like this:

Step 1: Expand the supported math - completed for four scheduling components

Phase 3 compiles the real scheduling arithmetic and Boolean expressions directly from compact traveler and alternative columns. Phase 5 adds categorical assignments, inequalities, and different timetable owners. Stateful timetable functions remain precomputed primitives, and other ActivitySim components may require more operations.

Step 2: Build a configured ActivitySim backend - completed for the tour-scheduling family

Phase 4 adds an explicit setting, preserves ActivitySim-owned random draws and timetable updates, and falls back to ActivitySim for unsupported tracing or estimation paths. Phase 5 reuses that contract for four scheduling components.

Step 3: Handle large destination choices - completed for the prototype model

Phase 6 packs 30 real ragged destination segments into one launch and proves a transfer-inclusive GPU win. Phase 7 batches live preprocessing across purposes and adds the 21-mode nested-logit GPU reducer. Much larger regional alternative sets remain future work.

Step 4: Prove correctness on public benchmark data - completed for the prototype scheduling family

Phase 2 checks all six prototype MTC mandatory-scheduling batches and records every numerical tolerance. Phase 5 verifies byte-identical final files after four scheduling components. Phase 6 checks all 30 destination batches and hashes all six A/B model trials. Phase 8 repeats the proof at 50,000 households on current ActivitySim. Phase 9 repeats the destination test across 1,454 public MTC zones and disables the scheduling path when its exactness gate fails. A differently structured public model is still needed before making a general claim.

Step 5: Run fair performance comparisons - completed for four scheduling components

Phase 5 compares three cached-Sharrow and three transfer-inclusive ChoiceForge scheduling runs. Phase 6 strengthens the design with alternating A/B processes and proves both component and whole-model gains. Other models and hardware still need the same treatment.

Step 6: Publish a reproducible evidence package

Include exact commands, pinned dependencies, raw timing samples, correctness reports, benchmark data instructions, and results from more than one GPU. Let other researchers reproduce or challenge the claim.

13. Why faster travel modeling could matter

Faster computation does not automatically make a better transportation decision. But it can change what analysts have time to study.

More scenarios

Instead of testing only a few projects, an agency could examine more combinations of fares, services, land use, road pricing, and policy assumptions.

Better uncertainty analysis

No forecast is certain. Faster models can support many runs with different population growth, fuel prices, remote-work rates, or random seeds. The range of results may be more informative than one forecast.

Faster feedback

Modelers could find data or configuration problems sooner, shorten development cycles, and respond more quickly to community questions.

Lower intermediate memory pressure

Fused calculations that avoid giant temporary tables may allow larger problems or reduce memory bottlenecks.

Important cautions

Speed cannot repair biased survey data, missing travel options, unrealistic assumptions, or unfair project goals. GPUs use electricity and require specialized hardware and skills. Public decisions still need transparency, equity analysis, local knowledge, and human judgment.

14. Frequently asked questions

Does the GPU predict travel differently?

It should implement the same model rules. The goal is to calculate them faster, not invent a different behavior model.

Does zero choice mismatch mean every number is identical?

Not always. Some kernel-level logsums differ by tiny amounts because floating-point arithmetic can be ordered differently in parallel, even when the selected alternative is unchanged. In Phases 5, 6, and 8, however, all seven substantive final CSV files were byte-for-byte identical in the reported A/B trials.

Why use 32-bit numbers?

They use half the memory of 64-bit numbers and are usually faster on GPUs. Whether they are accurate enough must be tested for each model. ActivitySim-compatible final sampling also has a 64-bit path in the adapter.

Why not move the entire model to the GPU immediately?

ActivitySim contains varied expressions, tables, joins, tracing, estimation tools, and irregular choice sets. A small, well-tested component creates a safer foundation than a large rewrite whose errors are hard to locate.

Is a GPU always faster?

No. Small workloads, frequent transfers, branching logic, or underused cores can make a GPU slower. The isolated sampling benchmark demonstrated this directly.

Which speedup is the headline?

Use the number for the boundary being discussed. Phase 3's scheduling kernel is 3.83x, Phase 7's smaller nested-logit reducer is 4.030x, and Phase 8's much larger nested-logit replay is 36.988x. At the 50,000-household scale, Phase 8 makes the four complete scheduling components 1.270x faster, complete trip destination 1.417x faster, and the complete model 1.117x faster. Phase 9's 1.400x destination ratio uses much larger 1,454-zone geography, but is only one matched pair. Only a carefully repeated full-model number should be called a whole-model speedup.

15. Mini glossary

Term	Plain-English meaning
ActivitySim	An open-source framework for activity-based travel demand models.
Amdahl's law	The idea that total speedup is limited by parts of a job that were not accelerated.
Alternative	One possible choice, such as walk, bus, or car.
Availability	A rule saying whether an alternative is possible.
Benchmark	A controlled experiment that measures speed and correctness.
Chooser	The person, household, tour, or trip making a simulated decision.
Compiler	A program that translates checked instructions into another executable form.
Compact representation	Data stored without unnecessarily repeating the same facts.
CPU	A flexible general-purpose processor with a modest number of powerful cores.
CUDA	NVIDIA's platform for programming its GPUs.
Feature	An input fact, such as travel time, cost, income, or vehicle access.
Floating point	The computer's approximate way of storing non-whole numbers.
Fusion	Combining connected calculations so intermediate data is not repeatedly stored or moved.
GPU	A processor with many simpler cores suited to large parallel calculations.
GPU kernel	A small program executed by many GPU threads.
Logsum	A summary of the attractiveness of all available alternatives.
Nested logit	A choice calculation that groups related alternatives before combining them.
Probability	A number from 0 to 1 describing how likely an outcome is.
Shared memory	Small, fast GPU memory close to threads in one block.
Skim	A table of travel time, distance, cost, or related values between places.
Synthetic population	Anonymous simulated households and people resembling a real region.
Utility	A model score representing how attractive an alternative is.

16. The bottom line

ChoiceForge asks a focused question: can a common bundle of travel-choice calculations be redesigned as one correctness-first GPU kernel that avoids huge intermediate tables?

The answer is now more precise. The fused kernel can decisively beat a strong 24-core CPU when each chooser has enough alternatives and features, but it loses on small, narrow jobs. Phase 1 also found a multi-warp synchronization defect and two near-boundary choice differences, demonstrating why wider correctness tests and precision rules matter as much as speed.

Phases 3 through 9 advanced integration through full MTC geography. Phase 8 improved whole-model time with exact outputs. Phase 9 confirmed exact destination results across a far larger zone system and held scheduling back until precision improves. Phase 11 then repeated that full-geography destination test and made the outcome stronger. Phase 12 identified the utility-calculation opportunity. Phase 13 completed the exact CPU answer key. Phase 14 generated the

matching GPU evaluator and proved exact agreement on real public-model batches, while keeping it safely in shadow mode.

17. The latest evidence: replicated destination speedup and a safer path to GPU utilities

This section updates the earlier phases. It uses two labels deliberately:

- **Proven production result** means a complete public ActivitySim run finished, its final files matched exactly, and the timing was repeated.
- **Foundation or microbenchmark** means a smaller, controlled engineering test succeeded. It is useful evidence, but it is not yet a claim that a whole ActivitySim model got faster.

Phase 10: a guardrail, not a shortcut

The full-geography scheduling experiment revealed a subtle risk. A GPU and CPU can both follow the same-looking formula yet round a last digit differently. If a random draw lands exactly near the boundary between two choices, that last digit can change the selected alternative. One changed early choice can also change the controlled sequence of later random draws.

Phase 10 added a **shadow guard**. Think of it as a second answer key running beside an experimental route. The established CPU/Sharrow result remains authoritative. The GPU is allowed to report success only if its shadow comparison passes; otherwise the system deliberately returns the CPU result. This is slower than blindly using the GPU, but it prevents a silent model change while the arithmetic issue is investigated.

Phase 11: the current headline result

Phase 11 focused on the part that was already exact: trip destination. It used the public Prototype MTC full geography with 1,454 zones and 50,000 households. Three fresh, interleaved A/B pairs were run: regular ActivitySim, then ChoiceForge, with the order alternated to reduce the chance that unrelated computer activity picked the winner.

Measured boundary	Regular ActivitySim median	ChoiceForge median	Result
Trip destination	39.7 s	28.6 s	1.388x faster
All 34 model steps	202.492 s	190.380 s	1.064x faster

The complete run saved a median of 12.112 seconds. The three paired whole-model savings were 12.172, 12.112, and 7.596 seconds. A simple resampling check placed the median saving between 7.596 and 12.172 seconds for these three pairs. Every optimized run was faster than its paired baseline.

Most importantly, all six substantive final CSV output files were **byte-for-byte identical** in every pair. That is stronger than saying the choices merely looked similar: the saved files matched character for character. The experiment also pins the exact ActivitySim source revision, configurations, patches, and hashes, so another person can repeat the same setup.

This is the honest current headline: on this workstation, for this public 50,000-household full-geography workload, ChoiceForge made the complete model 1.064 times faster while reproducing its reported final outputs exactly. It does not mean every ActivitySim model, computer, or future version will gain the same amount.

Why the next bottleneck is utility calculation

Destination choice works in two broad stages. First, the model evaluates many utility equations: for each potential destination it combines travel time, cost, household facts, destination facts, and rules. Second, it reduces those scores into probabilities and a selected destination.

ChoiceForge already makes the second stage faster. But the Phase 11 profiler showed that preparing and sending the first stage's finished utility values involved about 12,564,936 rows, or roughly 2.111 gigabytes of data. Sending a huge finished answer sheet to the GPU is like carrying every completed math problem to a very fast calculator. The more ambitious improvement is to send the compact ingredients and calculate the utility scores on the GPU too.

Phase 12: proving the ingredients can be translated safely

ActivitySim utility equations are written as expressions such as "travel time times a coefficient, plus an income term, only when a condition is true." A GPU cannot safely accelerate these expressions merely by guessing what they mean. The order of arithmetic, special functions, missing values, and table lookups all matter.

Phase 12 therefore built three safety layers:

- 1 An **expression reader** turns the public trip-mode-choice formulas into a small checked tree of operations. It successfully parses and evaluates all 253 distinct expressions used in that configuration on both CPU and GPU test paths.
- 2 A **skim adapter** gives the GPU path access to travel-time and cost lookup tables in the same way the model asks for them. Its tests confirm the adapter returns the expected values.
- 3 A **shadow capture** runs the GPU calculation beside the trusted CPU/Sharrow calculation without changing the official model answer. It records any difference for investigation.

The good news is that the deterministic GPU kernel exactly matched the local ordered CPU expression evaluator in the shadow test. The caution is that a compiled path can use different input precision, operation grouping, reduction order, or optional speed transformations such as `fastmath`. Any of these can change rounding in the last few binary digits. The shadows found differences from tiny fractions up to 0.25 for extremely large "unavailable" scores. No GPU utility result is allowed to replace Sharrow's production result until one explicit cross-device policy controls both targets.

A controlled speed test: encouraging, but not a production claim

Phase 12 also tested a deliberately simple 250,000-row, 64-feature, 21-alternative utility-and-choice pipeline. The CPU NumPy reference took 235.215 milliseconds. The GPU lowering plus 21-mode reduction took 28.988 milliseconds, an **8.114x** pipeline speedup, while the logsum check stayed within 0.0000000001.

This result answers a narrow question: can the compact-ingredient design be fast enough to matter? Yes. It does **not** yet prove an 8.114x whole-model improvement, because the test is not the full ActivitySim destination component and the CPU reference is not Sharrow's compiled production evaluator.

The strict IR plan: one recipe, two kitchens

The project is now pursuing a more rigorous solution rather than trying to tune away discrepancies one at a time. A **strict intermediate representation**, or strict IR, is a precise shared recipe for an equation. Instead of separately interpreting an expression for the CPU and generating a similar-looking CUDA kernel for the GPU, both targets are built from the same checked list of operations and numeric rules.

An everyday analogy: two kitchens can make the same named dish but use different measuring cups and a different order of steps. A strict recipe states the ingredients, measurements, order, and rounding rules so both kitchens produce the same dish. For ChoiceForge, the two kitchens are the strict CPU evaluator and the generated CUDA target.

The strict IR has generated a canonical description of the public MTC trip-mode utility: 379 terms across 21 alternatives. Phase 13 completed the strict CPU target, including separate ordered multiply and add steps, exact comparison reports, and fail-closed policy checks. Phase 14 completed the CUDA target generated from the same IR. That milestone passed 95 tests; the completed Phase 52 repository passes 215, including exact cross-device edge cases, compact skim gathering, a device-resident handoff to the nested-logsum reducer, the Phase 16 FP32 arithmetic policy, Phase 17 plan reuse, Phase 18 GPU-native runtime checks, Phase 48 probability and random-state checks, Phase 49 inter-stage packet checks, Phase 50 compact input-generation checks, Phase 51 fused-ABI checks, and Phase 52 persistent tiling and release checks.

This remains a project implementation, not a completed upstream Sharrow feature. Phase 15 removed the qualification-only utility transfer but failed its 50,000-household scale gate. Phase 16 recovered a repeated large destination-component win. Phase 17 added persistent plans and trip-mode continuation, strengthening that component win and making the five-run whole-model median faster. The strict path remains opt-in because the whole-model interval still includes zero and replication on other hardware and models is unfinished.

How to read the project status today

Question	Best current answer
Is there a real, repeated, exact speedup?	Yes. Phase 11 achieved 1.064x for the complete 50k-household, 1,454-zone public run, with byte-identical outputs in three pairs.
Is the destination component itself faster?	Yes. Its Phase 11 median was 1.388x faster.
Has the large utility equation been connected to the real GPU path?	Yes. Phase 17 reuses generated plans for destination and trip-mode utilities, with Sharrow fallback and explicit telemetry.
Is the compact GPU utility approach promising?	Yes. It was 8.114x faster in a controlled microbenchmark with a tight numerical gate.
What prevents a bigger claim today?	The destination component wins, but whole-model timing noise still crosses zero; a second GPU and a second public model have not been tested.
What changed in Phase 16?	An explicit FP32 policy, compact inputs, and caching produce a repeated 1.025x large destination-component win. The whole-model gate still fails.
What changed in Phase 17?	Checked persistent plans and trip-mode continuation produce a five-pair 1.040x destination win and a 1.006x faster whole-model median; the strict whole-model gate still fails.

18. Phase 13: building the exact CPU answer key

Phase 12 discovered that saying "use the same formula" was not precise enough. Two compilers may change when a number is rounded or combine a multiplication and addition. The answers usually remain close, but a travel model can place a random draw very near a choice boundary. That makes the last few digits important.

Phase 13 turns the arithmetic rules into a written and executable contract. The strict CPU evaluator is like an official answer key. It requires:

- 64-bit arithmetic while an expression is being calculated;
- one conversion of the finished feature to a 32-bit number;
- 32-bit coefficients and utility totals;
- the original term order;
- a separate multiplication and addition for every term; and
- normal IEEE rounding, with speed shortcuts disabled.

If the IR version, policy, or identifying hash changes unexpectedly, the evaluator refuses to run. It also refuses unresolved coefficients and malformed arrays. These are useful failures: they stop a plausible-looking but undefined calculation from entering the model.

Did it cover the real model?

Yes. The canonical public MTC utility contains 379 terms for 21 travel-mode alternatives. Every term and alternative executes under the strict policy. An independent, simple scalar loop produced exactly the same utility bits. The completed Phase 52 repository passes 215 tests.

Phase 13 then ran the public full-geography model with 1,001 households. It observed 30 real trip-mode batches containing 85,126 rows. That meant comparing 32,262,754 individual feature values and 1,787,646 utility values. ActivitySim completed all 34 model steps normally in 95.511 seconds, and Sharrow remained the official source of every model answer.

What did the comparison teach us?

The new strict answer key and current Sharrow did not match exactly. They agreed on 99.9118 percent of feature cells and 66.8085 percent of utility cells. The largest feature difference was about 0.0000305. The largest utility difference was 0.25, occurring among very large scores where a float32 number has relatively wide spacing.

Those percentages do not mean Sharrow is "66.8 percent correct." Sharrow follows its existing compiled arithmetic, while the strict evaluator follows the newly published cross-device policy. Many different utility cells can result from one tiny feature difference or a different order of hundreds of additions. Phase 13's job is to expose and classify that fact, not silently declare one existing compiler's behavior universal.

Every observed difference now has a stage. The reports found 28,466 feature cells whose compiled-expression result did not follow the strict expression policy. Among 593,346 different utility cells, 183 were explained by those input differences alone, while 593,163 also reflected a different accumulation rule. The first feature difference was about 0.00000763; the first utility difference was one float32 step.

Why Phase 13 is successful even though Sharrow differs

The goal was to create one stable answer key for future CPU and GPU code, not to copy whatever rounding happens on this workstation. All 379 terms and 21 alternatives run under that answer key. The comparison gate identifies the first different row, expression, alternative, values, and arithmetic stage. It can operate in observation mode, where Sharrow remains authoritative, or exact mode, where any difference stops qualification.

Phase 14 has now generated CUDA from this same strict IR and matched the strict CPU arrays exactly. Current Sharrow remains the safe ActivitySim fallback. The project will attempt another complete-model performance claim only after the generated path is integrated without qualification overhead and passes repeated byte-identical trials.

19. Phase 14: giving the GPU the exact same recipe

Phase 13 made the official answer key. Phase 14 built a code generator that turns that same checked recipe into a CUDA program. This matters because two separately handwritten programs can look equivalent while making tiny different choices about types, rounding, or operation order. Generating both evaluators from one hashed recipe greatly reduces that ambiguity.

What does the generated GPU program do?

For each model row, one GPU work group evaluates all 379 travel terms. It stores those temporary feature values in fast shared memory. Then separate GPU workers calculate the 21 travel-mode scores. Each worker uses the coefficients in the original order and performs a separate 32-bit multiplication and addition for every term. The resulting 21 scores can remain on the GPU and go directly into the nested-logsum calculation instead of taking a round trip through ordinary computer memory.

Inputs are packed by their real meaning: decimal numbers, whole numbers, and true/false values are not silently mixed. The compiled program is cached using the recipe hash, input types, coefficients, generated source, and compiler rules. Change any important ingredient and ChoiceForge builds a different kernel instead of reusing a stale one.

A tiny-number rule discovered by testing

Phase 14 found a useful edge case. A 32-bit floating-point number can be so tiny that it enters a special range called **subnormal**. The NVIDIA runtime on this machine turns such numbers into signed zero during these calculations. This is called **flush to zero**.

Ignoring the behavior would leave the phrase "exactly the same arithmetic" with a hidden exception. Instead, strict IR version 3 publishes the rule. The CPU answer key now also turns float32 subnormals into signed zero after the same steps. Tests deliberately use subnormal numbers, NaN ("not a number"), and infinity. This makes the cross-device agreement a real contract rather than an agreement that only holds for convenient inputs.

Did it match on the public model?

Yes. The same public 1,001-household, full-geography ActivitySim workload produced 30 trip-mode batches and 85,126 rows:

Exact Phase 14 check	Result
Batches	30 of 30
Feature values	32,262,754 of 32,262,754
Utility values	1,787,646 of 1,787,646
Largest feature difference	0.0
Largest utility difference	0.0
Terms and alternatives per batch	379 and 21

Every checked GPU bit matched the strict CPU answer key. ActivitySim still used Sharrow's established output as the official answer, so the experiment could not alter model behavior. That is the safe way to qualify a new backend.

Is it faster?

The generated kernel's diagnostic median was about 5.868 milliseconds. Preparing and transferring its inputs took about 116.157 milliseconds, and downloading the large comparison results took about 1.593 milliseconds. These numbers are not a fair speed contest. Phase 14 intentionally captures every one of the 379 feature columns so it can prove exactness term by term, while production code would keep compact inputs and useful outputs on the GPU.

Therefore the honest result is: **Phase 14 proves exact generated CPU/GPU semantics, not a new end-to-end speedup.** A separate test already confirms that the generated 21-column utility matrix can stay on the GPU through the nested-logsum reducer with no utility download and no reducer re-upload.

What does this change?

Before Phase 14, different arithmetic rules blocked the more ambitious GPU utility path. That obstacle is now resolved for this IR and hardware stack. That was the remaining Phase 15 problem: remove shadow-only transfers, connect generated utilities to the real path, and repeat public-model A/B runs. The next section explains what happened when those tests were completed.

20. Phase 15: putting the exact GPU recipe into the real model

Phase 14 proved that the CPU and GPU could follow the same recipe. Phase 15 asked a harder question: can the GPU's answer actually replace Sharrow's answer inside ActivitySim, stay on the graphics card for the next calculation, and make the real model faster?

The first connection worked, but exposed a data problem

The first candidate made all 379 ingredients as ordinary columns before the GPU kernel ran. Imagine copying every number from a library of road-time maps onto one enormous worksheet, sending the worksheet to a calculator, and then throwing it away. The arithmetic was fast, but preparing the worksheet was not.

At 1,001 households, the first repeated combined test looked faster than plain ActivitySim. A fairer test then compared Phase 15 directly with Phase 11, so both sides received the same destination batching and GPU nested-logsum work.

That test showed the first strict utility path itself was slower. Measuring the parts found the cause: looking up and materializing skim columns took about 15.35 seconds across three candidate runs, while the useful GPU work took much less time.

The compact skim design

ChoiceForge was changed so the generated kernel receives:

- the shared road-time and cost map cubes;
- the already-checked origin, destination, and time positions for each row;
- the smaller ordinary chooser columns; and
- the strict coefficients and recipe.

The kernel now looks up a needed map value directly. It does not build a giant 379-column feature table. Identical map cubes are uploaded only once, even when ActivitySim creates new temporary wrapper objects. After trip destination is finished, the cache is released so later model steps do not pay for unused GPU memory. If any part fails, ChoiceForge uses Sharrow instead.

Did compact GPU gathering stay exact?

Yes. The final qualification covered 30 real public-model batches:

Check	Exact result
Rows	85,126
Feature values	32,262,754 of 32,262,754
Utility values	1,787,646 of 1,787,646
Largest strict CPU/GPU difference	0.0
Utility downloads before nested logsum	0 bytes
Nested-logsum utility re-uploads	0 bytes

Every modeled trip decision also matched. A printed diagnostic called `destination_logsum` differed by at most 0.000008 at this scale. That field is not used as a later decision in this output check. The difference comes from the published strict arithmetic versus current Sharrow arithmetic; the strict CPU and GPU answers still match bit for bit.

The direct repeated speed result

Three fresh pairs compared Phase 11 with compact Phase 15 at 1,001 households. The destination component improved from a median 11.3 seconds to 10.3 seconds, or **1.097 times faster**. Every Phase 15 destination run was faster than every Phase 11 destination run. All 90 strict utility batches stayed on the GPU for the nested calculation, and all modeled decisions matched.

The complete model did not pass the stricter promotion rule. Its median was 84.631 seconds for Phase 11 and 85.445 seconds for Phase 15. Other small model steps varied enough to hide the one-second destination gain. ChoiceForge therefore claims component superiority at this qualification scale, not a new whole-model record.

What happened at the large scale?

The final diagnostic used 50,000 households, all 1,454 zones, and 4,188,312 utility rows. Decisions still matched, but compact Phase 15 took 33.3 seconds for trip destination versus 28.4 seconds for Phase 11. The complete model took 197.871 versus 192.519 seconds. This is a clear rejection, so running three more expensive pairs would not create a responsible speed claim.


The reason is now narrower and useful. The GPU kernel gives one block to each row and makes many scattered reads from large skim cubes. At 50,000 households, that access pattern is slower than Sharrow's compiled expression path. The next compiler should process tiles of nearby rows, combine repeated map reads, reuse a gathered value across multiple terms, and keep the exact ordered arithmetic rules.

What should a high-school reader conclude?

An experiment can succeed even when the final answer is "do not deploy this version." Phase 15 produced working, exact, fail-safe GPU integration and a repeated component win. It also used a public large benchmark to prove where the design stops winning. The production headline therefore remains Phase 11's repeated 50,000-household result. Phase 15 is the tested bridge to a better future compiler, not an excuse to replace a faster working system today.

21. How the complete calculation fits together

The earlier sections introduce each idea separately. This is the whole trip destination calculation in one view:



Phase 11 accelerates the preparation and nested-logsum part of this pipeline. Phase 15 additionally tries to calculate the 379 utility terms on the GPU and hand their 21 results directly to the GPU nested-logsum calculation. Keeping that handoff on the GPU avoids downloading a large utility table to ordinary memory and immediately uploading it again.

What exactly is a skim cube?

A **skim** is a lookup table containing travel time, distance, cost, or another measure between two zones. Imagine a map grid: choose an origin row and a destination column, and the cell tells you the travel time. A model needs many such maps for different travel modes and times of day. Stacking the maps makes a three-dimensional **skim cube**.

For every possible trip, the kernel receives the already-checked origin, destination, and time positions. It uses them like coordinates to read the needed cells. At small scale this avoids building a giant worksheet and is faster. At large scale, millions of rows may ask for far-apart cells. Those scattered reads are like repeatedly fetching books from unrelated library shelves: the arithmetic workers spend too much time waiting for data.

The next design should group nearby requests into **tiles**, reuse a fetched skim value when several utility terms need it, and arrange memory reads so neighboring GPU workers fetch neighboring data. This is a data-movement problem, not evidence that the equations are too difficult for a GPU.

22. Three different meanings of "the answers match"

The word **exact** can describe different boundaries. A responsible report must say which boundary it tested.

Level	Plain-English question	Phase 15 result
Arithmetic	Did the strict CPU and generated GPU follow the same numeric recipe bit for bit?	Yes: all 32,262,754 checked feature cells and 1,787,646 utility cells matched exactly.
Decisions	Did the model choose the same alternatives?	Yes: every modeled trip decision matched in the reported runs.
Saved files	Were complete output files identical byte for byte?	All six non-trip substantive files were byte-identical. Every modeled field in <code>final_trips.csv</code> matched. The reported <code>destination_logsum</code> diagnostic was bounded, not byte-identical.

The small `destination_logsum` difference is between current Sharrow arithmetic and the newly published strict arithmetic. It was at most 0.000008 for 1,001 households and 0.00001 for 50,000 households, below the 0.0001 gate. It is not a disagreement between the strict CPU and GPU, which matched bit for bit.

This distinction also explains why Phase 11 has the strongest production replication statement. In its repeated 50,000-household trials, all substantive final CSV files were byte-identical. Phase 15 has exact strict cross-device arithmetic and exact decisions, but one diagnostic column intentionally exposes the known difference from Sharrow's arithmetic order.

23. The rule for promoting an experiment

One fast run can be luck. The computer may be warmer, another program may have used a processor core, or a file may already be cached. ChoiceForge therefore uses fresh processes and alternates conditions: A1, B1, A2, B2, A3, B3. A is the established implementation and B is the candidate. The **median** is the middle of three timings, so one unusually fast or slow run has less influence.

For automatic promotion, all of these conditions must pass:

- 1 Run at least three fresh, interleaved A/B pairs.
- 2 Preserve every modeled decision and pass the complete-output correctness gate.
- 3 Make the median trip-destination time faster.
- 4 Make the median complete-model time faster.
- 5 Achieve complete separation: every candidate time must beat every corresponding baseline time.

Phase 15 qualification gate at 1,001 households	Result
Three fresh interleaved pairs	Pass
Exact modeled decisions	Pass
Faster destination median	Pass: 11.3 to 10.3 seconds
Every candidate destination run faster	Pass
Faster complete-model median	Fail: 84.631 to 85.445 seconds
Every candidate complete-model run faster	Fail

Phase 15 qualification gate at 1,001 households	Result
Final promotion	Rejected

The 50,000-household test used one pair. One pair cannot prove superiority, but it can reveal that a candidate is clearly unpromising: Phase 15 was slower at both the destination and complete-model boundaries. Repeating that losing configuration three more times would spend hours without fixing its identified memory-access problem. A redesigned kernel must return to the full three-pair gate before any new speed claim is made.

24. Hardware, memory, and the boundary of the claim

The numbers in this guide describe one measured system, not every computer.

Reproduction item	Recorded value
CPU	AMD Ryzen Threadripper PRO 5965WX, 24 cores and 48 threads
System memory	63.9 GB usable for the benchmark workstation
GPU	NVIDIA RTX A4000, 16 GB, compute capability 8.6
NVIDIA driver	571.59
Python	3.11.14, 64-bit
ActivitySim source	Commit 16ab11180a26912987eb902daf945e268f3efc11
Public geography	Prototype MTC Extended, 1,454 zones

The 50,000-household Phase 15 candidate reached 9,196 MiB of GPU memory. It did not run out of the A4000's 16 GB. The observed slowdown instead points to the one-block-per-row design and scattered skim reads. Another GPU with different memory bandwidth, cache sizes, or scheduling could behave differently, so the claim is deliberately limited to this hardware and software stack until other systems reproduce it.

The full public population contains 2,875,192 households. The upstream multi-process instructions call for hundreds of gigabytes of system memory, far beyond this workstation. A 50,000-household sample with all 1,454 zones is therefore the local public scale gate. It is large enough to expose the failed Phase 15 access pattern, but it is not a substitute for a full-population run on a high-memory server.

25. A practical reproduction checklist

Reproduction means more than downloading the code and seeing a program finish. Another researcher should be able to identify the same source, data, configuration, hardware, commands, outputs, and decision rule.

Pinned evidence

Item	Identifier recorded by the run
ChoiceForge Phase 15 evidence commit	7f40ce227e46a16801acd6abc264544457c97a7b
ActivitySim commit	16ab11180a26912987eb902daf945e268f3efc11
Prototype MTC data_full.tar.zst SHA-256	b402506a61055e2d38621416dd9a5c7e3cf7517c0a9ae5869f6d760c03284ef3
ActivitySim integration patch SHA-256	aa2eed95d99fe4853365b9ed49427722ac467ef928013ccf7944fd9eda0e04e2
Python environment lock SHA-256	84b97738100cd4b8c405af50ff823cdddf4a33f5f203deab7839e4a8739a3adc

Set up and verify the software

The detailed setup is in docs/activitysim-integration.md. In PowerShell, the important verification steps are:

```
git clone https://github.com/ActivitySim/activitysim.git tmp\activitysim-phase8-source
git -C tmp\activitysim-phase8-source checkout 16ab11180a26912987eb902daf945e268f3efc11
git -C tmp\activitysim-phase8-source apply ../../integration/activitysim-current-choiceforge.patch
uv venv --python 3.11 .venv-phase8
uv pip install --python .venv-phase8\Scripts\python.exe -e tmp\activitysim-phase8-source -e "[gpu, test]"
.\.venv-phase8\Scripts\python.exe -m pytest -q
```

Place the public full-geography data at benchmark-data/phase9-mtc-full/prototype_mtc_extended/data_full. Check its downloaded archive against the SHA-256 value above before extracting it. The benchmark scripts refuse to overwrite an existing run, so use a new short RunTag when repeating an experiment.

Run the exactness and performance gates

First run one qualification candidate. Then run three direct Phase 11 versus Phase 15 pairs:

```
.\scripts\run_phase15_candidate.ps1 -Households 1001 -RunTag reproduce-gate -MaxCandidateRows 2000000
.\scripts\run_phase15_incremental_ab.ps1 -Households 1001 -Repetitions 3 -RunTag reproduce-r3 -MaxCandidateRows 2000000
```

The candidate is explicitly enabled with CHOICEFORGE_STRICT_CUDA_CANDIDATE=1; it is off by default. The row policy CHOICEFORGE_STRICT_CUDA_MAX_ROWS can reject an ineligible batch before a large allocation. Any code-generation or GPU-reduction failure falls back to Sharrow, which remains authoritative.

Compare a new run with these machine-readable evidence files:

- benchmark-results/phase15-candidate-summary.json
- benchmark-results/phase15-p15finalr3-summary.json
- benchmark-results/phase15-p15compact50-summary.json

The manifests record source, patch, environment, data and configuration hashes, GPU and driver, Python version, individual timing samples, memory measurements, correctness results, and the final promotion decision. Matching the published median alone is not enough; the correctness and promotion fields must also pass.

26. The Phase 16 plan that was tested

Phase 15 showed that the next major opportunity was not to add more arithmetic to the existing one-block-per-row kernel. The approved Phase 16 plan was to change how data reaches that arithmetic and test each idea honestly:

- 1 Add measurement for skim-cache hits, memory transactions, and time spent waiting on gathers.
- 2 Sort or tile rows by origin, destination, and time without changing their final model order.
- 3 Coalesce neighboring skim reads and reuse one gathered value across every term that needs it.
- 4 Fuse tiled gathering, strict ordered utility accumulation, and nested-logsum reduction without writing a global feature or utility matrix.
- 5 Preserve strict IR version 3, typed inputs, default-off activation, bounded diagnostics, and Sharrow fallback.
- 6 Pass unit tests and exact CPU/GPU edge cases before running ActivitySim.
- 7 Qualify all 30 public-model batches and exact final decisions at 1,001 households.
- 8 Require three interleaved 50,000-household pairs and the full promotion rule before replacing Phase 11.
- 9 Repeat on at least one different GPU and one differently structured public ActivitySim model.
- 10 Only then propose the backend for upstream Sharrow or ActivitySim integration; until every gate passes, Phase 11 remains supported and Phase 15 remains opt-in research evidence.

27. Phase 16: the GPU finally wins the large target component

Phase 15 taught us that moving a calculation to a GPU does not automatically make it faster. Phase 16 used that failure as a map. It measured where time was going, tried several designs, rejected the slow ones, and found one policy that was both reproducible and faster for the targeted part of the public model.

Four ideas in ordinary language

First, constants no longer become giant repeated columns. If every row uses the same number, the GPU receives that number once. This is like writing a classroom rule on the board instead of printing the same rule on every student's paper.

Second, two names that point to the exact same column share one stored copy. Third, travel-time maps and compiled equation plans are cached so the program does not rediscover identical information for every batch. Fourth, an optional policy uses 32-bit floating-point arithmetic for the expressions instead of 64-bit arithmetic.

The last change matters on this NVIDIA RTX A4000. It can do ordinary 32-bit GPU arithmetic much faster than 64-bit arithmetic. Sharrow already stores each completed feature as a 32-bit number before combining features with coefficients, so Phase 16 publishes a separate 32-bit recipe and a separate CPU answer key for it. The older strict 64-bit recipe is still available and is not silently changed.

What is a numeric policy?

A **numeric policy** is a written recipe for how the computer handles numbers. It says how many bits a number uses, when rounding happens, whether operations may be rearranged, and how unusual values such as infinity are treated. Two programs can use the same equation but get slightly different last digits if their numeric policies differ.

ChoiceForge now has two named policies:

Policy	Expression arithmetic	Why it exists
Strict IR version 3	64-bit, then one 32-bit feature rounding	strongest continuation of the Phase 13-15 arithmetic contract
Phase 16 FP32	32-bit expression arithmetic and 32-bit feature storage	mirrors the deployed intermediate precision more closely and runs efficiently on this GPU

The run manifest records which policy was selected. The generated source hash also changes, so evidence from the two policies cannot be accidentally mixed.

Did the new CPU and GPU answers match?

Yes. On the real 1,001-household qualification:

Exact qualification question	Result
Real batches checked	30 of 30
Rows checked	85,126
Feature cells exactly equal	32,262,754 of 32,262,754
Utility cells exactly equal	1,787,646 of 1,787,646
Largest CPU/GPU difference	0.0
Fallback batches	0

The tests also include a simple equation where the 64-bit answer is 1 and the 32-bit answer is 0 because a very large 32-bit number cannot represent a change of just 1. That test proves the two policies really are different. It then proves that the published 32-bit CPU and GPU recipes agree exactly.

The large repeated performance proof

The public scale gate uses 50,000 households, all 1,454 zones, and 4,188,312 utility rows. It launches fresh processes in the order baseline 1, candidate 1, baseline 2, candidate 2, baseline 3, candidate 3.

Trip destination	Run 1	Run 2	Run 3	Median
Phase 11 baseline	28.5 s	28.5 s	28.8 s	28.5 s
Phase 16 FP32 GPU	28.4 s	27.7 s	27.8 s	27.8 s

Every candidate time beat every baseline time. The median speedup is 1.025 times, and the paired savings are 0.1, 0.8, and 1.0 seconds. Every modeled trip decision matched in all three candidate runs. The largest printed diagnostic logsum difference was 0.000010, below the 0.0001 limit.

This means the **component promotion gate passes**. It is a practical, repeatable GPU win for the part of the model Phase 16 was designed to replace.

Why is the whole model not promoted?

The full model contains many steps Phase 16 does not change. Their timings move slightly from run to run because of operating-system scheduling, file caches, CPU activity, and other noise. Whole-model medians were:

Complete model	Median
Phase 11 baseline	193.014 s
Phase 16 candidate	194.312 s

That is a 0.993-times result, which is slower, not faster. One candidate run was faster than its paired baseline, but two were slower. The **whole-model gate fails**. Reporting both gates prevents a real kernel success from being hidden and prevents that success from being exaggerated into an application-wide claim.

Slow ideas that were tested and rejected

Phase 16 did not simply keep the first plausible code:

- Loading 149 skim maps cooperatively for tiles of nearby rows was exact but synchronization and unnecessary eager loads made it slower.
- The model has only about 419-465 nonzero coefficients out of 7,959 positions. Skipping zeros sounds ideal, but one version created bulky control flow and another caused scattered shared-memory reads. Both were slower.
- Strict FP64 compaction reduced host packing from about 2.4 to 1.1 seconds and utility time from about 4.5 to 4.1 seconds, but destination still lost 32.5 to 30.1 seconds at large scale.
- FP32 reduced the large generated utility work to roughly 1.2-1.6 seconds and produced the repeated component win.

The lesson is that counting arithmetic operations is not enough. GPUs care about memory order, groups of threads following the same instruction, code size, and the kinds of number hardware they execute efficiently.

What Phase 17 needed to do

Phase 17 had three jobs: stop rebuilding the same GPU program, prevent that optimization from moving a delayed Sharrow cost into trip mode choice, and repeat the public benchmark enough times to separate a real destination win from ordinary timing noise. The next section explains what was implemented and what the measurements do and do not prove.

28. Phase 17: turn a kernel into a reusable backend

A fast kernel is like a fast kitchen appliance. It is not enough for the blender blades to spin quickly if someone must unpack, assemble, and wash the whole machine for every smoothie. Before Phase 17, ChoiceForge repeatedly did small pieces of setup around an already-fast generated kernel.

A compiled plan is a checked instruction packet

Phase 17 creates a **compiled plan**. The plan remembers:

- the exact equation recipe, identified by a cryptographic hash;
- which inputs are whole columns and which are single numbers;
- the data type and storage slot for every input;
- how skim lookups are indexed;
- the compiled CUDA kernel; and
- the coefficient matrix already stored on the GPU.

Reusing a plan is safe only when the next batch has the same shape of meaning. The row values may change. A price or coefficient value may change. But a column cannot quietly become a scalar, a floating-point value cannot quietly become an integer, and two different columns cannot quietly acquire a new shared-memory layout. If any of those rules changes, ChoiceForge refuses that plan and builds or selects a matching one. This is called **fail closed**: when the program is uncertain, it does not guess.

There is a subtle scalar rule too. Imagine two constants both happen to equal 2 today. Combining them into one slot would be wrong if tomorrow one becomes 3 and the other becomes 7. Persistent mode therefore gives semantic scalar inputs stable slots even when their current values are equal.

Why trip mode choice had to join the GPU path

An early experiment made destination faster but made the later trip-mode step about 2.4 seconds slower. The code had not created work from nowhere. It had stopped Sharrow from warming up during destination, so Sharrow paid its first compilation cost later during mode choice. This is a **displaced cost**: one timer looks better because the cost moved to another timer.

Phase 17 adds a continuation bridge. The same generated FP32 CUDA utility plans now serve trip mode choice. ActivitySim still owns the nested-logit probabilities, random numbers, final choices, and result tables. ChoiceForge replaces only the utility calculation. If the generated path encounters an unsupported case, it records the reason and calls the original Sharrow path. The reported runs had zero fallbacks.

The small exactness gate

On 1,001 public households, the final Phase 17 qualification checked:

Check	Result
Real destination batches	30 of 30 exact
Rows evaluated	85,126
Feature values	32,262,754 of 32,262,754 exact
Utility values	1,787,646 of 1,787,646 exact
Reused destination plans	20 calls
Reused trip-mode plans	10 calls

Check	Result
GPU fallbacks	0
Changed modeled decisions	0

Two saved columns are diagnostics rather than later decision inputs. Their last decimal places can differ because the declared FP32 GPU policy does not promise identical decimal text to Sharrow:

Diagnostic	Largest difference	Allowed gate
Destination logsum	0.000008	0.0001
Mode-choice logsum	0.00000191	0.00001

The output checker gives tolerance only to those two named columns. Every trip destination, trip mode, and other modeled output still has to match exactly.

The five-pair 50,000-household result

The clean proof used five fresh baseline/candidate pairs and fixed both BLAS thread settings at 16. A prior 24-thread attempt suffered a native OpenBLAS failure in a baseline process, so that incomplete series was discarded.

Measure	Baseline median	Phase 17 median	Result
Trip destination	28.4 s	27.3 s	1.040x faster
Complete model	191.474 s	190.307 s	1.006x faster

Every one of the five destination runs was faster than every baseline destination run. The five paired destination savings were 1.8, 0.8, 0.9, 0.8, and 1.9 seconds. A bootstrap calculation put the middle saving's 95% interval between 0.8 and 1.9 seconds. That is strong component-level evidence.

The complete model is more complicated. Its paired savings were 1.845, -0.042, 1.360, -0.038, and 2.595 seconds. Three won and two lost by less than one twentieth of a second. The median improved by 1.167 seconds, but the 95% bootstrap interval ran from -0.042 to 2.595 seconds. Because that interval still touches a small slowdown, the strict whole-model superiority gate does not pass. "The median was faster" is true; "whole-model superiority is proven" would be too strong.

Reusable buffers: useful, but still experimental

Creating GPU arrays also costs time. Phase 17 implements an optional workspace that lets a checked plan reuse its device input and output storage. A first version used pinned host memory. It uploaded faster but packed columns more slowly, so it was rejected. The retained version keeps fast NumPy column packing and reuses only device allocations.

That version passed the exact gate. In one 50,000-household diagnostic it reduced measured destination packing plus upload time from about 1,471.5 to 1,364.1 milliseconds, a saving of roughly 107 milliseconds. However, one pair cannot prove a whole-model effect, and unrelated initialization varied by more than two seconds in that run. Reusable buffers therefore remain an opt-in experiment, not part of the five-pair headline proof.

29. What the result means and what comes next

Here is the honest Phase 17 conclusion:

- the generated FP32 GPU backend exactly matches its published FP32 CPU answer key on every qualified feature and utility cell;
- all reported modeled decisions match the reference;

- repeated trip-destination superiority is proven on this RTX A4000 and public MTC workload;
- the five-run whole-model median is faster, but the strict whole-model proof is not yet complete; and
- the result has not yet been replicated on a second GPU or a different public ActivitySim model.

The fastest path to a stronger claim is to collect more independent, controlled 50,000-household pairs with the same locked software, hashes, and 16-thread setting. That narrows the uncertainty around the tiny whole-model losses. The most important scientific step after that is external replication: run the same exactness and A/B protocol on a second NVIDIA GPU and then on a second public ActivitySim model. A result that survives different hardware and different travel equations is much more useful to the ActivitySim and Sharrow communities than a single impressive number from one workstation.

30. Phase 18: can a chain of model steps live on the GPU?

After Phase 17, the obvious ambitious question was: what if the GPU did not calculate one isolated piece and then hand everything back to the CPU? What if household data entered the GPU once, several dependent model steps happened there, and only the finished answers came back?

Phase 18 builds the first honest version of that idea. It is not a whole travel model yet. It is a **vertical slice**, meaning a narrow but complete path from real input rows through multiple calculations to final outputs. Imagine building one fully working elevator before constructing every floor of a skyscraper. The elevator proves the important machinery can work together, but it does not mean the building is finished.

Does “GPU-only” really mean the CPU does absolutely nothing?

No. A regular computer must use its CPU to start Python, read a CSV file, read configuration, tell the GPU which kernel to launch, handle errors, and write a result file. Even a CUDA program is launched by host code.

In this project, **GPU-native modeled execution** has a precise meaning:

- The CPU may read input and upload one partition before modeling starts.
- The CPU may launch kernels and wait for them.
- The CPU may download final outputs after modeling finishes.
- The CPU may not calculate a utility, random choice, logsum, or modeled total.
- The program may not secretly use NumPy or pandas when a GPU operation is missing.

The runtime closes a gate called `seal_ingress`. After that gate closes, a host array entering a modeled stage, an intermediate modeled result leaving the GPU, or a CPU fallback causes a hard error. This behavior is called **fail closed**. If the program cannot honor its promise, it stops instead of quietly producing a misleading benchmark.

What is GPU state?

A travel model changes tables over time. A household first gets an auto ownership result. A person later gets a work location. Tours and trips are then created. Later steps depend on earlier answers. All those current tables are the model's **state**.

Phase 18 adds a `DeviceTable`. Its columns must be CUDA arrays, which means the actual numbers live in GPU memory. All columns must describe the same number of rows. A `GpuNativeRuntime` owns these tables and records what happens at the CPU/GPU boundary. It counts input bytes, output bytes, modeled transfers, CPU fallbacks, and kernel stages.

Why random numbers are part of correctness

Travel models use random draws to turn probabilities into simulated choices. Suppose household 42 gets random number 0.63 in one run. If changing the batch size gives it 0.18, the model can change even though no travel assumption changed. That would make scaling unsafe.

The new GPU random generator uses three stable labels:

```
entity ID + project seed + stream ID -> the same random draw
```

The entity ID identifies the household or person. The project seed identifies the run. The stream ID identifies the decision, such as first choice or second choice. A hash thoroughly mixes those integers, and the GPU turns part of the result into a number between zero and one.

The important fact is what the formula does **not** use: it does not use the row's position inside a partition. Household 42 therefore gets the same bits whether it is processed in one table of 2.8 million households or a small table of 250,000. A separate NumPy implementation checks the GPU generator bit for bit.

Why adding numbers can be nondeterministic

Floating-point addition is not perfectly associative. In exact mathematics, $(a + b) + c$ equals $a + (b + c)$. On a computer, rounding after each addition can make the last bits different.

Many fast GPU totals use **atomics**: many threads race safely to add their values to one total. The final answer can depend on which thread arrives first. Phase 18 instead sorts group IDs and gives one thread responsibility for each group. That thread adds rows in a fixed order. It may not be the fastest possible future method, but its behavior is easy to explain and reproduce.

What the public-data vertical slice does

The benchmark reads all 2,875,192 households from the public Prototype MTC table. It then performs this chain after the GPU boundary closes:

- **Stage 1:** Build eight household features, such as normalized household size, workers, automobile count, income, and zone information.
- **Stage 2:** Generate the first stable random stream on the GPU.
- **Stage 3:** Make a fused choice among 21 alternatives.
- **Stage 4:** Put that first answer into the inputs of a second choice.
- **Stage 5:** Generate a different stable stream and make the dependent second choice.
- **Stage 6:** Sort households by traffic analysis zone, or TAZ, and total the first choices within each zone.
- **Stage 7:** Download the final answers for checking.

The word **dependent** matters. The second choice really uses the first choice. This is not a collection of unrelated kernels that happen to run one after another.

Which parts are real, and which are invented for the test?

The households and their fields come from the real public MTC benchmark. The 21 alternatives and eight-feature calculation have the shape of travel-model choice work. However, the coefficients are fixed synthetic test values. Nobody estimated them from a survey, so these choices must never be interpreted as a forecast of real behavior.

The data transform also publishes its assumptions:

- Negative income codes mean missing for this systems test and become zero.
- Income above \$250,000 is capped before normalization.

- Very large household, worker, and automobile counts are capped before normalization so rare coding outliers do not dominate the test.
- Household ID is the stable random key.
- The calculation uses a declared 32-bit floating-point policy.

These choices make a stable systems benchmark. They do not claim to be the right behavioral specification for a transportation agency.

31. Phase 18 results: fast, reproducible, and carefully bounded

The full-table benchmark used the local NVIDIA RTX A4000, driver 571.59, Python 3.11.14, nine measured repetitions, and a fused parallel Numba CPU comparison. Compilation warm-up happened before the measured repetitions.

Full public household table	Median time
Parallel fused Numba CPU	0.457023 seconds
GPU modeled work only	0.031357 seconds
GPU including one upload and final downloads	0.055327 seconds

That is **14.575 times faster** for modeled computation and **8.260 times faster** even after the allowed boundary transfers.

This is a much larger speedup than the whole ActivitySim results reported in earlier phases because the scopes are different. Phase 18 measures a compact GPU-native vertical slice where almost every operation is parallel. It does not include CSV parsing, all ActivitySim components, checkpoint work, or output writing. A person comparing the numbers must keep that difference in mind.

What exactly reproduced?

The same GPU calculation ran in two arrangements:

- all 2,875,192 households together; and
- consecutive partitions containing at most 250,000 households each.

Every final GPU choice was identical. Every final GPU logsum had identical bits. This is the Phase 18 **replication guarantee**: scaling the workload into deterministic partitions does not change the GPU result.

Runtime telemetry also showed:

Boundary check	Result
Modeled CPU fallbacks	0
Host-to-device modeled bytes after ingress closed	0
Device-to-host modeled bytes before final output	0
Sampled active GPU allocation peak	about 452.4 MiB

The memory number samples active CuPy arrays after stages. A very short-lived temporary allocation could be higher, so it is a measured lower bound rather than a perfect hardware high-water mark.

Why are CPU and GPU answers not all bit-identical here?

The Numba CPU and CUDA GPU both implement the same mathematical model, but they do not promise the same lowest-level versions of exponential and fused-multiply-add operations. Across all 2,875,192 rows:

CPU/GPU comparison	Observed difference
First choices	1 row
Dependent second choices	3 rows
Largest dependent logsum difference	0.007143
Largest zone-total difference	1.0

One first-stage boundary choice can affect the second-stage feature, which is why the number grows from one to three. The observed rates are 0.348 and 1.043 differences per million rows. They pass the published limits of one and two per million.

This is not the same guarantee as Phases 14-17, where a specially defined CPU recipe and generated CUDA utility recipe matched feature and utility bits exactly. Phase 18 deliberately compares against a strong normal Numba choice implementation. Its cross-architecture claim is **numerical equivalence within written bounds**, while its GPU partition claim is **bit-exact**.

During development, a first full-table run used `log1p(income)`. NumPy and CUDA rounded that function differently by one float32 unit on some rows, and one row landed close enough to a choice boundary to change. The project did not hide the failure or simply rerun until it disappeared. It replaced the feature with capped linear normalization, documented the arithmetic policy, reran the full table, and still reported the remaining CUDA-versus-Numba boundary cases.

What fits in this GPU's memory?

The local RTX A4000 reports 16,376 MiB of GPU memory. The public MTC skim file contains 826 datasets. If every dataset is counted as its raw uncompressed array, the collection needs 13.389 GiB. Earlier real 50,000-household ActivitySim integration runs already peaked near 8.4 GiB.

Therefore the complete future model should not load every skim and every state table at once. Phase 18 proposes four planned pools:

Planned use	Budget
CUDA, driver, and failure reserve	2 GiB
Frequently used, or "hot," skims	4 GiB
Persistent model state	2 GiB
Largest component workspace	3 GiB
Remaining safety and partition room	about 4.99 GiB

A **hot skim cache** keeps the travel-time maps needed soonest on the GPU and evicts others safely when space is needed. A **population partition** processes a stable set of households and everything belonging to them, then moves to the next set. These are design budgets, not yet measured full-model limits. The maximum safe production partition cannot be claimed until the missing model components have real memory high-water measurements.

What does Phase 18 prove?

- A sealed GPU state boundary can be enforced rather than merely promised.
- Stable GPU random draws survive changes in partition size.
- Dependent choices and an ordered group total can stay on the device.

- This vertical slice can process the entire public household table at once.
- The GPU is substantially faster than fused parallel Numba for this work, including the boundary transfers.
- Every declared performance, numerical, partition, and fallback gate passes.

What does it not prove?

- It is not a complete ActivitySim replacement.
- Its synthetic coefficients do not make calibrated travel predictions.
- It does not prove every person, tour, trip, timetable, shadow-price, and skim table fits together in 16 GB.
- It does not yet provide restartable GPU checkpoints.
- It does not accelerate file parsing or report writing.
- It does not prove the same speed on another GPU or model.
- It does not claim that ordinary CPU and GPU math is bit-identical.

32. The practical path from this slice to a whole GPU model

The next-generation project is possible on this GPU if it grows in dependency order and treats memory and replication as design rules.

Phase A: finish the GPU table toolbox

Implement indexed joins, filters, stable sorting, scatter and gather, category encoding, group operations, and missing-value policies. Each operation needs a small readable CPU oracle, adversarial tests, and an explicit arithmetic rule.

Phase B: build the hot skim cache

Track which skim arrays each component needs. Upload them asynchronously, retain frequently reused arrays, and evict only after every dependent CUDA event finishes. Record cache hits, misses, transferred bytes, and memory peaks.

Phase C: port one calibrated household-person chain

Replace synthetic coefficients with a real public specification. Preserve stable entity channels and compare every intermediate column, not only the final output. A missing GPU operation must stop qualification rather than use a hidden CPU fallback.

Phase D: make restart behavior safe

Long models must recover after a failure. Either serialize device tables at declared checkpoints or make a documented host checkpoint boundary. Restoring a checkpoint must reproduce random streams and every downstream table.

Phase E: add tours, trips, timetables, destinations, and shadow prices

These are harder because they create variable numbers of rows, use large skim lookups, and update shared state. Port them one dependency layer at a time. Measure high-water GPU memory after every component and reduce partition size before an out-of-memory failure becomes possible.

Phase F: run the real whole-model proof

Only after the calibrated chain is complete should the project run fresh, interleaved CPU/GPU whole-model trials. The proof must include table hashes, random-stream checks, fallback counters, memory peaks, transfer bytes, component times, whole-model time, software hashes, and a second-machine replication.

The implication is exciting but specific. This workstation has enough GPU to build and test a serious next-generation model and to run large partitions very quickly. It does not have enough memory to hold the public model's entire skim collection plus every future state object without caching and partitioning. Success therefore comes from a GPU-native **system** - state, randomness, memory, checkpoints, and kernels together - not from writing one spectacular kernel.

33. Phase 19: replace the invented equations with a real calibrated chain

Phase 18 answered an engineering question: can several dependent calculations stay on the GPU and run quickly? Its answer was yes, but its choice equations used invented test coefficients. Phase 19 asks the harder scientific question:

> Can the GPU run real published travel-model equations and reproduce the > choices that ActivitySim already made?

The answer for the first calibrated household-to-person chain is **yes**.

What does “calibrated” mean?

A travel-model equation combines facts such as income, household size, age, location, and travel time. Each fact gets a number called a **coefficient**. A coefficient says how strongly that fact moves a choice up or down.

For example, this made-up equation is only an illustration:

```
score for owning two cars
= 0.4 × number of workers
+ 0.2 × income in thousands
- 0.3 × transit accessibility
```

Real modelers estimate coefficients from observed travel and household data. When a model uses those estimated numbers, it is **calibrated**. Phase 19 uses the real coefficients published with Prototype MTC Extended. It does not invent new behavior.

Which two real decisions run on the GPU?

The first model is **auto ownership**. Each household chooses one of five answers: zero, one, two, three, or four-or-more automobiles.

The second model is **mandatory tour frequency**. A *tour* is a journey that starts at home, visits one or more places, and returns home. A mandatory tour goes to work or school. A person can choose among these five answers:

- one work tour;
- two work tours;
- one school tour;
- two school tours; or
- both a work and a school tour.

Only people whose earlier daily-activity choice says they have a mandatory activity enter the second model. In the public checkpoint, that is 78,900 of 132,536 people.

The chain is truly connected:

```
50,000 households
    |
    v
GPU auto-ownership model
    |
    | join each answer to people in that household
    v
78,900 mandatory-person choosers
    |
    v
GPU mandatory-tour-frequency model
```

The second model uses the **new GPU auto answer**. It does not secretly copy ActivitySim's saved auto answer into the calculation. The saved answer is used only afterward as an answer key.

What is a checkpoint replay?

A long video game lets you save at a checkpoint instead of starting from the beginning after every test. ActivitySim can also save its tables after model steps. Phase 19 starts from public saved tables immediately before the chosen components.

This is powerful because the input state is fixed. The CPU reference, the GPU, and ActivitySim's saved output all begin from the same households, people, zones, work locations, school locations, accessibility values, and earlier daily-activity answers.

It also limits the claim. Phase 19 does **not** recalculate school location, work location, or the earlier daily-activity model on the GPU. Those are frozen inputs. This is a two-component calibrated replay, not yet a whole-model run.

34. How Phase 19 makes a real choice

Both components use a method called **multinomial logit**, or MNL. The name is less important than its four steps.

Step 1: evaluate expressions

An **expression** turns input columns into a useful model feature. Examples in the real auto-ownership file include:

- whether the household has exactly two drivers;
- household income, capped within a range;
- retail accessibility by car or transit;
- whether the home is in San Francisco County; and
- estimated automobile time savings per worker.

The auto model has 29 active expressions. The mandatory-tour model has 98. Together, the GPU evaluates 127 published expressions.

The safe expression reader understands only operations that were reviewed and tested, such as addition, comparison, Boolean “and/or,” clipping a value to a range, and choosing one of two values with where. If an unfamiliar operation appears, the run stops. It never sends the expression to unrestricted Python code and never quietly falls back to the CPU.

Step 2: calculate utilities

For every alternative, each expression value is multiplied by its calibrated coefficient. The products are added to form a **utility**. Utility is a model score, not money or electrical service. A higher utility means the alternative is more attractive to the model.

With 50,000 households and five auto alternatives, the first component makes 250,000 utility values. With 78,900 people and five tour-frequency alternatives, the second makes 394,500 utility values.

Step 3: turn utilities into probabilities

MNL converts the five utilities in a row into five probabilities that add to one. A numerically safe version first subtracts the largest utility, then uses the exponential function, and finally divides each result by the row total.

Suppose the probabilities were:

alternative:	A	B	C	D	E
probability:	0.10	0.25	0.40	0.20	0.05

These probabilities describe ranges on a line from zero to one. A random draw of 0.52 lands in C's range, so C is selected.

A **logsum** is the logarithm of the total exponentiated utility. Modelers use it as a summary of how attractive the complete choice set is. Phase 19 keeps logsums on the GPU and downloads them only as final diagnostic outputs.

Step 4: use ActivitySim's exact random draw

Matching probabilities is not enough. Two programs can have the same probabilities but choose different answers if they use different random draws.

ActivitySim creates a stable seed from four items:

```
run seed + channel name + model-step name + household/person ID
```

It then uses an algorithm called **MT19937**, from NumPy's older RandomState, to create the draw. Phase 19 implements the needed part of MT19937 directly in a CUDA kernel. The GPU produces the exact same 64-bit floating-point draw as NumPy for every tested household and person.

Phase 19's kernel supported the first draw in a model step, called offset zero. These two public components need exactly that. Phase 20, explained below, has since implemented and proved later offsets as well.

How does a person find the right household answer?

Tables do not always have rows in matching order. A person's row contains a `household_id`. The GPU must find the household row with that ID.

Phase 19 adds a GPU **keyed join**. It sorts source IDs, searches for every requested ID, checks that no ID is missing, and gathers the matching values. The same tool also looks up land-use and accessibility facts by zone ID.

This is a basic database operation performed on the GPU. It is essential for a whole travel model because household, person, tour, trip, and zone tables constantly refer to one another by IDs.

35. How we proved the result and what it means

The proof uses three separate answer layers.

- 1 A plain NumPy version independently calculates expressions, utilities, probabilities, ActivitySim random draws, and choices.
- 2 That CPU version must exactly reproduce ActivitySim's saved checkpoint choices.
- 3 The GPU is compared with every important CPU intermediate and separately with ActivitySim's saved final choices.

This avoids circular reasoning. The GPU is not declared correct merely because another function that shares all its code agrees with it.

Correctness results

Public-checkpoint test	Result
Households	50,000
All people in input state	132,536
Mandatory-person choosers	78,900
CPU auto choices different from ActivitySim	0
GPU auto choices different from ActivitySim	0
CPU tour-frequency choices different from ActivitySim	0
GPU tour-frequency choices different from ActivitySim	0
Expression feature differences	0
Random-draw bit differences	0
Largest utility difference	about 0.00000000000000018
Largest probability difference	about 0.00000000000000044
Choice differences across nine repeated GPU runs	0

Those tiny utility and probability differences come from low-level floating-point library arithmetic. They are many orders of magnitude below the written limits and did not change a single choice.

Speed results

The program first warmed up compilation. It then measured nine complete CPU and GPU replays on the RTX A4000.

Measured boundary	Median time	Relative speed
Independent CPU replay	0.458724 seconds	1.000×
GPU modeled work	0.025713 seconds	17.840× faster
GPU with input upload and final download	0.037997 seconds	12.073× faster

The transfer-inclusive number is the most practical result for this checkpoint replay. It includes moving the input tables to the GPU once and bringing the final choices and logsums back. It does not include reading Parquet files from disk or running all the upstream and downstream ActivitySim components.

Did the GPU secretly use the CPU?

No modeled fallback was recorded. After the GPU boundary closed:

Forbidden event	Recorded amount
CPU modeled fallbacks	0
Late host-to-GPU modeled transfers	0 bytes
Early GPU-to-host modeled transfers	0 bytes

The CPU still launches kernels and checks a few true/false error flags. That is control work, not a second model calculation.

Does Phase 19 make the first 18 phases useless?

No. It makes Phase 19 the best **headline**, but it could not stand alone as a trustworthy project.

Earlier phases built and tested the choice kernels, strict expression language, float rules, ActivitySim adapters, public checkpoints, failure policies, device-state runtime, and transfer counters. They also studied destination and scheduling components that Phase 19 does not run. Phase 18 proved that a dependent GPU state chain could work before real calibrated behavior was added.

The right summary is:

- Phase 19 replaces Phase 18's synthetic behavioral demonstration for this household-to-person boundary.
- Phases 1-18 remain the engineering foundation, audit trail, and evidence for other component types.
- Phase 19 still does not replace a full fresh ActivitySim CPU/GPU comparison.

36. Phase 20: make a real tour table on the GPU

Phase 19 ended with one frequency answer per mandatory person. Phase 20 asks a different kind of computer question: how do we create a table when different input rows create different numbers of output rows?

One person may choose one work tour. Another may choose two school tours. A third may choose one work tour and one school tour. This is called **variable-length table expansion**.

Imagine students lining up for a school photo. Some need one chair and some need two. Before assigning chairs, the organizer counts how many each student needs. A **prefix sum** adds the counts from left to right:

```

chairs needed:      1  2  1  2
starting chair:     0  1  3  4
total chairs needed:                6

```

The GPU uses the same idea. It counts each person's tours, calculates the starting output position, allocates exactly 81,983 rows, and runs one fused kernel that writes all columns.

What is in a tour row?

A tour row needs much more than “work” or “school.” It contains:

- the tour's own stable ID;
- the owner person's and household's IDs;
- whether it is a work or school tour;
- which tour of that type it is;
- which mandatory tour should be scheduled first;
- how many mandatory tours that person has;
- the home origin and work or school destination;
- its category; and
- its participant count.

That is 12 columns. Phase 20 compared every generated value, not just a sample or the row count, with ActivitySim's public saved table. Every column had zero mismatches across all 81,983 rows.

A subtle work-and-school rule

ActivitySim physically stores work rows before school rows. But for a person who is not classified as a worker and has both tours, school must be scheduled first. The row order stays the same while the schedule number is swapped.

This distinction sounds tiny, but getting it wrong changes the person's timetable and can affect later choices. ChoiceForge has a focused test for the rule and also checks it across the complete public table.

37. Why IDs must be boring and predictable

A random-looking but stable tour ID lets later model steps find the same tour, attach a repeatable random stream, and compare two runs.

The public model has 41 possible tour labels when mandatory, optional, joint, and at-work tours are considered together. The mandatory labels occupy fixed positions:

Tour label	Position among 41 labels	ID rule
First school tour	31	person ID $\times$ 41 + 31
Second school tour	32	person ID $\times$ 41 + 32
First work tour	39	person ID $\times$ 41 + 39
Second work tour	40	person ID $\times$ 41 + 40

For example, if person 100 has a first work tour, its ID is $100 \times 41 + 39 = 4,139$.

ChoiceForge checked the formula against every public mandatory tour. More importantly, the generated ID set exactly equaled the set of tour IDs consumed by the next scheduling component. That proves the two stages connect, instead of merely producing two separately correct-looking results.

38. What mandatory-tour scheduling decides

After making a work or school tour, the model must decide when it starts and ends. A **tour departure-and-duration alternative**, shortened to **TDD**, is one possible time window. Examples might be “leave at 7, return at 17” or “leave at 9, return at 15.” The public model has 190 base TDD alternatives.

Not every window is possible for every tour. A later tour cannot overlap a time already occupied by an earlier tour. Its scores can depend on:

- the person's worker or student type;
- income and household size;
- travel time to work or school;
- whether the destination is downtown;
- the start, end, and duration of the proposed window;
- how attractive the available travel modes are at those times;
- the previous tour's end time; and
- open blocks in the person's timetable.

This is why scheduling is a harder downstream test than another five-answer frequency model.

The full public capture

Phase 20 resumed the preserved 50,000-household ActivitySim pipeline exactly after mandatory-tour frequency, ran only mandatory scheduling, and captured:

- 81,983 real tour choosers;
- six first/second work, school, and university groups;
- 15,242,743 feasible chooser-time rows;
- published coefficients and expressions;
- real time-dependent mode-choice logsums;
- real timetable facts;
- ActivitySim's random draws; and
- ActivitySim's selected time positions.

The compact format does not save a huge table with every expression already calculated. It saves shared ingredients once and lets the generated CPU and GPU programs evaluate the expressions. The six compressed files total about 12 megabytes.

39. The gate failed first, and that was success

The first full run did **not** pass. Three GPU scheduling choices and one CPU reference choice differed among 81,983 tours.

This was not dismissed as “only a few.” Every mismatch occurred where the random draw was extraordinarily close to the boundary between two cumulative probabilities. The investigation found that the older kernel changed the arithmetic recipe:

- ActivitySim keeps its random draw as a 64-bit floating-point number.
- The old kernel rounded that draw to 32 bits.

- ActivitySim normalizes its 32-bit probability weights before subtracting them from the draw.
- The old kernel compared the rounded draw with unnormalized weights.

The corrected kernel keeps the draw at 64 bits, creates normalized 32-bit probabilities, and subtracts them in the same alternative order as ActivitySim. The independent CPU answer key was corrected to use NumPy's exact 32-bit probability-reduction behavior too.

A new regression test places a draw just one 64-bit step above 0.5. Rounding it to 32 bits changes the answer, so the old bug cannot quietly return.

After the correction:

Proof check	Result
CPU schedule choices different from ActivitySim	0 of 81,983
GPU schedule choices different from ActivitySim	0 of 81,983
CPU choices different from GPU choices	0 of 81,983
GPU TDD values different from public checkpoint	0 of 81,983
Largest CPU/GPU logsum difference	about 0.00000381
Choice differences across nine GPU repeats	0

The small logsum difference is reported honestly. The CPU and GPU add many 32-bit weights using different reduction trees, so their last few bits need not match. The behavioral choice output is exact.

40. Phase 20 speed results

Compilation was warmed before nine measurements on the RTX A4000.

Work measured	CPU median	GPU median	Speedup
Create tour table, data already on GPU	0.006904 s	0.000601 s	11.496×
Create tours including upload and ID download	0.006904 s	0.001101 s	6.272×
Schedule kernel, compact data already on GPU	0.173748 s	0.009601 s	18.097×
Schedule kernel including compact transfers	0.173748 s	0.059196 s	2.935×

These are meaningful wins at their named boundaries. They do not mean that the complete travel model is now 17.8 times faster.

The scheduling timer starts after ActivitySim has prepared the mode-choice logsums, feasible time alternatives, and timetable facts. Those preparation steps remain on the CPU. In simple language, the GPU now cooks the scheduling recipe very quickly, but ActivitySim still gathers and measures several major ingredients.

41. Random streams can now continue and restart

A repeatable simulation sometimes needs the first random number for an entity, then the second or the 313th. Phase 19 implemented only the first number and correctly refused anything else.

Phase 20 implements any nonnegative **random offset** for ActivitySim's MT19937 generator. Tests check offsets 311, 312, and 313 because that crosses an internal state-refresh boundary where a careless implementation often fails. Every tested GPU double matches NumPy bit for bit. The common first-draw case still uses its faster specialized kernel.

A **random ledger** records how many draws each table and model step has used. Saving that ledger prevents a restarted model from accidentally reusing the first random number and changing its scientific result.

Phase 20 also writes an audit manifest. For each tour column it records the row count, data type, and SHA-256 fingerprint. It records completed components, the selected TDD fingerprint, the random offset, and the capture fingerprint. This makes silent changes detectable.

It is not yet a fully self-contained restart file because the compact scheduling-preparation arrays live beside it rather than inside it.

42. What did Phase 20 prove, and what remains?

Phase 20 proves that a calibrated GPU choice can create a correctly shaped dependent table, give its rows ActivitySim-compatible identities, and feed all of them into a real downstream calibrated choice kernel. This is more than a standalone arithmetic demonstration.

It still does not prove a complete GPU-only ActivitySim model. The frozen upstream location and CDAP inputs from Phase 19 remain. Scheduling preparation still uses the CPU. Other tour and trip components remain outside this chain. Only one NVIDIA GPU architecture and one public model have been used for this new proof.

The next major phase should move scheduling preparation itself:

- 1 Keep the needed travel-time and travel-cost skim arrays in a managed GPU cache.
- 2 Calculate time-dependent mode-choice logsums on the GPU.
- 3 Represent each person's timetable as device state.
- 4 Construct the feasible TDD alternatives from that timetable.
- 5 Schedule first tours, update the timetable, then schedule later tours.
- 6 Feed those arrays directly into the already qualified scheduling kernel.
- 7 Compare the entire scheduling component, including preparation, with a fresh ActivitySim run.
- 8 Repeat on another NVIDIA architecture and a second public model.

The long-term lesson is that raw GPU arithmetic was never the only problem. The difficult and valuable work is preserving identities, table growth, random streams, time state, maps, arithmetic rules, restart evidence, and exact behavior while data crosses a chain of decisions. Phase 20 closes the table- growth and downstream-kernel parts of that problem. Scheduling preparation is now the clearest next frontier. Phase 21, explained next, completes that frontier at an honest compact-cache boundary and proves the raw-skim CUDA side separately.

43. Phase 21: prepare real schedules on the GPU

Phase 20 began with a prepared list of possible times. That was useful, but it left a fair question: what if preparing the list takes much longer than choosing from it?

Phase 21 moves that preparation work. For each person, the GPU stores a small 21-slot timetable. Think of it as a row of boxes, one box for each modeled hour. A box can be empty, mark the start of a tour, mark its end, or mark time in the middle of a tour.

For every tour, the GPU considers 190 possible combinations of leaving and returning. It asks whether each combination collides with something already on the person's timetable. It keeps the feasible combinations and throws away the colliding ones. This creates a list of different length for every tour.

The compact computer format for those different-length lists is called **CSR**, short for compressed sparse row. You do not need to memorize the name. It is simply:

```
one long list of all feasible time choices
+ a small list saying where each person's part begins and ends
```

CSR matters because the GPU does not need to reserve 190 output rows for every tour after it knows many choices are impossible.

Why tours must be scheduled in order

A person may have two mandatory tours. The second tour cannot be checked until the first tour has occupied part of the timetable. Phase 21 therefore uses a careful mixture of parallel and sequential work:

- thousands of people and alternatives are processed together inside a batch;
- the first-tour batch finishes and updates the timetable; and
- only then does a later-tour batch begin.

This is not a weakness. It is the causal rule of the model. Running later tours too early would be fast but scientifically wrong.

44. How 190 time choices become a 25-number cache

Network travel times do not change separately for every modeled hour. The public MTC model groups hours into five broad skim periods, such as early, morning, midday, afternoon, and evening.

A tour has an outbound period and an inbound period. Five possibilities in each direction make a 5-by-5 table, or 25 slots. Because a tour cannot return before it leaves, only 15 slots are used for a first tour.

This lets Phase 21 replace millions of repeated mode-logsum values with one small cache per tour. A **logsum** is a single number summarizing how attractive all the travel modes are for that outbound and inbound period pair. It is not a simple average: attractive modes contribute more, and related modes are grouped by a nested-logit formula.

The cache builder does not assume repeated values are equal. It checks their 32-bit patterns. If two hourly alternatives mapped to the same slot contain different values, the build stops with an error.

The original captured prepared arrays used 518,909,174 bytes in memory. The Phase 21 primitive arrays use 12,688,620 bytes, **40.896 times smaller**. The six compressed files on disk total about 6.3 megabytes.

45. What exactly the Phase 21 speed test includes

The stopwatch begins with the compact 5-by-5 cache and per-tour facts. It includes all six real scheduling batches and all of these operations:

- 1 checking the timetable for collisions;
- 2 counting feasible alternatives;
- 3 building the CSR list;

- 4 deriving the previous tour's end time;
- 5 calculating all seven timetable-dependent ActivitySim facts;
- 6 gathering the right mode logsum;
- 7 calculating calibrated scheduling utilities and probabilities;
- 8 using ActivitySim's original 64-bit random draw;
- 9 selecting the TDD; and
- 10 updating the timetable before the next batch.

Here, **TDD** means a tour departure-and-duration alternative. It identifies a start time and end time from the public 190-row alternative table.

Nine measurements on the RTX A4000 produced:

Qualified work	Compiled CPU median	GPU median	GPU advantage
Data already resident	0.214878 s	0.021069 s	10.199x
Primitive transfer included	0.214878 s	0.024755 s	8.680x

Including primitive transfers added about 3.7 milliseconds to the GPU median and reduced its advantage from 10.199x to 8.680x. This is why the report keeps resident and transfer-inclusive boundaries separate instead of hiding data movement. Both measured GPU boundaries were much faster than the same compiled CPU reference.

The benchmark also lists ActivitySim's 23.338-second scheduling-component time for context. It does **not** divide that number by 0.021 seconds to claim a thousand-fold improvement. ActivitySim's timer includes raw mode-logsum work, pandas tables, and workflow management that are outside this stopwatch.

46. How Phase 21 proves the answer is the same

Speed is accepted only after the answer passes independent checks.

The public workload contains:

Item	Count
Households	50,000
Persons with timetable rows	78,900
Mandatory tours	81,983
Sequential scheduling batches	6
Possible TDD alternatives	190 per tour
Feasible rows actually regenerated	15,242,743

A readable CPU version and a compiled parallel CPU version serve as answer keys. The GPU must match the captured ActivitySim CSR offsets, feasible TDD IDs, eight row values, and adjusted chooser values. It then must match every selected TDD. A second GPU run must produce the same result again.

The final counts are simple:

- CPU preparation mismatches: **0**;
- GPU preparation mismatches: **0**;

- CPU TDD mismatches: **0 of 81,983**;
- GPU TDD mismatches: **0 of 81,983**;
- CPU-versus-GPU TDD mismatches: **0**; and
- differences on the repeated GPU run: **0**.

The evidence file also stores the machine and software versions, all nine raw timing samples, source fingerprints, input fingerprints, per-batch counts, and a restart/checkpoint fingerprint. A fingerprint is a SHA-256 hash: a long label that changes if the bytes change.

The GPU-only method rejects a NumPy or pandas modeled array. This prevents a hidden CPU fallback from accidentally being timed and described as GPU work.

47. The live raw-skim gate and the one-choice mystery

The compact-cache benchmark is the strongest speed result, but where does the cache come from? Phase 21 also answers that question with a real ActivitySim integration run.

The live path reads the public network skim tensors. A **tensor** is a multi-dimensional array. Here its coordinates can include origin zone, destination zone, and time period. The generated GPU program evaluates the public tour-mode specification for 21 travel alternatives and then performs the nested-logit reduction.

Six live calls processed 1,210,124 rows. Every call used CUDA, none fell back to the CPU, and the utility matrix stayed on the GPU when it entered the nest reducer.

The gate did not pass on the first try.

First, the compiler did not understand expressions asking for a reverse trip, the maximum of forward and reverse distance, and special round-trip skim directions. It stopped safely instead of guessing. Phase 21 added those exact operations and tests that prove a reverse lookup swaps origin and destination without uploading a second copy of the network.

Next, all GPU calls ran, but one of 81,983 schedules changed. A fresh CPU-only control matched the frozen reference, so the GPU arithmetic needed more work.

The cause was a tiny but real recipe difference:

- Sharrow's utility calculation uses 32-bit numbers and permits a fused multiply-add operation.
- The strict ChoiceForge proof compiler used separate multiplication and addition and deliberately forbade fusion.
- ActivitySim then promotes the utilities to 64-bit numbers for a particular sequence of nested exponent, sum, and logarithm operations.

A **fused multiply-add**, often written FMA, calculates $a \times b + c$ as one hardware instruction with one final rounding step. Separate multiplication and addition round twice. Usually the difference is far too small to matter. This one tour had a random draw extremely close to a probability boundary, so it did matter.

Phase 21 did not weaken the old strict contract. It added a separate, explicitly named Sharrow-compatible fused policy for this live path. It also added a CUDA nest reducer that follows ActivitySim's real mixed-precision tree. An attempted all-32-bit nest policy changed 1,203 schedules and was rejected.

After the correct policies were combined, the live result was:

Live proof check	Result
CUDA calls	6
CUDA mode-logsum rows	1,210,124

Live proof check	Result
CPU fallbacks	0
Utility download/re-upload before nesting	0 bytes
TDD differences	0 of 81,983
Start-time differences	0
End-time differences	0

This failed-first history is part of the evidence. It shows why "almost every choice" was never treated as enough.

48. What Phase 21 means, and what still needs to be done

Phase 21 proves two major facts:

- from a compact time-period logsum cache, the GPU can regenerate the complete calibrated mandatory-scheduling preparation and choice boundary exactly and about ten times faster than compiled parallel CPU code; and
- from real public raw skims, the generated CUDA utility and nest engine can feed ActivitySim without changing any mandatory schedule.

These are complementary proofs, not yet one continuous production timer. The live ActivitySim API still creates pandas tables and materializes the compact cache between the raw-skim GPU work and the standalone device-resident scheduler. Joining those two already qualified halves is the next major step.

After that, a larger program of work remains:

- 1 checkpoint the live GPU timetable so a stopped run can restart exactly;
- 2 port non-mandatory, joint, and at-work scheduling;
- 3 build a managed hot cache for the 13.389-GiB public skim collection;
- 4 partition the population without changing identities or random draws;
- 5 repeat the proof on another NVIDIA GPU;
- 6 repeat it on a second public activity-based model; and
- 7 run repeated complete-model CPU/GPU pairs before claiming whole-model superiority.

Phase 21 does not make Phases 1 through 20 unnecessary. Those phases created the public captures, independent CPU answers, random-number rules, stable tour IDs, expression compiler, device handoff, precision tests, and fail-closed habits that made Phase 21 provable. A large success that cannot be reproduced would be a demonstration. This project is trying to build evidence.

49. Phase 22: join the two fast halves

Phase 21 ended with two strong but separate results. One result started with raw road and transit data and calculated mode-choice logsums on the GPU. The other started with a small logsum cache and prepared and chose schedules on the GPU. Phase 22 connects them.

The connected path now works like this:

- 1 ActivitySim supplies a batch of tours, their stable IDs, and network skim lookups.
- 2 Generated CUDA code calculates 21 mode utilities for each representative time combination.

- 3 A CUDA nested-logit reducer turns those utilities into one logsum per combination.
- 4 The logsum stays on the graphics card and is placed into a small cache for its tour.
- 5 The GPU timetable rejects schedule choices that collide with an earlier tour.
- 6 The GPU evaluates the scheduling rules, uses ActivitySim's exact random draw, chooses a departure-and-duration alternative, and marks those hours busy before the next batch begins.
- 7 Only final schedule labels return to ActivitySim for the normal path.

There are six ordered batches because first tours must be scheduled before second tours. Phase 22 completed all six for 81,983 mandatory tours.

50. Why one answer was still incredibly hard

The first connected GPU run was wrong for exactly one tour. The reference chose TDD 169, while the GPU chose 168. This was not ignored as a harmless rounding difference.

That tour's random number was only about **0.00000000085** beyond a probability boundary. Imagine drawing a line with a thick marker and asking whether a dust speck is on the left or right edge. Several mathematically reasonable computer recipes can round the line by a few billionths and put the speck on different sides.

The investigation found several hidden parts of ActivitySim's recipe:

- Sharrow builds 65 32-bit features, including two temporary features whose coefficients are zero.
- It performs a 32-bit dot product. A **dot product** multiplies matching pairs and adds all the products.
- Because of one ActivitySim safety setting, it exponentiates the utilities without first subtracting the largest utility.
- NumPy's CPU exponential and sum do not promise the same final bits as CUDA's GPU exponential and parallel sum.

All of those methods are numerically sensible. But “sensible” is not the same as “guaranteed to make the identical simulated choice.” A simulation is a chain: one changed schedule can change a person's later availability and then change more decisions.

Hard-coding 169 would have been cheating. It would only memorize the test. The solution instead asks the GPU how close every random draw is to the nearest choice boundary. If the distance clears the guard tested on this public model, the GPU result is accepted. If it is very small, the row is marked ambiguous and the exact ActivitySim/Sharrow recipe decides it. A different model or GPU must test the guard again; it is not a universal law of floating-point math.

51. Is Phase 22 completely GPU-only?

No - and that distinction matters.

Most of the expensive work is on the GPU, and the bulk modeled logsums are not downloaded. But 57 of 81,983 tours were close enough to a probability boundary to use the exact CPU/Sharrow adjudicator. That is **0.0695%** of the tours. The adjudicator downloaded 11,400 bytes, about eleven kilobytes, per full run.

This is a hybrid with a tiny, measured correctness exception. Calling it “absolutely no CPU” would be inaccurate because ActivitySim still organizes the workflow, owns the random numbers, writes output tables, and resolves those 57 ambiguous choices.

Why keep this exception? Because removing it before CUDA and Sharrow share an identical definition for dot products, exponentials, sums, and rounding would trade a true exact result for a marketing slogan. A future Sharrow GPU backend or a dedicated expression compiler can remove the resolver, but only after it passes this boundary test and every other saved proof gate.

52. The final Phase 22 evidence

The team ran three fresh CPU/GPU pairs from the same public 50,000-household checkpoint. Each run included raw network skim setup, ActivitySim orchestration, all six scheduling batches, and output writing for the resumed component.

Pair	CPU time	GPU-connected time	CPU / GPU
1	42.358 seconds	36.599 seconds	1.157x
2	40.389 seconds	31.963 seconds	1.264x
3	40.250 seconds	32.030 seconds	1.257x

The GPU path was faster in every pair. The middle paired speedup was **1.257x**. Comparing the middle CPU and GPU times gives **1.261x**. These are smaller than the 8.680x-to-10.199x compact-kernel results because the live timer includes shared setup and ActivitySim work that still runs on the CPU. Both measurements are useful; they answer different questions.

Every live GPU run produced the same proof facts:

- 6 of 6 batches joined;
- 1,210,124 raw-skim rows used generated CUDA;
- 0 CUDA fallbacks;
- 57 exact boundary checks;
- 11,400 boundary-transfer bytes;
- 0 changed random draws;
- 0 changed TDD choices out of 81,983;
- 0 changed start times;
- 0 changed end times; and
- a restart record containing fingerprints of the selected TDDs and final timetable.

The result means this machine can make the complete resumed mandatory-tour scheduling component faster, starting from raw public skims and ending with exact schedules. It does not yet mean the whole ActivitySim model is faster or GPU-only. Non-mandatory, joint, at-work, trip, and destination components still need the same connected proof. A second GPU and a second public model are also needed before claiming that the result is portable.

Most importantly, Phase 22 shows why the earlier phases were not wasted. The one-in-81,983 boundary problem was found because the project preserved random draws, stable identities, CPU references, public checkpoints, arithmetic policies, and fail-closed tests. Those are the pieces that turn “the GPU looked fast” into a result another reviewer can challenge and reproduce.

53. Why the 1.257-times result was not the GPU's real limit

The Phase 22 timer measured a whole ActivitySim component. That was the right test for compatibility, but most of its roughly 32 seconds were not spent in the scheduling kernel. ActivitySim had to restart a saved pipeline, create and join pandas tables, organize model calls, synchronize state, validate results, and write output. Making a small kernel even faster could not remove that surrounding work.

This is an example of **Amdahl's law**. It says that a system cannot become much faster by improving a part that occupies only a small fraction of total time. If a one-second calculation sits inside 30 seconds of unchanged setup, making that calculation instantaneous still leaves about 30 seconds.

The escape is to change who owns the workflow. Phase 23 uploads the needed state once and lets several dependent model stages share it on the GPU.

54. What a device-resident runtime is

Think of a school group project stored in a shared online folder. If every student downloads the entire folder, changes one sentence, emails it back, and waits for someone else to upload it, most time is spent moving and organizing files. A better system keeps the current document in one shared place and records which person changed which section.

Phase 23 does the same for model data. A **device table** is a set of columns in GPU memory. The runtime gives every table a name and version. Every model stage must declare which tables it reads and writes. It follows these rules:

- input tables may be uploaded before the runtime is sealed;
- after sealing, modeled outputs must be GPU arrays;
- a host NumPy array or hidden CPU fallback causes a hard failure;
- all outputs of a stage are checked before any become official;
- replacing an existing table must be requested explicitly;
- temporary tables are released after their final user;
- only named final columns may be published; and
- checkpoints record enough actual data to restart, not only a success label.

This still needs a CPU process to read files and launch CUDA kernels. "Device resident" means the modeled state and arithmetic stay on the GPU between the declared entry, publication, and checkpoint boundaries.

55. The six connected modeled stages

The public Phase 23 graph begins with 50,000 households and 132,536 persons. It performs six connected kinds of work:

- 1 The calibrated auto-ownership model chooses the number of cars for every household.
- 2 The mandatory-frequency model uses those new car choices and decides which mandatory people make one or two work or school tours.
- 3 The GPU creates the variable number of tour rows.
- 4 Stable tour IDs connect every generated tour to its scheduling row.
- 5 Six ordered scheduling batches test time choices, calculate probabilities, and select TDDs.
- 6 The timetable is updated so later tours cannot overlap earlier tours.

The graph ends with 78,900 mandatory-person choices and 81,983 scheduled tours. The compact 5-by-5 mode-logsum caches are uploaded inputs in this phase. Phase 22 proved that CUDA can generate them from raw skins, but joining that producer inside the sealed Phase 23 graph remains future work.

56. Two optimizations that only a resident runtime can use well

The first optimization is a **compiled join map**. Households repeatedly look up land-use data by zone, and persons repeatedly look up household data. The older code sorted and searched those keys every run. Phase 23 validates the relationships once and keeps the resulting row numbers on the GPU. Later runs use direct array indexing.

The second optimization is a **fused fixed-choice compiler**. The mandatory frequency model has 98 published expressions and five possible answers. The old GPU path created 98 feature columns, multiplied them by coefficients,

then ran separate probability and choice operations. The new generated CUDA kernel does all of this for one person in one thread:

- read the person's permanent fields;
- insert the auto-ownership answer made by the previous GPU stage;
- evaluate 98 expressions;
- accumulate five utility scores;
- calculate probabilities and a logsum; and
- use the exact ActivitySim random draw to select the answer.

It changed zero of 78,900 frequency choices. Its largest logsum difference from the independent dense CPU calculation was less than 0.0000000000000001.

57. The Phase 23 proof

The team started three independent Python and CUDA processes. Each process ran nine measured CPU repetitions and nine measured resident-GPU repetitions after both compilers were warmed.

Result	Process 1	Process 2	Process 3	Middle result
CPU modeled time	0.7673 s	0.7605 s	0.7677 s	0.7673 s
GPU resident time	0.0313 s	0.0353 s	0.0315 s	0.0315 s
Resident speedup	24.516x	21.555x	24.405x	24.405x
Setup-inclusive speedup	1.356x	1.314x	1.360x	1.356x

"Setup-inclusive" charges the one-time input upload, scheduler creation, device join-map compilation, one modeled run, and final publication to only one run. It excludes file reading and compiler warm-up on both CPU and GPU. Optional checkpoint writing is reported separately and took about 0.34 seconds.

If the same graph is run repeatedly, the fixed setup cost is shared. Arithmetic using the measured medians gives 9.055x over ten repeated runs and 20.868x over one hundred. These are amortized calculations, not measurements of different policy scenarios; measured parameter-batch scenarios are still future work.

Every proof gate passed in all three processes:

- 0 changed auto-ownership choices;
- 0 changed mandatory-frequency choices;
- 0 differences in all 12 generated tour columns;
- 0 changed tour IDs or TDDs;
- 0 changed timetable cells;
- 0 repeat differences across 27 measured GPU runs;
- 0 post-seal modeled transfers;
- 0 CPU fallbacks; and
- 0 schedule differences after checkpoint restore.

58. What this success means - and what it does not mean

Phase 23 proves that the GPU's large advantage was hidden by the old CPU-owned workflow. When several calibrated components share resident state, the modeled chain is more than 20 times faster in every independent process on this machine. Even a single run remains faster after paying the measured setup and publication costs.

It does not prove that the entire ActivitySim model now runs in 0.03 seconds. Earlier workplace/school location and CDAP state are frozen inputs. Scheduling mode-logsum caches are also inputs. Destination choice, non-mandatory tours, joint tours, at-work subtours, trips, shadow pricing, and normal pipeline output still need device-resident implementations.

Phase 24 completed skim-cache management, and Phase 25 has now bound the real utility/logsum producer to that cache. Remaining work is to create the dense chooser rows inside the versioned runtime, connect the generated caches to its timetable stage, represent sampling and shadow-pricing state as versioned tables, then port the other model components. The proof must eventually be repeated on another GPU and another public model.

The important change is that the project now has both pieces: a compatible ActivitySim path that proves exact integration, and a resident runtime that proves the architecture can deliver large speedups. Compatibility explains today's result; residency shows the path to the next generation.

59. Phase 24: keeping the useful road-network facts on the GPU

A travel model repeatedly asks questions such as, "How long is the drive from zone 42 to zone 900 in the morning?" The answer lives in a **skim**. You can picture a skim as a giant spreadsheet: origins are rows, destinations are columns, and each cell contains a time, distance, toll, fare, or other network measurement. A time-dependent skim is a stack of five spreadsheets for early morning, morning peak, midday, afternoon peak, and evening.

The public MTC file stores 826 matrices. As uncompressed 64-bit numbers, they would occupy about 13.39 GiB and nearly fill this 16 GiB card before model state and workspace are added. Loading everything would be fragile.

Phase 24 builds a **hot cache**: the frequently used data. It reads the reviewed expression recipe, called the strict IR, and discovers required skim names automatically, avoiding errors in a hand-typed list.

The 315-term recipe contains 209 logical bindings. "Logical" means a meaning in an equation, such as outbound time or the reverse-direction inbound time. One reverse binding can use the same physical cube with its origin and destination swapped. Shared directional views reduce this to 149 physical allocations.

60. Why 6.38 GB fits when 13.39 GiB did not

The original matrices use **float64**, which stores each number in 8 bytes. The qualified GPU equations use **float32**, which stores each number in 4 bytes. Float32 has less precision, so a model may use it only when the numeric contract and choice tests allow it; it is not a free change.

The required source values total 12,757,865,000 bytes. Their float32 GPU form is exactly half: 6,378,932,500 bytes, or about 5.94 GiB. Phase 24 declares an 8 GiB budget first. Too much data, a missing matrix, or a wrong shape causes an error instead of silent data loss.

The .omx file is only about 0.734 GB because HDF5 compression saves disk space, not expanded working memory. Comparing GPU allocation only with the compressed file size would be misleading.

The cache uploads each physical cube once. The runtime then attaches those existing GPU arrays without making a second copy, gives them a named table and version, and seals entry. No new modeled host array can enter afterward.

61. How every raw read was checked

The proof uses all six real mandatory-scheduling logsum batches. They contain 1,210,124 tour-and-time rows. A small number refer to destination zone 0, ActivitySim's missing-destination marker. Those cannot be legal matrix positions, so the proof reports and excludes 5,530 of them. That leaves 1,204,594 valid public OD/period rows.

For each valid row, both sides read all 209 bindings: 251,760,146 values per run. Instead of downloading that giant table, each feeds every exact 32-bit pattern through two 64-bit **hashes**, or compact fingerprints. A changed input is overwhelmingly likely to change them; source-file and matrix hashes add protection against accidental corruption.

Three new Python/CUDA processes each ran five measured repetitions:

Result	Process 1	Process 2	Process 3	Middle result
CPU time	5.691 s	5.662 s	5.669 s	5.669 s
GPU resident time	0.029 s	0.069 s	0.028 s	0.029 s
Resident speedup	193.114x	82.323x	205.639x	193.114x
Upload-inclusive speedup	1.813x	1.851x	1.802x	1.813x

GPU timing varied with memory clocks and system state, so all values are reported. Even the slowest was 82.323 times faster. All output fingerprints matched; 15 GPU repetitions were exact; post-seal modeled transfers and CPU fallbacks were zero.

The huge 193.114-times number is for this isolated memory-access and hashing job. It must not replace Phase 23's 24.405-times calibrated chain result or Phase 22's 1.257-times live ActivitySim component result. Different timing boundaries answer different questions.

62. What Phase 24 changes, and the honest next step

Before Phase 24, a bounded hot-skim cache was only a plan. It is now measured on public data, protected by a hard budget, and registered inside the sealed runtime. The raw values are bit-exact, fast to access, and small enough to leave space for later model state.

It does **not** yet calculate the real 315 mode-choice terms from those cubes. It does not perform the 21-mode nested-logit calculation, produce the 5-by-5 scheduling logsum cache, or replace that precomputed input in Phase 23. The all-binding hash program is a proof instrument, not a travel-choice model.

This was the seven-job plan for the next connected phase:

- 1 Create the real chooser and OD/period row indices on the GPU.
- 2 Give these resident cube pointers to the already-qualified strict CUDA expression plan.
- 3 Evaluate 315 terms for 21 travel modes.
- 4 Reduce those utilities through ActivitySim's nested-logit tree.
- 5 Scatter the resulting logsums into each tour's 5-by-5 scheduling cache.
- 6 Define reproducible arithmetic for the 57 choices near probability boundaries, or retain and measure an explicit adjudication boundary.
- 7 Replace Phase 23's saved logsum input and rerun every proof gate.

Phase 25 completed jobs 2 through 5 and removed the bulk saved-logsum input from its new producer. Job 1 remains because ActivitySim still prepares dense rows and coordinates. Job 6 remains as the explicit 57-row adjudication

boundary. Job 7 is partly complete: the replacement producer is proven, but the historical Phase 23 timetable benchmark has not yet been rewired to consume it directly. The next section explains the measured result.

63. Phase 25: connecting the road facts to the real equations

Phase 24 was like placing all the needed road maps on a workbench. It proved that the GPU could find and read the right cells quickly, but it did not yet use those cells to predict a travel mode.

Phase 25 adds the missing calculator. For every tour-and-time possibility, it does the following:

- 1 Read facts about the traveler and the tour.
- 2 Read time, distance, toll, fare, and transit facts from the resident skims.
- 3 Evaluate 315 real expressions from the public MTC model.
- 4 Turn those expression values into scores for 21 travel modes.
- 5 Combine related modes using the model's nested-logit tree.
- 6 Produce one logsum that summarizes how attractive the available modes are.
- 7 Place that logsum in the correct cell of a 5-by-5 time-period cache used by tour scheduling.

The six programs represent work, school, and university tours for the first and second mandatory tour. Together they process 1,210,124 rows. That means 381,189,060 expression evaluations and 252,915,916 logical skim reads in one complete replay.

64. What "sealed and resident" means here

The first time a batch appears, ActivitySim has already prepared its traveler fields and origin, destination, and time-period numbers. Phase 25 freezes those inputs in GPU memory. It also keeps the compiled equation program, coefficients, output space, and raw skim arrays there.

There was one subtle trap. To place each new logsum in a scheduling cache, the computer needs a map saying, for example, "source row 12,345 belongs to tour 700 and morning-to-evening cache cell 9." An early passing version rebuilt that map on the CPU during every replay. The values stayed on the GPU, but the operation was not truly resident.

The final version compiles this map once. About 19.36 MB of device positions remain on the GPU, along with reusable output caches. During a measured replay, the CPU does not rebuild the map, upload it, or receive the modeled logsums. It only launches the GPU work.

The device state for this boundary includes:

Resident item	Approximate size
149 unique skim arrays	6.20 GB
Dense equation inputs	348.5 MB
Origin/destination/time coordinates	154.9 MB
Compiled cache-placement map	19.4 MB

The six programs share the 149 skim arrays. They are not six separate 6.20 GB copies. Counting GPU pointers across the whole process proves that only 149 physical arrays exist.

65. How the Phase 25 result was proved

One fast run can be luck. The proof starts three fresh Python and CUDA processes. Each process first runs the real ActivitySim mandatory-scheduling component, captures the six actual programs, and checks final four times against the frozen public answer. Then it warms the resident path once and measures five complete replays.

Result	Process 1	Process 2	Process 3	Middle result
Resident time	0.169039 s	0.169739 s	0.169423 s	0.169423 s
Speedup versus initial live CUDA setup/execution	10.389x	9.655x	9.475x	9.655x
Changed logsum bits across five replays	0	0	0	0
Changed final TDD/start/end values	0	0	0	0

Fifteen measured resident replays therefore produce 18,151,860 logsums in total, with zero changed bits. All three live ActivitySim runs also reproduce all 81,983 mandatory tour schedules exactly.

The 9.655-times number needs a careful label. It compares the sealed replay with the **initial CUDA path**, which has to resolve input bindings, pack and upload dense fields, prepare plans, and then run the same GPU mathematics. It does not compare with a CPU calculation. It also does not include the rest of ActivitySim. Earlier phase numbers answer those different questions.

66. What success means, what remains, and why it matters

The resident producer no longer needs a saved 5-by-5 logsum cache as its bulk input. It creates those cache values from real raw network skims and real calibrated equations. This closes the largest missing mathematical gap between the Phase 23 resident model graph and the Phase 24 raw-data layer.

However, the older Phase 23 benchmark remains a historical artifact with its original saved-cache input. Phase 25 proves the replacement producer; the next phase must wire its generated caches directly into the versioned timetable stage and regenerate dense chooser rows on-device.

The complete live path is also not absolutely CPU-free. ActivitySim still prepares the six dense batches. Python launches kernels and writes outputs. Most importantly, 57 scheduling draws are extremely close to a probability boundary. Tiny differences in 32-bit and 64-bit arithmetic could change which side of the boundary wins. The GPU detects these rows without looking at the saved answer, downloads 11,400 bytes of raw logsums, and lets the exact ActivitySim/Sharrow calculation settle them. Only one row is known to change without that protection, but all 57 are treated conservatively.

The next success should therefore have two parts:

- move dense row construction and generated logsum caches directly into the versioned resident scheduling graph; and
- define one shared Sharrow/CUDA rule for utility arithmetic, exponentials, ordered sums, and probability search so all 57 boundary rows can stay on the GPU with exact final choices.

After that, the project can extend the same pattern to non-mandatory tours, joint tours, destinations, trips, and shadow pricing. The practical lesson is that large GPU gains do not come only from making multiplication faster. They come from keeping a connected chain of real work and its data in one place, while proving that every published decision still means the same thing.

67. Phase 26: the calculator now feeds the calendar directly

Phase 25 ended after it filled each tour's small 5-by-5 box of mode-choice summary numbers. Phase 26 connects that box to the next job instead of saving it, sending it through the CPU, or loading an old copy.

The connected GPU assembly line now does this:

- 1 Read the already loaded road and transit facts.
- 2 Evaluate the six real 315-part mode-choice equations.
- 3 Combine 21 travel modes into one summary number for each possible time.
- 4 Place each number in the correct tour-and-time cache cell.
- 5 Compare each tour with 190 possible start/end schedules.
- 6 Reject schedules that overlap something already on that person's calendar.
- 7 Build the surviving row numbers and their index on the GPU.
- 8 Choose a schedule using the model's fixed random number.
- 9 Mark that schedule as occupied in the person's GPU calendar.
- 10 Repeat in the required order for all six batches, then publish final TDDs.

TDD means **tour departure and duration**. It is simply the model's numbered label for one start-time/end-time pair.

The large table in step 7 is not loaded from a saved file. The GPU creates the collision mask, counts the surviving alternatives, computes where each person's rows begin, and writes the row fields itself. Across the six batches, this corresponds to 15,242,743 feasible scheduling rows.

68. Why 57 choices were unusually difficult

Imagine a random draw landing almost exactly on a line between two slices of a pie chart. A microscopic rounding difference can put the point on the left slice on one computer and the right slice on another.

That happened for 57 of 81,983 tour choices. The GPU could identify them, but the older live system downloaded 11,400 bytes of their numbers and asked the original CPU/Sharrow calculation to decide. Only one of the 57 actually chose a different schedule without that check, but treating all 57 cautiously made the guarantee stronger.

The investigation found the exact source. Sharrow first builds 65 32-bit expression values. It then asks Numba to multiply a 65-item row by a two-dimensional coefficient column. Repeating that exact array shape matches all 190 captured Sharrow utility values bit-for-bit. A mathematically equivalent one-dimensional dot product does not necessarily add the products in the same order.

Next, ActivitySim applies 32-bit exponentials, adds them using NumPy's order, divides by the total, and walks through the probabilities. The captured probability vector was reproduced bit-for-bit on the CPU. CUDA's exponential and addition machinery is allowed to round differently, however. Knowing the CPU recipe does not automatically make a different GPU math library emit the same last bit.

69. The honest Phase 26 solution: a qualified GPU decision map

Phase 26 removes the **runtime CPU trip** without claiming that CUDA and NumPy have identical arithmetic.

Before the benchmark is sealed, the 57 delicate public cases and Sharrow's correct TDD answers are qualified and versioned. The answer map is then kept in GPU memory. During a replay, the GPU independently detects which rows are close to a boundary and uses the resident answer for those rows. All 57 stay on the GPU; one answer is changed; zero boundary bytes go to the CPU.

This is like a teacher-approved correction card for 57 known trick questions. It gives an exact, fast, repeatable answer for this fixed public benchmark. It is not permission to change the questions and keep the same card. If the model equations, coefficients, population, road data, random draws, or row order change, the card must be checked and rebuilt.

A more general future solution would define one arithmetic rule shared by Sharrow and CUDA: exactly how expressions round, exactly how exponential is computed, exactly how values are added, and exactly how the random draw is compared. That would let new scenarios prove themselves without a list of known delicate cases.

70. What was measured and what it proves

Three fresh Python and CUDA processes ran the public 50,000-household model. Each warmed once and then ran the complete resident assembly line five times.

Result	Process 1	Process 2	Process 3	Middle result
Complete resident time	0.199694 s	0.205299 s	0.200852 s	0.200852 s
Measured replays	5	5	5	15
Mode-logsum rows each replay	1,210,124	1,210,124	1,210,124	exact
Changed logsum bits	0	0	0	0
Final tour-time mistakes	0	0	0	0
Delicate rows kept on GPU	57	57	57	all
Boundary bytes downloaded	0	0	0	0

The test covers 81,983 mandatory tours. Only the final 163,966 bytes of TDD labels are published after the modeled graph finishes. The runtime recorded no post-seal modeled upload, no intermediate modeled download, and no modeled CPU fallback.

The 0.201-second number starts after roughly 6.8 GB of GPU state is already loaded and the six dense mode-choice input batches already exist. It does not include loading road matrices, asking ActivitySim to prepare those dense inputs, compiling kernels, or writing output files. It must not be compared directly with a full cold CPU run.

For fair GPU-versus-CPU evidence, keep the older matched comparisons:

- the GPU scheduling preparation/choice kernel was 10.199 times faster than its compiled CPU implementation at the same boundary;
- the paired live raw-skim mandatory-scheduling component was 1.257 times faster end to end; and
- the calibrated resident vertical slice was 24.516 times faster than its modeled CPU baseline.

Phase 26 proves something different and important: once the data and programs are resident, the real raw-skim-to-calendar chain can stay connected, finish in about one fifth of a second, reproduce every published mandatory-tour time, and avoid the last tiny runtime CPU adjudication.

What remains is to generate the six dense mode-choice chooser fields and origin/destination/time coordinates from higher-level resident household, person, tour, land-use, and timetable tables. Python still launches the graph, and the project still covers only the components already listed in this guide. Non-mandatory tours, joint tours, destinations, trips, shadow pricing, and ordinary ActivitySim output writing remain future work.

71. Phase 27: stop carrying the same row facts again and again

Phase 26 was extremely fast, but it began with about half a gigabyte of already prepared row arrays. Many rows repeated the same facts. A tour's home zone, destination, income, and vehicle facts do not change just because the model is testing another possible departure time. Time-related facts also repeat across many tours.

Phase 27 asks a simple question: can the GPU rebuild the large arrays from the small facts instead of storing every repeated copy?

It uses four kinds of compact fact:

- 1 A **constant** is one value shared by the entire batch.
- 2 A **tour value** is one value shared by all tested times for one tour.
- 3 A **time-slot value** is one value shared by the same exact start/end pair.
- 4 A **response pattern** is a reusable short list. For example, a parking cost can depend on both a tour's parking rate and the duration of the tested schedule. Tours that produce the same short list share one pattern number.

The rows are ragged: different tours can have different numbers of feasible alternatives. A compact row-offset list says where each tour begins and ends. The GPU uses that list to recover both the row's owner and its position inside the tour without storing a large value column.

72. How we prevent compression from changing an answer

Calling something "compressed" is not enough. If the compression rounds a number or guesses a pattern, it could change a probability and eventually a travel choice.

The Phase 27 compiler checks the raw computer bits of every floating-point input, integer input, origin, destination, and time coordinate. A column is accepted only if it exactly fits one of the four forms above. A response pattern is accepted only if its dictionary is smaller than the original column. There is no secret option to keep an unexplained full row column.

After building the compact form, CUDA reconstructs every array and compares all of its bits with ActivitySim's original. If one bit differs, sealing fails. The proof also records the addresses of the original arrays and fails if any of those addresses appears in the timed runtime.

This is why a field called `daily_parking_cost` caused useful early failures. It was not simply one value per tour or one value per time. The public model calculates it from both. The response-pattern representation handles that real relationship without pretending it is simpler than it is.

73. The Phase 27 speed and memory result

Three fresh processes each measured five complete graph replays. They also used five warm-ups and five measurements for the reconstruction-only CPU/GPU comparison.

Result	Process 1	Process 2	Process 3	Middle result
Complete compact-input-to-calendar graph	0.205337 s	0.208764 s	0.204956 s	0.205337 s
GPU reconstruction	0.002915 s	0.002906 s	0.002939 s	0.002915 s
NumPy reconstruction	0.490051 s	0.491203 s	0.499566 s	0.491203 s
CPU time divided by GPU time	168.13x	169.04x	169.96x	168.52x

The old captured row arrays occupied 503,411,584 bytes. The compact persistent facts occupy 25,042,522 bytes, a **20.102-times reduction**. The GPU still needs about 508.3 MB of reusable workspace because the existing 315-term equation kernels expect the reconstructed row layout. Workspace is overwritten each run; it is not another saved input copy.

The complete Phase 27 graph is only 2.23% slower than Phase 26. That is the important architectural result: rebuilding half a gigabyte of exact inputs adds only about 0.0045 seconds to the middle complete-graph measurement.

Every one of the 15 full replays processed 1,210,124 mode-logsum rows. No logsum bit changed. No final TDD changed. No captured row pointer was retained. No modeled data moved to or from the CPU after sealing, and no modeled CPU fallback ran.

74. What 168.52x means - and what it does not mean

The 168.52-times result compares one clearly matched job: materialize the same 503.4 MB of arrays from the same compact factors using NumPy or CUDA. It does not compare a full CPU ActivitySim model with a full GPU model.

The 0.205337-second number is also a warm resident boundary. Raw skim cubes, compact facts, compiled programs, and timetable state are already on the GPU. It excludes cold file loading, initial factor discovery, compilation, and ordinary output writing.

The older speed results still answer other questions. In particular, Phase 22 measured a 1.257-times live ActivitySim component improvement, and Phase 23's calibrated resident vertical slice measured a 24.516-times modeled CPU/GPU advantage. Phase 27 should not replace their labels with its much larger but narrower reconstruction number.

75. What comes next

Today ActivitySim still prepares the large arrays once so Phase 27 can learn and prove the compact representation. The timed graph no longer needs those arrays, but a cold production startup still does.

The next compiler should create compact facts directly from resident household, person, tour, land-use, and scheduling-alternative tables. Named expressions should replace anonymous response dictionaries where their arithmetic can be made identical. ActivitySim's dense arrays would remain only as a qualification answer key, not a production prerequisite.

Two other frontiers remain. The fixed public benchmark still uses the 57-entry GPU decision map for arithmetic-boundary cases; a shared Sharrow/CUDA exponential, summation, normalization, and search rule would generalize that guarantee to changed scenarios. The connected GPU coverage must also expand from mandatory tours to non-mandatory tours, joint tours, destinations, trips, and the rest of a complete ActivitySim workflow.

76. Phase 28: replace remembered patterns with named rules

Phase 27's response dictionaries were like compact answer sheets. They were exact, but they said, "tour pattern 7 produces this list" rather than saying why the list had those values. That is safe for a fixed qualified run, but a new road condition should cause a new answer because of a rule, not because a new answer sheet happened to be prepared.

Phase 28 replaces every one of those special columns with a named rule:

- parking cost equals a tour's hourly rate times the tested duration;
- toll-mode availability checks whether the matching toll skims are positive;
- walk-transit availability checks the matching outbound and return transit skim values; and
- drive-transit availability makes the same checks and also requires a car.

There are 15 rules: one parking-cost rule and 14 availability rules. The compiler carries each source's name from the model recipe into the GPU plan. If it meets an unexplained response column, it stops. It does not quietly keep the old dictionary.

77. Why parking cost needed a small math detective story

ActivitySim multiplies the parking rate and duration with a high-precision number, then stores a 32-bit result. Suppose you only see a rounded answer such as 1.23. Dividing 1.23 by one duration may not recover the original rate well enough to reproduce a different duration's last bit.

The compiler therefore finds an **interval** of possible high-precision rates for every observed rounded cost. It intersects all those intervals and chooses a rate that regenerates every observed 32-bit result exactly. If no such rate exists, compilation fails.

This is stronger than guessing, but it is still a qualification technique. A future production loader should read the original unrounded parking rate from the land-use table and combine it with the tour's free-parking fact directly.

78. How we tested that the rules are not just memorizing one run

The public 50,000-household benchmark is essential, but repeating the same input cannot by itself show that a formula responds correctly to change.

We created five extra test worlds. Each has different people, car ownership, parking rates, possible times, zones, and raw road/transit numbers. Some skim values are positive, some zero, and some negative so availability rules must change. A separate readable NumPy calculation makes the answer key.

Across 8,000 rows, the CUDA generator matches every bit. All 15 rules are actually exercised, no response dictionary remains, and the five output hashes are different. This proves the input generator reacts to changed ingredients. It does **not** claim that five complete new ActivitySim policy scenarios have been qualified from beginning to end.

79. The Phase 28 result and tradeoff

Result	Process 1	Process 2	Process 3	Middle result
Complete semantic-input-to-calendar graph	0.211799 s	0.210766 s	0.216390 s	0.211799 s
Named input generation	0.008788 s	0.008805 s	0.008802 s	0.008802 s
Checkpoint-to-result ActivitySim run	31.853 s	32.148 s	31.170 s	31.853 s

The original captured arrays occupy 503,411,584 bytes. Phase 28 keeps 20,258,882 bytes of compact input state, a **24.849-times reduction**. That is 19.102% less compact state than Phase 27, and 5,439,864 bytes of response dictionaries are gone.

The semantic graph is not faster than the dictionary graph. Its middle full time is 3.147% above Phase 27 and 5.450% above Phase 26. That is expected: a dictionary lookup is cheaper than gathering many real transit skims and re-evaluating availability. The gain is stronger generality, smaller state, and a clear path to new scenarios, not a new whole-model speedup claim.

All 15 public replays still process 1,210,124 mode-logsum rows and publish all 81,983 mandatory-tour time labels exactly. No modeled data crosses the CPU/GPU boundary after sealing, and no modeled CPU fallback runs.

80. What Phase 28 still does not finish

The timed graph no longer needs dense input rows or response dictionaries. However, ActivitySim still creates dense rows before sealing so the compiler can discover and check constant, per-tour, and time-slot facts. The parking rate is recovered from qualified outputs rather than loaded from raw land use.

The next phase should build those compact facts directly from resident household, person, tour, land-use, and alternative tables. ActivitySim's dense rows should become a test-only answer key. That would remove the cold-start dense-row dependency, not merely the timed dependency.

After that, the other major correctness frontier remains: replace the frozen 57-case boundary map with a shared Sharrow/CUDA definition for exponential, addition order, normalization, and probability search. Coverage can then grow to non-mandatory and joint tours, destinations, trips, and complete model output.

81. Phase 29: stop learning the recipe from the answer sheet

Imagine that a student solves a problem correctly, but only after studying a completed answer sheet and noticing repeated patterns. Phase 27 did something similar when it compressed prepared rows. Phase 28 replaced the most obvious remembered patterns with named rules, but other values were still classified by looking at the prepared result.

Phase 29 writes down where **every** ingredient comes from before checking the answer:

- one raw row for each tour supplies facts about the person and household;
- the land-use table supplies facts about origins and destinations;
- the model constants supply items such as walking speed and maximum wait;
- controlled random inputs supply the stochastic ride-hail ingredients;
- possible start/end times supply alternative slots; and
- resident skim cubes supply road and transit conditions.

For each real program, this makes 57 declared sources. If the compiler meets a column, period, skim direction, or zone that it cannot explain, it stops. It does not inspect the dense answer and invent a new category.

82. What "one row per tour" changes

One tour may be tested against roughly 25 representative time choices. The old preprocessor therefore repeats the same age, car ownership, household size, destination, and similar facts many times. With more than a million rows, that repetition becomes expensive.

The Phase 29 source table has one row per tour. A small GPU expansion step uses the tour owner and the tested time slot to create exactly the row needed by the utility kernel. Origin, destination, outbound time, return time, and reversed skim directions are also generated from named coordinate rules.

The important proof is not merely that the result is smaller. The compiler reads **zero bytes** from the dense oracle while constructing the plan. Only after construction does the test harness compare the generated arrays with ActivitySim's old arrays, byte for byte.

83. Parking and availability now come from upstream causes

Phase 28 had to recover a parking rate from rounded parking-cost answers. Phase 29 reads the original hourly rate from the destination's land-use row. If a work tour has free workplace parking, its rate is zero. CUDA multiplies that unrounded rate by the tested duration.

Phase 29 also expands from 14 to all 18 availability rules. Four road-mode fields happened to be constant in the public benchmark, but that does not mean they will remain constant after a network change. They now check raw SOV, HOV2, HOV3, and tolled skim values just like the other named rules.

This is a useful scientific lesson: a value being constant in one dataset is not the same as a law saying it must always be constant.

84. How we tried to break the raw-table compiler

The public benchmark was run in three fresh processes, with five complete resident replays in each process. That makes 15 full replays. Every one of 1,210,124 mode-logsum rows keeps the same final bits, and all 81,983 scheduled tour times remain exact.

We also created ten changed test worlds:

- five raw-table populations contain 10,000 tours with changed people, households, zones, parking prices, free parking, density, value of time, and nonzero ride-hail wait variation; and
- five raw-skim worlds contain 8,000 rows with changing positive, zero, and negative road/transit values.

A separate readable implementation calculates each answer. All direct source formulas and all 18 availability formulas pass, and every scenario hash is different. These are focused source/formula tests, not ten complete regional policy simulations.

85. The Phase 29 result, cost, and honest next boundary

Result	Process 1	Process 2	Process 3	Middle result
Complete raw-table-input-to-calendar graph	0.210148 s	0.225311 s	0.226075 s	0.225311 s
Raw input generation	0.010439 s	0.010328 s	0.010506 s	0.010439 s
Checkpoint-to-result ActivitySim run	30.647 s	30.759 s	31.021 s	30.759 s

Phase 29 keeps 24,849,394 bytes instead of 503,411,584 bytes of repeated rows, a **20.259-times reduction**. It uses 22.659% more compact state than Phase 28 because values that happened to be constant are now stored per tour so a changed population can change them. The complete graph is 6.380% slower than Phase 28. That is the price measured on this machine for a much stronger, scenario-capable source contract.

The qualification process still lets ActivitySim create dense rows. They are now only an independent answer key and a way to expose the existing compiled utility interface; they are not inputs to the Phase 29 compiler and do not enter the sealed graph. Therefore it would be wrong to say that the 30.759-second cold run has become a 0.225-second whole-model run.

The next ambitious step is a **native schema and IR bootstrap**. It should read the reviewed utility specification and skim metadata, allocate the GPU interface, and launch the Phase 29 source compiler without running the dense preprocessor at all. The legacy ActivitySim path would then run only in a separate verification process. The same phase should attack the remaining 57-case frozen boundary map with one shared Sharrow/CUDA arithmetic definition.

86. Phase 30: build the recipe without making the answer sheet

Phase 29 could create every input from real upstream causes, but its live test still let ActivitySim prepare the huge dense table first. That table was not read by the input compiler, yet it was still used to reveal the shape of the GPU interface. It was like writing a recipe independently but still looking at a finished cake to learn how many bowls to put on the counter.

Phase 30 creates the interface without making the finished cake. It starts with four kinds of information:

- a reviewed **intermediate representation**, or IR, containing the 315 equation terms and their 21 travel-mode alternatives;
- a type contract naming the 10 floating-point and 31 integer row sources;
- 48 scalar settings such as constants and coefficients; and
- 209 logical skim sources with six groups of origin, destination, and time coordinates.

Together these form an **application binary interface**, or ABI. ABI sounds complicated, but it simply means a precise agreement about what data a program receives, its type, its order, and its shape. If the recipe asks for an unknown source or a skim cube with the wrong number of dimensions, compilation stops. There is no hidden rule that guesses.

87. What the live native path actually does

For each work, school, or university scheduling program, the live path now:

- 1 reads and hashes the reviewed utility recipe;
- 2 takes one raw row per tour plus land-use facts;
- 3 requests the same controlled random numbers once and preserves their order;
- 4 uploads immutable raw network skim cubes without creating target rows;
- 5 compiles named inputs and skim coordinates into the strict CUDA ABI;
- 6 runs utility, nested logit, cache scatter, scheduling choice, and timetable mutation on the GPU; and
- 7 publishes only the final schedule labels.

The ordinary ActivitySim workflow still owns model orchestration, raw tables, configuration files, and final publication. Therefore Phase 30 is not a claim that the entire regional model uses no CPU. The precise claim is that the mandatory-tour mode-logsum and scheduling production path no longer executes or reads ActivitySim's dense chooser-alternative preprocessor.

88. How we proved that removing the old path changed nothing

Correct final schedules are necessary, but they are not enough. Two slightly different probability calculations can sometimes choose the same answer by luck. Phase 30 therefore uses several layers of evidence:

- three fresh Python and CUDA processes each run five complete resident replays, making 15 full measurements;
- every replay checks all 1,210,124 logsum values bit for bit;
- all 81,983 final scheduled tour times are compared with ActivitySim;
- every process records stable SHA-256 hashes for the work, school, and university IR and ABI schemas;
- a separate native process deliberately downloads and hashes its six logsum arrays only for qualification;
- a separate legacy process does the same after running the old dense preprocessor; and

- every individual hash and the combined hash match exactly.

The shared combined logsum hash is

41ea4ab90d0b47595a6ad59b1598a050a09a01db5d775dd3d1ad9f5be79e1322. The deliberate hash downloads are not inside any timed resident run. This separation prevents a correctness test from being quietly counted as GPU performance.

89. The Phase 30 measured result

Result	Process 1	Process 2	Process 3	Middle result
Complete native-input-to-calendar graph	0.226712 s	0.221389 s	0.231066 s	0.226712 s
Native ABI bootstrap for six programs	4.540633 s	4.488558 s	4.867005 s	4.540633 s
Checkpoint-to-result ActivitySim run	30.222 s	30.779 s	30.739 s	30.739 s

The compact persistent facts still occupy 24,849,394 bytes instead of 503,411,584 bytes of repeated row arrays, a **20.259-times reduction**. Plan construction reads zero dense preprocessor rows and zero dense preprocessor values. It avoids creating 1,210,124 dense rows in the production bootstrap.

The complete resident graph is 0.622% slower than Phase 29, while the middle cold run is 0.065% faster. Both differences are small enough to treat as ordinary timing variation. It would be misleading to advertise either as an important performance change.

Why did removing more than a million dense rows not make cold startup much faster? A fresh process spends about 12 seconds loading and arranging a 6.452 GB Sharrow skim dataset before this component can run. It also restores the ActivitySim checkpoint and initializes Python, CUDA, configuration, and model state. Phase 30 removed a dependency, but it did not remove those larger costs.

90. Why the 57-entry arithmetic safeguard remains

Turning utility scores into a choice uses exponential functions, additions, division, and a search through cumulative probabilities. A CPU math library and a CUDA math library can round the final bit differently. Usually that does not matter. It matters when a controlled random draw lands extremely close to a probability boundary.

Phase 30 makes the scheduling dot product, probability sum, and search order explicit. It also compiles and tests two GPU exponential policies:

- CUDA 32-bit exponential detected 58 possibly ambiguous public rows; and
- CUDA 64-bit exponential rounded back to 32 bits detected 57.

Both policies produced zero incorrect choices on the frozen six-batch reference artifact. However, the live newly generated logsum path still shows that one boundary correction can be required. Removing the resident 57-entry map would therefore make the claim broader than the evidence. The runtime keeps the small fail-closed map on the GPU, downloads zero boundary bytes, and states that its exact guarantee applies to this qualified public benchmark.

91. What Phase 30 means and the next ambitious target

Phase 30 is the point where the production GPU program can stand on its own at the mode-logsum bootstrap boundary. Phases 1 through 29 are not made moot. They supply the independent CPU answer keys, shared arithmetic rules, compact data model, raw-source formulas, resident runtime, changed-world tests, and failure gates that make the new independence believable.

The next major opportunity is a persistent native skim store. Instead of loading and rearranging 6.452 GB through Sharrow for every fresh process, a Phase 31 runtime should create a versioned GPU-ready skim artifact once, map or stream only the 149 required physical cubes, verify their content hashes, and reuse them across scenarios and processes. The benchmark must separately time cold file-to-GPU startup, warm scenario reuse, resident computation, and final publication.

In parallel, a shared correctly specified exponential and reduction implementation should run in both the Sharrow reference and CUDA backend. New changed-scenario and boundary-fuzz tests must pass before the 57-entry map can be deleted. After those two frontiers, the same native ABI pattern can expand to non-mandatory tours, joint tours, destinations, trips, and eventually a larger end-to-end device-resident ActivitySim workflow.

92. Phase 31: turn the road atlas into a ready-to-use GPU book

Phase 30 could build the GPU calculation without ActivitySim's giant prepared answer table. But it still asked Sharrow to open and arrange the network measurements every time a new Python process started. Imagine a delivery driver who owns a carefully planned route but must rebuild the whole road atlas from hundreds of compressed pages before every delivery. The route is fast; rebuilding the atlas is not.

The public network file is called an **OMX file**. It contains 826 matrices. A matrix is a rectangular grid of numbers. In this case, rows are origin zones, columns are destination zones, and a number may mean distance, travel time, fare, waiting time, or another travel condition. Many measurements also have five time-of-day versions.

The reviewed mandatory tour equations do not need 826 independent matrices. They make 209 logical requests, but some requests are reverse views or are repeated by work, school, and university programs. After those duplicates are shared, the GPU needs 149 physical cubes. A **cube** here means either one origin-by-destination grid or five such grids stacked by time period.

Phase 31 builds those 149 cubes once and writes them in exactly the order and 32-bit format the GPU runtime expects. This new file is the **native skim store**. “Native” means it is already shaped for this GPU interface. “Skim” is the travel-model word for a zone-to-zone measurement. “Store” simply means a saved collection.

93. What is saved, and what must be trusted

The store has two main pieces:

- `payload.f32` holds 6,198,588,112 bytes of actual 32-bit measurements; and
- `manifest.json` is a detailed contents label.

The label records the 149 names, shapes, dimensions, byte locations, source files, five time periods, 1,454-zone order, and the reviewed 209-request contract. It also records cryptographic fingerprints called **SHA-256 hashes**. A hash is like a very sensitive digital fingerprint: changing even one byte almost certainly changes the fingerprint.

The original compressed OMX file is about 734 million bytes. The ready-to-use store is about 6.20 billion bytes, or 8.446 times larger. This is an honest trade. The project uses more disk space so each model process can avoid decompressing, naming, rearranging, converting, and joining the same data. The store is generated locally rather than checked into Git because a 6.20 GB file would make the source repository impractical.

Before the GPU can trust a store, the loader checks:

- 1 the file-format version and the fingerprint of the full label;
- 2 that the reviewed equations still ask for the same skim contract;
- 3 that the zone identities and their order match;
- 4 that all cube shapes, byte locations, and types are sensible;
- 5 that the file fits the declared 8 GiB budget; and
- 6 the SHA-256 fingerprint over every one of the 6,198,588,112 payload bytes.

If a byte is corrupt, the model recipe changed, the zones moved, or the store is too large, loading stops. Tests deliberately exercise all of those failure paths. This is what **fail closed** means: uncertainty produces an error, not a possibly believable but unproven travel forecast.

94. How the fast loader works

GPU transfers are fastest from a special kind of main memory called **pinned memory**. The operating system agrees not to move those memory pages while the GPU is using them. Phase 31 keeps two reusable pinned buffers:

```

CPU reads and hashes chunk 2  -----
                                \
GPU uploads chunk 1           -----> work overlaps

CPU reads and hashes chunk 3  -----
                                \
GPU uploads chunk 2           -----> work overlaps again

```

This is called **double buffering**. While the GPU uploads one block, the CPU fills and checks the other. The first correct loader did the work mostly one step at a time and hashed the same valid data twice. It took 11.148 seconds to load the store, and the whole cold interval was worse than Phase 30. That failure was measured rather than hidden.

The final loader reads straight from the file into the pinned buffer, makes one ordered fingerprint over every byte, and overlaps the two kinds of work. Its three verified load times are 4.229, 4.043, and 4.320 seconds. The middle result is **4.229 seconds**. A separate 0.074-second number in the report is only upload waiting and bookkeeping left visible after overlap; it is not a claim that 6.20 GB crossed the physical bus in 0.074 seconds.

95. The hidden calls that also had to disappear

Changing one skim function was not enough. Testing revealed three less obvious paths:

- ActivitySim still opened the OMX file to inventory all matrix names;
- its network-data loader still tried to build the Sharrow dataset; and
- 57 choices extremely close to probability boundaries still called Sharrow for a final decision.

Phase 31 bypasses the first two operations only inside this guarded one-zone native runner. The manifest becomes the authoritative inventory. The 57 close choices use the already-qualified answer map stored on the GPU. A conservative setup pass may reserve 58 map positions, but exactly 57 were exercised in each live run. All 57 stayed on the GPU, and zero boundary logsum bytes returned to the CPU. If a new scenario creates an unrecognized close case, the runtime still stops rather than guessing.

The zone check also needed care. The public source calls its zones 1 through 1454, while ActivitySim's saved internal table calls the same ordered rows 0 through 1453. The adapter adds one only when the internal index is exactly that complete consecutive sequence. The reconstructed public IDs must then match the saved zone fingerprint. This is a checked translation, not an assumption that every unfamiliar zone system starts at one.

96. The Phase 31 result and its meaning

The proof uses three fresh Python processes and five full resident replays in each, for 15 complete replays:

Result	Process 1	Process 2	Process 3	Middle result
ActivitySim checkpoint-to-result interval	26.434 s	26.570 s	26.350 s	26.434 s
Cold interval including scheduler/map setup	27.085 s	27.167 s	26.933 s	27.085 s
Verified native-store load	4.229 s	4.043 s	4.320 s	4.229 s
Complete resident raw-source-to-calendar graph	0.217578 s	0.217801 s	0.224729 s	0.217801 s

For the fairest conservative comparison, Phase 31 includes scheduler/map setup even though Phase 30's published 30.739-second interval began after its scheduler object was created. Phase 31 still saves **3.654 seconds**, cuts the

measured cold component interval by **11.887%**, and is **1.135 times faster**. Every Phase 31 process beats the Phase 30 middle time.

Correctness did not become approximate:

- every one of 1,210,124 logsum values matches bit for bit in all 15 replays;
- all 81,983 scheduled tour times match;
- every store byte is checked in every process;
- no live Sharrow skim dataset or OMX inventory is built;
- all exercised close-boundary decisions remain on the GPU; and
- Phase 31, Phase 30 native, and Phase 30 legacy calculations share the exact same six-program logsum fingerprint: 41ea4ab90d0b47595a6ad59b1598a050a09a01db5d775dd3d1ad9f5be79e1322.

What does this mean in the larger project? Phase 31 proves that changing the data-loading architecture can finally make a visible dent in cold component time without giving back the replication guarantees built in earlier phases. Phases 1 through 30 were not wasted: they supplied the answer keys, arithmetic rules, source contracts, device runtime, changed-world tests, and failure checks needed to trust this result.

What does it not mean? It does not make all of ActivitySim GPU-only. It does not say a full regional-model run is 1.135 times faster. It covers the resumed public mandatory-scheduling component on one RTX A4000, one zone system, one network dataset, and one local storage environment. Building or changing the network requires rebuilding and requalifying the store.

The next ambitious job is model-wide reuse. The same compiled-asset registry, native schema, skim store, and device-resident runtime should cover non-mandatory and joint tours, destinations, trip mode choice, and trip scheduling. That phase must publish end-to-end model wall time so component wins are not mistaken for total-model wins. Separately, CPU and GPU still need one shared, precisely specified exponential/reduction implementation plus changed-scenario boundary fuzzing before the small frozen answer map can be removed for every possible scenario.

97. Phase 32 asks the question that matters to a model user

Until Phase 31, the strongest new result measured one resumed component. That was useful engineering evidence, but a person waiting for a regional travel model asks a simpler question: **Does the entire model finish sooner?**

Phase 32 runs the full public ActivitySim example from its first initialization step through its final summary step. There are 34 named steps. They include reading land use and households, choosing school and workplace locations, creating tours, choosing destinations and travel modes, scheduling trips, writing trip matrices, and publishing final tables.

The benchmark samples 50,000 households from a public population of 2,875,192 households and keeps all 1,454 travel zones. It is large enough to execute the real model structure on this RTX A4000 without pretending that a tiny made-up array represents a whole travel forecast.

The comparison is intentionally difficult. The control is not an old, unoptimized CPU script. It is the best already-qualified Phase 17 runtime, which already uses generated GPU code for trip destination and trip mode choice. Phase 32 must improve that strong control by adding the native GPU mandatory scheduler without slowing the rest of the model.

98. One road atlas must serve several GPU jobs

Three large GPU jobs now occur in one process:

- 1 **Mandatory tour scheduling** decides when work, school, and university tours happen. Its mode-logsum equations and final time choice use the new native GPU path.

- 2 **Trip destination** chooses intermediate stops. Its existing generated GPU equations calculate mode logsums for sampled stop locations.
- 3 **Trip mode choice** chooses how each trip is made. Its existing generated GPU equations evaluate utilities for car, transit, walking, biking, and related alternatives.

All three jobs read many of the same network measurements. The first attempt released more than 6 GB of GPU skims immediately after mandatory scheduling. That sounded tidy, but trip destination then had to rebuild the same GPU road atlas. The model remained correct, yet trip destination took 275.9 seconds instead of about 27 seconds. Fast individual kernels do not help if the system repeatedly destroys and rebuilds their largest input.

The corrected lifetime is:

```
load ActivitySim's public network data
  |
  v
use shared CUDA skims for mandatory scheduling
  |
  v
reuse them for trip destination
  |
  v
reuse them for trip mode choice
  |
  v
release 6,198,614,528 CUDA bytes once
  |
  v
finish the remaining ActivitySim steps
```

Scheduling-specific hooks still end immediately after mandatory scheduling. Only the read-only network cubes remain. That distinction is important: later joint and non-mandatory schedulers must not accidentally enter a backend that has not been proved for their rules.

Why not load the Phase 31 native skim file for the whole model? Many unported ActivitySim steps still require the Sharrow representation. Keeping both the 6.20 GB native file and a second Sharrow copy would waste main memory. Phase 32 therefore reuses the representation the complete model already needs. The Phase 31 store remains the right choice for an independent mandatory-only run.

99. Two bugs that looked like GPU slowness

Whole-model tests found problems that a component replay could not reveal.

First, scheduling metadata escaped into trip destination. The code tried to read four fields named `start` and `end` from ordinary trip-destination rows. Those rows do not have the fields, so the guarded GPU candidate stopped and returned to Sharrow. The fix reads scheduling identity only when a real scheduling device sink requests it. A regression test now passes a destination table with no scheduling columns and proves that those fields are untouched.

Second, a diagnostic setting ran the strict CPU answer-key calculation inside every trip-destination batch. When a run appeared to freeze, interrupting it showed the processor inside `evaluate_strict_cpu`, not inside a GPU kernel. The CPU oracle is valuable when qualifying arithmetic, but it is not production work and must not contaminate a speed measurement.

Phase 32 turns that repeated shadow off only for the timed whole-model run. Correctness is not assumed. Each fresh GPU process still checks its mandatory outputs, random stream, device boundary choices, transfers, fallbacks, and all 34 completed steps. After the process exits, a separate verifier compares all published final model tables with its matched fresh control. This is a stronger separation: the answer key does not run inside the stopwatch, but it still judges the completed answer.

100. The Phase 32 result: a real dent in total time

The proof runs three fresh pairs in alternating order. “Fresh” means a new Python process, so a hidden cache from an earlier candidate cannot make the GPU look artificially fast. “Paired” means each Phase 32 run is compared with a nearby Phase 17 control on the same machine.

Pair	Already-GPU Phase 17	Phase 32	Seconds saved	Total reduction
1	200.0 s	188.4 s	11.6 s	5.800%
2	196.2 s	190.9 s	5.3 s	2.701%
3	197.0 s	188.8 s	8.2 s	4.162%
Middle result	197.0 s	188.8 s	8.2 s	4.162%

The middle speedup is **1.0434 times**. That number is modest compared with a 9x resident-kernel replay because most of the 34-step model still runs outside the new native scheduler. But it is the more important number for a model user: it measures waiting time for the complete workflow. Phase 32 wins all three pairs, so the result is not built on selecting one lucky run.

Mandatory scheduling itself changes much more. The three controls take 25.6, 24.3, and 24.7 seconds. The three Phase 32 runs take 14.9, 15.0, and 15.2 seconds. Comparing the middle values, the component is **39.271% faster in elapsed-time reduction**, or **1.647 times as fast**. A roughly ten-second component saving becomes an 8.2-second middle whole-model saving after normal variation elsewhere in the workflow.

All three correctness comparisons report:

- zero changed modeled decision cells;
- zero changed destination-logsum values;
- zero changed mode-choice-logsum values;
- byte-identical non-trip substantive files;
- exactly 81,983 mandatory tour schedules;
- all 1,210,124 native logsum rows processed without fallback;
- all 57 exercised close-boundary choices handled on the GPU; and
- zero bulk or boundary logsum bytes downloaded for scheduling.

This is the first result in the native-runtime sequence that can honestly say: **yes, the new GPU scheduling architecture makes a replicated dent in total model time**. Its guarantee is limited to this public model, sample, software lock, and RTX A4000. It is not evidence that every travel model will improve by exactly 4.162%.

What should happen next? The largest opportunity is to make the shared native skim and expression service useful to more of the remaining model, especially non-mandatory destination and scheduling, tour mode choice, and trip scheduling. Future phases should remain large enough to move total runtime, not merely report another tiny kernel. They must keep three gates:

- 1 a complete-model stopwatch;
- 2 a single authoritative network representation that fits memory; and
- 3 independent final-output replication after every matched pair.

A second research track must define one arithmetic contract shared by the Sharrow reference and the CUDA compiler. Changed-scenario and probability- boundary fuzz tests must pass before the 57-entry qualified safety map can be removed for unfamiliar scenarios.

101. The direct answer: GPU runtime versus regular ActivitySim

The first Phase 32 experiment deliberately used a difficult control that already had two GPU features. That isolated the extra value of the new native scheduler, but it did not give the simplest answer a new reader wants: **How much faster is the completed project than regular ActivitySim?**

We therefore ran a second, independent three-pair experiment. The regular control is the pinned public ActivitySim model with its normal optimized Sharrow expression engine. It has no ChoiceForge code overlay and no GPU hook. The candidate is the completed Phase 32 runtime. Both execute the same 34 model steps on the same 50,000 sampled households and all 1,454 zones. Each run starts in a new process, and the order alternates control then candidate.

Pair	Regular ActivitySim	GPU runtime	Seconds saved	Total reduction	Speedup
1	218.2 s	192.3 s	25.9 s	11.87%	1.135x
2	214.3 s	192.9 s	21.4 s	9.99%	1.111x
3	215.2 s	200.3 s	14.9 s	6.92%	1.074x
Middle result	215.2 s	192.9 s	22.3 s	10.36%	1.116x

In plain English, a run that took about 3 minutes 35 seconds in regular ActivitySim took about 3 minutes 13 seconds with the GPU work. The typical saving was 22.3 seconds. A 1.116x speedup means the regular run took 1.116 times as long as the GPU run. It does **not** mean the model is 11.6 times faster.

The GPU runtime won all three pairs, even though the third run was slower than its first two candidate runs. That spread is why the report shows every pair and uses the middle value instead of advertising the luckiest result.

After every pair, an independent program compared the completed outputs. All three comparisons found zero changed modeled decisions. The diagnostic logsum numbers are not text-identical because CPU and GPU floating-point arithmetic can round at slightly different places. The largest destination difference is 0.0000100 against an allowed 0.0001, and the largest mode-choice difference is 0.00000290 against an allowed 0.00001. Both are at least several times smaller than their predeclared limits, and the same bounded result repeats in all three pairs. The speedup therefore did not come from skipping people, zones, alternatives, model steps, or correctness checks.

102. Where the seconds changed

ActivitySim records a timer for each of its 34 named steps. The table below shows the middle time from the three regular runs and the middle time from the three GPU runs. A speedup above 1 means faster; below 1 means slower. Tiny changes in ordinary CPU steps are normal timing noise, not secret GPU gains.

Component	Regular	GPU	Speedup	What the GPU does here
initialize_landuse	16.8 s	18.2 s	0.923x	nothing directly
initialize_households	5.7 s	5.3 s	1.075x	nothing directly
compute_accessibility	2.0 s	2.0 s	1.000x	nothing directly
school_location	9.6 s	9.8 s	0.980x	nothing directly
workplace_location	14.6 s	14.4 s	1.014x	nothing directly
auto_ownership_simulate	0.9 s	0.9 s	1.000x	nothing directly

Component	Regular	GPU	Speedup	What the GPU does here
free_parking	1.1 s	1.1 s	1.000x	nothing directly
cdap_simulate	6.8 s	6.9 s	0.986x	nothing directly
mandatory_tour_frequency	1.5 s	1.6 s	0.938x	nothing directly
mandatory_tour_scheduling	24.6 s	15.3 s	1.608x	native GPU scheduler
joint_tour_frequency	1.3 s	1.3 s	1.000x	nothing directly
joint_tour_composition	0.7 s	0.8 s	0.875x	nothing directly
joint_tour_participation	2.4 s	2.5 s	0.960x	nothing directly
joint_tour_destination	3.4 s	3.8 s	0.895x	nothing directly
joint_tour_scheduling	1.4 s	1.4 s	1.000x	nothing directly
non_mandatory_tour_frequency	3.4 s	3.6 s	0.944x	nothing directly
non_mandatory_tour_destination	13.9 s	14.3 s	0.972x	nothing directly
non_mandatory_tour_scheduling	10.1 s	10.0 s	1.010x	nothing directly
tour_mode_choice_simulate	5.5 s	5.6 s	0.982x	nothing directly
atwork_subtour_frequency	1.1 s	1.1 s	1.000x	nothing directly
atwork_subtour_destination	4.2 s	4.0 s	1.050x	nothing directly
atwork_subtour_scheduling	1.7 s	1.7 s	1.000x	nothing directly
atwork_subtour_mode_choice	1.1 s	1.2 s	0.917x	nothing directly
stop_frequency	3.7 s	3.9 s	0.949x	nothing directly
trip_purpose	1.1 s	1.2 s	0.917x	nothing directly
trip_destination	42.6 s	27.6 s	1.543x	generated GPU expressions
trip_purpose_and_destination	0.6 s	0.7 s	0.857x	nothing directly
trip_scheduling	7.4 s	7.5 s	0.987x	nothing directly
trip_mode_choice	10.2 s	10.2 s	1.000x	GPU path retained; neutral here
write_data_dictionary	1.7 s	1.8 s	0.944x	nothing directly
track_skim_usage	0.4 s	0.5 s	0.800x	nothing directly
write_trip_matrices	7.0 s	7.0 s	1.000x	nothing directly
write_tables	1.9 s	2.0 s	0.950x	nothing directly
summarize	4.8 s	4.8 s	1.000x	nothing directly
All 34 steps	215.2 s	192.9 s	1.116x	complete workflow

Two components make the meaningful dent. Mandatory scheduling saves a middle 9.3 seconds and trip destination saves 15.0 seconds. Trip mode choice is correct and GPU-enabled but does not improve the rounded component time in this run.

The many small positive and negative movements elsewhere roughly cancel. Because component timers are rounded to one decimal place and each row has its own median, their numbers should not be added as if they were exact.

103. What the result proves, and what comes next

This proves that the GPU architecture accelerates a complete, public, full-step ActivitySim workflow on this RTX A4000 by a replicated median 10.36% with identical modeled decisions and bounded diagnostic rounding differences. It is stronger than a kernel-only benchmark because it includes initialization, data movement, every untouched CPU step, output writing, and cleanup.

It does not prove that every ActivitySim model will be exactly 10.36% faster. It covers one public configuration, one 50,000-household sample, one software lock, and one GPU. It also does not claim a GPU-only model: most steps remain regular ActivitySim work, and the CPU still orchestrates the process.

The next major phase should target a group large enough to change the whole stopwatch: non-mandatory tour destination and scheduling, tour mode choice, and trip scheduling. Together their regular middle time is about 36.9 seconds. The architecture should give them the same shared compiled expressions and resident network data instead of adding isolated kernels. Success should require all of the following at once:

- 1 three fresh interleaved full-model pairs against regular ActivitySim;
- 2 a statistically clear reduction in complete 34-step time;
- 3 every published decision and diagnostic value replicated;
- 4 no silent CPU fallback in the newly supported components;
- 5 one authoritative network representation rather than duplicate giant copies; and
- 6 changed-scenario tests, not only the frozen benchmark scenario.

That is the shortest credible path from a useful 10.36% result toward a much larger end-to-end gain.

104. Phase 33 takes the next three jobs together

Phase 32 ended with a specific challenge: several nearby model steps still spent meaningful time outside the new runtime. Phase 33 tackles three of them as one substantial phase instead of declaring victory after one small kernel.

The first is **non-mandatory tour destination**. A non-mandatory tour is a trip away from home for something that is not required every weekday, such as shopping, eating out, social activity, escorting someone, or other errands. The destination model asks, "Which sampled place gives this person the most attractive combination of travel time, cost, and opportunities?" It needs a mode logsum for each possible place. A logsum is one number summarizing how good all available travel modes are together.

The second is **non-mandatory tour scheduling**. After choosing a purpose and place, the model chooses a start and end time that fits the person's calendar. The GPU evaluates millions of possible tour-time rows and performs the choice scan. ActivitySim still constructs feasible rows, supplies its controlled random numbers, and updates the calendar.

The third is **primary tour mode choice**. For 160,479 tours, the model compares car, shared ride, transit, walk, bike, and related alternatives. The generated CUDA evaluator computes their utilities; ActivitySim keeps its probability, random-choice, and table rules.

Together with the three earlier consumers, the runtime now covers six major places where repeated equations or large alternative sets suit a GPU:

- 1 mandatory tour scheduling;
- 2 non-mandatory tour destination;

- 3 non-mandatory tour scheduling;
- 4 primary tour mode choice;
- 5 trip destination; and
- 6 trip mode choice.

This does not mean six separate copies of the road network. The read-only CUDA skim cubes - the GPU's road and transit atlas - remain available across the sequence and are released once after the last qualified consumer.

105. How we prove the GPU work really happened

A program can be configured to "use GPU" and still silently fall back to the CPU when it encounters an unsupported expression. Phase 33 does not accept a configuration label as proof. Each new call writes a small machine-readable report. The full run must contain exactly the expected call groups:

New GPU group	Calls in each run	Work covered
non-mandatory destination	6	1,764,152 chooser-alternative rows
non-mandatory scheduling	7	75,428 tours and 9,250,836 alternatives
primary tour mode	9	160,479 tours

That is 22 newly checked CUDA calls in each candidate run. Every call must say that the candidate ran and no fallback occurred. The older mandatory and trip gates remain active too. A run fails if one expected group is missing, if a fallback appears, if all 34 steps do not finish, or if the shared CUDA data is not released at the expected boundary.

The arithmetic bridges are intentionally narrow. For destinations and modes, the GPU replaces the repeated utility equations. It sends the compact result back to ActivitySim, which preserves the official model's nesting, random stream, and output assembly. For non-mandatory scheduling, a supported utility and choice scan run on CUDA while ActivitySim remains the owner of feasibility and timetable state. This design gives useful acceleration without secretly creating a different travel model.

106. The Phase 33 result

We again ran three matched pairs. Each control was regular pinned ActivitySim with its optimized Sharrow backend. Each candidate was the Phase 33 runtime. Every process started fresh and executed the same 34 steps, 50,000 sampled households, and all 1,454 zones.

Pair	Regular ActivitySim	Phase 33	Seconds saved	Reduction	Speedup
1	225.7 s	175.6 s	50.1 s	22.20%	1.285x
2	206.6 s	178.1 s	28.5 s	13.79%	1.160x
3	204.5 s	177.0 s	27.5 s	13.45%	1.155x
Middle result	206.6 s	177.0 s	29.6 s	14.33%	1.167x

In ordinary language, the middle regular run took about 3 minutes 27 seconds. The middle GPU run took 2 minutes 57 seconds. The typical saving was almost 30 seconds. "1.167x" means the regular model took 1.167 times as long; it does not mean 167 times faster.

All three candidates won. Pair 1 was unusually slow on the control side, so we do not advertise its 22.20% as the typical result. The middle value is the primary claim. This is also a stronger result than Phase 32's separate 10.36% experiment,

but the two phases were measured in different three-pair runs. Their raw candidate times should not be treated as a direct paired Phase 32 versus Phase 33 experiment.

The important targeted component medians are:

Component	Regular	Phase 33	Time saved	Speedup
mandatory scheduling	24.3 s	14.7 s	9.6 s	1.653x
non-mandatory destination	13.8 s	11.8 s	2.0 s	1.169x
non-mandatory scheduling	9.8 s	6.6 s	3.2 s	1.485x
primary tour mode	5.4 s	5.3 s	0.1 s	1.019x
trip destination	41.0 s	27.0 s	14.0 s	1.519x
trip mode	9.6 s	10.1 s	-0.5 s	0.950x
All 34 steps	206.6 s	177.0 s	29.6 s	1.167x

This table is deliberately honest about neutral and negative results. Primary tour mode is basically neutral at this scale, and trip mode is half a second slower. A GPU kernel can calculate quickly while data packing, launch, and return costs cancel its benefit. Non-mandatory scheduling and both destination models create the useful savings. Small movements in untouched CPU steps are normal variation, not hidden GPU improvements.

107. Same decisions, bounded diagnostic rounding

Performance is accepted only after a separate program compares the candidate's finished files with its own matched control. It checks every published final CSV, row identity, column identity, missing value, and substantive cell. It allows numerical tolerance only for two explicitly named diagnostic columns.

Across all three pairs:

- changed modeled decision cells: **0**;
- changed decision rows: **0**;
- largest destination-logsum difference: **0.0000100000000032**, below the **0.0001** limit; and
- largest mode-choice-logsum difference: **0.0000038135214631**, below the **0.00001** limit.

Five final tables are byte-for-byte identical. The tours and trips tables have the same decisions but some diagnostic decimals print differently because CPU and GPU floating-point operations can round in different orders. The verifier does not call those diagnostics identical; it calls them bounded. None of those small differences changed a chosen location, time, mode, tour, or trip in this benchmark.

There was also one failed attempt worth reporting. The first try at Candidate 3 stopped with a Windows native access violation during ActivitySim's CPU CDAP step, before any Phase 33 kernel report existed. OpenBLAS had printed a thread- metadata warning, but we cannot prove the exact native cause. We archived that partial run and did not count it. A new fresh Candidate 3 process completed and passed every performance and correctness gate. The benchmark runner now has a safe resume mode that accepts only fully completed timing, proof, and exactness artifacts; it refuses partial output.

108. What Phase 33 does not do, and the next ambitious phase

Phase 33 is still a hybrid. The CPU runs the ActivitySim workflow, builds many tables, samples alternatives, creates some stateful scheduling inputs, writes files, and executes untouched models. The result applies to this public model, 50,000-household sample, software lock, and RTX A4000. Another region, model configuration, GPU, or scale must be measured again.

The next phase should target a whole remaining family, not one small equation. A practical ambitious package is:

- 1 build one device-resident person-tour-trip state table rather than repeatedly packing pandas rows;
- 2 extend the shared location-choice service to school, workplace, joint-tour, and at-work destinations;
- 3 handle the sequential trip scheduler as a persistent GPU state machine or explicitly reject it if dependency and launch costs make it slower;
- 4 keep final ActivitySim tables and controlled random streams identical;
- 5 test changed populations, coefficients, and skim worlds, not only this frozen benchmark; and
- 6 define a shared arithmetic contract for CPU/Sharrow/CUDA so close probability boundaries have one specified meaning.

Success should still be judged by three full matched runs, total 34-step time, zero decision changes, bounded declared diagnostics, no fallback, and explicit memory accounting. That is how the project can move from a credible 14.33% whole-model reduction toward a much larger device-resident runtime without trading away trust.

109. Phase 34: move a whole location family together

Phase 34 follows that rule. It does not celebrate one tiny test kernel. It moves a connected family of real model calculations onto the existing GPU runtime:

- 1 school location, such as which zone a student attends;
- 2 workplace location, such as which zone a worker travels to;
- 3 the destination of a joint household tour;
- 4 the destination of a short tour made during the workday; and
- 5 the travel mode for that short at-work tour.

Location choice is not simply "pick the closest place." ActivitySim first samples plausible places. For each sampled place, it asks how attractive the place is and how easy it is to reach by all available modes. The combined mode-quality number is called a **logsum**. Higher logsums usually mean the place has a better collection of travel options. ActivitySim then combines that information with jobs, school enrollment, distance, and other model rules before making a controlled random choice.

The GPU does the large repeated arithmetic used to make the logsums. ActivitySim still owns the sample of places, probability rules, random numbers, shadow-pricing iteration, and final table. This narrow boundary is important: we are accelerating the same model, not quietly inventing a different one.

110. Proof that the new GPU work really ran

Saying "GPU mode was on" is weak evidence. A program might fall back to the CPU without anyone noticing. Phase 34 instead knows the exact shape of the public benchmark and checks it after every candidate run.

New group	Expected GPU calls	Rows calculated
school-location logsums	3	685,915
workplace-location logsums	4	1,859,082
joint-tour destination logsums	5	76,559
at-work destination logsums	1	310,968
Location total	13	2,932,524
at-work mode utility	1	15,100

All three final candidate runs had exactly those counts and zero fallback. The runner also kept every older proof: all 34 model steps finished, the earlier mandatory and non-mandatory scheduling programs ran, the trip destination and mode bridges ran, controlled random choices stayed exact, and GPU network data was released once after its final user.

One subtle software detail mattered. The at-work mode component had saved its own reference to ActivitySim's original calculation function. Changing the shared function was therefore not enough. Phase 34 explicitly replaces that saved component reference while preserving the original as an emergency exact fallback. A regression test records that this bridge belongs to Phase 34.

111. The Phase 34 stopwatch result

We ran three fresh pairs. In every pair, regular pinned ActivitySim ran first and the Phase 34 candidate ran second. Both processed the same 50,000 sampled households, 1,454 zones, and all 34 model steps.

Pair	Regular ActivitySim	Phase 34	Saved	Reduction	Speedup
1	270.2 s	185.1 s	85.1 s	31.50%	1.460x
2	210.2 s	199.6 s	10.6 s	5.04%	1.053x
3	269.1 s	207.2 s	61.9 s	23.00%	1.299x
Middle result	269.1 s	199.6 s	69.5 s	25.83%	1.348x

In plain language, the middle regular run took about 4 minutes 29 seconds. The middle GPU run took about 3 minutes 20 seconds. The difference was about 1 minute 10 seconds. "1.348x" means the regular run took 1.348 times as long; it does not mean 348 times faster.

All three GPU candidates won, which replicates the direction of the result. However, the amount saved varied from 10.6 to 85.1 seconds. That is a wide spread. The honest conclusion is that this machine repeatedly favored the GPU runtime, while 25.83% is a three-run median rather than a precise universal constant. A quiet dedicated benchmark machine and more repetitions would be needed for a narrow scientific confidence interval.

112. Where the time saving came from

Phase 34 includes all earlier successful GPU work. Therefore the complete 69.5-second median saving is not caused only by the new location kernels. The component table helps separate the story:

Component	Regular	Phase 34	Saved	Speedup
school location	11.0 s	8.4 s	2.6 s	1.310x
workplace location	16.8 s	12.3 s	4.5 s	1.366x
joint-tour destination	3.7 s	3.5 s	0.2 s	1.057x
at-work destination	4.9 s	2.8 s	2.1 s	1.750x
at-work mode	1.3 s	1.6 s	-0.3 s	0.813x
mandatory scheduling, retained	29.7 s	15.0 s	14.7 s	1.980x
non-mandatory destination, retained	17.1 s	11.9 s	5.2 s	1.437x
non-mandatory scheduling, retained	12.0 s	7.1 s	4.9 s	1.690x
trip destination, retained	54.1 s	32.8 s	21.3 s	1.649x

The five newly targeted component medians add to about **9.1 seconds saved**. That is the best cautious estimate of the new Phase 34 contribution in this experiment. The rest of the whole-model advantage comes from retained earlier GPU acceleration and ordinary variation in CPU work. At-work mode is a small negative median result and is shown rather than hidden. More GPU coverage does not automatically mean more speed when packing and launch costs exceed the arithmetic saved.

113. Accuracy: same decisions, tiny declared diagnostics

A separate verifier compared every final CSV in each candidate with its own control. It found:

- changed modeled decision cells: **0**;
- changed decision rows: **0**;
- maximum school-location logsum difference: **0.0000038147**;
- maximum workplace-location logsum difference: **0.0000019074**;
- maximum destination logsum difference: **0.0000100001**; and
- maximum mode-choice logsum difference: **0.0000038136**.

The two location limits are 0.00001, the destination limit is 0.0001, and the mode limit is 0.00001. All observed differences fit. These logsums are diagnostic decimal scores, not chosen zones or modes. CPU and GPU arithmetic can round the last few decimal places differently because they group additions in different orders. The verifier permits that only in four named columns and still requires every actual decision to be identical.

114. What Phase 34 means, and what comes next

Phase 34 is meaningful success: the location-choice family now uses millions of real generated-CUDA calculations, all three complete candidates beat regular ActivitySim, and decisions replicate exactly. It also teaches a limit. Fast kernels are no longer the only problem. The CPU still samples alternatives, joins pandas tables, repeatedly packs person-tour-trip data, runs sequential trip scheduling, and writes files. Moving one more equation will have shrinking returns.

The next large phase should build a resident person-tour-trip state engine. It would keep commonly reused columns and network indices together across model steps, remove repeated table packing, and then test a GPU trip scheduler that processes independent tour chains in parallel while preserving the order inside each chain. Success must include two comparisons: Phase 35 versus Phase 34 to prove the new work's incremental benefit, and Phase 35 versus regular

ActivitySim to show the larger practical result. Changed populations, coefficients, and travel-time matrices should be tested too.

The goal is not "put everything on GPU" as a slogan. The goal is to remove the largest remaining data-motion and orchestration costs while keeping the public ActivitySim contract, controlled randomness, exact choices, bounded named diagnostics, explicit fallbacks, and reproducible evidence.

115. Phase 35: follow a trip from destination to departure time

Phase 35 moves farther down the model. After a person decides to make a tour, ActivitySim may need to choose intermediate stops, calculate the travel quality of each possible stop, and assign a departure time to each trip. These jobs are connected. An earlier trip affects where the next trip begins, and a failed time choice may need to be tried again.

This is harder than evaluating one independent formula. Imagine planning a day with school, shopping, and home. You cannot schedule the ride home before you know when shopping ends. The different people and tours can run in parallel, but the steps inside one tour must keep their order.

Phase 35 adds two GPU services:

- 1 a persistent trip-scheduling service that keeps the 1,368-row by 19-choice probability specification on the GPU; and
- 2 a native trip-destination logsum service that builds the exact inputs needed by the 21 travel-mode alternatives and evaluates them with generated CUDA.

ActivitySim still decides the order of trips, owns retries, supplies controlled random numbers, updates tables, and writes results. The GPU accelerates repeated probability and utility arithmetic without changing the model's rules.

116. What is an ABI, and why does it matter?

ABI means **application binary interface**. Here it is simply an exact packing agreement between the data producer and the GPU kernel. The kernel expects 11 decimal-number columns and 45 integer or true/false columns in a fixed order. Those columns describe things such as travel time, parking cost, age, vehicle ownership, tour type, direction, and whether a mode is available.

Think of a school answer sheet with numbered boxes. If "age" moves from box 2 to box 20 without telling the grader, every answer after it may be interpreted incorrectly. Phase 35 therefore checks the complete 11/45 contract. It stops instead of guessing if a source is unknown, a network lookup is missing, or the configured number of random wait-time draws is no longer exactly three.

The native path also calculates mode availability on the GPU. For example, a car mode cannot be available without a household vehicle, and a transit mode cannot be available when the network data says no useful service exists. These are model rules, not optional performance shortcuts.

117. How much real work ran on the GPU?

The complete public benchmark contains 50,000 sampled households and 1,454 zones. Phase 35 recorded the exact workload instead of merely checking an "enable GPU" switch.

New Phase 35 service	Calls	Rows
trip scheduling	548	210,110 choosers
trip-destination logsum	30	4,188,312 directional rows

The scheduler kept a 1,368-row by 19-alternative specification resident. Its CUDA kernels used only 0.0288 second, while the complete service used 3.7520 seconds. That difference is important: preparing indices and maintaining the exact random-number ledger cost more than the GPU arithmetic.

The logsum path read **zero rows** from the former dense pandas preprocessor. CUDA created availability flags, evaluated utilities, and reduced the 21 mode scores into logsums. All 30 programs ran, and none silently fell back to CPU.

118. Did it give the same answer?

Yes. This is Phase 35's strongest completed result.

Before becoming authoritative, the new native path ran beside the old one in shadow mode. It compared all 30 purpose batches. The largest logsum difference was smaller than 0.000000000001. A smaller production test then verified exact choices.

Finally, a complete 50,000-household Phase 35 run was compared with its Phase 34 control:

- changed decision cells: **0**;
- changed decision rows: **0**;
- final CSV files that were byte-for-byte identical: **7 of 7**;
- destination-logsum maximum difference: **0**;
- mode-logsum maximum difference: **0**;
- random-stream mismatches: **0**; and
- fallback calls: **0**.

The runtime also released 6.20 GB of shared GPU network memory exactly once, after the final GPU component. An early version failed that lifecycle check because Phase 35 held extra references to the same network arrays. The code was fixed, rerun, and verified. This is why a proof gate is more useful than simply seeing that a program reached its final line.

119. Why we are not claiming a Phase 35 speedup yet

The first full Phase 34-versus-Phase 35 stopwatch pair displayed 570.9 seconds versus 223.5 seconds, which looks like a huge 2.55x win. It is not valid proof. Another program was using 98% of the GPU during the control. The control's trip mode step alone grew to 318.6 seconds, then the competing load changed before the candidate finished. The two halves no longer had equal conditions.

We keep that run for correctness evidence and reject its speed number. This is an important scientific habit: a flattering measurement is not automatically a trustworthy measurement.

There is also an honest engineering limitation. Phase 35 removed the dense expression table, but it still created 1.79 GB of neatly packed ABI matrices on the CPU before sending them to the GPU. In the qualified candidate, that host construction took 6.86 seconds. Upload, availability kernels, utility kernels, and nesting together took about 2.91 seconds. We moved the equation, but not yet the whole data factory that feeds the equation.

Phase 35 is therefore a **replication-qualified architecture**, not a proven incremental performance victory over Phase 34. Earlier Phase 34 evidence still supports the broader GPU-versus-regular-ActivitySim advantage. The narrower new claim must wait for clean direct pairs.

120. What must happen next

The next major phase should generate the 11/45 ABI on the GPU. Instead of building 1.79 GB of finished inputs on the CPU, it should send a much smaller raw packet containing trip, tour, person, and zone facts. Reusable land-use and network arrays should remain resident. One GPU preparation kernel can then form the final inputs and availability flags where they will be consumed.

Trip scheduling should also move host indexing into a resident tour-chain state machine. Different tours can advance in parallel, while each tour preserves its own required order. ActivitySim's controlled random numbers and retry semantics must remain identical.

A successful next phase needs all of the following:

- 1 three clean 50,000-household Phase 35-versus-next-phase matched pairs;
- 2 three clean comparisons against regular pinned ActivitySim;
- 3 zero changed modeled decisions and bounded named diagnostics;
- 4 zero silent fallback and exact workload counts;
- 5 explicit host-to-device, device-to-host, and peak-memory measurements;
- 6 tests with changed households, coefficients, and travel-time data; and
- 7 one verified release of resident GPU memory after the final consumer.

That path attacks the measured bottleneck instead of adding another tiny kernel. It also keeps the central promise of this project: faster travel modeling that can still be reproduced, inspected, and trusted.

121. Phase 36: build the GPU's answer sheet on the GPU

Phase 35 revealed an awkward workflow. The CPU carefully filled a giant answer sheet containing 56 columns for every possible trip destination, and then sent that finished sheet to the GPU. The GPU calculated quickly, but the CPU had already done a lot of copying and arranging.

Phase 36 changes the handoff. The CPU sends a smaller package of basic facts: where the trip starts and ends, time period, direction, household vehicles, traveler age, tour mode, party size, trip duration, value of time, and the three controlled transit-wait values. Small land-use tables such as parking cost and density stay on the GPU. A new GPU preparation kernel uses those facts to fill the complete answer sheet where the next GPU kernel will read it.

This is like sending ingredients and a recipe to a fast kitchen instead of assembling 56 labeled plates in a distant warehouse and trucking every plate to the kitchen.

122. What exactly became smaller?

The old finished input used 428 bytes for each possible destination row. The new raw packet uses 84 bytes:

Packet part	Bytes per row	Examples
14 whole-number facts	56	zones, period, modes, age, vehicle count
2 decimal facts	16	duration and value of time
3 controlled wait values	12	transit waiting times
total	84	compact Phase 36 packet

In the full public model there were 4,188,312 such rows. The old method would construct 1,792,597,536 bytes, about 1.79 GB. Phase 36 transferred 351,818,208 bytes, about 351.8 MB. That removes 1,440,779,328 bytes from this boundary, an **80.37% reduction**.

This does not mean the whole computer uses 80% less memory or the whole model is 80% faster. It means one precisely measured data handoff became 80% smaller. That distinction keeps the result useful and honest.

123. How do we know the smaller package did not change the model?

First, a 500-household shadow run built both the Phase 35 and Phase 36 utility inputs. It compared the GPU results before accepting the new path. Then a second 500-household run used only Phase 36. Finally, the complete public model ran with 50,000 households and 1,454 zones.

The large run completed all 34 ActivitySim steps and all 30 trip-purpose programs. It processed all 4,188,312 utility rows on the new path, never fell back to the CPU evaluator, and released 6.20 GB of shared GPU network memory after the final GPU consumer.

An independent program compared the completed Phase 36 output with Phase 35:

- changed modeled decision cells: **0**;
- changed decision rows: **0**;
- maximum destination-logsum difference: **0**;
- maximum mode-logsum difference: **0**; and
- byte-for-byte identical published CSV files: **7 of 7**.

Several safety rules support that result. Values too large for the compact whole-number format stop the run. Unknown inputs stop the run. Missing CUDA support stops the run. Paired missing numerical values are compared correctly in shadow mode. The system does not quietly switch to another evaluator and pretend the test passed.

124. Did Phase 36 make the whole model faster?

It almost certainly removed useful work, but this qualification cannot make a new scientific stopwatch claim. Another program was actively using the GPU. Phase 36 happened to finish the complete model in 174.0 seconds versus 224.6 seconds in an older Phase 35 artifact, and its trip-destination step happened to take 22.6 versus 28.9 seconds. Those runs did not occur under matched quiet conditions, so the differences are observations, not proof of speedup.

What *is* proved is stronger than a hopeful microbenchmark: Phase 36 produces the exact same public-model answers while eliminating a counted 1.79 GB dense CPU construction and replacing it with a 351.8 MB compact transfer. A supplied test script can run three fresh alternating Phase 35-versus-Phase 36 pairs when the machine is quiet.

The next ambitious step is a resident trip-state engine. Today the compact packet is still rebuilt for each purpose. Phase 37 should upload common trip, tour, person, and household facts once, reuse them across all 30 programs, and fuse preparation with utility calculation when safe. It should also let many independent tour schedules advance on the GPU while preserving the required order of trips inside each tour. Changed populations, coefficients, land use, and travel-time matrices must all pass before calling that engine general.

125. Phase 37: stop writing the GPU answer sheet at all

Phase 36 built the 56-column answer sheet on the GPU. That removed the giant CPU factory, but it still wrote the completed sheet into GPU memory. The very next GPU kernel read it and threw it away after calculating utilities.

Phase 37 fuses those two jobs. Imagine a cook reading ingredients and a recipe, calculating each needed quantity mentally, and putting only the finished meals on the counter. The cook no longer fills 56 temporary bowls first.

One generated GPU kernel now reads the same compact 84-byte packet. It looks up resident parking, terminal, density, topology, and travel-time facts; derives all availability rules and temporary inputs in fast local variables; evaluates the 379 utility expressions; and writes only 21 final mode scores. The later nested-logit step still turns those scores into the model's combined logsum.

126. How much temporary GPU work disappeared?

For each possible destination row, Phase 36 temporarily stored 404 bytes of utility inputs and 64 bytes of skim coordinates. Phase 37 stores neither.

Removed device intermediate	Bytes in the full public run
11-float and 45-integer input arrays	1,692,078,048
grouped origin/destination/time coordinates	268,051,968
total removed	1,960,130,016

That is about 1.96 GB of temporary allocations and writes removed across the 30 programs. It does not mean 1.96 GB less data is uploaded: the compact 351.8 MB packet is still built and uploaded. It also does not automatically prove the whole model is faster. A fused kernel does more work in each GPU thread and can use more registers, so a fair stopwatch and hardware counters are still needed.

127. How did we prove the fusion was real and correct?

The ordinary compatibility input was deliberately reduced to **one dummy row**, only 468 bytes. If the new kernel secretly read an old full-size array, the test could not succeed safely. The generated CUDA text was also inspected by the program itself: any remaining row read from a legacy input or coordinate array stopped compilation.

This safety design found a real bug. Two input names shared one numbered slot. The first compiler version replaced one name but left the alias pointing to the old input. The one-row setup exposed it. The generator was fixed to replace every alias, and all qualification runs were repeated.

The evidence then climbed four steps:

- 1 source-generator and one-row-bootstrap unit tests;
- 2 a 500-household shadow run where Phase 36 and Phase 37 utilities were compared inside the same process;
- 3 a 500-household production run using only minimal legacy state; and
- 4 the complete public run with 50,000 households and 1,454 zones.

The large run completed all 34 ActivitySim steps and fused all 30 utility programs over 4,188,312 rows. It used zero CPU fallback. A separate verifier compared it with Phase 36: changed decision cells were **0**, destination and mode-logsum maximum differences were **0**, and all seven published CSV files were **byte-for-byte identical**.

128. Is Phase 37 faster, and what should happen next?

Phase 37's 30 fused utility kernels used 2.157 seconds in the recorded full run. Building their compact packets took 4.471 seconds, uploading took 0.081 second, and nested logit took 0.079 second. The complete set of ActivitySim model steps took 216.9 seconds.

Those are honest measurements, but not a new speed claim. The GPU was not quiet enough for a controlled Phase 36-versus-Phase 37 pair, and an older Phase 36 run happened to finish sooner. Comparing unmatched runs could falsely call the change either a win or a loss. The proved result is exact behavior plus removal of 1.96 GB of device intermediates.

The next major phase should attack the remaining 4.47-second packet factory. Stable trip, tour, person, and household facts should be uploaded once into a normalized resident store. Each of the 30 programs would then send only row selection, facts that truly changed, and the controlled wait draws. The test plan should include changed households, coefficients, land use, and skim data; GPU peak-memory and occupancy counters; and three alternating quiet timing pairs with exact final-output checks.

After that, the larger wall-time opportunity is a resident trip scheduler. Different tours can advance in parallel, but trips inside one tour must keep their exact order, retry behavior, and controlled random numbers. That is the ambitious path from a proven fused component toward a genuinely device-resident travel-model runtime.

129. Phase 38: stop repeating the same trip facts

Phase 37 stopped writing a giant temporary answer sheet on the GPU, but its small input packet still repeated the same facts for every possible destination. Imagine a student answering 50 multiple-choice questions and writing their name, age, address, and school at the top of every answer. Most of that writing is unnecessary.

Phase 38 separates what changes from what stays the same. Every possible destination row now carries only three small whole numbers: origin, destination, and a pointer to the correct trip state. Stable facts such as age, household size, car ownership, tour mode, duration, and value of time are written once per trip direction. The GPU follows the pointer when it needs them.

Why "per direction" instead of simply "per trip"? ActivitySim gives outbound and inbound calculations their own three controlled wait draws. Those values can differ even for the same trip. Combining them would save a little more memory but could change the result, so Phase 38 deliberately keeps two states.

130. How much smaller did the input become?

Full 50,000-household public run	Phase 37	Phase 38
possible destination rows	4,188,312	4,188,312
bytes sent/stored as the compact input	351,818,208	64,171,392
repeated bytes removed	-	287,646,816

That is an **81.76% reduction**. The Phase 38 total consists of 50,259,744 bytes of row-specific coordinates and pointers plus 13,911,648 bytes for 183,048 outbound/inbound states belonging to 91,524 trips.

Five reusable GPU work areas grow when necessary and are refilled for later jobs. They were reused 120 times during the 30 utility programs. Phase 38 also keeps Phase 37's larger structural win: 1.96 billion bytes of temporary GPU answer-sheet and coordinate arrays are still gone.

131. How do we know the smaller representation did not cheat?

Before normalizing anything, the program checks that every candidate belonging to a trip agrees on all facts that are supposed to be stable. It separately checks each direction's controlled wait draws. A changed value, missing half, or unexpected ordering stops the run. There is no quiet CPU fallback.

The proof then used several independent checks:

- 1 unit tests deliberately changed stable facts and wait draws and confirmed that the run stopped;
- 2 a 500-household shadow run calculated utilities both old and new ways and found zero differences;
- 3 the complete 50,000-household, 1,454-zone model used the new path for all 30 programs and all 4,188,312 rows with zero fallback;
- 4 an independent verifier found zero changed decisions, zero logsum difference, and all seven published CSV files byte-for-byte identical; and
- 5 all 158 repository tests passed.

One early small run used the wrong-size reference file. The model and shadow comparison completed, but the reporting gate correctly rejected it. We did not count that run; we corrected the reference and repeated the proof.

132. Is Phase 38 faster, and what should we build next?

It unquestionably moves less data. Its recorded packet-building work took 1.362 seconds instead of 4.471 seconds in the Phase 37 artifact, and upload took 0.012 instead of 0.081 second. Those observations are encouraging, but the two runs were not alternating trials under identical quiet conditions. The GPU kernel itself also varied from 2.157 to 2.467 seconds. It would be bad science to turn unmatched measurements into a new speedup claim.

The proved result is exact behavior plus an 81.76% smaller packet. The supplied test script can run three alternating Phase 37-versus-Phase 38 pairs when the machine is quiet. Those pairs are the next timing evidence to collect.

The next *engineering* phase should be much larger. Instead of making another small table change, keep trip scheduling, failed-trip retries, destination iteration, and utility evaluation inside one resident GPU service. Different tours can run in parallel, but trips within one tour must remain in order and must use exactly the same controlled random numbers. That could remove pandas frame construction and CPU/GPU handoffs that still dominate the 32.6-second trip-destination step. It also needs mutation tests for changed people, land use, coefficients, and skims, hardware memory/occupancy measurements, exact complete outputs, and quiet matched timings before promotion.

133. Phase 39: ask the GPU to score every possible trip destination

Before ActivitySim picks a small set of possible destinations for a trip, it first considers every zone in the region. In this public benchmark there are 1,454 zones. Doing that for 91,524 trip records means scoring 133,075,896 trip-and-zone combinations.

Phase 39 built a CUDA kernel for that large score sheet. It ran the reviewed 15-part public scoring recipe for all 133 million combinations without first building a 133-million-row pandas table. It used the GPU for the repetitive arithmetic but deliberately left ActivitySim in charge of turning scores into probabilities, drawing its controlled random numbers, choosing zones, combining duplicate picks, and retrying failed trips.

This was a good scientific boundary because we could compare it carefully. It was not yet a good performance boundary, as the measurements showed.

134. Why can two correct computers produce slightly different decimals?

A score is made by multiplying several facts by coefficients and adding the products. Decimal computer numbers have limited precision. Addition is not perfectly associative: rounding after $(a + b) + c$ can differ by a few final bits from rounding after $a + (b + c)$. A GPU and Sharrow's CPU program may add the same 15 products in different groupings.

Most tiny score changes cannot change a choice. But ActivitySim chooses by placing a controlled random number on a cumulative-probability line. If that number is extremely close to a boundary between two zones, even a tiny score difference might matter.

Phase 39 therefore calculates a safe error range around every GPU score. It asks: "Would every value in these ranges still put all 30 controlled draws in the same zones?" If yes, the GPU answer is safe. If not, the complete row of 1,454 zones is calculated again with live Sharrow, using the same random draws. This happened for 10,721 of 91,524 chooser rows, or 11.71%.

We also wrote a separate small CPU evaluator that copies the exact shape and dot-product operation generated by Sharrow. In a full shadow check, all 15,588,334 guarded utility cells matched live Sharrow bit for bit. Unsupported equations or data layouts stop instead of quietly falling back.

135. Did Phase 39 preserve the model's answers?

Yes, within the project's stated claim boundary. Three fresh experiments each ran the complete 34-step, 50,000-household model twice: first with Phase 38 and then with Phase 39. An independent verifier compared the published tables after every pair. Every modeled decision matched in all three pairs.

We also learned why end-to-end checks matter. A tighter experimental error range passed a comparison of all 133,075,896 GPU utility cells and reduced the number of exact rechecks. Yet a later diagnostic number called `destination_logsum` differed by as much as 0.0002125, above our 0.0001 rule. The actual choices still matched, but we rejected the tighter version and kept the conservative guard. Proving one intermediate calculation is not enough; we must also check what later model steps do with it.

136. Was it faster?

No. This is the most important honest result of Phase 39.

Complete public benchmark	Phase 38	Phase 39
middle whole-model time from three pairs	160.8 s	165.6 s
middle trip-destination time	18.8 s	23.5 s
pairs won	-	1 of 3

The whole model was 2.99% slower in the middle result. The trip-destination step was 25.0% slower. Therefore Phase 39 failed its performance promotion gate and Phase 38 remains the version to use.

The surprising part is that the GPU kernels were fast: together they took only about 0.71 to 0.81 second. The surrounding hybrid work was expensive. About 532 megabytes of dense utility scores came back from the GPU. CPU probability, choice, and exact-guard work then consumed roughly 8.16 to 8.70 seconds, with the exact guard alone taking about 5.66 to 6.18 seconds. We accelerated the arithmetic but left the bottleneck around it.

This is not a failed project. It is a successfully rejected design. The tests prevented a fast-looking sub-second kernel from being advertised as a faster model when the complete model was actually slower.

137. What should Phase 40 build?

Phase 40 should move the **whole sampler**, not merely its scoring kernel. The scores should remain on the GPU while it turns them into probabilities, uses a replica of ActivitySim's controlled random-number rules, performs inverse-CDF selection, combines duplicates, and creates the probability correction used by later destination calculations. Only the much smaller sample should return to ActivitySim. That would remove the 532-megabyte download and attack the eight seconds surrounding the sub-second kernel.

There is one hard requirement: Sharrow's CPU evaluator and CUDA must agree on a canonical arithmetic recipe. A shared expression compiler can define exactly which precision and addition order both backends use. Alternatively, an upstream Sharrow GPU backend can own the same contract. Phase 40 must then test changed people, coefficients, land use, skims, and random draws; compare every utility during qualification; prove the random ledger and all final decisions; measure GPU memory and occupancy; and win at least three fresh matched Phase 38/candidate runs. Only that larger device-resident boundary has a credible path from a fast kernel to a meaningfully faster travel model.

138. Phase 40: keep the whole destination sampler on the GPU

Phase 39 sent the giant score sheet back to the CPU. Phase 40 stops doing that. The GPU now keeps all 133,075,896 trip-and-zone scores in its own memory while it performs the next three jobs:

- 1 turn scores into probabilities;
- 2 place ActivitySim's controlled random numbers on the probability line and select 30 possible destinations for each trip; and
- 3 combine repeated selections of the same destination.

ActivitySim still creates the random numbers. This matters because its random system is like a carefully labeled deck of cards: every household, person, tour, and trip must receive the same card in every reproducible run. Phase 40 uploads those exact cards to the GPU instead of inventing a new random stream.

The GPU returns the small list of selected destinations, their probabilities, and duplicate counts. It does not return the enormous table of every score and every error range.

139. What does "device resident" mean?

The GPU is often called a device, and ordinary computer memory is called host memory. "Device resident" means data stays in GPU memory while several related operations use it.

Imagine cooking in one kitchen but carrying every ingredient across the street between each recipe step. Even if the stove is extremely fast, the walking can erase the advantage. Phase 39 did the computer version of that: it calculated scores on the GPU and carried about 532 megabytes of scores back to the CPU.

Phase 40 keeps the ingredients beside the GPU stove. Because it also keeps the error table there, it avoids 1,064,607,168 bytes of dense downloads - a little over one billion bytes. It downloads about 36.2 megabytes of compact results instead. The random-number upload is about 22.0 megabytes.

140. How does probability sampling work here?

Each destination has a utility score. A larger score means the destination is more attractive under the model's rules. The sampler exponentiates the scores and divides each result by the total:

`probability = exp(utility) / sum of exp(all utilities)`

The probabilities form adjacent sections of a line from zero toward one. A controlled random draw lands somewhere on that line, and the section containing the draw identifies the chosen zone. This is inverse cumulative distribution sampling.

Phase 40 sorts each trip's 30 draws for an efficient walk across the 1,454 zones, then restores the original draw order exactly as ActivitySim does.

The same destination can be drawn several times. ActivitySim keeps its first appearance and records how many times it was picked. A second CUDA kernel now performs that duplicate-counting rule.

141. Why does Phase 40 still need an exact checker?

GPU and CPU programs can round the same decimal calculation in slightly different ways. Most tiny differences do nothing. A difference matters when a random draw is very close to the boundary between two destinations.

Phase 40 surrounds every GPU utility with a conservative error range. It asks whether each draw selects the same zone for every value inside those ranges. A row that cannot prove this goes to a small exact evaluator that copies Sharrow's 15-number dot-product shape and arithmetic rules. Phase 39 had already compared this evaluator with live Sharrow and found every guarded utility bit for bit identical.

There is a second subtlety. The chosen zone can stay the same while its reported probability changes slightly. ActivitySim later uses the logarithm of that probability in a diagnostic called `destination_logsum`. Phase 40 therefore also calculates a safe selected-probability risk. Rows above 0.00009 use the exact evaluator even when their selected zone is stable.

142. What went wrong during development, and how was it fixed?

The first full Phase 40 attempt selected every destination correctly, but one trip's `destination_logsum` differed by 0.000238 . Our allowed difference is 0.0001 , so the run failed. This is why checking final outputs matters.

One attempted fix used an extremely strict probability rule. It made every published file byte identical, but it sent 91,518 of 91,524 chooser rows back to exact CPU arithmetic. The trip-destination step grew to 75.2 seconds. We rejected it.

A much looser CDF range reduced the guard to 5,951 rows and the sampling work to 7.05 seconds, but another diagnostic outlier appeared at 0.000213 . A rule for same-zone work samples fixed the output but guarded 23,185 rows and became slower. We rejected both variants too.

The final rule combines the conservative choice-boundary proof with the 0.00009 selected-probability risk. It exact-checks 13,607 rows, or 14.87% of the workload. This is a useful lesson: the fastest-looking internal experiment is not the winner unless the complete model still meets its accuracy promise.

143. What did the final full run prove?

Question	Result
Did all 34 ActivitySim steps finish?	yes
How many destination-sampling GPU programs ran?	30
How many trip chooser rows were covered?	91,524
How many trip-and-zone scores were evaluated?	133,075,896
How many controlled random draws were preserved?	2,745,720
How many utility bytes returned to the CPU?	0
How many dense download bytes were avoided?	1,064,607,168
How many rows needed exact arithmetic?	13,607 (14.87%)
How many fallback calls occurred?	0
How many modeled decision cells changed?	0
Largest destination-logsum difference	0.000012
Allowed destination-logsum difference	0.0001
Mode-choice-logsum difference	0

Six final files unrelated to trip diagnostics were byte-for-byte identical. The remaining diagnostic difference was 8.3 times smaller than the allowed limit. The clean run passed every proof gate.

144. Was Phase 40 faster?

There are two honest answers because there are two comparisons.

Against the already GPU-accelerated Phase 38 runtime, no. In one fresh matched pair, Phase 38 took 162.8 seconds for all model steps and the Phase 40 candidate took 169.0 seconds. The trip-destination step took 19.0 versus 24.4 seconds. The final code later recorded 165.0 seconds overall and 23.0 seconds for trip destination, but that was not another alternating pair. Phase 38 therefore remains the version to use when the only goal is the lowest measured time.

Against the regular pinned ActivitySim CPU baseline already measured in three pairs, the whole GPU project is still substantially faster. The CPU middle result was 206.6 seconds overall and 41.0 seconds for trip destination. Compared descriptively with the final Phase 40 run, that is about 1.25 times faster for the whole model and 1.78 times faster for trip destination. Those gains belong to all the GPU phases together, not to Phase 40 alone.

145. Why is Phase 40 still a major result if it did not beat Phase 38?

It proves that the whole full-zone sampler can live on the GPU without changing modeled decisions. More than one billion bytes of unnecessary dense movement are gone. The random ledger, probability order, duplicate rules, and complete model outputs remain under explicit tests.

It also identifies the remaining obstacle precisely. The GPU utility, choice, and duplicate kernels together take only about 2.44 seconds. Exact arithmetic for the 13,607 guarded rows takes about 7.33 seconds. Transfer is no longer the main problem; identical CPU/GPU arithmetic is.

146. What should the next generation build?

The next project should give Sharrow and CUDA one shared arithmetic language. A dedicated expression compiler could specify the data type, multiplication, addition order, exponential, probability sum, and boundary comparison once, then generate both CPU and GPU programs. An upstream Sharrow GPU backend could provide the same guarantee inside Sharrow itself.

If ordinary CUDA rows become bit-identical to the CPU authority, the 7.33-second exact guard can disappear. That is the clearest remaining route to making the device-resident architecture faster than Phase 38, not merely smaller and more GPU-native.

The promotion test must remain strict: altered people, land use, coefficients, skims, and random draws; complete output comparison; zero fallback; bounded diagnostics; and at least three fresh alternating Phase 38/candidate pairs that the candidate wins. Phase 40 gives that future compiler a proven resident sampler to plug into.

147. Phase 41: make the CPU and GPU follow the same arithmetic recipe

Phase 40's GPU was already fast, but it had to send 13,607 uncertain rows to the CPU. That CPU check consumed about 7.33 seconds. Phase 41 removes it by answering a deceptively simple question: **in exactly what order does Sharrow add 15 numbers?**

Addition with rounded computer numbers is not perfectly associative. In school math, $(a + b) + c$ equals $a + (b + c)$. In float32 computer math, either route can round at a different point and produce a last-bit difference. Those last bits can matter when a random draw is almost exactly on the boundary between two destinations.

We discovered that Sharrow's 15-number dot product was not a plain left-to-right loop. Numba sent it to an OpenBLAS matrix-vector routine. On this machine, that routine adds three groups of four products from left to right, then adds the final three products one by one. Phase 41 puts that recipe in a named, versioned arithmetic contract and generates the CUDA statements from it.

148. How did we prove the utility recipe was correct?

First, a controlled experiment tried the candidate addition recipes on 200,000 random rows. The four-at-a-time left-associated recipe had zero bit mismatches. Ordinary sequential, tree, and lane recipes had tens of thousands.

Then we used the real public travel model. The GPU and live Sharrow separately calculated every trip-and-zone utility:

Test	Result
real utility cells compared	133,075,896
cells with even one different bit	0
largest difference	0

That is much stronger than saying the answers were "close." For this reviewed program and environment, every utility bit matched.

149. Why were exact utilities still not enough?

The first guard-free experiment still changed 15 trip destinations. The reason was the next formula:

```
probability = exp(utility) / sum(exp(all 1,454 utilities))
```

CUDA's ordinary `expf` function and NumPy's float32 exponential use different approximations. NumPy also adds a long row with a recursive pairwise tree, while the first GPU version used a sequential total. Both answers looked very close, but a few controlled random draws sat close enough to a boundary to notice.

We therefore added a second shared contract. It copies NumPy 2.4.6's AVX2 float32 exponential polynomial and generates NumPy's exact 1,454-item pairwise sum tree. On 5,000 synthetic full rows (7,270,000 exponentiated values), the final GPU totals matched NumPy bit for bit on every row. After this change, the full travel model returned to zero changed decisions.

150. What did the final Phase 41 run prove?

We ran three fresh pairs. Every pair ran Phase 40 first and Phase 41 second on the same 50,000-household public benchmark. Every Phase 41 candidate had:

- 30 GPU sampling programs;
- 91,524 trip chooser rows;
- 133,075,896 utility cells;
- 2,745,720 ActivitySim-controlled random draws;
- zero CPU guard rows;
- zero fallback calls; and
- zero changed modeled decision cells.

All six non-trip output files were byte-for-byte identical. The largest `destination_logsum` difference was 0.000012, safely below the 0.0001 limit, and the mode-choice-logsum difference was zero.

151. How much faster is Phase 41?

Complete public benchmark	Phase 40	Phase 41	Result
middle trip_destination time	23.1 s	15.1 s	8.0 s saved; 1.53x
middle whole-model time	165.8 s	158.3 s	7.5 s saved; 1.047x
matched pairs won	-	3 of 3	promoted

Inside the sampler, the median complete GPU-to-compact-result boundary was about 1.71 seconds. Phase 40's final clean boundary was 10.48 seconds. That is about 6.12 times faster at the directly changed boundary. The old 7.33-second CPU guard section now takes roughly 0.00012 second because it has no rows and does no arithmetic work.

Compared only as context with the older regular ActivitySim middle result, 158.3 versus 206.6 seconds is about 1.31 times faster for the whole model, and 15.1 versus 41.0 seconds is about 2.72 times faster for `trip_destination`. Those larger figures describe the accumulated project, not Phase 41 alone.

152. What are the assumptions, implications, and next step?

The exact contracts are deliberately specific. They were qualified on an NVIDIA RTX A4000, the public 1,454-zone model, NumPy 2.4.6's AVX2 path, and the reviewed SciPy/OpenBLAS and Numba versions. A different CPU instruction set, NumPy release, OpenBLAS kernel, alternative count, or model equation may need a new generated schedule and a new proof. Unsupported cases stop instead of quietly falling back.

This does not make Phases 1 through 40 pointless. Those phases supplied the expression compiler, strict references, RNG tests, output verifier, resident skims, compact trip state, full-zone utility kernel, and resident sampler that made Phase 41 possible. Phase 41 replaces one bottleneck inside that system; it does not replace the system.

The larger implication is important: CPU/GPU replication does not always need a slow CPU safety pass. Arithmetic library behavior can be made explicit, versioned, generated, and tested. The next ambitious project is to generalize this from one 15-expression program to a real Sharrow GPU backend or dedicated expression compiler that can emit matching CPU and CUDA programs for many ActivitySim models. It should cache contracts by hash and automatically rerun mutation tests for coefficients, skims, land use, people, and random draws.

153. Phase 42: from one exact recipe to a small compiler

Phase 41 knew one important recipe: 15 utility terms and 1,454 alternatives. Phase 42 changes the recipe into data that a compiler can read. In plain English, the policy says:

- how many numbers are multiplied;
- exactly which partial products are added together and in what order;
- which float32 exponential recipe to use;
- how the long probability total is split into a tree; and
- whether arithmetic shortcuts such as fused multiply-add are allowed.

The compiler checks the policy, gives it a SHA-256 fingerprint, writes a CPU reference evaluator, and writes matching CUDA source. SHA-256 is like a very sensitive digital fingerprint: changing the addition order or even the number of alternatives changes the fingerprint. A cached kernel is reused only when that fingerprint matches.

This does not mean the compiler understands every possible mathematical expression. It means the exact utility-reduction and probability-normalization parts that caused our replication problem are now reusable instead of being handwritten for one model shape.

154. How did we test that “general” is more than a name?

We compiled and ran the policies on the real RTX A4000 for several shapes:

GPU test	Shapes tested	Answer
utility addition	1, 3, 5, 15, 17, and 31 terms	zero different bits
exponential probability total	1, 7, 8, 129, 257, and 1,454 choices	zero different bits

Each shape used 257 random rows. We compared the 32 individual bits of every float32 result, not just rounded decimal text. We also made two deliberate mutations. Changing grouped addition to ordinary left-to-right addition changed the compiler fingerprint. Changing 1,454 alternatives to 1,453 did too.

This is useful evidence, not infinite proof. It establishes that one generator works across small, boundary, and public-model sizes. A brand-new CPU math library or expression type still needs its own reviewed policy and test.

155. The other Phase 42 idea: stop rebuilding giant tables

The arithmetic kernel was no longer the only cost. ActivitySim visits trip destination three times: roughly first stops, second stops, and third stops. For each visit it handles ten trip purposes. Older code repeatedly built a large outbound-plus-inbound pandas table even though most columns repeated.

Phase 42 passes a compact view instead: one base sampled table plus two small coordinate arrays. Before CUDA uses it, the code checks that the stable trip information really is stable, that each directional row has exactly three controlled random draws, and that equivalent directional states have the same draws. Only the small representative state is then packed for the GPU.

Think of sending one class roster plus two seating charts instead of printing the entire roster again for every seating arrangement.

156. What does “compile once and reuse” mean here?

There are ten purpose models, and the same ten models are used in each of the three trip-number passes. In every final Phase 42 run we observed:

Item	Built on first pass	Reused later	Total uses
purpose/logsum contract	10	20	30
strict expression document	10	20	30
native CUDA ABI source	10	20	30
final simulation specification	10	20	30
compact directional bundle	-	-	30

The cache stores instructions and compiled structure, not travelers' answers. Trips, probabilities, destinations, and random draws are still calculated for the current run. ActivitySim still owns and advances its keyed random-number ledger, which is essential for replication.

157. What did the complete public benchmark prove?

We ran three new pairs. Every pair started Phase 41 and Phase 42 in fresh processes and completed all 34 ActivitySim steps for 50,000 households and 1,454 zones. The trip workload contained 91,524 trip rows and 2,094,156 retained destination-sample rows.

Every Phase 42 report passed its runtime gates. Three separate output verifiers found:

- zero changed modeled decision cells;
- zero changed modeled decision rows;
- zero destination_logsum difference; and
- zero mode_choice_logsum difference.

Phase 42 was faster in all three matched pairs, not only in the middle result. That matters because whole-model timings naturally move by a few tenths of a second as the operating system and storage do other work.

158. How much faster is Phase 42?

Complete public benchmark	Phase 41	Phase 42	Result
middle trip_destination time	14.9 s	10.7 s	4.2 s saved; 1.393x
middle whole-model time	156.7 s	151.8 s	4.9 s saved; 1.032x
matched pairs won	-	3 of 3	promoted
changed decision cells	-	0 in every pair	exact

The trip component used 28.19% less time. The complete model used 3.13% less time because trip destination is only one part of 34 steps. In context, the latest 151.8-second model is about 1.36 times faster than the older regular ActivitySim median of 206.6 seconds. The latest 10.7-second trip destination is about 3.83 times faster than the older 41.0-second regular result. Those larger ratios describe the accumulated project, not Phase 42 alone.

159. Where does the remaining trip time go?

The code directly measured the middle Phase 42 trip boundary:

Work inside trip destination	Middle time
full-zone GPU sampling	1.674 s
prepare compact inputs	1.540 s
run remaining preprocessor work	1.045 s
calculate trip mode logsums	3.165 s
make final sampled-destination choices	1.812 s
complete measured boundary	9.312 s

ActivitySim's outer component timer reads 10.7 seconds because it includes framework work around this measured boundary. The table tells us why another tiny kernel is unlikely to transform the whole model: the remaining time is spread across preparation, preprocessing, logsums, final simulation, Python orchestration, and other model components.

160. Did Phase 42 meet every goal, and what comes next?

No. We set ambitious stretch goals of trip destination below 8 seconds and the whole model below 150 seconds. The qualified medians are 10.7 and 151.8 seconds, so both stretch-goal flags remain false. We did not lower the goals after seeing the result. Phase 42 is still a successful promotion because it won all three pairs, preserved exact results, made the target component 28% faster, and converted a special arithmetic trick into reusable infrastructure.

The next large phase should use the shared expression IR to compile the sparse final-simulation and preprocessor work, retain selected probabilities and logsums on the GPU, and replace more pandas group/merge work with compact indexed arrays. The measured targets are about 1.045 seconds of preprocessing, 1.812 seconds of final simulation, and parts of 1.540 seconds of preparation.

Success must still mean more than a fast stopwatch. The next phase needs fresh matched whole-model pairs, one fingerprinted policy per compiled program, fail-closed handling of unsupported expressions, ActivitySim-owned random draws, and independent exact comparison of every substantive output. That is how this grows toward an upstream Sharrow GPU backend without trading away trust.

161. Phase 43 starts by measuring the last choice step

After Phase 42, we did not guess which little piece to accelerate. We timed the final step that chooses one sampled destination for each trip. It took about 1.83 seconds. Roughly 0.58 seconds calculated how attractive destinations were, but 0.72 seconds was spent asking ActivitySim for keyed random numbers and turning them into choices. Probability math itself took only 0.02 seconds.

A keyed random number is like a reproducible raffle ticket attached to one trip ID. If we run the same model again, that trip receives the same sequence of tickets. This is essential: without it, a faster program could make different choices merely because it asked for random numbers in a different order.

162. Why millions of sample rows did not need millions of tickets

The benchmark has 91,524 intermediate trips. The destination sampler keeps 2,094,156 possible destination rows because each trip can have many options. The three wait-time disturbances and final choice ticket belong to the trip, not separately to every possible destination.

ActivitySim safely expanded trip-level random values over all those option rows. That is convenient for general pandas code, but our compact GPU runtime already knows which options belong to each trip. Phase 43 therefore asks ActivitySim for random state on unique trips and keeps it compact. Across the complete run it:

- keeps 183,048 directional trip rows for outbound and destination-to-end directions;
- avoids 4,005,264 repeated directional option rows;
- creates one final choice ticket for each of 91,524 trips; and
- replaces 66 small purpose-level requests with nine trip-batch requests.

ActivitySim still owns the generator and seeds. We changed the shape of the request, not the values or meaning of the raffle tickets.

163. How can batching random requests give identical answers?

Inside one trip-number pass, each trip belongs to exactly one purpose, such as work, school, shopping, or escort. Phase 43 first proves that option rows are contiguous, trip order matches, and no trip ID appears in two purpose groups. Only then may it combine the ten groups.

For each of three trip-number passes, it asks for outbound normals, destination-to-end normals, and final uniform choice tickets. That is three requests per pass and nine total. Because ActivitySim's ledger is keyed by trip ID, one combined request advances the same per-trip channel as the trip's one original purpose request. If these assumptions fail, the compact path stops instead of guessing.

The final-choice adapter also keeps ActivitySim's probability checks, error handling, exact `choice_maker` function, table indexes, and output types. It only supplies the ticket that ActivitySim already generated. A mismatched index immediately uses the original path.

164. What did Phase 43 achieve?

Measured work	Phase 42	Phase 43	Result
final destination simulation	1.817 s	1.152 s	1.578x faster
complete measured trip boundary	9.261 s	8.526 s	1.086x faster
ActivitySim trip_destination component	10.6 s	9.9 s	1.071x faster
all 34 model steps	152.4 s	152.1 s	1.002x; within noise

We ran three fresh Phase 42/43 pairs. The measured trip boundary was faster in all three, and final simulation was faster in all three. In every pair, all seven final tables were byte-for-byte identical. There were zero changed decision cells and zero destination- or mode-logsum differences. The full test suite passed 179 tests.

The whole model did not show a strong new speedup: one pair was faster overall and two were slower because unrelated steps varied by more than 0.7 seconds. The median moved by only 0.3 seconds. So the honest claim is a replicated 6.6% trip-destination component improvement, not proof that the entire model is now meaningfully faster.

165. Does Phase 43 add a new GPU kernel?

Not by itself. Earlier phases put sampling, logsums, scheduling, and other heavy math on the GPU. Phase 43 makes that GPU runtime better by removing redundant CPU table expansion and random-manager calls at its boundary. This matters because a fast GPU kernel can still be hidden behind slow preparation, copying, and orchestration.

This distinction is important when reading performance claims. The 1.578x final-simulation improvement comes from compact exact orchestration around the GPU-resident pipeline, not from pretending a CPU operation is a GPU kernel.

166. What remains for the next ambitious phase?

About 0.58 seconds still evaluates the 14 sparse final destination utility terms. More time constructs joined interaction tables, groups options by trip, pads different-length option lists into rectangles, and maps a selected position back to a destination.

The next large project should compile those 14 terms from the same shared expression language, represent each trip's options with compact start/end offsets, and fuse utility, exact probabilities, and selection without giant pandas tables. It should first run in shadow mode beside ActivitySim, compare every utility and choice, fingerprint its arithmetic rules, and reject unknown expressions. Only after exact shadow results should it become authoritative and repeat three matched complete-model runs.

That could remove most of the remaining 1.15-second final simulation and part of preparation. It cannot by itself transform a 152-second model, because trip destination is now only about 6.5% of total time. A large whole-model dent will require applying the compact compiled runtime to several other costly components, especially initialization, location choice, tour scheduling, and trip mode choice.

167. Phase 44 builds the compact final-choice machine

Phase 43 ended with one last general-purpose ActivitySim step. For every trip, it combined traveler facts with roughly 23 sampled destinations, calculated their attractiveness, padded the unequal option lists into a rectangle, made one reproducible random choice, and translated the winning position back into a zone number. The math was not the only cost. Building and reshaping large pandas tables also took time.

Phase 44 replaces that table plumbing with a small, purpose-built runtime. It does this for all 30 purpose-and-trip-number calls, covering 91,524 trips and 2,094,156 sampled destination rows. It does not change the model question or remove any legal options.

168. What is a ragged array, and why does it help?

"Ragged" means the rows do not all have the same length. One trip may have 18 sampled destinations while another has 27. A spreadsheet-like rectangle must add blank cells so every row has the same width. A ragged representation keeps all real options next to each other and stores a short list of starting points:

```
starts: [0, 18, 45, 68]
options: [18 for trip A][27 for trip B][23 for trip C]
```

Trip B's options are positions 18 through 44. This offset list replaces a group-by operation that repeatedly asks pandas which rows belong to which trip. Phase 44 first proves that rows are contiguous and in chooser order. A Numba-compiled loop then creates ActivitySim's required rectangle directly, using exactly the same -999 value for fake padded options.

169. What is the 16-slot expression program?

An expression is a small formula, such as "distance from the trip origin" or "log of destination size." The public final-choice file has 16 ordered slots. Two slots temporarily store distances; the other 14 contribute real utility terms. Together they describe destination size, unavailable zero-size zones, directional distance, closeness to tour endpoints, correction for sampling, and mode-choice logsums.

"Utility" is only a score: a higher score means a destination is more attractive to the model. ActivitySim turns all scores into probabilities and uses a controlled random ticket to choose one destination.

Phase 44 checks the exact text, order, count, and coefficient shape before it runs. If a future model changes one expression, the optimized path stops instead of silently assuming the old meaning. Sharrow still evaluates the formula using its trusted float32 CPU compiler. Phase 44 makes the data around that evaluator smaller and more direct.

170. Is Phase 44 itself a new GPU kernel?

No, and that distinction matters. Earlier phases already run the enormous full-zone destination sampling and trip logsum calculations on the GPU. Phase 44 improves the last boundary around that GPU pipeline. Its new padding loop is Numba-compiled CPU code, and the 16-slot evaluator remains Sharrow CPU code so its arithmetic stays identical.

This is still useful. A GPU can finish heavy math quickly and then sit idle while Python builds tables around the result. Removing that overhead makes the whole GPU system faster. But we do not label CPU work as GPU work merely because it is located next to a GPU kernel.

171. How did we prove the result?

We ran three fresh matched pairs. In each pair, Phase 43 was the control and Phase 44 was the candidate. Both used the same public Prototype MTC extended model, 50,000 households, 1,454 zones, GPU stack, Sharrow requirement, and random-number rules.

Measured work	Phase 43	Phase 44	Result
compact final-choice boundary	1.092 s	0.822 s	1.328x faster
trip_destination component	10.1 s	9.5 s	1.063x faster
all 34 model steps	158.1 s	156.4 s	1.011x faster
matched pairs won	-	3 of 3	repeatable direction
changed decision cells	-	0 in every pair	exact

The boundary saved about 0.270 seconds, the complete trip-destination step saved 0.6 seconds, and the median full model saved 1.7 seconds. The unusually large 7.2-second gain in the first whole-model pair included unrelated steps also becoming faster, so we do not credit all of it to Phase 44. The stable boundary measurement and three component wins are the stronger evidence.

For every pair, an independent checker compared all published results. All seven CSV files were byte-for-byte identical. It found zero changed modeled decisions, zero destination-logsum difference, and zero mode-logsum difference. Every runtime proof gate passed, and the complete test suite passed 183 tests.

172. What do these results mean, and what comes next?

Phase 44 proves that exact GPU acceleration is a systems problem, not only a kernel problem. Fast kernels, compact data layouts, controlled randomness, expression contracts, and complete-output tests all have to work together. It also turns the remaining final-choice work into five clearly measured pieces: about 0.093 seconds to build narrow frames, 0.582 seconds for Sharrow utility evaluation, 0.028 seconds for padding, 0.045 seconds for probabilities, and 0.057 seconds for choice mapping.

The next direct target is the 0.582-second Sharrow evaluator. A future GPU backend cannot simply calculate similar numbers; one tiny rounding change can move a random choice across a probability boundary. The safe next phase is to translate the reviewed 16-slot program into the shared expression language, generate a strict CPU reference and CUDA version from one arithmetic policy, compare every utility bit and final choice in shadow mode, and reject unknown expressions. Only then should the GPU version become authoritative.

Even perfect removal of that 0.582 seconds cannot transform a 156-second model by itself. A major whole-model gain requires reusing the same compact compiled runtime in several expensive components, such as location choice, tour scheduling, and trip mode choice. Phase 44 supplies both a proven pattern and an honest measurement of the remaining ceiling.

173. Phase 45 takes the destination sampler model-wide

Phase 44 improved one late trip-destination choice. Phase 45 asks a larger question: can one GPU destination machine serve several parts of ActivitySim? It now serves five families:

- where students attend school;
- where workers work;
- where a household's joint tour goes;
- where non-mandatory tours such as shopping or social trips go; and

- where a short tour made during work goes.

These five families invoke 19 sampling programs. At 50,000 households they cover 201,390 decision makers and 274,223,637 possible person-zone score cells. The GPU calculates those dense scores, normalizes them, processes 6,041,700 reproducible random tickets, and removes duplicate sampled zones. A compact shared adapter then handles the 4,696,676 sampled rows used for final choice.

174. What does destination sampling mean?

The model has 1,454 zones. Fully calculating every expensive downstream detail for every person and every zone would waste time. Sampling is a two-stage shortcut that preserves statistical meaning:

- 1 Give every legal zone a simple preliminary attractiveness score.
- 2 Use that score to draw 30 candidate zones for each person or tour.
- 3 Calculate richer logsums and details only for those sampled candidates.
- 4 Correct for the sampling probability so the final model is not biased toward zones that happened to be sampled more often.

Phase 45 accelerates the enormous first two steps. It does not reduce the number of zones, change the 30 draws, or replace ActivitySim's random-number ledger.

175. What formulas does the GPU evaluate?

The preliminary score uses a small reviewed vocabulary. Distance is divided into pieces: the first mile can have one penalty, miles 1-2 another, miles 2-5 another, and so on. Larger destinations are generally more attractive, so the model uses $\log(1 + \text{size})$. School and workplace models add shadow prices, which are balancing adjustments that help simulated workers and students fit available jobs and school places. Workplace choice also lets high-income travelers respond differently to distance.

The program has 7 terms for tour destinations, 9 for school, and 11 for work. Phase 45 rejects an unknown formula instead of pretending it understands it. That fail-closed rule is essential: a fast answer to a different formula is not an optimization.

176. Why was exact arithmetic difficult?

Computers store decimal-looking numbers as binary floating-point values. When several values are multiplied and added, changing the grouping of additions can change the last bit. Sharrow's CPU program groups its dot product like the installed OpenBLAS library. The GPU therefore uses the same grouped-left schedule rather than a mathematically equivalent but differently rounded sum.

The $\log(1 + \text{size} \times \text{shadow price})$ feature caused a remaining one-bit mismatch. It depends only on the 1,454 destinations, not on every traveler. Phase 45 now lets NumPy calculate those 1,454 values once using the original column types, then uploads them. This replaces hundreds of millions of repeated logarithms and makes the largest 38,525,184-cell workplace utility shadow bit-for-bit identical to Sharrow.

177. Why is there still a tiny CPU guard?

Identical scores do not quite guarantee identical random choices. NumPy and CUDA can round the exponential and division used for probabilities one bit differently. Usually this changes nothing. It matters only when a random ticket is extremely close to the boundary between two zones.

We compared the GPU and NumPy paths and measured the largest choice-changing boundary displacement as 0.000000217. Production uses a wider 0.0000005 gate. Only 7,313 of 201,390 rows - 3.63% - enter it. For those rows, the

program downloads their already-exact scores and asks NumPy to repeat only final normalization and choice. It does not rebuild a pandas interaction table or rerun Sharrow. The remaining 96.37% stay on the GPU path.

This guard is not a hidden fallback. It is a measured adjudication rule with a reported row count and an independent end-to-end output test.

178. What did the three official comparisons show?

Each pair started a fresh Phase 44 control, then a fresh Phase 45 candidate, on the same public 50,000-household, 1,454-zone, 34-step ActivitySim model.

Pair	Phase 44	Phase 45	Time saved	Whole-model speedup
1	151.5 s	150.3 s	1.2 s	1.008x
2	152.5 s	149.2 s	3.3 s	1.022x
3	152.1 s	149.7 s	2.4 s	1.016x
Middle result	152.1 s	149.7 s	2.4 s	1.016x

Model component	Phase 44	Phase 45	Result
school location	8.0 s	6.7 s	1.194x faster
workplace location	11.7 s	11.3 s	1.035x faster
joint-tour destination	3.3 s	3.5 s	0.943x; 0.2 s slower
non-mandatory-tour destination	11.3 s	10.2 s	1.108x faster
at-work subtour destination	2.8 s	2.6 s	1.077x faster
all five targets together	37.1 s	34.5 s	1.075x faster

The small joint-tour program cannot fully pay back its first-program setup cost, so it became 0.2 seconds slower. We report it because honest performance work includes regressions. The five targets together improved in every pair, and the complete model improved in every pair.

179. How accurate are the Phase 45 results?

An independent checker compared every published modeled decision in all three pairs. It found zero changed decision cells each time. Mode-choice logsums were identical. Maximum differences were about 0.00000191 for school and workplace location logsums and 0.000003 for tour-destination logsums, below the declared 0.00001 and 0.0001 limits.

This means the locations, destinations, schedules, modes, and trips chosen by the model are the same. A few diagnostic decimal strings differ in their last digits because selected probabilities differ by at most one float32 unit.

180. How much faster is the project than regular ActivitySim now?

The latest fresh causal comparison is Phase 44 versus Phase 45: 152.1 versus 149.7 seconds. An older established regular ActivitySim/Sharrow experiment had a 206.6-second middle result. Comparing those numbers only as context, the latest runtime is about 1.380 times faster, saves 56.9 seconds, and is 27.5% lower. Because those runs were conducted in different phases, this is not a fresh matched CPU-versus-Phase-45 experiment and must not be presented as one.

181. What should the next major phase do?

The next phase should build a persistent model-wide destination service rather than another tiny component patch. It should:

- compile the 7-, 9-, and 11-term GPU programs before timed model steps;
- retain destination features, coefficients, and scratch buffers across all 19 calls;
- batch small joint-tour segments so setup cost is shared;
- move the final sampled-choice utility and probability path onto the same strict GPU arithmetic system; and
- keep ActivitySim's keyed randomness, fail-closed formula checks, exhaustive shadows, three fresh matched pairs, and independent final-output verification.

That attacks both the one regressing small component and the remaining CPU boundary. It also creates an upstreamable backend shape: one versioned expression compiler and resident service, rather than a collection of model- specific shortcuts.

182. Phase 46 turns many GPU calls into one service

Phase 45 had the right GPU math, but each destination call still behaved a bit like renting a new workshop: find the tools, reserve worktables, do one job, and clean up. Phase 46 creates one workshop at the start and keeps it open for the whole model.

The service handles 19 destination-sampling calls for school, work, joint tours, non-mandatory tours, and at-work tours. It compiles the four legal formula shapes before the timed model steps, grows its GPU work area only when it sees a larger job, and then reuses that memory. Its largest measured work area is about 391 megabytes. That is enough for 26,496 decision makers times 1,454 possible zones, plus choices, probabilities, duplicate counts, and random-number state.

Why do this? Starting tiny jobs and repeatedly asking the computer for large blocks of memory can take longer than the arithmetic itself. A persistent service lets later calls reuse work already paid for.

183. How can a GPU make the same random numbers as ActivitySim?

"Random" in a simulation does not mean uncontrolled. ActivitySim gives each person or tour a seed, which is like a recipe number, and an offset recording how many random tickets have already been used. If those tickets change, later choices can change even when every probability is identical.

Phase 46 implements NumPy's MT19937 ticket machine on CUDA. For each row it:

- 1 reads ActivitySim's seed and offset;
- 2 rebuilds the same MT19937 state on the GPU;
- 3 skips exactly the already-used tickets;
- 4 produces the requested 64-bit random values; and
- 5 advances ActivitySim's official offset by the same amount.

Each qualified run generated 6,243,090 values for 402,780 chooser rows in 38 GPU calls. A live shadow compared all 705,270 tickets for 23,509 gradeschool students and found zero bit differences. ActivitySim still owns the ledger; the GPU is an exact ticket printer, not a new source of randomness.

184. What does "compute each exponential once" mean?

The model first gives all 1,454 zones scores called utilities. It turns a score into a positive weight using an exponential, roughly "raise the number e to this score." Earlier GPU choice code could repeat that expensive calculation while it summed weights and searched for a chosen zone.

Phase 46 writes each weight once into reusable GPU memory. It then adds the 1,454 float32 weights in the same pairwise pattern NumPy uses and searches the resulting probability line for all 30 random tickets. This is important for both speed and accuracy: changing the order of additions can change the last binary digit.

The four legal preliminary programs are still checked exactly: 9 terms for school, 11 for workplace, and two 7-term tour forms. An unfamiliar expression stops the optimized path. The service cannot quietly run a fast version of the wrong model.

185. Why did the first full Phase 46 run fail?

The first attempt completed school and workplace choice, but one student's school zone changed. That one change altered later data, and the model's fail-closed scheduling check stopped the run. This was a useful failure.

We compared the person step by step. The GPU random tickets were bit-for-bit correct. The problem was a memory-saving shortcut. Probabilities are unchanged in pure mathematics if every utility has the row's largest utility subtracted. Phase 46 performed that subtraction directly inside the only utility buffer. Phase 45 kept the original utilities for its exact boundary checker.

Those two routes were algebraically equal but not arithmetically identical in float32. A ticket near a probability boundary sampled zone 659 instead of 660. The richer final choice then selected zone 583 instead of 585.

The fix keeps the original utility surface. A short reusable vector stores one maximum per row, and the weight kernel subtracts it while making weights. The exact checker still sees the untouched values. The next complete run returned to zero differences. This is why an agent must run and verify code: a plausible edit and a passing small test did not prove the full simulation.

186. How did we make the exact safety path fast?

About 7,313 of 201,390 rows have a random ticket close enough to a probability boundary that NumPy should decide them. That is 3.63%. Phase 45 used a plain Python loop for these rows, looking across as many as 1,454 zones.

ActivitySim already contains a Numba-compiled version of the same preserved- order sampling rule. "Numba-compiled" means Python-like code is turned into fast machine instructions. We tested both implementations and found identical zone choices and identical selected-probability bits. In a focused test the ActivitySim version was about 160 times faster. Phase 46 compiles it during prewarm, and the benchmark charges that compile time.

Phase 45 also asked pandas to sort a large sample table twice. Phase 46 knows that every decision maker has at most 30 draws, so it sorts those tiny rows directly and builds an already ordered table. A 26,496-row test was exact and 4.53 times faster for the packing step.

187. What did the official three comparisons show?

Every comparison started a fresh Phase 45 control and then a fresh Phase 46 candidate. Both ran the same public Prototype MTC extended data, 50,000 households, 1,454 zones, 34 model steps, and strict output checker.

Pair	Phase 45 lifecycle	Phase 46 lifecycle	Time saved	Speedup
1	150.4 s	146.935 s	3.465 s	1.024x
2	148.3 s	147.211 s	1.089 s	1.007x
3	148.2 s	146.088 s	2.112 s	1.014x
Middle result	148.3 s	146.935 s	1.365 s	1.009x

Phase 46's lifecycle number includes 1.69 to 1.73 seconds of cold prewarm. We did not hide setup outside the clock.

Directly targeted component	Phase 45	Phase 46	Result
school location	7.9 s	5.9 s	1.339x faster
workplace location	10.1 s	8.7 s	1.161x faster
joint-tour destination	3.3 s	3.7 s	0.892x; 0.4 s slower
non-mandatory-tour destination	9.5 s	8.7 s	1.092x faster
at-work subtour destination	2.5 s	2.3 s	1.087x faster
all five together	33.0 s	29.4 s	1.122x; 10.9% lower

The joint-tour groups are small, so they do not fully repay the persistent GPU service overhead. We keep that regression visible. The five targets together improved in all three pairs, and the complete lifecycle also improved in all three.

188. How strong is the replication guarantee?

For each candidate, an independent verifier compared the seven published CSV files: accessibility, households, joint-tour participants, land use, persons, tours, and trips. All seven files were byte-for-byte identical in every pair. Changed modeled decision cells were 0, 0, and 0.

The runtime also proved that all 19 calls used the persistent service, all 274,223,637 dense utility cells stayed on the reviewed CUDA path, all 38 random calls covered the expected workload, the largest workspace stayed under one gibibyte, and every older GPU proof gate still passed. The complete repository test suite passed 192 tests.

This is stronger than saying totals look similar. It says the published files are the same bytes for this benchmark, software environment, configuration, and hardware. It is not a promise that any future formula or GPU will work without requalification; unknown programs still stop, and a new environment must rerun the same tests.

189. What did Phase 46 achieve, and what remains?

Phase 46 proves a model-wide GPU service can own memory, probability sampling, duplicate handling, and controlled random generation while preserving exact results. Target work is 10.9% lower, but it is only part of a 147-second model, so the full lifecycle improves about 0.92%.

It is not GPU-only: ActivitySim orchestrates, the 3.63% boundary uses NumPy/Numba, and richer final choice among 4,696,676 samples remains Sharrow CPU work.

Next: compile strict final-choice expressions on GPU, batch small joint tours, and requalify with cold costs and byte-exact outputs.

190. What did Phase 47 try to change?

After Phase 46, the GPU could quickly make a shortlist of possible zones, but the CPU still performed the richer final scoring of that shortlist. Imagine a school decision. The shortlist may contain up to 30 zones. The final score can include distance pieces, school size, a shadow-price correction that balances zone capacity, a travel-mode score, and a correction for how the shortlist was sampled.

Phase 47 built one strict GPU compiler for those final formulas. It covers five kinds of decisions: school location, workplace location, joint-tour destination, non-mandatory-tour destination, and at-work subtour destination. Across the public benchmark, that is 19 calls, 201,390 decision makers, and 4,696,676 shortlisted alternatives.

"Strict" matters. The compiler recognizes exactly four reviewed formula shapes and four measured shortlist widths: 21, 25, 29, and 30. A new or changed formula stops the fast path. It cannot quietly calculate a similar but wrong answer.

191. What is a final-choice utility?

A utility is a score, not money or electrical power. A zone gets points or penalties based on its features. A simplified example is:

`score = distance penalty + size benefit + mode score + sampling correction`

The real formulas split distance into ranges such as 0-1, 1-2, 2-5, 5-15, and over 15 miles. This lets one extra mile matter differently on a short trip than on a long trip. Workplace formulas also let distance affect higher-income groups differently. School and work models include shadow prices, which nudge choices toward places with available capacity. Tour models include the mode-choice logsum, a summary of how attractive all available travel modes are.

Finally, the sampling correction prevents a zone from gaining or losing an unfair advantage merely because the shortlist process happened to include it more or less often.

192. What does the GPU do now?

For each shortlisted zone, a CUDA kernel reads the distance from the already resident skim, builds every required feature, multiplies features by model coefficients, and adds them in the same reviewed float32 order as Sharrow on the CPU. It writes directly into a compact table organized by decision maker.

The Phase 46 service supplies reusable GPU memory and exactly the same MT19937 random tickets as ActivitySim. A second GPU kernel turns utilities into weights, searches the probability line, and proposes the selected shortlist position. This avoids building a large pandas interaction table and running the Sharrow CPU utility compiler during production.

Phase 47 is still not "no CPU." The compact utility table returns to NumPy for ActivitySim's exact probability normalization and logsum. Normally the GPU selection is accepted. If a ticket is extremely close to a probability boundary, NumPy decides that row exactly. ActivitySim still coordinates the 34 model steps and owns the official random ledger.

193. How did we prove the GPU arithmetic was identical?

We did not trust a few example rows. We enabled an expensive shadow mode for a complete 50,000-household model run. Every one of the 19 final programs ran twice: once through the new CUDA compiler and once through Sharrow's CPU compiler. The checker compared the binary float32 representation of every utility.

The result was zero differing bits across all 4,696,676 shortlisted alternatives. It also compared GPU and exact final positions. Only 7 of 201,390 rows were close enough to a probability boundary to enter the safety guard, and there were zero unguarded choice disagreements.

We then ran three normal production comparisons. An independent verifier found zero changed decision cells and seven byte-for-byte identical published CSV files in every pair. That includes the final households, persons, tours, and trips - not just a total or average.

194. Why replace the Numba safety helper?

Phase 46 used ActivitySim's Numba-compiled preserved-order sampler for the safety rows. It was fast after compilation, but a fresh process spent about 1.7 seconds compiling it before the model began. Phase 47 needs only a tiny number of final safety rows, so that cold cost was wasteful.

The replacement uses NumPy's float64 cumulative sum and `searchsorted`. A focused test compared it with ActivitySim's compiled helper across realistic and full 1,454-zone shapes. Choices and selected float32 probabilities were identical. Phase 47's combined preliminary-runtime prewarm fell to about 0.012 seconds, while its four final CUDA programs added roughly 0.003 seconds after the first CUDA cache was populated. The first fresh cache paid about 0.054 seconds. All cold costs are included in the lifecycle results.

195. How fast is the new final-choice boundary?

The Phase 46 final-choice runtime took a 1.786-second median. Phase 47 took 0.362 seconds. That is 4.94 times faster and saves about 1.424 seconds directly at this boundary.

Measurement	Phase 46	Phase 47	Result
final-choice runtime	1.786 s	0.362 s	4.94x faster
all five target components	29.8 s	27.7 s	1.076x faster
complete 34-step lifecycle	144.600 s	143.815 s	1.0055x faster

The complete lifecycle improvements in the three pairs were 1.933, 1.785, and 0.485 seconds. Phase 47 won every pair, but the median whole-model saving is only 0.785 seconds, or 0.54%.

196. Why can 4.94x become only 1.0055x for the whole model?

Suppose a two-hour trip includes a one-minute ticket line. Making that line five times faster is useful, but it cannot make the entire trip five times faster. Before Phase 47, this final-choice boundary was only about 1.2% of the already accelerated Phase 46 lifecycle.

Most of the 34-step model was unchanged. It still loads data, creates people and tours, schedules activities, makes modes and trips, writes matrices, and produces reports. Amdahl's law is the name for this limit: total speedup is bounded by the fraction of total time that was improved.

This is why the report keeps three numbers separate: kernel or boundary speedup, target-component speedup, and whole-model speedup. Mixing them would make a real result sound much larger than it is.

197. Which components improved, and which did not?

Component	Phase 46	Phase 47	Result
school location	6.1 s	5.2 s	1.173x faster
workplace location	9.2 s	8.7 s	1.057x faster
joint-tour destination	3.4 s	3.5 s	0.971x; 0.1 s slower
non-mandatory-tour destination	8.4 s	8.4 s	tied at timer precision
at-work subtour destination	2.5 s	2.2 s	1.136x faster

The five components together improved by 1.8 to 2.3 seconds in every pair. Joint tours are still too small to hide per-call setup, so batching small segments remains important. Reporting that regression prevents an average from hiding where the next work belongs.

198. What can we honestly claim now?

We can claim that, on the pinned public Prototype MTC extended benchmark and this RTX A4000 environment, a strict GPU compiler made the final sampled-choice boundary 4.94 times faster while producing byte-identical published outputs in three fresh comparisons. We can also claim a replicated 1.0055x complete lifecycle improvement over the already GPU-accelerated Phase 46 runtime.

We cannot claim that ActivitySim is entirely GPU-only, that every possible ActivitySim expression is compiled, or that every computer will get the same speed without rerunning qualification. Exact logsums and normalization still cross a compact CPU boundary. Unknown formulas fail closed. New hardware, software, data, or formulas must rerun shadows and output checks.

199. What should the next ambitious phase build?

The next phase should join several adjacent operations into one resident destination graph instead of adding another isolated kernel. It should:

- keep compact samples, mode logsums, final utilities, probabilities, and selected results on the GPU across neighboring stages;
- compile an exact GPU normalization and logsum ABI, with CPU shadows proving its addition, exponential, division, overflow, and zero-probability rules;
- batch the small joint-tour segments so they share uploads and launches;
- retain random state by stable chooser identity instead of reconstructing it at each call;
- expose the compiler and numeric policy behind a Sharrow or ActivitySim backend interface rather than only a benchmark runner; and
- preserve fail-closed formulas, exhaustive shadows, three fresh matched pairs, cold-cost accounting, and independent final-table verification.

The likely direct prize is removal of the remaining compact utility download and Python call boundaries. The larger prize is architectural: a reusable upstream backend in which a whole chain of destination work is device resident. Only that wider graph can turn another large boundary speedup into a meaningful whole-model reduction.

200. What did Phase 48 build?

Phase 47 sent a short list of final destination scores back to the CPU. The CPU turned those scores into probabilities, selected one destination, and calculated a logsum. Phase 48 keeps that chain on the GPU.

The new chain is called a **resident destination graph**. "Resident" means the working numbers stay in GPU memory between related calculations. "Graph" means the output of one operation becomes the input of the next operation: score, weight, total, probability, random choice, and logsum. It is not a picture or a road network in this context.

The graph covers all five final sampled-destination families: school, workplace, joint tours, non-mandatory tours, and at-work subtours. The public run contains 19 calls, 201,390 decision makers, and 4,696,676 shortlisted places.

201. How does a score become a probability?

Suppose three places have final scores of 1, 2, and 3. A model should prefer the higher score, but a score is not yet a percentage. A common travel-model rule uses the exponential function, written `exp`:

```
weight = exp(score - highest score)
```

Subtracting the highest score keeps the numbers safe without changing their relative chances. Our example becomes approximately 0.135, 0.368, and 1.000. Their total is 1.503. Dividing each weight by that total produces probabilities of about 9%, 24%, and 67%. They add to 100% apart from tiny rounding effects.

The model then places those probabilities end to end on a line from zero to one. A random ticket lands somewhere on that line. If the ticket is 0.40, the third place wins because the first two places cover only the first 0.33 of the line. This is called inverse cumulative-distribution selection. The long name just means "walk along the probability line until the ticket is passed."

202. What is a logsum, and why keep one CPU operation?

A logsum summarizes how attractive a whole group of alternatives is. It rises when there are more good choices and falls when choices are poor. Later travel model steps can use it as an accessibility or mode-quality signal.

After stabilizing scores, its essential calculation is:

```
logsum = highest score + log(sum of exponential weights)
```

The GPU now calculates the weights and their total. This computer's CUDA `log` operation does not always produce the identical final bits as NumPy's `log`, so Phase 48 returns one small total per decision maker and lets NumPy perform the final scalar logarithm. That is an intentional exactness boundary, not an oversight. Returning one number per person is far cheaper than returning every destination score.

203. Why can two correct exponential functions disagree?

Computers store most decimal-like values in **floating-point** form. A 32-bit floating-point number, or `float32`, has only a fixed number of binary digits. Many real numbers must be rounded to the nearest available pattern.

Functions such as `exp` are not calculated from an infinite mathematical definition every time. Fast libraries use carefully designed approximations. NumPy on this Windows CPU and CUDA on the RTX A4000 can both be very accurate while rounding the last bit differently. Usually that difference is harmless. Here, a one-bit probability change could put a rare random ticket on the other side of a boundary and select another zone. Exact replication therefore means matching the bits, not merely printing the same six decimal places.

204. How did we prove the exponential rule instead of guessing?

A `float32` value has exactly 2^{32} , or 4,294,967,296, possible bit patterns. Phase 48 includes a scanner that visits every one. Some patterns mean infinity or "not a number"; 4,278,190,080 patterns are finite values that can be compared.

The scan compared the CUDA approximation with NumPy 2.4.6. Across all finite magnitudes it found 2,180,536 differences, mostly in extreme ranges that the destination model does not accept. The backend deliberately accepts ordinary exponential inputs only from -80 through 80. There are 2,235,564,034 finite bit patterns in that range, and exactly 73 produced a different result.

The project stores those 73 exact input and output bit pairs in a tiny sorted correction table. A hash acts like a tamper-evident fingerprint for the table. The reproducible scan found exactly the same 73 pairs and the same fingerprint. Inputs outside the declared range stop the fast path, except for the reviewed -999 padding marker used for an empty shortlist slot.

This is stronger than testing millions of random examples: it is complete for every possible `float32` value inside the declared rule.

205. What does "fail closed" mean here?

A fast backend can be dangerous if it silently guesses what to do with a new formula. Phase 48 publishes a versioned contract. It recognizes four reviewed formula shapes and shortlist widths of 21, 25, 29, or 30. It also declares its number types, random-number rule, exponential range, logsum rule, and hashes.

If a formula, width, number range, or random-state relationship is outside that contract, required mode stops with an error. It does not return a plausible but unproved answer. Test mode can run both paths and compare them. This design follows the useful idea in ActivitySim's Sharrow backend: compiled speed is optional during development, but a production run can require it and refuse silent fallback.

206. How can a random number stay on the GPU and still be the same?

ActivitySim does not use one uncontrolled stream of random numbers. It assigns repeatable random state to stable chooser identities, such as a tour ID. This lets a model reproduce decisions even when work is divided differently.

Phase 46 generated the 30 shortlist-sampling draws with the same MT19937 rule as NumPy and kept each generator's internal state on the GPU. Before final choice, ActivitySim's intervening logsum work advances each chooser's official ledger by six draws. Final rows may also be reordered.

Phase 48 maps each final row back to its stored chooser identity, advances that state by exactly six values, and uses the next value as the final ticket. It resumes only when the chooser sets and ledger changes pass strict checks. All 19 calls resumed successfully in every qualified run. Tests also reseed NumPy from scratch, generate the long sequence independently, and compare every 64-bit random value after reordering and skipping six draws.

207. What data transfer did Phase 48 remove?

Phase 47 downloaded a dense compact utility table containing one score for every shortlisted alternative. Across 19 calls that table occupied 23,759,764 bytes, about 23.76 megabytes.

Phase 48 leaves it on the device. It transfers 2,618,966 bytes in total: chosen positions, guard flags, selected probabilities, row totals, maxima, and other small orchestration results. In simple terms, the system now returns mostly answers instead of returning the worksheet used to calculate the answers.

This is a roughly 9.1-to-1 reduction relative to the dense utility download, although it is not zero transfer. ActivitySim still needs compact results to continue its CPU-controlled model.

208. How did the complete live shadow work?

The expensive shadow ran the entire 50,000-household public model. For every one of the 19 calls, it calculated the final chain through both the new GPU graph and the authoritative CPU answer path. It compared the binary results at six checkpoints:

- final utility;
- exponential weight;
- row total;
- normalized probability;
- selected alternative; and
- logsum.

Across all 4,696,676 shortlisted alternatives, every mismatch count was zero. The actual unshifted utilities ranged from about -54.528 to 22.047, inside the declared -80 to 80 contract. Seven of 201,390 choice rows were close enough to a cumulative-probability boundary to enter the independent safety guard. None disagreed before that check.

Finally, an output verifier compared accessibility, households, joint-tour participants, land use, persons, tours, and trips. All seven CSV files were byte-for-byte identical to the Phase 47 reference.

209. How fast was Phase 48?

Three fresh-process matched pairs compared Phase 47 with Phase 48. Each pair ran a control and candidate on the same machine and independently checked the published outputs.

Measurement	Phase 47	Phase 48	Honest result
resident final boundary, median	0.4002 s	0.3182 s	1.258x faster; 20.5% lower
direct boundary pairs won	-	3 of 3	replicated
all five target components	28.1 s	27.9 s	too close to timer noise
complete model, median	147.2 s	144.717 s	1.017x observed
complete-model pairs won	-	2 of 3	not replicated
changed decision cells	-	0, 0, 0	exact

The direct boundary saved 0.0363, 0.0934, and 0.0568 seconds in the three pairs. Its median saving was 0.0820 seconds. That is a real repeated boundary gain over an already accelerated Phase 47 control.

210. Why are we not claiming a whole-model victory?

The complete runs changed by -0.617, +2.982, and +2.483 seconds. The median favored Phase 48, but the first candidate was slower. A roughly 145-second run contains operating-system scheduling, file output, memory management, and 34 model steps. Their normal variation is much larger than an 0.082-second local saving.

The correct conclusion is therefore narrower: Phase 48 made its directly measured boundary faster in all three pairs and preserved exact outputs. It did not make a large enough dent to prove whole-model superiority. Reporting the attractive 1.017x median without the two-of-three fact would exaggerate the evidence.

This is Amdahl's law again. Once a boundary is only about 0.3 to 0.4 seconds, even removing it completely could improve a 145-second run by only about one quarter of one percent.

211. What is the larger achievement?

The main achievement is an exact reusable **backend contract**, not merely an 82-millisecond saving. The project now has:

- a compiled destination formula shared by CPU proof and CUDA execution;
- device-resident utility, exponential, normalization, probability choice, and compact logsum reduction;
- exact continuation of ActivitySim's identity-keyed random state;
- a complete float32-domain exponential qualification method;
- a fail-closed version and arithmetic fingerprint;
- a full live dual-run shadow at every important numerical checkpoint; and
- four independently verified complete runs: one shadow and three production candidates, each with seven byte-identical output files.

That is the kind of evidence an upstream ActivitySim or Sharrow backend needs. It separates a portable interface from machine-specific arithmetic evidence and makes unsupported work visible.

212. What should the next major phase do?

The next phase must be large enough to remove work measured in seconds, not another few hundredths of a second. It should build an **inter-stage destination supergraph**.

Today, ActivitySim still constructs pandas tables and crosses Python function boundaries between sampling, mode-choice logsums, final destination choice, and later table updates. The next graph should keep the compact sample and mode logsums resident across those stages, batch the many small joint-tour and at-work calls, and return only final zones plus required diagnostics.

The practical order is:

- 1 measure a detailed CPU/Python/GPU timeline for all five destination families and identify at least one second of removable orchestration;
- 2 define one versioned input/output schema for the entire destination chain;
- 3 compile and batch sampling, logsum consumption, final utility, choice, and required aggregation inside one reusable service;
- 4 expose that service behind a real Sharrow or ActivitySim backend interface;
- 5 keep Phase 48's arithmetic hashes, state-ledger checks, fail-closed rules, complete dual-run shadow, and byte-identical output verifier; and
- 6 require three matched wins at the wider component boundary before testing a whole-model performance claim.

The target should be a repeated multi-second reduction across the five destination components. Only a boundary that large has a fair chance to rise above whole-model noise and make a meaningful dent in total ActivitySim time.

213. What did Phase 49 actually build?

Phase 49 built the first cross-stage part of that destination supergraph. A mode-choice logsum is a single score summarizing how easy it is to reach one possible destination using all available travel modes. The GPU already made that score. Phase 48 then brought every score back to the CPU, put it in a pandas table, and sent the same information back to the GPU for final choice.

Phase 49 keeps that vector on the graphics card. It hands the result directly from the mode-logsum calculation to the final destination formula. Think of two workers at the same workbench: previously the first worker mailed a part to another building, where it was labeled and mailed back. Now the first worker passes it across the bench.

214. How does the system know it passed the right rows?

Travel tables often repeat an ID because one person or tour has several possible destinations. Merely checking that two arrays have the same length is not enough. The sequence [7, 7, 9] is different from [7, 9, 9], even though both have three rows.

Each Phase 49 packet carries the complete repeated ID sequence. The final stage must present exactly the same IDs in exactly the same order. A missing packet, reordered duplicate, changed number of rows, stale packet, or packet left at the end stops the model. The fast path never guesses which rows match.

Tests deliberately try missing, reordered, malformed, and unconsumed packets and require them to fail.

215. Why do school and workplace need a special return path?

The three tour-destination families only need the final chosen zone. School and workplace also publish the chosen destination's mode-choice logsum for later analysis. Returning no logsums would therefore change a required output.

Phase 49 first chooses the destination on the GPU. It then gathers just one original 64-bit logsum per selected person and returns those compact values. Stable person IDs put them back into the correct output rows. Across the public run, this returns 107,794 values, or 862,352 bytes, instead of returning every logsum for 2,543,709 school and workplace alternatives.

The final utility uses a 32-bit copy, but the published value comes from the original 64-bit device vector. That distinction preserves the output bits.

216. How much data movement did Phase 49 remove?

Across all 19 calls, the former bridge downloaded 37,573,408 bytes of 64-bit logsums and uploaded 18,786,704 bytes of 32-bit copies. Phase 49 eliminates both full transfers, then returns 862,352 bytes of selected school/workplace output.

The net reduction is 55,497,760 bytes per complete run. In friendlier units, that is about 55.5 megabytes of copying removed from a small boundary.

Transfer	Phase 48	Phase 49
all logsum values from GPU to CPU	37.57 MB	0
selected published values	included above	0.86 MB
final-utility values from CPU to GPU	18.79 MB	0
net bytes avoided	-	55.50 MB

217. Was the handoff really faster and exact?

Yes at the directly measured boundary. A qualifier used the exact row counts of all 19 public calls and repeated both paths 31 times. The old path copied GPU-to-CPU, converted to 32-bit numbers, and copied back. The new path converted on the GPU and returned only selected published outputs.

Measurement	Phase 48 bridge	Phase 49 bridge	Result
median time	13.265 ms	1.226 ms	10.822x faster
time reduction	-	90.759%	repeated 31 times
32-bit result mismatches	-	0	bit-identical
alternative rows covered	4,696,676	4,696,676	complete workload shape

This is a strong boundary result. It is not evidence that the whole travel model became ten times faster; the boundary was only about 13 milliseconds to begin with.

218. Did it make the complete ActivitySim model faster?

Three fresh Phase 48/49 pairs ran the full 34-step, 50,000-household public model. All three pairs produced zero changed decision cells in seven published files.

The whole-model median changed from 144.100 to 143.517 seconds: an observed 0.583-second, 0.405%, or 1.004x improvement. The candidates changed by +0.284, +0.583, and -0.016 seconds relative to their paired controls. Because the third candidate was very slightly slower, we do not call this replicated whole-model superiority.

The apparent 0.583-second median is also far larger than the approximately 0.012-second boundary saving. That tells us normal lifecycle variation, not the handoff alone, dominates the whole-model number. The honest promoted claim is 10.822x for the handoff, exact full outputs, and no proven large lifecycle dent.

219. What remains for the ambitious supergraph?

Phase 49 removed an unnecessary trip across the CPU/GPU boundary, but it did not remove the much larger work that creates dense logsum inputs. The current 19 calls still make large pandas rows, resolve formula bindings, pack dozens of columns, and upload them before the mode utility kernel runs.

The audit found about 1.14 seconds of binding resolution, 0.77 seconds of host packing, 0.31 seconds of upload, and 2.65 seconds of mode-utility kernels, with about 1.95 billion bytes of dense inputs. The next phase should describe each person or tour once, keep sampled destination IDs compact, keep land-use and skim tables resident, and generate the 10 floating and 31 integer inputs on the GPU.

That next boundary is hard because waiting-time formulas, six controlled random draws, availability flags, time periods, skims, and exact arithmetic all have to match. But it is measured in seconds rather than milliseconds. It is the credible route to a repeated component gain and a visible reduction in total ActivitySim time.

220. What did Phase 50 build?

Phase 50 built the larger input half of the destination supergraph. Before a GPU can score a possible destination, it needs facts such as the traveler's age, car ownership, value of time, trip period, destination parking cost, local density, and whether each mode is available.

ActivitySim used to make one large spreadsheet row containing those facts for every possible sampled destination. The same traveler facts were copied again and again. Phase 50 instead sends one compact record per person or tour, plus a short list of destination zone numbers. A CUDA kernel looks up land-use and travel-time facts and creates the needed inputs on the graphics card.

The reviewed contract has 10 decimal-number inputs, 31 whole-number or true/false inputs, and six groups of origin, destination, and time coordinates. The travel-mode formula then uses 315 terms to score 21 ways of traveling.

221. Why does compact data matter?

Imagine 25 students each comparing 30 colleges. Their age, household, and car ownership do not change across the 30 choices. Writing those facts 30 times is easy to understand, but expensive at millions of rows.

The public run contains 201,390 people or tours and 4,696,676 sampled destinations. The old boundary packed 192,563,716 values and moved about 1.95 gigabytes to the GPU. Phase 50 sends about 58.26 megabytes of compact state and creates the repeated values there. It avoids 1.896 gigabytes, or about 97% of that input transfer.

This does not mean the whole model uses 97% less memory or becomes 33 times faster. It means one specific input boundary moves 97% fewer bytes.

222. What happened to the random waiting times?

Taxi and ride-hail waiting times use six repeatable random values for each chooser: origin and destination values for taxi, solo ride-hail, and shared ride-hail. These are not casual random numbers. ActivitySim gives every stable person or tour ID its own numbered sequence so the model can repeat exactly.

ActivitySim normally makes the six values once and then repeats them over all sampled destinations. Phase 50 asks the same official random channel for the same six values on the unique IDs, but skips the repetition. The random ledger advances exactly as before. The first full implementation used ActivitySim's ordinary repeated form and passed; the compact version then passed the entire model proof again.

223. How much faster are the destination parts?

Three fresh pairs compared Phase 49 with Phase 50. Each control and candidate started in a new process, and every output was checked independently.

Measurement	Phase 49	Phase 50	Plain-English result
five destination components, median	28.4 s	16.9 s	11.5 s saved; 1.680x
reduction in those components	-	40.5%	repeated in all 3 pairs
complete 34-step model, median	148.700 s	136.518 s	12.182 s saved; 1.089x
complete-model pairs won	-	2 of 3	median win, not every run
changed modeled decision cells	-	0, 0, 0	exact decisions

The targeted destination work improved in every pair: 28.8 to 19.2 seconds, 27.9 to 16.9 seconds, and 28.4 to 16.8 seconds. This is the first recent phase whose direct improvement is many seconds rather than milliseconds.

224. Why did one whole run still lose?

The first pair's whole model took 150.2 seconds for Phase 49 and 150.7 seconds for Phase 50, even though its destination stages saved 9.6 seconds. Unchanged work elsewhere became slower: for example, daily-activity and scheduling steps varied with operating-system, CPU, and memory conditions.

The other candidates won by about 12.2 and 11.9 seconds. The median whole-run gain is useful evidence that the optimization now makes a visible dent, but a careful report does not hide the loss. The strongest replicated statement is: all three destination-component pairs won by a large amount.

225. Did the answers stay the same?

Yes. Against Phase 49, all seven published CSV files are byte-for-byte identical in all three pairs. Phase 50 changes how inputs are created, not the arithmetic policy already established by Phase 49.

Against regular ActivitySim, the independent checker still finds zero changed decision cells across accessibility, households, joint-tour participants, land use, persons, tours, and trips. Some stored logsum diagnostic decimals differ in the last few digits because the qualified GPU arithmetic follows its declared floating-point policy. The largest destination-logsum difference is 0.00001, ten times below its 0.0001 limit. The largest mode-choice-logsum difference is about 0.000003814, below its 0.00001 limit. For that cumulative comparison we say **exact decisions and bounded diagnostic logsums**, not that every byte in every CSV is identical.

226. What assumptions does Phase 50 make?

The fast path is intentionally strict. It is qualified for the public Prototype MTC extended model, its 1,454 zones, its reviewed mode formula, its five time periods, its current skim directions, and its current lack of separate transit-access subzones. Sampled destinations for one chooser must be together, and owner facts must not change within that group.

The land-use table is fingerprinted. A new input, changed zone layout, different period name, reordered owner group, unsupported direction, or chunked call stops with an error. That can sound unfriendly, but it protects research results: unsupported work is visible instead of silently switching to a path with different arithmetic.

227. What exactly is still not on the GPU?

ActivitySim still controls the 34 model steps, owns pandas result tables, and runs many non-destination models and file-writing tasks on the CPU. Phase 50 also still creates the 10-plus-31 dense input arrays in GPU memory. The next kernel reads those arrays immediately to compute utilities, so there is one remaining write-then-read stage on the device.

That is why this is a major improvement, not a claim that the entire model is GPU-only. The graphics card now performs the expensive destination input and choice mathematics, while the CPU remains the workflow conductor.

228. What should Phase 51 do?

Phase 51 should compile a **fused compact-source utility kernel**. Instead of first writing every dense input and then reading it, one generated kernel would look up a compact value exactly when a formula term needs it and add the term directly to one of the 21 mode scores. Those scores would flow straight into nested logit and final destination choice.

The phase should also precompile the ten purpose-specific programs and retain compact owner tables across related calls. A successful result should reduce the five destination components below 12 seconds, win all three fresh matched pairs, preserve zero changed modeled decisions, keep diagnostic logsums within their declared limits, and use no fallback. That is ambitious enough to seek another roughly five-second direct saving rather than polishing a tiny boundary.

229. How fast is the latest system versus regular ActivitySim?

We ran a second experiment specifically for that question. Three fresh pairs compared regular pinned ActivitySim, with its normal Sharrow acceleration required, against the complete ChoiceForge system through Phase 50. This is a serious baseline, not a deliberately slow teaching program.

Complete 50,000-household model	Regular ActivitySim	Phase 50	Result
pair 1	209.3 s	138.918 s	1.507x
pair 2	209.3 s	134.118 s	1.561x
pair 3	208.6 s	133.916 s	1.558x
median	209.300 s	134.118 s	75.182 s saved; 1.561x

The median reduction is 35.92%, and Phase 50 wins all three pairs. Every pair has zero changed modeled decision cells. Four output files are byte-identical; the differing logsum diagnostic decimals stay within the small declared limits explained above.

This is a **cumulative** result. It includes the work from earlier phases: faster scheduling, GPU trip destination, resident skims and random state, compiled final choice, the Phase 49 handoff, and Phase 50 compact inputs. The separate Phase 49/50 experiment is what tells us Phase 50 itself removes about 11.5 median seconds from the five destination components. Keeping those two comparisons separate prevents us from crediting one phase for the whole project's improvement.

230. What did Phase 51 build?

Phase 51 joined two GPU jobs that used to be separate. Phase 50 first created a very large set of temporary input rows on the GPU. A second GPU kernel then read those rows and calculated how attractive 21 travel modes were.

Phase 51 uses one generated kernel. For each sampled destination it reads the small person-or-tour record, looks up destination and travel facts, rebuilds the 10 decimal inputs and 31 whole-number or yes/no inputs in a tiny shared workspace, and immediately evaluates the 315-term travel-mode formula. The huge temporary input rows never exist.

This happens for all 19 reviewed destination-logsum calls, covering 201,390 people or tours and 4,696,676 sampled destinations. The nested-logit and final destination-choice stages remain on the GPU after the fused calculation.

231. What is a fused kernel, in everyday language?

Imagine a cafeteria where one worker fills 4.7 million trays, puts every tray in a giant warehouse, and a second worker immediately retrieves each tray to make a meal. Phase 50 removed the truck trip to an outside warehouse, but it still used a giant warehouse inside the GPU.

A fused kernel is more like a cook with a small workbench. The cook takes the compact ingredients for one order, prepares the few facts needed, and makes the meal immediately. There is no giant pile of finished trays between jobs.

The GPU still runs many orders at once. One short int32 owner number connects each sampled destination row to its person or tour. Threads cooperate to fill shared memory, which is a small, fast workspace visible to threads in the same block. They then compute the same formula terms and the same 21 ordered utility sums as before.

232. How much GPU memory disappeared?

The result is large and repeatable in every qualified run:

Device-memory item	Phase 50	Phase 51	Change
dense row allocations across 19 calls	1,953,817,216 B	0 B	eliminated
row-owner maps across 19 calls	0 B	18,786,704 B	one int32 per sampled row
net allocation volume across 19 calls	-	-	1,935,030,512 B removed
largest call's live row state	295,698,208 B	2,843,252 B	292,854,956 B removed
compatibility storage per call	large row arrays	416 B	tiny bootstrap only
compact host upload	58,259,296 B	47,787,016 B	17.98% lower

These 19 calls happen one after another, so 1.94 GB is the **sum of allocation volume avoided**, not 1.94 GB that was all live at once. The best estimate of the peak capacity improvement is the largest-call row: about 292.85 MB less live temporary state.

"ABI" means the agreed shape and type of information passed between two pieces of code. Here it was a large set of row-shaped arrays. Removing it is important because GPU memory has a hard limit. Saving about 1.94 billion bytes of allocation work across a run reduces allocator and memory-pool pressure; removing about 292.85 MB from the largest call can let a larger population or choice sample fit without chopping it into smaller batches.

This does not mean the whole model now uses only 416 bytes. Skim cubes, land-use data, utilities, random state, and other model tables still occupy memory. The number describes this specific temporary boundary.

233. Did Phase 51 itself make Phase 50 faster?

Not convincingly on this GPU. Three fresh Phase 50/51 pairs gave these medians:

Incremental measurement	Phase 50	Phase 51	Honest conclusion
generator plus utility / fused kernel work	1.723 s	1.739 s	approximately neutral; 0.9% slower
complete instrumented logsum service	8.570 s	8.943 s	4.35% slower
five destination components	16.4 s	17.0 s	3.66% slower; 1 of 3 pairs won
complete 34-step model	136.300 s	135.018 s	median 0.94% better, but only 2 of 3 pairs won

The whole-model median happened to improve, while the directly targeted components became slower. That is a warning that unrelated operating-system, CPU, memory, and compilation variation can dominate a one-second difference. We therefore do **not** claim that Phase 51 is faster than Phase 50.

The original goal in Section 228 was to push the five destination components below 12 seconds and win all three pairs. Phase 51 did not reach that speed goal. Its successful result is the memory breakthrough and exact fused execution. Keeping the missed target in the explainer is part of making the research reproducible rather than turning every experiment into a victory.

234. Why was the service slower if it removed so much memory traffic?

Removing an array also changes how work is scheduled. Phase 50's small input kernel handled a group of destinations belonging to the same person, so it could reuse that person's data. Its utility kernel then read the finished rows. Phase 51 exposes one block per destination row for high parallelism, but those blocks must repeatedly find and read their compact owner state.

The actual fused GPU pipeline is nearly tied with the old two-kernel pipeline. The clearer penalty is new first-use compilation: the Phase 51 source must be compiled once before its first call, adding roughly 0.30 seconds at the median relative to Phase 50. Compact packet preparation still takes about four seconds and is not removed by fusion.

Several alternatives were run and rejected. Fully inlining every lookup did too much repeated work. One block per owner reused data but serialized rows. A 128-thread block was slower than 256 threads. A redundant synchronization was removed. Non-aliasing compiler hints did not help. The retained int32 packet is protected by range checks and reduces capacity even though it did not create a clear speed win.

235. Did the travel answers stay the same?

Yes. Three incremental Phase 50/51 pairs and three regular-ActivitySim/Phase 51 pairs were checked independently. Every comparison found zero changed modeled decision cells.

In the incremental experiment, all seven published CSV files were byte-for-byte identical in every pair. That is stronger than merely checking that the chosen destination or mode was the same.

The fused compiler uses the same hashed strict intermediate representation, the same 315 feature order, and the same ordered 21-alternative floating-point accumulation. It also checks that exactly 10 float sources, 31 integer sources, and six skim-coordinate groups are present. Unknown or changed inputs stop the run. There is no hidden CPU or generic fallback.

236. How fast is the complete project now versus regular ActivitySim?

That is a different and much stronger comparison. Three fresh pairs compared regular pinned ActivitySim with Sharrow required against the complete ChoiceForge system through Phase 51:

Complete 50,000-household run	Regular ActivitySim	Phase 51	Speedup
pair 1	202.6 s	135.117 s	1.499x
pair 2	202.6 s	135.318 s	1.497x
pair 3	202.8 s	137.817 s	1.472x
median	202.600 s	135.318 s	1.497x

The median saves 67.282 seconds, or 33.21% of total model time. Phase 51 wins all three pairs. The five destination components fall from a 42.8-second median to 16.9 seconds: 25.9 seconds saved, 60.51% lower, or **2.533x**.

This is the accumulated advantage of Phases 1-51, including scheduling, mode-choice, trip-destination, resident state, exact GPU random draws, compact inputs, and the fused memory path. It would be wrong to say Phase 51 alone created the 67-second gain. The Phase 50/51 experiment above isolates what Phase 51 changed.

237. What assumptions protect this result?

The fast path is qualified for the public Prototype MTC Extended model at the reviewed 50,000-household scale, with 1,454 dense zero-based zones, the current 315-term/21-mode formula, current five time periods and skim directions, and sampled destinations grouped contiguously by stable chooser ID. It assumes no separate transit subzone access-distance column and no chunking of these location-logsum calls.

The land-use table is fingerprinted. Owner facts must stay constant inside a group. Compact whole-number fields must fit in int32; the code checks this before narrowing them and widens them back to 64 bits at the strict formula boundary. A changed schema, period, direction, land-use table, owner order, or integer range raises an error.

These restrictions are not a promise that every ActivitySim model already works. They are replication guarantees for the public model actually tested.

238. What should Phase 52 do?

The next phase should keep the 1.935-GB aggregate allocation saving and 292.85-MB largest-call live-state saving, then attack the two costs Phase 51 exposed.

First, compile the fused program before component timing and cache it by the hashes of the formula and compact ABI, including across fresh processes. That targets the roughly 0.30-second first-use penalty.

Second, retain compact owner packets across related purpose calls and generate a multi-row owner tile. Two or four rows should share cached owner facts while independent warps calculate their rows in parallel. This seeks both owner locality and enough GPU work at once, instead of choosing only one.

Phase 52 should not pass merely because memory stays low. It should require all of the following together: no dense ABI, no fallback, zero changed decisions, three-of-three wins in synchronized kernel and complete service time, and three-of-three wins across the five destination components. That is the practical path from a capacity breakthrough to a provable incremental speed improvement.

239. Did Phase 52 meet that harder goal?

Yes. Phase 52 is the first phase after the large Phase 51 memory improvement to show a repeated speed win against the immediately previous GPU system at every measured level.

Three fresh Phase 51/52 pairs produced these medians:

What was timed	Phase 51	Phase 52	Improvement
fused GPU kernel pipeline	2.030 s	1.657 s	1.225x; 18.36% lower
complete instrumented destination service	9.069 s	7.846 s	1.156x; 13.48% lower
five destination model components	16.9 s	16.2 s	1.043x; 4.14% lower
all 34 model steps, including prewarm	139.0 s	136.527 s	1.018x; 1.78% lower

Phase 52 won all three pairs on every row of this table. That matters because a single computer run can be accidentally fast or slow. Repeating a matched control-then-candidate pair three times makes the claim much harder to fake with ordinary timing noise.

240. What does “persistent service” mean?

Think about doing 19 school assignments that reuse the same textbook and calculator. A nonpersistent approach cleans the desk, finds the book, checks the instructions, and gets a new sheet of paper for every assignment. A persistent service leaves the trusted book, instructions, calculator, and reusable workspace ready for the next assignment.

Phase 52 remembers two safe plans:

- the **semantic plan** says which model facts each formula name means;
- the **native ABI plan** says the exact order, type, and shape of values sent to the CUDA kernel.

It also reuses GPU memory for compact input packets, row-owner numbers, and utility results. Buffers grow when a larger call arrives, then serve smaller calls without another allocation. In each qualified run, the system reused 9 semantic plans, 9 native plans, 16 utility buffers, 152 packet buffers, and 16 row-owner buffers.

“Persistent” does not mean memory is kept forever. After the nineteenth and last destination call, the service releases these objects at an explicit safe boundary.

241. What is prewarming, and why is the source hash important?

GPU source code cannot run directly. It must first be translated into a binary program for the particular GPU. That compilation is useful work, but doing it inside the first timed destination call makes the model look slower and makes run-to-run timing less predictable.

Phase 52 stores the qualified CUDA source in the repository, reads it before the destination timers begin, and asks CuPy to compile it. CuPy keeps a content-addressed binary cache on disk, so a fresh Python process can reuse the compiled result. The three qualified prewarm calls took about 0.010 to 0.023 seconds on this machine.

The exact source has this SHA-256 fingerprint:

```
599a9704be0992d2863320390cbca0028c7a578ecacf72d69de2e658a5d79906
```

A hash is like a very sensitive digital fingerprint. Changing even a small part of the source produces a different fingerprint. Every Phase 52 call records the expected hash, and the qualification rejects mismatches. This connects the benchmark claim to the actual checked-in program instead of to an unknown program left in a compiler cache.

242. What is a four-row GPU tile?

Each modeled person or tour considers several possible destinations. Nearby rows in the sample often belong to the same owner and therefore share facts such as age, household size, vehicle ownership, and tour purpose.

Phase 51 gave one GPU thread block one sampled row. Phase 52 gives a block four adjacent rows. The block loads shared owner facts once into a small, fast on-chip workspace. It then loads each row's destination and time values and calculates the same 315 formula terms for all 21 travel modes independently.

This is like bringing four worksheets about the same student to one table: the student's record is opened once, while four separate destination answers are still calculated. Two-row tiling was built and tested first. Four rows ran faster on this GPU, so four became part of the checked-in proof contract.

The tile does **not** mix decisions between rows. Each row keeps its own destination facts, arithmetic order, utility outputs, probabilities, random draw, and final answer.

243. Did Phase 52 keep the Phase 51 memory win?

Yes. Across the 19 calls, Phase 52 still eliminates 1,953,817,216 bytes of dense GPU input allocation and avoids the same amount of dense host packing. It uploads only 47,787,016 compact bytes, for 1,906,030,200 net bytes of upload avoided. Its row-owner maps total 18,786,704 bytes, exactly one four-byte int32 number per sampled row. Compatibility bootstraps total only 7,904 bytes.

These numbers are totals across sequential calls, not all memory live at one instant. The distinction from Section 232 still applies.

Phase 52 also fixes a lifecycle issue discovered during development. One experimental run completed every new destination call, but Windows later refused a small CPU allocation while ActivitySim was building a trip table. The final runtime now releases unused Phase 52 buffers, runs garbage collection, and returns unused CUDA allocator blocks immediately after the last destination consumer. It freed 101,630,976 bytes at that boundary in every qualified run while preserving the shared skim arrays needed later.

244. Are the answers still the same?

Yes. Six independent comparisons were made: three Phase 51/52 pairs and three regular-ActivitySim/Phase 52 pairs. Every verifier found:

- zero changed modeled decision cells;
- zero rows with changed modeled decisions;
- all declared floating diagnostic differences below their limits;
- no hidden fallback to a generic generator or CPU evaluator.

In every incremental comparison, five published files were byte-for-byte identical. Some published person and tour files contain floating logsum diagnostics. Their text can differ by a few millionths because the GPU follows the qualified float policy. The largest destination-logsum difference was about 0.000003, far below its 0.0001 limit; school and workplace maxima were about 0.00000191, below their 0.00001 limits. These tiny diagnostics changed no choice.

This is why the honest claim is **exact modeled decisions with bounded logsum diagnostics**, not that every decimal character in every output file is always identical.

245. How fast is Phase 52 versus regular ActivitySim?

Three new public 50,000-household comparisons gave:

Pair	Regular ActivitySim, all steps	Phase 52, all steps	Whole-model speedup	Regular destination	Phase 52 destination	Destination speedup
1	204.8 s	142.427 s	1.438x	43.6 s	16.7 s	2.611x
2	209.0 s	139.086 s	1.504x	43.2 s	15.8 s	2.734x
3	205.4 s	138.028 s	1.488x	43.0 s	16.3 s	2.638x
median	205.4 s	139.086 s	1.477x	43.2 s	16.3 s	2.650x

The median complete run saves 66.314 seconds, or 32.29%. The destination group saves 26.9 seconds, or 62.27%. Phase 52 wins all three pairs, and all three output comparisons have zero changed decisions.

The 1.477x complete-model result belongs to the accumulated ChoiceForge system through Phase 52. Phase 52 alone contributes the smaller but now replicated Phase 51/52 gains in Section 239. Keeping those two comparisons separate prevents us from giving one phase credit for years of earlier work.

246. What exactly is proven, and what is assumed?

The proof covers the public Prototype MTC Extended model on this machine at 50,000 households and 1,454 zones. Each run contains 34 timed model steps. The new service covers 19 destination-logsum calls, 201,390 people or tours, and 4,696,676 sampled destination rows.

It assumes the reviewed 315-term, 21-mode formula; the current five time periods and skim directions; dense zero-based zone mapping; contiguous groups of sampled rows belonging to one stable owner; no chunking inside these calls; and current land-use and compact integer ranges. The program checks important parts of these assumptions and stops on a mismatch.

The benchmark uses three fresh-process matched pairs in control-then-candidate order. That design controls workload and catches repeated regressions, but it does not eliminate every temperature, operating-system, antivirus, or cache effect. Results on another GPU, another ActivitySim model, or a different sample size must be measured again. Phase 52 is a strong replicated result, not a universal law that every model will be exactly 1.477 times faster.

247. What is the next big phase?

Do not spend Phase 53 merely changing the tile from four to another small number. The qualified tiled kernel takes only about 1.657 seconds, while building compact packets on the CPU still takes roughly 4.3 seconds inside the service. The larger prize is removing that remaining host preparation path.

Phase 53 should build a **model-wide device-resident destination data plane**:

- 1 Destination sampling should leave canonical owner IDs, sampled destination IDs, and group offsets on the GPU.
- 2 Permanent GPU tables should hold reusable person, tour, household, land-use, time-period, and skim facts with versioned fingerprints.
- 3 The tiled utility service should gather those facts directly, without first constructing pandas-shaped compact packets on the CPU.
- 4 Logsum, probability, keyed random choice, and selected-output gathering should remain connected on-device.
- 5 Fixed-shape segments should be tested with CUDA Graph replay to reduce the many small Python launch and orchestration costs.

A successful Phase 53 must keep Phase 52's source hashes, exact decisions, bounded diagnostics, zero fallback, memory lifecycle, and three-pair testing. It should eliminate most of the roughly 4.3-second packet-preparation cost, win all three matched pairs, and aim to bring the five destination components below 14 seconds. That would move the project from "a faster persistent kernel service" to "an end-to-end resident destination pipeline."

248. What did Phase 53 actually build?

Phase 53 built the first upstream compact-owner destination data plane. That name sounds complicated, but the idea is simple: stop copying the same person or tour facts onto every possible destination before the GPU can use them.

Imagine one student considering 25 colleges. The old path made 25 copies of the student's age, household, car access, and other facts, attached one college to each copy, and then scanned the copies to recover the one student record. Phase 53 keeps one student card plus a list of 25 college IDs. The GPU combines the card and each college only while doing the calculation.

ActivitySim has two public places where it builds destination logsums. Phase 53 intercepts both places before their large joins. For school and workplace it joins one sampled row per person to the original person table. For tour destinations it joins one sampled row per tour to the original person/tour facts. It then passes this compact owner table and the full destination vector to the already proven Phase 52 CUDA service.

249. Why is this more than deleting a safety check?

Phase 52 received a huge table whose owner facts were already repeated. It had to scan each repeated column to prove that all rows in a group really belonged to the same unchanged owner. Simply skipping those scans would be faster but unsafe.

Phase 53 changes where the data comes from. It constructs the compact table inside ActivitySim's public logsum functions from the authoritative one-row-per-person or one-row-per-tour source. It still checks that:

- sampled rows for an owner are next to one another;
- every compact owner ID exactly matches the corresponding sampled group ID;
- every tour-to-person relationship exists;
- required fields, zones, periods, integer ranges, formula hashes, and shapes match the qualified contract.

If any check fails, the model stops. There is no generic or CPU fallback. The speedup comes from avoiding unnecessary data construction, not from trusting unverified repeated data.

250. How much faster is Phase 53?

Three matched comparisons used the public 50,000-household, 1,454-zone model. Each comparison covered all 19 destination-logsum calls, 201,390 owners, and 4,696,676 sampled destinations.

What was measured	Phase 52	Phase 53	Improvement
build the compact packets	4.312 s	1.300 s	3.318x; 69.86% lower
complete destination service	7.846 s	4.989 s	1.573x; 36.41% lower
five destination model components	16.2 s	12.6 s	1.286x; 22.22% lower

Phase 53 won all three pairs on every row. It also beat the Phase 52 goal of bringing the five components below 14 seconds.

The difference between the rows matters. "Packet preparation" is the host work most directly changed. "Destination service" adds formula planning, GPU utility work, nesting, and handoff. "Five components" also includes destination sampling, pandas work, ActivitySim orchestration, and final choice. The speedup naturally gets smaller as unrelated work is included.

251. Are the travel choices still the same?

Yes. All verified school zones, workplace zones, and tour destinations are exactly the same as the reference. No modeled destination decision changed.

Some stored logsum diagnostics differ by a few millionths because of the already qualified GPU floating-point policy. The largest observed difference was 0.000001907. The school/workplace limit is 0.00001, and the tour-destination limit is 0.0001. The observed difference is more than five times below the tighter limit and more than fifty times below the tour limit.

The honest replication statement is therefore: **exact destination decisions, bounded logsum diagnostics, and zero fallback**. It is not a claim that every printed decimal character is identical.

252. Why was the qualification split into two pieces?

This Windows computer repeatedly refused to allocate a 79-megabyte pandas array during unchanged mandatory tour scheduling. At the time, Windows reported more than 24 gigabytes of virtual memory available. The same failure occurred after the seven new school/workplace calls and before the later tour destination calls, so it was not caused inside the Phase 53 code.

Instead of hiding the problem or reporting a partial run as complete, the qualification used two deterministic pieces:

- 1 initialization through workplace location, covering seven destination calls;
- 2 a resume from a previously verified mandatory-scheduling checkpoint through at-work destination, covering the other twelve calls.

The two pieces together reproduce exactly 19 calls, 201,390 owners, and 4,696,676 sampled rows. Six formal shards - two per matched pair - passed all runtime and output gates. A separate resumed downstream run finished the rest of the model with zero changed modeled decisions.

This proves the measured destination improvement. It does **not** turn the two pieces into a measured monolithic whole-model time.

253. What is the current whole-model comparison?

The latest successful monolithic Phase 52 median is 136.991 seconds. Replacing only its five measured destination-component times with the Phase 53 measured times gives a conservative projection of 133.391 seconds.

Regular ActivitySim's measured median is 205.4 seconds. The Phase 53 projection is therefore about **1.540x faster**, saving about 72 seconds or 35.06%.

The word "projection" is essential. The Phase 52 and regular ActivitySim times are measured complete runs. The 133.391-second Phase 53 number is a component-substitution estimate necessitated by the unrelated Windows allocation failure. It should become a measured number after the host memory problem is fixed or the benchmark is repeated on a clean machine.

254. Did Phase 53 make the GPU kernel itself faster?

Not directly. Phase 52 had already made the four-row CUDA kernel fast. Phase 53 made the road into that kernel much shorter.

This is an important systems lesson: a fast engine does not make a car fast if workers spend most of the trip unloading and repacking the luggage. GPU projects often stall because data is copied, expanded, converted, and checked in CPU-shaped tables between kernels. Phase 53 removes a large repeated table before it exists, so the proven kernel receives useful compact data sooner.

The CUDA formula, four-row tile, random-stream rules, nested logit, and final choice arithmetic stay the same. The same checked-in source fingerprint is recorded in every call:

599a9704be0992d2863320390cbca0028c7a578ecacf72d69de2e658a5d79906

255. What assumptions and limits should a reviewer remember?

The proof applies to the public Prototype MTC Extended model on this machine, with 50,000 households, 1,454 zones, the reviewed 315-term and 21-mode program, five time periods, current skim directions, unchunked destination calls, and the current ActivitySim person/tour relationships.

Another model might name fields differently, use noncontiguous samples, add a new zone layer, or need another time system. It must qualify those contracts before enabling this route. Three matched pairs reduce normal timing noise but do not prove the exact speed on every computer or GPU.

Most importantly, Phase 53 is not a “GPU-only whole model.” ActivitySim still controls workflow, tables, many model components, and the authoritative random stream. Phase 53 makes the destination path substantially more resident and compact while preserving compatibility and proof.

256. What should the next ambitious phase do?

Phase 54 should make the destination packet genuinely device-owned from sampling onward.

The sampling service already knows owner IDs, group offsets, and sampled destination IDs. Instead of rebuilding those vectors through pandas and NumPy, it should publish versioned GPU buffer leases directly to the logsum service. The GPU should also transform each owner’s six controlled normal draws into taxi and ride-hail wait times using a numerically qualified CUDA routine and resident land-use density bands.

The target should be ambitious and testable:

- all 19 public calls and 4,696,676 sample rows;
- exact destination decisions and the existing logsum limits;
- zero hidden CPU or generic fallback;
- three matched wins over Phase 53;
- destination service below four seconds;
- a successful measured monolithic complete run once the host allocation issue is removed.

That phase would move the project from “compact before the kernel” to “retain the destination packet on the device across sampling, logsum, probability, and choice.”

257. What did Phase 54 build?

Phase 54 built the device-owned destination packet proposed above. When the GPU sampler finishes making a destination list, it now gives the next GPU stage a short-lived, numbered claim ticket called a **lease**. The ticket points to the exact group boundaries and destination IDs already in GPU memory. The next stage uses them directly instead of copying the list back through CPU tables and rebuilding it.

The GPU also makes the six controlled bell-curve random numbers needed for each person or tour and converts them into expected taxi and ride-hail waiting times. Sampling, random inputs, wait construction, destination utility, probability, and choice are now connected much more closely.

258. What is a buffer lease, in ordinary language?

Think of checked luggage being transferred between two flights. The old system unloaded every bag, wrote a new list on paper, and then loaded the same bags again. A buffer lease is a secure electronic transfer slip: “bags 1 through 25 belong to traveler A, these are their exact storage locations, and this slip is valid only for transfer number 17.”

The consumer checks the transfer number, source object, row count, owner count, index, column name, and owner sequence. A stale or mismatched ticket causes the run to stop. The lease is deliberately short-lived and synchronous; it is not permission to keep using memory after another task changes it.

259. How can GPU random numbers match old CPU random numbers?

A pseudo-random generator is a deterministic recipe. Give it the same seed, advance it the same number of steps, and use the same conversion recipe, and it produces the same bit pattern.

ActivitySim relies on NumPy's legacy RandomState behavior. Phase 54 implements its MT19937 state update and polar Box-Muller normal conversion on CUDA. Tests try several seeds, offsets, and output lengths and require every generated 64-bit normal value to match NumPy exactly, not merely approximately. The implementation follows NumPy's published legacy source and compatibility policy: [legacy distribution source](#) and [random compatibility policy](#).

This does not mean every possible random API is supported. It means the exact bounded legacy stream used by this public model is implemented and checked.

260. What are the wait tables and why generate them on the GPU?

The travel formula needs an estimate of how long a taxi or transportation- network-company vehicle might take to arrive. The estimate depends on a controlled normal draw and the destination's density band. Previously the CPU made a table of those values and uploaded it.

Phase 54 sends the compact ingredients to a CUDA kernel. The kernel applies the same float32 rules and creates the owner-by-density-band wait table where the destination calculation will use it. Tests compare the entire table with the host reference exactly.

261. How much faster is Phase 54 itself?

Three matched Phase 53/54 comparisons cover the same 19 calls, 201,390 owners, and 4,696,676 sampled destinations:

Measured boundary	Phase 53	Phase 54	Improvement
packet preparation	1.300 s	0.306 s	4.252x; 76.48% lower
complete destination service	4.989 s	3.713 s	1.344x; 25.58% lower
five destination components	12.6 s	11.3 s	1.115x; 10.32% lower

Phase 54 won every comparison at every boundary. The shrinking speedup is normal: each lower row includes more unchanged ActivitySim work outside the new packet path.

262. Did the full model finally run, and how fast is it?

Yes. The earlier failure was traced to running many host math workers during a large unchanged pandas allocation. The GPU benchmark now limits OpenBLAS, OpenMP, MKL, and Numba host work to one thread. Three fresh complete runs took 134.267, 133.264, and 133.714 seconds, including startup bookkeeping. The median is **133.714 seconds**.

The separately measured regular ActivitySim CPU median is 205.4 seconds. The latest system is therefore **1.536x faster**, saves about 71.7 seconds, and uses **34.90% less total time**. It is also 2.39% lower than the measured 136.991-second Phase 52 median.

The incremental destination result is the strongest experiment because Phase 53 and Phase 54 are directly matched. The CPU and Phase 52 figures are older measured medians on the same machine and workload, not interleaved pairs with these three full runs.

263. Are accuracy and reproducibility still protected?

Yes. All three complete runs pass every proof gate. Every published school, workplace, and tour destination stays exactly the same. Every controlled normal stream is generated on CUDA, every lease is both published and consumed, every wait table is made on CUDA, and no generic or CPU fallback is used.

The largest observed cached floating diagnostic difference is 0.000005722. It remains below the declared limits. The honest statement is still **exact modeled decisions with bounded floating diagnostics**, not that two different processors must print every intermediate decimal identically.

264. What did the speed cost in memory and complexity?

The persistent GPU service now uses at most 462,283,900 bytes, about 441 MiB. It keeps separate random state for sampling and normal generation plus reusable packet workspaces. That is modest on the tested 16 GiB GPU, but it is not free and smaller GPUs must be qualified.

The lease is safe under the current immediate producer-to-consumer workflow. It is not a permanent cache, and code must not mutate the sample between publishing and consuming it. Different sample layouts, random APIs, expression programs, or models need their own checks.

265. What should Phase 55 do?

The next ambitious step is a versioned **device entity store** plus ahead-of-time linked execution plans. Today, Phase 54 still assembles some one-row-per-person or one-row-per-tour facts on the CPU and Python still plans individual calls. Phase 55 should keep authoritative person, tour, household, and land-use columns on the GPU across destination families and launch a hash-addressed packet-to-choice graph without rebuilding Python objects.

The success bar should be meaningful: exact decisions, existing diagnostic limits, zero fallback, three matched wins, destination service below 2.5 seconds, five destination components below 10 seconds, and a measured complete model median below 130 seconds. That would attack orchestration and remaining owner-state movement, not merely tune an already fast inner kernel.

266. What did Phase 55 actually build?

Phase 55 built two connected pieces.

First, it made an ahead-of-time plan book for the ten destination calculations used by this public model. Earlier versions opened ActivitySim expression files, interpreted their meaning, worked out every input, and generated GPU code when a run started. Phase 55 does that planning once in a special developer capture run. A normal production run reads the reviewed plan book and goes straight to the already-defined GPU inputs.

Second, it made a CUDA sample-compaction step. After the GPU draws 30 possible destinations for each traveler, many draws may repeat the same place. The new kernel counts repeats, orders the unique places exactly as ActivitySim expects, and lets the next GPU stage reuse the compact destination buffer through the lease system built in Phase 54.

267. What does “ahead of time” or AOT mean?

Imagine a school play. A just-in-time system reads the script, assigns roles, and plans every scene after the audience arrives. An ahead-of-time system does that preparation during rehearsal and brings a checked stage plan to the performance.

Here, the “script” is a travel-model expression such as “driving time times a coefficient, plus parking cost, plus a household effect.” The AOT plan records which columns, constants, coefficients, and travel-time arrays the GPU needs. This reduces repeated Python planning. It does not change the mathematics.

268. What are the plan atlas and a cubin?

The **plan atlas** is a typed JSON file containing ten reviewed destination plans. “Typed” means it remembers whether a value was a list, tuple, map, integer, or decimal instead of guessing when it is loaded.

A **cubin** is compiled GPU machine code for a particular NVIDIA GPU architecture. It is similar to giving a processor an executable file instead of source code. This project builds an `sm_86` cubin for the RTX A4000 from the same checked CUDA source used since Phase 52.

The cubin is faster to start and easier to review, but it is less portable. A different GPU architecture needs its own compiled and tested binary. Phase 55 stops with an error on the wrong architecture rather than secretly compiling a different program.

269. How can we trust files made ahead of time?

Phase 55 uses SHA-256 fingerprints. A fingerprint is a long number calculated from all bytes in a file. Even a tiny edit produces a different number.

The production runtime checks the atlas fingerprint, each relevant model configuration fingerprint, the CUDA source fingerprint, the cubin fingerprint, the GPU architecture, the kernel name, and the number of plans. A missing plan or mismatch stops the run. Production cannot silently enter capture mode or fall back to a generic CPU calculation.

This is a replication guarantee about the reviewed program. It does not prove that every future travel model is supported. A changed model must build and qualify a new atlas.

270. What was learned from the first failed design?

The first atlas stored nearly the entire expression document for every plan. It was about 30 MB and made the supposedly faster startup slower. That version was rejected.

The second design stores only what execution needs: bindings, coefficients, scalar values, source order, alternative count, and proof metadata. It is about 3.4 MB. The loader also originally requested the same travel-time cube many times when several expressions referred to it. The final version resolves each immutable cube once and shares it across plans.

This is why benchmarking matters. A design can sound efficient and still lose when measured.

271. What were the final speed results?

The final experiment ran three fresh Phase 54 controls and three fresh Phase 55 candidates. Every candidate ran the complete 34-step public ActivitySim model for 50,000 households.

Measured boundary	Phase 54 median	Phase 55 median	Improvement
cold two-process destination shards	3.713 s	2.737 s	1.356x; 26.3% lower
monolithic 19-call Phase 55 service	not reconstructed	2.249 s	below 2.5 s target
five destination components	10.7 s	10.1 s	1.059x; 5.61% lower
complete 34-step model	138.2 s	135.909 s	1.017x; 1.66% lower

Phase 55 won the destination-component total in all three full pairs. It also won complete model time in all three pairs, saving 3.791, 1.291, and 2.691 seconds. The improvement becomes smaller as more unchanged CPU work is added, which is expected.

The five-component median missed the under-10-second stretch target by 0.1 seconds. The full model missed the under-130-second stretch target by 5.909 seconds. Those goals remain open even though the replicated comparison passed.

272. Did the new compaction kernel make sampling faster?

No, not by itself. This negative result is important.

The Phase 54 sampler median was 2.182 seconds. The final Phase 55 sampler median was 2.584 seconds. ActivitySim still demands a pandas DataFrame on the CPU, so the compact results must still be copied to host memory and arranged as a table. GPU compaction also adds its own sorting work.

Why keep it? The compact GPU buffer can be handed directly to the next GPU stage without uploading the destinations again. It is a bridge toward a truly device-native workflow. In Phase 55, the AOT savings and downstream reuse are larger than the local sampler cost, so the five-component and complete-model totals still improve. It would be misleading to call the compaction kernel alone a speedup.

273. How accurate and repeatable are the results?

All three final full candidates passed every inherited and Phase 55 proof gate. Each had exactly 19 atlas hits, 19 expression-code-generation bypasses, 19 CUDA sample compactions, 19 lease transfers, and zero fallback calls.

Independent verification compared the published model outputs with the fixed public reference. All modeled decision cells were identical. Floating logsum diagnostics stayed inside limits declared before the run. As before, the honest claim is **exact decisions and bounded floating diagnostics**, not identical text for every intermediate decimal.

The final-code shard experiment also ran six separate processes and passed all gates. This protects against accidentally timing a warm Python process as if it were a fresh run.

274. What assumptions limit the claim?

The proof applies to the tested RTX A4000, its sm_86 instruction set, this software environment, 50,000 public benchmark households, 1,454 zones, ten known destination plans, 21 mode alternatives, 315 expression terms, five time periods, and the current ActivitySim ordering and random-stream rules.

Phase 55 is not a whole model running only on the GPU. ActivitySim still uses the CPU for workflow, pandas tables, initialization, reporting, and many model steps. The reported 1.017x is an incremental gain over Phase 54, which was

already heavily GPU accelerated. It should not be confused with the larger GPU-versus-regular-ActivitySim comparison established by earlier phases.

275. Why is a 1.017x whole-model gain still useful?

It is small, but it is measured against a strong GPU baseline rather than an unoptimized CPU prototype. It repeats in all three pairs, keeps every modeled decision, and removes runtime compiler work. That makes production behavior more predictable as well as faster.

It also identifies the new limiting factors. The destination service itself is now only about 2.25 seconds. Most of the roughly 136-second model is elsewhere. Making that service infinitely fast could save only a few more seconds. Large future gains require removing CPU-shaped tables and accelerating other major components, not polishing the same inner kernel forever.

276. What should Phase 56 do?

Phase 56 should replace the host sample-table compatibility boundary with a general device-native entity and sample store. Person, household, tour, destination, probability, repeat-count, and random-state fields should remain in versioned GPU columns from sampling through logsums, probability, and final choice. ActivitySim should receive a host table only when a public model output must be published.

The next proof should require:

- a versioned schema that supports more than these ten hard-coded plans;
- reproducible binary bundles for each supported GPU architecture;
- a sampler that beats Phase 54 in all three matched pairs;
- five destination components below 10 seconds median;
- a full-model median below 130 seconds;
- exact decisions, bounded diagnostics, zero fallback, and three matched wins.

That would address the real lesson of Phase 55: the next large gain is no longer inside the destination formula. It is in removing the CPU table boundary around the formula.

277. What did Phase 56 actually choose to do?

It changed priorities after measuring the whole program again.

The earlier plan was to build an even broader GPU entity store. That is still a useful idea. But the new profile found a larger and easier-to-prove delay: every fresh run rebuilt the same giant travel-time data structure. That setup took about 14 to 19 seconds. Meanwhile, the destination service improved through Phases 45-55 was already only about 2.25 seconds.

Even a magically instant destination service could not save as much time as removing the repeated setup. Phase 56 therefore solved the measured largest removable problem first.

278. What is a skim, and why is it so large?

A **skim** is a travel lookup table. Give it an origin zone, a destination zone, a time of day, and a travel measure, and it can return something such as driving time, transit time, distance, or cost.

This benchmark has 1,454 zones, many travel measures, and five time periods. The computer rearranges the compact source file into arrays that formulas can read quickly. The resulting Sharrow memory image is 6,452,305,336 bytes, about 6.45 GB. Rebuilding it is like unpacking the same large reference book before every exam.

279. What is a memory-mapped file?

A **memory-mapped file**, or `mmap`, is a file that a program can use as if it were a large array in memory. The operating system loads the needed pages when the program touches them. Another fresh process can reopen the already arranged data instead of decoding and copying the source again.

Phase 56 builds the 6.45 GB image once. Later model runs attach it in about 0.6 seconds. The file stays on this machine and is not placed in GitHub because it is large and can be reproduced from the public inputs.

280. Does reusing a cache risk using old or wrong data?

Yes, unless the cache has a strict identity check. A fast wrong answer would be a failure.

Phase 56 calculates SHA-256 fingerprints for the complete 6.45 GB cache, its layout description, the original `skirts.omx`, `land_use.csv`, and `network_los.yaml`. It also fingerprints a manifest that records those facts. Before a normal run, it rechecks the sources, metadata, file size, modification time, contract version, and manifest fingerprint. Any mismatch stops the run.

The slow complete cache fingerprint is checked when the artifact is built and formally qualified. Normal runs use strong source checks plus cache identity facts so they do not spend several seconds rereading all 6.45 GB. This rule is written down and tested; it is not an informal promise that "the file looks right."

281. What bug did Phase 56 find in Sharrow?

Sharrow already supported memory-mapped data, but its version in this public environment had a small ownership bug. After writing the large NumPy array, it asked Python whether the whole array was "true" or "false." An array with billions of values has no single sensible yes-or-no answer, so NumPy stopped with an error.

The project added a narrow bridge. When, and only when, that argument is a NumPy `mmap`, the bridge supplies Sharrow's documented ownership flag. Sharrow then reopens the same file and continues its normal reconstruction. The bridge is removed after the run, and a unit test proves the special case. This is a compatibility repair, not a change to travel-model mathematics.

282. How fast is Phase 56?

Three independent Phase 55 controls were each followed by a fresh Phase 56 candidate:

Trial	Phase 55	Phase 56	Time saved	Speedup
1	127.100 s	118.836 s	8.264 s	1.070x
2	128.300 s	119.039 s	9.261 s	1.078x
3	128.400 s	118.137 s	10.263 s	1.087x
median	128.300 s	118.836 s	9.464 s	1.080x

The median complete-model time is 7.38% lower. Initialization alone falls from 14.1 to 0.6 seconds, a 23.5x speedup. Every Phase 56 candidate finishes below 120 seconds, so the old under-130-second goal is cleared three times.

For the larger project picture, regular pinned ActivitySim was separately measured at a 205.4-second three-run median on this machine and workload. The latest 118.836-second result is therefore a cumulative **1.728x speedup**, 42.14% less time, and 86.564 seconds saved. That older CPU experiment was not interleaved with Phase 56, so use the Phase 55/56 table to describe what Phase 56 alone contributed.

283. Is the comparison fair?

The Phase 56 number includes about 0.528 seconds of cache validation and the inherited Phase 55 prewarm cost. Those costs are not hidden before the timer. The one-time cache construction is reported separately because it is not repeated for every model run. Building ActivitySim through initialization takes about 23 seconds; hashing and fully rechecking the new artifact takes about 4.7 more seconds.

All runs use the same public 50,000-household, 1,454-zone model and the same RTX A4000. Each candidate is a fresh process. Each output is compared with the fixed public reference before large temporary results are removed.

284. Did speed change the answers?

No published modeled decision changed in any of the three candidates. Every inherited GPU proof gate also remains active: exact controlled random streams, bounded floating diagnostics, the reviewed Phase 55 plan atlas and cubin, device sample leases, and zero silent fallback.

The final qualification has nine additional gates. They prove the cache build, manifest fingerprint, shared artifact identity, three runtime validations, three sub-three-second initializations, three exact comparisons, three complete model wins, and a median below 130 seconds. All nine pass.

285. Is Phase 56 a GPU kernel speedup?

Not by itself. It is a whole-system speedup around an already GPU-accelerated model. The CUDA kernels from earlier phases still do the accelerated travel calculations, but Phase 56 stops the CPU process from rebuilding invariant travel data before those kernels can work.

This distinction matters. The honest claim is **Phase 56 makes the complete GPU-enabled model 1.080x faster than Phase 55 on this benchmark**. It would be wrong to say that a new arithmetic kernel became 1.080x faster. Good performance engineering improves the measured bottleneck, even when the bottleneck is data preparation rather than multiplication.

286. Does Phase 56 make Phases 1-55 unnecessary?

No. Phase 56 reuses data faster, but it does not replace the GPU execution, exact random-number rules, expression compiler, native buffers, destination service, output verifier, or replication harness built earlier. Without those pieces, attaching skims quickly would only start a slower CPU-shaped model a few seconds sooner.

The phases answer different questions. Earlier phases prove that important travel decisions can run on the GPU without changing published choices. Phase 56 proves that the complete accelerated system can launch efficiently and repeatably.

287. What should the next major phase do?

Phase 57 should aim for a complete lifecycle below 105 seconds. It should use the Phase 56 verified launch path and attack whole remaining components, not shave microseconds from the already small destination kernel.

The largest median components now include mandatory tour scheduling at about 14 seconds, trip mode choice and trip destination near 10 seconds each, and several 6-7-second frequency, scheduling, and output steps. The next system should keep a general versioned person-household-tour entity store on the GPU, fuse one complete component from input columns through final choice, and create a CPU table only at the required publication boundary.

Success should again mean exact decisions, bounded diagnostics, no silent fallback, three fresh matched wins, and a large whole-model improvement. The ambitious target is not "another phase exists." It is roughly 14 more seconds removed from a model that now takes about 119 seconds.

288. What changed in Phase 57?

Phase 57 changes how the model prepares possible departure and return times. Previously, the CPU created more than 15 million rows: one row for each tour and available time choice. It then repeatedly merged and shortened these tables before asking the GPU to calculate travel attractiveness.

The new kernel reads the current timetable directly. Thousands of GPU workers check time conflicts at once and keep the first available choice for each departure-period and return-period pair. Only this compact list returns to the CPU. The model then calculates all six batches of travel-attractiveness scores again using the current tours, current skims, and controlled random numbers.

This removes 15,242,743 large host-table rows and supplies 1,210,124 smaller representative rows. It does not reduce the set of actual scheduling choices. The scheduling kernel still considers the available hourly alternatives and updates each person's timetable after choosing a time.

289. Why is taking one representative time legitimate?

The public MTC model already groups hours into five travel periods. For example, several morning departure times may use the same morning travel conditions. Its configuration chooses representative times for these period pairs when computing logsums. Phase 57 implements that existing rule more directly.

It preserves the order of the first available alternative in every pair. Order matters because the model's controlled random draws and downstream tables must remain aligned. Integer minimum and integer counting operations make this new reduction independent of which GPU worker happens to finish first. They introduce no floating-point summation approximation.

The numerical utility and logsum calculations still use the earlier reviewed GPU arithmetic. The inherited near-boundary choice checks remain necessary; Phase 57 does not remove the project's existing numerical limitations.

290. What happened to the faster-looking cache experiment?

An early experiment saved old logsum values, trip matrices, and summary tables, then reused them when input fingerprints matched. Its checked decision columns were unchanged and one run took about 108.48 seconds. That is evidence about repeating an unchanged scenario, not evidence that the GPU calculates a new scenario faster. Its hash checks also did not establish a complete model-wide dependency or semantic equivalence proof.

We therefore removed that result-replay implementation from the release candidate. The released Phase 57 path recalculates all six logsum batches and regenerates the trip matrices and summary tables. Its formal benchmark uses this live-computation path. Early cache measurements remain development evidence and must not be presented as Phase 57 GPU-kernel speedups.

Reusing invariant input skims, as Phase 56 does, remains part of the baseline and candidate. Reusing input data and replaying a model's calculated answers are different operations and must be described separately.

291. How do we test whether it really works?

First, an independent CPU loop checks every combination of timetable conflict codes on randomized inputs. Tests change the timetable to empty and occupied states and require identical GPU results. Another test changes representative times in the configuration and checks that the new calculation responds.

Second, the full public model runs from a fresh process. It must execute all 34 timed steps, use six live GPU logsum computations, preserve the complete representative-row count, and match every checked published decision column. The separately declared logsum diagnostics retain their numerical tolerances.

Third, fresh Phase 56 controls and Phase 57 candidates run in three matched pairs. Both sides include validation of their input-skim cache and program prewarm. The comparison also records launch-to-exit wall time. The commonly reported sum of model-step times is smaller than process wall time because imports, startup, and other orchestration also take time.

Finally, a separate benchmark compares this reduction against compiled CPU code with several thread counts. It includes uploads, downloads, and allocation in the GPU measurement. That benchmark answers a kernel question; the full ActivitySim comparison answers the end-to-end question.

292. What does this let us claim?

The single development run of the live path reduced mandatory scheduling from 13.9 to 7.2 seconds and the measured model total from 118.34 to 111.73 seconds. Every checked decision cell matched. These are preliminary single-pair figures; the final replication record below is the authority for a repeated claim.

The new kernel uses live inputs, but the complete application still contains benchmark-specific scheduling artifacts and sparse reference decisions inherited from earlier phases. It is a qualified prototype for this 50,000-household public workload, not a general GPU implementation of every ActivitySim model.

The target remains a measured total below 105 seconds. A positive speedup does not automatically meet that target. A report must state both the improvement we measured and any remaining gap. Faster kernels also do not prove that the underlying travel model predicts real people more accurately; they preserve the model's calculations while taking less time.

293. What did the finished replication actually show?

We ran three complete Phase 56 controls and three complete Phase 57 candidates. Each candidate produced exactly the same checked travel decisions as the reference. All correctness gates passed, and all 238 automated repository tests passed. Some attractiveness-score diagnostics still differ slightly within their declared tolerances; "same decisions" does not mean every number is bit-for-bit identical.

What we measured	Previous GPU system	New GPU system
Whole-model median, including validation and prewarm	142.08 seconds	128.16 seconds
Launch-to-exit wall-clock median	151.23 seconds	138.26 seconds
Mandatory tour scheduling median	16.5 seconds	8.0 seconds

The observed whole-model improvement is 1.109x, or 9.80% less time. The targeted scheduling step takes 51.52% less time. All three pairs improved, with whole-run savings of 11.69, 19.65, and 12.72 seconds.

The final runs were slower overall than the earlier development test. Other, unchanged steps also varied, and each pair ran the old version first. That means we cannot assign all 13.92 seconds of median savings to the new kernel alone. Alternating which version runs first is an important follow-up. These are honest measurements on this computer, not a promise that another computer or every new run will have the same times.

The under-105-second goal was **not met**. The replicated median remains 23.16 seconds above it. The machine-readable result explicitly says `replicated_improvement_target_not_met`. A working optimization and an unfinished larger target can both be true.

294. Did the GPU beat a strong CPU implementation too?

Yes, for this particular reduction. We wrote compiled CPU code for the same job and tried 1, 4, 8, 24, and 48 CPU threads. A thread is a worker that can carry out part of a program. More workers can help, but their coordination also costs time, so we measured rather than assumed which count was best.

Across seven trials, the fastest tested CPU configuration used 48 threads and took a median of 21.73 milliseconds for all six batches. The GPU took 8.25 milliseconds, including copying inputs in and results out and allocating memory. The GPU was **2.63x faster** and returned exactly the same integer answers in every trial. A millisecond is one thousandth of a second.

These measurements reuse already initialized compiled functions. First use also takes time: about 1.32 seconds for the CPU and 0.59 seconds for the GPU in that process, with existing disk compiler caches possibly reused. The complete model measurements include their own initialization costs.

Here is the crucial interpretation: the GPU's advantage over this compact CPU calculation saves only about 13.5 milliseconds. Avoiding the old giant tables is what saves whole seconds in ActivitySim. A good CPU rewrite could also avoid those tables. We have demonstrated both a better data layout and a local GPU advantage, but not that the GPU alone caused every second of the full-model improvement or that no better CPU algorithm could ever exist.

295. What remains before the bigger ambition is fulfilled?

First, make the experiment stronger: alternate run order and repeat a fresh regular-CPU ActivitySim comparison. The old 205.4-second CPU result is useful history, but it is not a newly matched comparison with these slower formal runs. Do not divide unrelated historical timings and call that a new proof.

Second, tackle complete remaining steps. Trip mode choice takes about 11.3 seconds, trip destination about 10.7 seconds, and non-mandatory tour frequency about 8.9 seconds in the latest candidate medians. Keeping person, tour, and trip information in a shared GPU store could remove repeated copying and table construction. We must measure those costs before promising the savings.

Third, broaden correctness beyond this fixed benchmark. The full prototype still relies on earlier benchmark-specific scheduling contracts and a small set of reference decisions near numerical boundaries. Those dependencies need live replacements and tests with changed inputs and seeds. The new kernel's changed-timetable tests do not by themselves prove that the entire model works for every city or scenario.

The next success is not another impressive tiny-kernel ratio. It is a lower complete-run time, including all required work, with repeatable travel choices and a clearer path for another researcher to reproduce and extend it.

296. What is Phase 58 trying to change?

Imagine planning an afternoon: leave school, visit a shop, then return home. You cannot arrive home before leaving the shop. A travel model has to respect these dependencies for hundreds of thousands of trips. A **chain** is simply this ordered set of related trips.

Earlier versions often sent a trip's result back to the CPU before computing the next dependent trip. Phase 58 lets one GPU worker finish the short chain for one person-tour, while many other workers handle other tours. Intermediate departure times stay on the GPU. The CPU receives the completed iteration's answers and decides which failed trip groups need another attempt.

This is not the same as making the entire model GPU-only. The CPU still handles tables, some random-number arithmetic, failed-group retries, and final reports. We have moved a larger connected calculation onto the GPU, not eliminated the CPU.

A scheduling "failure" means a draw did not find a supported departure time inside the permitted window, not that a simulated vehicle crashed. ActivitySim retries the affected trip group and, on its final attempt, applies its existing correction rule. The new kernel preserves that rule. In the first formal pair, both versions consumed 210,110 scheduling draws and recorded the same 66,899 failed attempts, including repeated attempts for some trips. The new version needed 100 iteration calls instead of 548 smaller scheduling calls.

297. How can a random model reproduce exactly the same trips?

The model does not use uncontrolled randomness. Each trip has an identity and a reproducible stream of numbers. Think of a numbered deck of cards assigned to each trip. Its **seed** chooses the deck; its **offset** says which card to draw next. Rearranging the trips in a table must not change their decks.

The new shared service generates uniform random numbers on the GPU, using the same algorithm as ActivitySim. Uniform means every equal-sized interval between zero and one has the same probability. The CPU still owns the official record of seeds and offsets. The GPU reads that record and advances it exactly once for each required draw. Later parts of the model must see the same next card.

An **epoch** means the currently active model step. A deck from destination choice cannot silently be reused as if it belonged to mode choice. The service checks the active step and restores its temporary hooks when the step finishes. Tests change seeds, shuffle trip order, select subsets, and compare later draws, not just the first answer. This catches errors that otherwise appear much later.

298. What changes when choosing between car, transit, walking and other modes?

A **utility** is the model's numerical score for an option. A **probability** turns those scores into chances. Related options are grouped in **nests**: two similar car options are more alike than driving and walking. The model combines probabilities through this family tree, then uses a random draw to pick a mode.

Before Phase 58, GPU-generated scores came back to the CPU for much of that probability and choice work. Now the scores feed the next GPU calculation directly. Only final mode labels and small diagnostic results need publication. For the public development run, this avoids downloading a 37.2-megabyte table of mode scores. It also avoids constructing the large CPU probability table.

Floating-point numbers are rounded approximations. Two processors can disagree by a tiny amount, especially after logarithms and exponentials. Usually such a difference does not change a choice. Near the dividing line between two choices, however, it might. The new selector flags draws within one billionth of a probability boundary and checks those rows with the actual upstream CPU calculation, using the same already-drawn number. It never draws a second card or looks up a saved benchmark answer.

That one-billionth guard is a tested engineering precaution, not a mathematical proof covering every imaginable input. Full runs must still pass an independent comparison of all published travel decisions. Diagnostic scores have explicit small tolerances; exact decisions do not mean every decimal is identical.

299. What did the stronger CPU comparison teach us?

System	What it tells us
Regular CPU ActivitySim	How the whole accelerated project compares with the original application
Accelerated stack with compact CPU scheduling	Whether the improved compact layout also works well on the CPU
Accelerated stack with compact GPU scheduling	What using the GPU adds at that same small boundary

In two new runs, regular CPU ActivitySim took 300.2 and 299.8 seconds for its model steps. The other two systems both took about 111.4 seconds, including their charged validation and warm-up work. Compact CPU and GPU scheduling were effectively tied for the whole application. The earlier small GPU kernel really was faster, but its advantage saved milliseconds rather than many seconds.

The lesson is not that GPUs are unhelpful. It is that avoiding giant tables is an algorithm improvement that can help both processors. We must measure larger connected GPU work to demonstrate a meaningful additional application benefit. The full accelerated stack contains many other GPU changes; its total advantage cannot be assigned to this one scheduling kernel.

The new regular CPU timing is slower than the old 205.4-second historical result. We report the actual new observations instead of choosing whichever old number makes a preferred story. Computer load and caches can change absolute times. The regular CPU run uses the pinned single-process configuration with one numerical worker. We have not shown that it is the fastest possible CPU setup. The separate 48-worker CPU control changes only the small feasibility kernel.

300. Which tempting shortcut did we reject, and why?

Not all random numbers have a uniform distribution. A **normal distribution** is the familiar bell-shaped curve, where values near the middle are more common. The model uses such numbers while constructing some travel-time inputs.

We tried extending an older GPU normal-number generator to all new trip random calls. More demanding tests found 12 differences among 3,168 values. These were tiny rounding differences, sometimes just one bit, even though the next-card positions were correct. The older generator had matched its earlier fixed benchmark, but that did not prove identical arithmetic for every seed and offset.

We rejected the broader shortcut and kept these ordinary normal draws on the CPU. The audit and a script to reproduce it are saved. We did not loosen the tests or call the differences harmless without checking how they would be used. There is also a tested shortcut for a truly zero-variance expression, where the random part cannot affect the answer; the full development run did not use it, so it gets no credit for the speed improvement.

This is an important distinction: a failed extra optimization does not erase the separately verified improvements. It tells us exactly where a stronger arithmetic implementation or proof is still needed.

301. How is the final speed experiment made harder to fool?

We run the previous GPU version and the new version in fresh processes. For one pair, the old version goes first; for the next, the new version goes first. Six pairs therefore give each version three turns in each position. This reduces the risk that whichever runs second always benefits from a warmed-up computer.

The source files receive digital fingerprints, called **hashes**. If code changes during or between matched runs, the qualification rejects the timing series. Every run must complete all 34 steps and pass the independent output checks. We keep two clocks: model work plus required validation/prewarm, and the entire process from launch to exit. Both must improve in every qualified pair.

Repeated processes can still reuse files and previously compiled code, so these are not completely cold-computer measurements. No other heavy tests run alongside them. Component times are recorded in tenths of a second, so tiny differences in unchanged steps are not evidence that those steps were optimized.

The six runs use the same public population and seed. They test repeatability of speed, not six different cities. Changed-input and changed-seed unit tests cover the new kernels; the earlier benchmark-specific mandatory-scheduling artifacts still prevent a claim of unrestricted whole-model scenario support.

The 105-second target refers to the charged model-work clock. It does not mean the entire application must disappear from the screen within 105 seconds. Launch-to-exit wall time includes additional startup and bookkeeping and is reported separately. We also perform a final comparison of saved output files after the timed process finishes; that comparison is an experiment check, not a travel-model calculation hidden inside one version's timing.

302. What did the finished Phase 58 experiment achieve?

All six matched pairs improved both clocks. Every candidate preserved every checked published travel decision, and all 264 automated repository tests pass. The checker reports `replicated_improvement_target_met`: the charged model median is below 105 seconds, without waiving correctness or repetition.

What we measured	Previous GPU version	Phase 58
Charged complete-model median	110.97 seconds	103.81 seconds
Launch-to-exit median	118.27 seconds	111.02 seconds
Trip destination median	9.75 seconds	9.00 seconds
Trip scheduling median	6.20 seconds	3.65 seconds
Trip mode choice median	9.85 seconds	6.05 seconds

The full model takes 6.45% less charged time, saving about 7.16 seconds. Adding the three targeted component medians gives 25.80 versus 18.70 seconds, a 27.52% reduction. Nearly all the observed full-model saving is therefore visible in the components we actually changed. The six candidate totals range from 103.23 to 104.03 seconds. The smallest matched saving is 6.20 seconds; the largest is 8.19 seconds. There is no losing pair hidden by the average.

Against the newly measured regular CPU configuration, charged model time falls from 300.00 to 103.81 seconds: 2.89x as fast. Actual launch-to-exit time falls from 306.36 to 111.02 seconds: 2.76x as fast. The CPU controls were measured before the six old/new GPU pairs, not inside every pair. Those cumulative comparisons and the stronger paired incremental comparison are distinct evidence.

Some unchanged steps are not faster, and a few are slower than regular CPU ActivitySim. The technical report includes all 34 component rows, not just the winners. The complete application benefits from many earlier GPU kernels, better data layouts, and cached immutable inputs. We do not attribute every saved second to GPU hardware alone.

303. What remains after this successful phase?

We now have a repeatedly faster live trip runtime on this public benchmark. CPU retries, table publication and older benchmark-specific scheduling artifacts remain.

Next is a **versioned data store**: one shared home for person, tour and trip facts, with rules for marking copies outdated when a fact changes. Correct updates must be tested before removing further CPU/GPU handoffs. Ordinary normal arithmetic also needs a stronger shared implementation.

Broader support requires removing inherited benchmark dependencies and testing changed populations and seeds through the entire model. Today's checks do not prove a replacement for every ActivitySim city or configuration. Faster execution preserves the checked computation; it does not make the behavioral model more accurate about real people.

304. Why does Phase 59 test a different population and random seed?

A fast calculator is not very useful if it works only for the numbers used in its demonstration. Earlier phases established repeatable speed on one public population. Phase 59 asks a different question: can the connected model accept changed live inputs and still reproduce a newly run ordinary CPU model?

A **random seed** chooses a repeatable sequence of random numbers. Changing it can change people's simulated choices, even though their characteristics and the behavior equations are unchanged. A **population sample** changes which households participate. A **retry limit** changes how often the model may try again after failing to fit a trip into its available time window.

These changes exercise different paths through the program. For example, the 10,000-household, seed-991 test encountered 22 destination alternatives where the original run had used other widths. Two inherited guards accepted only a short list of familiar widths. We extended the reviewed arithmetic to every actual width from 1 through 32 and tested each width. We did not add fake alternatives to make the old shape fit: that could change rounding and choices.

Another old guard treated a new configuration directory as a changed model. The replacement still checks the original equations and coefficient files by hash. It permits only a narrowly defined seed/sample settings overlay. An unreviewed equation or constant is still rejected. Supporting these scenarios does not mean that every city's configuration is now supported.

305. How can scheduling stop depending on a saved demonstration?

Imagine teaching a calculator by giving it yesterday's completed worksheet. Reproducing that worksheet is a useful test, but the real calculator must also solve today's worksheet without looking at yesterday's answers.

The new mandatory-tour scheduler reads today's model equations, people, tours, available times and random draws. It builds its small table of travel-mode accessibility values from live GPU calculations. The scheduler's constructor does not load captured scheduling inputs. The experiment deliberately points the old artifact argument at a nonexistent directory to check this separation.

Seven timetable expressions and two temporary assignments form a supported vocabulary. Their meaning and order are checked before compilation. Every feasible time alternative must have a finite live accessibility value; a missing cache entry cannot silently become zero.

Rounding is still delicate near a choice boundary. Those few rows are solved again by the live CPU calculation, including the accessibility inputs, using all feasible times and the same already-consumed random values. Rechecking only the last equation with GPU-produced inputs is not sufficient: the seed-17 experiment caught two wrong corrections by that weaker guard. This does not read a saved answer or draw a second random number. An independent comparison with the CPU run occurs after the simulation. Test answers belong in the examiner, not in the calculator.

306. What does a complete GPU retry controller do?

A **tour** is a sequence of trips that returns to its starting point. A model might try to place several trips inside a person's available hours. If one leg fails, it may retry that leg. A leg that already succeeded stays finished.

Previously, the CPU repeatedly assembled the unfinished group, sent work to the GPU and collected results. The new controller gives each GPU thread one person-tour and lets it carry out the entire retry sequence. It returns the final departure choices and an account of how many random numbers each trip actually used. The CPU still prepares the initial compact inputs and publishes the final table.

There is a subtle rule: when an outbound leg finishes, later inbound retries must follow exactly the upstream model's treatment of that now-absent leg. Using a plausible but different lower time bound would change the algorithm. Tests compare the complete upstream retry loop, not just one easy iteration.

For speed, the GPU prepares enough random values for up to 100 attempts before it knows how many will be needed. Think of laying out a deck of cards but only marking the cards actually played as used. Unused values do not advance the model's random-number ledger. This costs extra memory; the report records the random buffer and generator-state sizes rather than pretending they are free.

The stronger compiled CPU control delivers an important warning: the compact retry calculation itself is faster on this CPU than on this GPU. The observed component benefit comes mainly from eliminating repeated Python and table processing. An algorithm improvement and a hardware advantage are not the same claim. We retain the negative CPU/GPU result alongside the successes.

307. What is a versioned device data store?

Suppose a class shares a spreadsheet. If somebody changes a departure time, another person must not unknowingly use yesterday's copy. A version number is a simple way to say which edition a calculation is allowed to use.

The device store keeps numeric columns for people, tours and trips, indexed by stable IDs rather than accidental row positions. Publishing equal data reuses its GPU copy. Publishing changed data creates a new version. Reordering or removing rows invalidates the previous layout. Deleting a table explicitly invalidates its leases, which are permission slips for reading one version. Even positive and negative zero are distinguished in the stored byte pattern.

A consumer must check its lease immediately before reading. A raw GPU pointer cannot revoke itself, so keeping such a pointer after mutation is outside the contract. CPU tables remain authoritative at publication boundaries. Snapshot checks and uploads cost time and memory, and that cost stays inside the run.

The retry controller uses the store's working columns. Other components still use their existing input paths. This is a tested foundation for sharing more data, not a claim that the whole model already lives exclusively on the GPU.

308. Why can writing the same matrices faster matter?

An **origin-destination matrix** is a grid: a row is a starting zone, a column is an ending zone, and a cell contains an amount of travel. The model publishes 115 such matrices in five standard OMX files. Every cell and every zone label must remain present logically, even if a different file layout is used.

Many cells are zero. The new writer divides a matrix into square tiles, leaves all-zero tiles at the file format's standard zero fill value, and compresses the remaining tiles with ordinary gzip and byte shuffling. Four CPU workers compress tiles; one writer places them in the file. Ordinary OMX readers see the same complete grid. Even a stored negative zero is preserved.

Tests tried several tile sizes. On an actual public morning-period output, 128-by-128 tiles took about 0.19 seconds to write, compared with 0.94 seconds for the upstream writer. That roughly 4.9-times ratio excludes the work of calculating and grouping the trips. It is also a CPU file-writing improvement, not a GPU speedup. The full model must be timed separately.

Profiling showed that putting the grouping operation on the GPU could save less than a second even if it were free. Writing compressed output was the larger opportunity. We therefore kept the existing grouping arithmetic and optimized storage. No matrices, reports or model steps were removed to make the stopwatch look better.

309. Why do cold caches and disk choice change the answer?

A compiler translates program instructions into machine code. Its first use can be slow; a **cache** keeps the translated code for later runs. A fresh process can reuse that cache, so fresh-process timing is not the same as a completely cold-computer measurement.

While the SSD had little available space, new diagnostic outputs were written to the larger hard disk without deleting earlier results. An empty compiler cache caused extensive first-use compilation. The successful 10,000-household changed-seed diagnostic took about 793 seconds. It proved output agreement, not good speed. That cost remains disclosed even though later runs reuse compiled code.

The SSD subsequently had enough free space for new experiments. Timed CPU and GPU comparisons use the same storage and cache policy, with no concurrent heavy benchmarks. We do not compare one system's hard-disk cold start with the other's SSD cached run and call the difference a GPU advantage.

310. How should we interpret the final Phase 59 experiment?

There are three separate questions. **Correctness** asks whether the tested CPU and accelerated runs make the same decisions and produce the same matrix values. **Repeatability** asks whether measured speed improvements recur in balanced fresh-process pairs. **Generality** asks how far the supported input domain extends beyond those tests.

Changing seeds, population size and a model setting strengthens the third claim, but three scenarios cannot establish a universal theorem. Small logsum diagnostics retain explicit numerical bounds; exact travel decisions do not mean every intermediate floating-point number is identical. Faster computation does not make the behavioral model itself a more accurate prediction of people.

The acceptance checker separates the engineering target of finishing in under 100 seconds from correctness. A faster, correct candidate that misses that target must be labeled as a target miss. The source version must match across the final paired runs and scenario runs. Missing outputs, exceeded numerical bounds, mixed source versions or an improperly ordered pair reject qualification.

311. What do the stronger CPU comparisons teach us?

Imagine comparing two kitchens. One has a faster oven, but it also uses a better recipe and spends less time carrying ingredients around. A faster dinner does not tell us how much credit belongs to the oven alone.

We therefore give both processors the same compact calculation. For actual public-model mode-choice inputs covering 442,682 trips, the GPU's final probability-and-choice reducer is about 1.60 times faster than the best tested compiled CPU version. The CPU was allowed up to 48 threads. The comparison does not include creating the input scores, moving data or starting a compiler. It measures one calculation, not an entire model run.

For the compact departure-retry calculation, the result reverses: the CPU is about 3.37 times faster. That is useful evidence, not something to hide. It suggests a future runtime should choose its processor based on measured cost, rather than insisting that every calculation run on the graphics card.

The GPU retry implementation also keeps extra random-number state and unused possible draws ready. In the 50,000-household test, those two allocations are about 217 MB and 70 MB, plus about 20 MB of working arrays and additional entity snapshots. Here MB means one million bytes. These figures do not measure every allocation in the complete application. Reducing table overhead can be useful while still spending extra memory; both sides of that tradeoff matter.

312. Are the published reports really the same?

The model does not only choose individual trips. It also writes 24 summary reports and 115 travel matrices in five files. Planners read those products, so checking only the internal choices would leave a gap.

The report check requires the same files, column names, row order and values. Four reports have one harmless difference: a departure hour can be written as 5 or 5.0. Both mean exactly five. The checker allows that spelling change only in those reports' departure-time column, with no rounding tolerance. Changing a trip count, losing a row or changing a

real time value still fails. It records both versions' file fingerprints and lists the spelling differences. Thus we claim exact report values, not identical text bytes everywhere.

Matrix checks are stricter about the stored numbers: every matrix's numeric bytes, shape, data type and zone labels must agree. The compressed file itself can have a different layout while presenting exactly the same numbers to a standard reader. Think of storing the same book on different shelves: the shelf arrangement can change, but no words may disappear.

313. Why can a major engineering change produce a modest speed gain?

A complete model is like a relay race. Making one runner much faster does not remove the time spent by everyone else. This phase removes repeated retry coordination and speeds up output writing, but also pays for version checks, data snapshots and a stronger live CPU check of difficult scheduling choices. Those new checks replace a narrower benchmark-dependent shortcut. Their cost belongs in the stopwatch, even though they make the result more trustworthy.

This is why there are two kinds of progress. A speed improvement makes a supported run finish sooner. A stronger correctness contract makes the result usable with more kinds of inputs. We want both, but should not count one as the other. Passing three changed scenarios is valuable even if it does not make the original scenario dramatically faster.

To make the next large difference, the next experiment should target the remaining expensive complete steps and their data preparation, not multiply small kernel speedups together. A CPU/GPU dispatcher could select the faster processor for a supported calculation. More consumers could use the versioned data store directly, avoiding repeated transfers and table rebuilding. Each change still needs exact decision and random-ledger checks, changed-scenario tests and a complete-model timing comparison. A proposed saving is not an achieved saving until that experiment is run.

314. What did the completed Phase 59 experiment achieve?

Complete model clock	Regular CPU	Previous accelerated version	Latest version
Launch to exit	314.90 seconds	112.52 seconds	109.84 seconds
Charged model work	308.40 seconds	104.82 seconds	102.69 seconds

These are medians, the middle of the measured values. The accelerated versions were measured in six pairs, alternating which ran first. The latest version won every pair on both clocks. The two new regular CPU controls ran afterward, so that comparison is useful but was not an interleaved CPU/GPU experiment. All runs used fresh processes and existing compiler caches on the same machine. The CPU baseline is ordinary pinned ActivitySim with numerical libraries using one thread, not the fastest imaginable rewritten CPU model.

In everyday terms, a run falls from about five minutes fifteen seconds to one minute fifty seconds compared with that regular CPU setup. This phase alone saves about 2.68 seconds compared with the previous accelerated version. The latest version still needs roughly ten more seconds of improvement to reach the 100-second goal. No candidate reached that target; none was omitted to make the result look better.

Selected work	Previous accelerated version	Latest version
Choosing work/school tour times	7.15 seconds	10.80 seconds
Scheduling individual trip departures	3.70 seconds	2.30 seconds
Writing travel matrices	6.60 seconds	2.50 seconds

The first row gets slower because we now check difficult choices using their complete live inputs instead of relying on the narrower previous arrangement. The other two rows get faster. These are component medians, so adding them does not

give an exact full-model saving. The separate technical comparison lists every one of the 34 steps, including those with no gain.

Correctness testing covers the original public sample and three changed cases: 10,000 households with seed 991, 50,000 households with seed 17, and a setting allowing seven departure attempts. All match independently generated regular CPU outputs. This establishes repeatable results for those tested inputs, not a proof for every city, seed, model equation or computer. Near-boundary checks are an explicitly tested engineering safeguard, not a universal theorem about floating-point arithmetic. Matching the reference means the optimization does not change the tested simulation; it does not prove that the simulation perfectly predicts real people.

The next large opportunity is to reduce complete-step data preparation and CPU/GPU handoffs while extending the live-input contract to more scenarios. The current result is a stronger, modestly faster hybrid system with clear evidence. It is not a finished GPU-only replacement for all of ActivitySim.

315. Why does Phase 60 work on the CPU around the GPU?

Imagine a school kitchen with a fast oven. If staff spend most of the lunch period repeatedly copying orders and sorting ingredients, buying a faster oven will not solve the queue. The GPU is similar: its calculations can be fast while the surrounding preparation still takes substantial time.

We profiled complete model steps to find where the time actually went. **Profiling** records how much time functions take. That recording adds work, so a profiled run is a diagnostic, not the number used to advertise speed. The experiment found repeated table construction, Python loops over millions of time-slot rows, repeated formatting of identical labels, and a compiled CPU calculation running with only one active thread.

Phase 60 addresses those measured costs while leaving the established GPU equations and live CPU boundary checks in place. This is an improvement to the combined system. It would be misleading to describe the resulting saving as proof that a new GPU kernel beats a CPU kernel.

316. How can preparation become faster without changing the model?

Several changes replace repeated work with a smaller equivalent operation:

- **Household activity rules:** CDAP coordinates whether family members stay home, do mandatory activities or do other activities. Its rules expand into a table of possible household patterns. The new constructor builds that same ordered table in arrays from today's coefficients, instead of repeatedly growing a table cell by cell. It does not load yesterday's choices.
- **Time-slot identities:** a pair of supported 16-bit time values can be represented by one unique 32-bit key. Sorting those keys gives exactly the same pair order, including negative values, with less work. Tests sample across the full numerical range, include its endpoints and check repeated pairs.
- **Timetable lookups:** use a small period lookup and gather the two required coordinates directly. Avoid copying a person's entire day for every possible time choice. Availability values and boundary rules stay unchanged.
- **Report labels:** if thousands of rows share a bin such as 0.00 - 1.00, format its label once and reuse it. Bin ranks, numeric conversion, row order and empty/missing-input behavior must remain compatible with the original.

The first tests caught an empty-report case that the new formatting code could not handle. The implementation now keeps the original path for empty or missing inputs. A faster common path does not excuse incorrect unusual inputs.

317. Why give the CPU a stronger chance?

A **thread** is a stream of work a processor can execute. The tested CPU has 24 physical cores and can handle 48 hardware threads. More threads are not always faster, but leaving a genuinely parallel calculation at one thread can make a hardware comparison misleading.

The non-mandatory tour-frequency calculation was already compiled by Sharrow to run different rows in parallel. We repeated its actual live input batches at 1, 4, 12, 24 and 48 CPU threads. Every resulting utility byte matched. The sum of the eight segment medians fell from 4.22 seconds at one thread to 0.27 seconds at 48 threads. Those numbers exclude preparing the inputs and publishing the model results; they are not full-step times.

The candidate therefore uses 48 threads for that one step and restores one afterward. Other numerical libraries remain limited to one thread. Changing the number of workers does not change the within-row arithmetic or random draws. Both comparison versions have the same thread-pool capacity, so creating that pool is not an unreported advantage for one version.

We also measure a stronger regular CPU ActivitySim configuration with 48 Numba threads throughout, alongside its ordinary one-thread setting. Numba is the compiler that turns selected Python calculations into machine code. This is still one model process, not 48 separate simulations. The stronger comparison must be reported even if it makes the GPU-based system's headline ratio smaller.

318. What makes the next speed claim credible?

The final test freezes the implementation, runs six old/new pairs in alternating order and compares every published output. Each candidate must finish in under 100 seconds from process launch to exit to meet the target. Three changed scenarios must also agree with their independently generated regular CPU outputs. The thread-sweep run is excluded because its repeated calculations deliberately make it slower.

All previous limits remain: warm compiler caches are not a cold computer; component medians are not additive full-model measurements; bounded diagnostic scores are not bitwise-identical intermediate arithmetic; and a finite set of scenario tests is not a proof for every city or every model configuration. The result must distinguish faster preparation, parallel CPU work and earlier GPU acceleration rather than assigning all credit to one processor.

319. What did Phase 60 actually achieve?

Complete model configuration	Median waiting time	What this comparison means
Regular ActivitySim, one Numba thread	296.18 seconds	The more limited CPU configuration, measured twice
Regular ActivitySim, 48 Numba threads	201.61 seconds	A stronger CPU control, measured twice
Previous accelerated system, Phase 59	98.42 seconds	Six fresh matched control runs
New accelerated system, Phase 60	90.16 seconds	Six fresh candidate runs

The latest complete system is 2.24 times faster than the stronger regular CPU configuration. A run that took about 3 minutes 22 seconds now takes about 1 minute 30 seconds: roughly 1 minute 51 seconds less waiting. Compared with the already accelerated Phase 59 system in the same paired experiment, this phase saves about 8.27 seconds, or 8.40%. These are different questions and must not be presented as the same gain.

All six candidates finish between 89.91 and 90.69 seconds. Each beats its matched control on both clocks, and every candidate beats both our 100-second goal and the harder 95-second goal. **Charged time**, which counts model steps plus recorded validation and prewarming, falls from 91.92 to 83.46 seconds. **Elapsed time** is the time from launching the model

process until it exits; it includes additional process overhead. Independent comparison of completed output files happens afterward and is outside both model clocks.

The result keeps every checked modeled decision, all 115 logical matrices and values in all 24 summary reports. The existing equivalent departure-label formatting rule and numerical limits for diagnostic scores still apply. The final implementation also passes 467 repository tests and the three changed-input scenarios described above. It is not a claim that every intermediate floating-point value is identical or that every possible future city configuration has been proven.

For each of the 34 model steps, including ones that remain on the CPU, see the [complete CPU and accelerated comparison](#). That table includes fresh one-thread CPU measurements too. Its per-step numbers are medians measured separately and rounded by upstream timers; they need not add exactly to the full-model median. A component's ratio is a useful description, not an independently randomized component experiment.

The biggest incremental step saving is deciding how many non-mandatory tours people make: 6.70 to 2.90 seconds, largely by using existing CPU parallelism. Mandatory scheduling falls from 10.30 to 8.80 seconds; household activity coordination from 6.30 to 4.90 seconds; and non-mandatory scheduling from 6.30 to 4.90 seconds. Some steps get slightly slower in these measurements. We keep those rows too, because success means less total waiting with correct outputs, not making every row in a table look better.

320. What does this mean for the larger project?

We have made a meaningful dent in the waiting time for a real, calibrated, complete public model on this workstation. We have also learned that improving the GPU's surrounding work can matter as much as writing another GPU kernel. The earlier GPU work is still used; this phase makes the combined system more efficient instead of replacing that foundation.

There are important boundaries. The experiment runs 50,000 households within the public 1,454-zone geography, not every household in the region. It uses previously populated compiler and filesystem caches, not a newly installed computer. More processor threads do not automatically accelerate every function. Exact reproduction of ActivitySim's outputs also does not prove that ActivitySim perfectly predicts what real people will do; behavioral validation is a separate transportation-science question.

The next substantial opportunity is to reduce repeated preparation and data handoffs in the remaining expensive complete steps, especially trip destination and live scheduling. Phase 60 did not add another user of the shared GPU entity store or remove CPU boundary checks. Those remain possible engineering directions, not completed features or promised speedups. Any next change must again beat a strong CPU alternative, improve the complete model, and pass the output checks on changed inputs. Separately, a clean-machine replication package and cold-start measurements would make this evidence easier for another researcher to reproduce.

321. Phase 61: making the whole team work better

Think of a restaurant with a very fast oven. Buying the oven helps, but dinner is still slow if the cooks repeatedly fetch the same ingredients, copy every order by hand, or send tiny jobs to the oven that a person could finish sooner. Our GPU is like that oven. Phase 61 works on the whole kitchen: the CPU, the GPU, the information they share, and the handoffs between model steps.

The starting point was Phase 60's roughly 90-second complete run. We aimed for a median below 75 seconds, with 70 seconds as a harder goal. A target is not a prediction and does not become a result just because we wrote it down. We keep the same travel behavior, random-number rules and output checks.

The model still represents a public sample of 50,000 households across 1,454 zones. It executes all 34 steps, including output files and reports. We did not replace the real calibrated model with a simpler synthetic demonstration.

322. What are skims, shared columns and generations?

A **skim** is a lookup table of travel information. If you know an origin, destination and time of day, it may tell you the travel time or cost. A three-dimensional skim is like several ordinary maps stacked together, one for each time period.

The original software provides a flexible way to find these entries. For the ordinary in-memory maps used here, we can use the already checked row and column positions directly. Special compressed, blended or unusual layouts still use the original route. Direct access does not mean ignoring zone identity, reversing a journey incorrectly, or assuming that morning and evening are interchangeable. Tests caught one time-index mistake before the final path was accepted.

A **column** is one field for many records: for example, every trip's departure time. A shared device store keeps useful numeric columns on the GPU. Actual tour- and trip-mode calculations now consume some of these stored columns. Trips can carry shared state from choosing a destination through scheduling and into choosing a mode of travel.

The store attaches a **generation number**, like an edition number, to its data. A consumer gets a permission slip called a lease for that edition. When relevant data changes, an old lease becomes invalid. We also compare the live CPU values and their exact data types before borrowing a stored column. A newly calculated field does not silently become yesterday's field. This protects against a fast but wrong simulation that uses stale information.

Sharing is not free. Checking, copying and moving information still take time. We count those costs in the complete run, rather than describing every reused column as a guaranteed speed improvement.

The measured new consumers are still a small subset: one participant-count column for tours and one direction flag for trips. The store holds more columns than these consumers currently use. This is a real working connection, not proof that the entire model now lives on the GPU or that all saved data movement pays for the store's overhead.

323. How can a timetable fit into a few numbers?

The model needs to know whether a person's proposed tour overlaps existing commitments. Its timetable contains a small vocabulary of states, including start/end boundaries and occupied periods. Two boundaries touching are not always the same thing as two activities overlapping.

A **bit** is a switch with two positions, zero or one. One 32-bit integer can hold 32 switches. We encode each relevant timetable state in its own row of switches. A few bitwise operations can then check exactly the same collision rules that originally inspected many period entries. Both a compiled CPU function and a GPU kernel use this representation.

The approved fast representation supports at most 32 periods and the existing state vocabulary. It does not claim to handle an arbitrary calendar. Every call reads the current timetable, including changes or rollbacks. Tests cover all 25 pairs of the five possible states, the highest bit, and changed windows. The live diagnostic also checks over 17.7 million availability queries against the original collision rules.

Here is the crucial lesson: a faster representation is not automatically a GPU victory. The CPU can use the same clever representation. If a GPU finishes the arithmetic quickly but needs extra delivery time for the data, the CPU may finish the real job first. Phase 61 retains both implementations and selects the CPU timetable path for the complete qualified candidate.

324. How do we make randomness faster without changing anyone's choices?

The model uses **pseudorandom** numbers: numbers produced by a repeatable recipe, not unpredictable physical events. Each record has a seed and an offset. The seed starts its private number sequence; the offset says how far along that sequence it has already moved. The record of these positions is the **random ledger**.

A **uniform** draw spreads numbers evenly across a range. A **normal** draw follows the familiar bell-shaped pattern: values near the middle are common, and extreme values are less common. They are different recipes, so matching one does not prove that the other matches.

Earlier model steps now use our already reviewed GPU uniform generator where supported. Normal draws use a new compiled CPU batch operation. This removes much of the slow one-record-at-a-time Python work while keeping each original seed, offset and requested draw count. ActivitySim still performs its existing scaling, repeated-record broadcasting and lognormal transformations. Unhandled request forms keep the original implementation.

An early normal-generator experiment was wrong. The compiler noticed that some skipped draws had unused values and removed them. Those draws were essential: skipping ahead changes the next random number. We rejected that version. The revised version returns and checks a checksum of the discarded draws, so advancing the sequence has an observable result and cannot simply be deleted.

This is why running the code matters. The revised normal generator matched the original NumPy bit patterns on 933,292 live row streams across 33 batches. Separate tests cover different seeds, long offsets, odd draw counts, repeated records and the final ledger. Matching a bell-shaped histogram alone would not have been enough: two different sets of numbers can have the same shape but make different people choose different trips.

325. What else changed, and why keep the CPU?

Tour mode choice asks whether a tour uses a car, transit, walking or another mode. Phase 61 connects it to the existing resident GPU pipeline for scoring and reducing the 21-mode choice tree. Very close decisions still receive the existing live CPU check, using the same already-consumed random draw. The check does not secretly read a saved answer or roll the dice again.

Other changes are ordinary CPU engineering. A report with thousands of rows but a handful of distance bins now formats one label per observed bin and reuses it. Trip scheduling can group integer person/tour identities without first constructing hundreds of thousands of Python row tuples. Their order and group identities remain the same.

We also tested how CPU worker threads wait between jobs. A waiting policy is a request to the runtime, not a different travel model. A CPU worker that keeps checking for work can compete with the CPU thread launching GPU jobs. But putting workers to sleep can also add wake-up costs. We measured this tradeoff rather than assuming that one setting must always be better.

The final system is intentionally **hybrid**. GPUs handle substantial parallel scoring and choice work; CPUs still handle coordination, selected calculations, numerical boundary decisions and file operations. A complete-model speed gain cannot honestly be credited only to GPU arithmetic when CPU algorithms and data handling also improved.

326. What does a fair CPU-versus-GPU test require?

There are three different comparisons in this project:

- **Regular CPU ActivitySim versus the complete accelerated system:** does a planner wait less for the same checked outputs? We include fresh regular CPU controls with both one and 48 Numba threads.
- **Previous accelerated version versus the new version:** did this phase improve the already fast system? Six matched pairs alternate which version runs first, reducing the risk of favoring one because of run order.
- **Equivalent CPU and GPU calculations:** which processor is better at a specific task when both get the same good algorithm and actual live inputs? We test several CPU thread counts rather than selecting a weak opponent.

The last comparison has two delivery boundaries. **Resident** means that the inputs are already on the processor that needs them. **Transfer-inclusive** adds the necessary uploads and downloads. Neither should be mistaken for a complete-model time. They omit other work that the actual application still has to perform, such as finding the correct records and checking live state.

The timetable and tour-mode controls disclose CPU wins instead of hiding them. The normal-generator gain is explicitly a CPU batching result. Earlier qualified GPU wins remain part of the project, but they do not prove that every new calculation should move to the GPU.

For the complete model we freeze source and configuration files throughout each measurement, run no competing benchmark jobs, retain all output checks, and distinguish development experiments from final repeated evidence. Instrumented runs that save input captures or collect detailed profiles are not eligible as performance evidence. A 75-second goal is a separate gate: correctness can pass and speed can improve while that goal still fails.

327. What did Phase 61 actually achieve?

The final repeated experiment, not the fastest development run, gives these complete-model results:

Configuration	Median elapsed time	Evidence
Regular CPU ActivitySim, one Numba thread	296.52 seconds	Two fresh runs
Regular CPU ActivitySim, 48 Numba threads	201.17 seconds	Two fresh runs
Previous accelerated system, Phase 60	91.33 seconds	Six matched controls
New accelerated system, Phase 61	79.22 seconds	Six matched candidates

Against the stronger regular CPU setting, this is 2.54 times faster: about 3 minutes 21 seconds becomes 1 minute 19 seconds, saving about 2 minutes per run. Against the already accelerated Phase 60, this phase saves 12.10 seconds, or 13.25%. Those are different comparisons. We do not pretend that all the earlier GPU work happened in Phase 61.

Every matched pair improves on both clocks. New elapsed times range from 78.75 to 79.68 seconds. The charged median, which includes model steps plus recorded validation and prewarming, falls from 84.67 to 72.82 seconds. Elapsed time additionally includes process overhead. The independent output audit occurs after the run and is outside these model clocks.

The 75-second median goal is still 4.22 seconds away. The more ambitious 70-second goal also fails. We publish the status as a replicated improvement with the wall-time target not met, rather than quietly changing the target or picking a different clock.

All checked travel decisions, all 115 logical travel matrices and values throughout all 24 summary reports agree. Four reports retain the documented equivalent departure-key spelling difference, such as 5 versus 5 . 0. Diagnostic scores still have their existing numerical limits; this is not a claim that every intermediate number has identical bits.

Three changed scenarios also pass against independently generated regular CPU answers: 10,000 households with seed 991, 50,000 households with seed 17, and the 50,000-household retry-7 configuration. These test a smaller sample, changed random numbers and changed retry behavior. They are not proof about every possible model configuration. The final repository suite passes 545 tests, before and after qualification.

The [complete 34-step comparison](#) retains every row, including slowdowns. The largest incremental reductions include tour mode choice (5.35 to 3.25 seconds), trip mode choice (5.65 to 3.90), non-mandatory scheduling (5.00 to 3.30), summary reporting (3.90 to 2.60), and mandatory scheduling (8.70 to 7.50). These separately rounded step medians need not add exactly to the total median. A step ratio is descriptive, not an independently randomized experiment for that step.

328. What does this mean, and what remains to do?

We can run this real, calibrated public model substantially faster on the existing workstation while reproducing its checked outputs. If someone ran 100 comparable scenarios, the measured saving against the stronger regular CPU setting would amount to roughly 3.4 hours of model execution. That is an illustration based on repeated similar runs, not a guarantee for every planning project or its total data-preparation time.

The project also has stronger evidence about where not to use GPU arithmetic. A compact CPU algorithm can beat a GPU once moving information is counted. This does not make earlier GPU wins disappear. It means we should choose a processor for the complete task rather than treating GPU use itself as the goal.

The next meaningful performance target is work that still consumes whole seconds: repeated flow/compiler setup for household activities and the live CPU scheduling reference, plus preparation and utility work in trip destination. A reusable compiled reference plan must reproduce the same arithmetic and preserve the independent live check. Replacing that check with saved choices is not an acceptable shortcut.

Shared data also needs broader, demand-driven use. The current new consumers read only two columns while the store publishes more. Future work should publish what programs actually need, connect additional useful live inputs, and prove that the saved copying exceeds the cost of managing the store. Trying a compiled CPU reducer inside the GPU utility pipeline is another measurable option, but its small isolated win does not promise a large complete-model improvement.

Finally, another researcher should reproduce the environment independently, measure cold startup and test more scales and configurations. This result uses warm caches, one workstation and a 50,000-household sample, not all regional households or every ActivitySim feature. Matching the model's answers establishes computational replication; proving that those answers accurately predict human behavior remains a separate scientific task.

329. Why does Phase 62 focus on preparation rather than another GPU kernel?

A fast restaurant needs more than a fast stove. Someone must read the order, find ingredients and wash the utensils. Making the stove twice as fast does not halve the customer's wait if these other jobs take most of the time. The same is true here: a GPU can finish arithmetic quickly while the CPU is still preparing the instructions and arranging the next inputs.

Phase 62 targets those remaining preparation costs. It keeps useful compiled programs, avoids rebuilding identical parts of calculation recipes, and publishes only the shared columns that a live calculation actually requests. It also tests running several scenarios in one long-lived process. These are application improvements, not a new claim that GPU arithmetic became faster. Earlier GPU kernels and the independent CPU checks still do their jobs.

An **expression** is a small formula, such as travel time multiplied by its importance. A **compiler** translates formulas into instructions the computer can execute. An **intermediate representation**, or IR, is an organized recipe between the written formula and those machine instructions. The same recipe can help generate both a CPU reference and a GPU implementation.

The new compiler reuses the parsed structure of an unchanged formula. It still reads the current coefficients and creates independent recipe objects. Changing a coefficient must change the calculation; reusing an old answer would be incorrect. Tests deliberately change coefficients and inputs to check this distinction. The shared recipe preserves the project's existing arithmetic rules rather than inventing a more convenient rounding policy.

The runtime also avoids repeated successful checks of the same writable compiler-cache directory. A cache is a place to keep reusable work. Its actual program identities, version checks and writes are not bypassed. Deleting or replacing the directory causes a new check. This scoped shortcut is for the supported sequential execution, not a guarantee against someone changing permissions while the model is running.

330. What does "only prepare what is needed" mean?

Think of a spreadsheet with dozens of columns. If a calculation uses two, copying every column to another computer wastes work and memory. Previously, the shared entity store published substantially more than its new consumers read. Phase 62 waits until a live program names the columns it needs.

In the development comparison, retained shared entity data fell from 42,055,332 bytes to 1,726,514 bytes, about 96% less. The same 19 segment consumers still use the two relevant columns: tour participant counts and whether a trip is outbound. This is the size of that particular store, not the whole model's GPU memory. Nor does a 96% smaller store imply a 96% faster model. The useful question is how much waiting the change removes overall.

The runtime checks the live table's identity, data type and exact host values before reuse. If the model changes a value, the next calculation must see the new value. It does not simply assume that an array at an old memory address still contains the same information.

Trip-destination preparation also stops compiling one base GPU program that the selected normalized production path never calls. That path uses its own fused program. The binding checks and recipe validation remain, and shadow or other paths that need the base program still compile it. Avoiding unused work is different from skipping a calculation the model needs.

331. How can a repeated-scenario worker save time safely?

A **scenario** is one complete model run with specified assumptions and a random seed. The seed determines a reproducible sequence of random draws. Starting a new Python process for every scenario repeats imports, input-file parsing and some program setup. A **persistent worker** is a process that stays open to run the next scenario.

Keeping a process open is easy; preventing yesterday's answers from leaking into today's calculation is harder. A timetable remembers occupied hours, and a random-number ledger remembers which draws belong to which entities. Neither should survive as the next scenario's mutable model state.

The selected worker keeps compiled CPU and CUDA programs with matching source identities. It rebuilds the ActivitySim and ChoiceForge application module graphs between scenarios. Each model invocation constructs a fresh workflow state, working tables, timetables and random-number ledgers. The reused programs receive the new scenario's live inputs. A recipe may be reused; its previous cooked meal may not.

The worker also keeps private snapshots of numeric raw input tables, before the model samples or changes them. Each scenario receives a deep copy: a separate set of values it can change without changing the stored snapshot. The snapshot budget is 1 GiB, roughly one billion bytes; the public benchmark uses 668,561,212 bytes. Unsupported input types use the original file reader. Both the regular CPU worker and the hybrid worker receive this same reuse opportunity, so the comparison does not give only the GPU side free setup.

Input files and generated program files are checked using SHA-256 hashes. A hash is a fingerprint of file contents; changed contents should produce a different fingerprint. Scenario-boundary checks reject detected changes. The supported assumption is that nobody edits the input files during the batch. Hashing at boundaries is not a file lock or protection against every possible concurrent edit.

332. Why was keeping network data on the GPU not selected?

Network **skims** summarize travel conditions between places, such as travel time or cost. Keeping them on the GPU sounds attractive because uploading them again costs time. Phase 62 implemented a content-checked, 2 GiB skim cache and tested that changed values could not reuse an old entry.

The working set was larger than that cache. An entry would be removed to make room, then needed again later. This is called **cache thrashing**. Development counts showed only four hits but 145 misses per scenario. Checking billions of bytes and maintaining the cache also cost time. The content-cached trial took 232.55 seconds for three scenarios; a

subsequent program-only trial took 216.62 seconds. These are development comparisons, not balanced final measurements proving an exact causal saving from this one feature.

The selected formal worker therefore does not keep GPU skim arrays between scenarios. It reuses programs and private raw input snapshots instead. The optional skim cache remains tested code, not the default or a claimed win.

An early cache trial also failed the old rule requiring a large GPU memory release after the last consumer. That failure was real: a cache deliberately retained some memory. The optional cache now has a separately named bounded ownership check. The selected no-cache path still obeys the original release rule. We do not relabel an old failed contract as a successful release.

333. How do we test whether scenario state really resets?

The batch test runs **A, then B, then A again**. A uses the default public 50,000-household setup. B changes the seed to 17. The second A must reproduce the original independent CPU reference, even after B changed the worker's recent execution history. All three runs get their own output directory.

There are four strategies: separate hybrid processes, one persistent hybrid process, separate regular CPU processes, and one persistent regular CPU process. The whole sequence is repeated with the strategy order reversed. That makes 24 complete model executions for the batch experiment alone. Both persistent strategies receive the raw-input and CPU program reuse opportunities. Every output is audited, not just the last A.

The fresh-process experiment is separate: six alternating Phase 61/62 pairs, two regular CPU runs at each of one and 48 Numba threads, and three changed scenarios. Together with the batch experiment, that is 43 full-model runs. Source and configuration files stay frozen throughout the timed series. Tests, detailed profiling and PDF rendering do not run beside the models.

The batch clock includes its first scenario, initial setup, input checks, private copies, state resets and process overhead. Independent output audits are outside the execution clock for every strategy. Dividing a three-run batch time by three describes an average within that batch; it is not the time a person should expect when launching one fresh process.

A-B-A, mutation tests and content checks provide concrete evidence for this specific worker and public benchmark. They do not prove every possible scenario, network, concurrency pattern or unlimited number of repetitions. In particular, rebuilding application modules does not prove that every third-party library releases every old object. A long-running service needs additional memory and lifecycle qualification.

334. What did the fresh-process experiment actually achieve?

The final six pairs and fresh CPU controls give these elapsed-time medians. A median is the middle measurement after sorting the results; for an even number of measurements, it is the average of the two middle ones.

Configuration	Launch-to-exit median	Repetitions
Regular CPU ActivitySim, one Numba thread	299.83 seconds	Two
Regular CPU ActivitySim, 48 Numba threads	204.84 seconds	Two
Phase 61 hybrid, remeasured in this series	80.63 seconds	Six
Phase 62 hybrid	76.81 seconds	Six

The complete accelerated model is 2.67 times faster than the stronger regular CPU setting: about 3 minutes 25 seconds becomes 1 minute 17 seconds. It saves 128.02 seconds, or 62.50%, per comparable fresh run on this workstation. Against

Phase 61 in this experiment, it saves 3.81 seconds, or 4.73%. The larger overall gain includes all the previous phases; this phase did not create that entire advantage from scratch.

Every pair improves on both the elapsed and charged clocks. Phase 62 elapsed times range from 76.25 to 78.14 seconds. Charged medians fall from 73.83 to 70.37 seconds. The elapsed clock includes the remaining process overhead; the charged clock includes the model and recorded validation/prewarming. Independent output audits happen afterward and are outside both clocks.

The fresh-process target was a median below 70 seconds. At 76.81 seconds, it is missed by 6.81 seconds; the 65-second stretch goal is also missed. Neither a faster development run nor a later warm batch scenario changes that result. The published status is a replicated improvement with the target not met.

All checked travel decisions, 115 logical travel matrices and values throughout 24 summary reports agree. Four reports have the documented equivalent departure-key text difference, such as 5 versus 5.0. Diagnostic intermediate scores keep their existing numerical bounds. The three changed scenarios pass too. This establishes replication of the specified outputs, not bit-identical values for every intermediate array.

The [34-step table](#) includes improvements and slowdowns. Household activity coordination falls from 4.85 to 3.55 seconds, mandatory scheduling from 7.50 to 6.90, stop frequency from 2.95 to 2.50, and trip destination from 7.90 to 7.55. For example, at-work subtour frequency instead rises from 0.60 to 0.80. Step timings are rounded and their separate medians do not have to sum to the whole-run median. We did not individually randomize every optimization, so these are component observations, not proof of each change's isolated causal contribution.

335. What did the repeated-scenario experiment achieve?

Each cell below is the median total for a complete A-B-A sequence, measured twice with the four-strategy order reversed. Both the initial run and process setup are included.

Engine	Three fresh processes	One persistent worker
Phase 62 hybrid	235.65 seconds	212.34 seconds
Regular CPU, 48 Numba threads	594.42 seconds	545.74 seconds

The hybrid worker saves 23.30 seconds, or 9.89%, across three scenarios. CPU reuse saves 48.68 seconds, or 8.19%. That CPU improvement matters: comparing the persistent hybrid only with the fresh CPU would give the GPU side an unfair setup advantage. With both persistent, the hybrid is 2.57 times faster, saving 333.40 seconds across this three-scenario workload.

All 24 batch-experiment outputs pass their independent audits, including both returns to A for every strategy. Subsequent hybrid scenarios reuse 23 source-keyed CUDA programs and 19 compiled CPU functions, plus unchanged Sharrow generated modules. Both engines reuse the three numeric raw input tables through private copies. No modeled choices are a cross-scenario cache.

The persistent hybrid total averages about 70.78 seconds per scenario over these three runs. This is not a fresh-process median and does not meet the under-70-second fresh-process goal by a different definition. Batch timing also charges per-scenario input hashing, which the separate fresh-model experiment does not add. Compare like with like.

Two repetitions per strategy establish a useful finite result, not a precise prediction of every future workload. A planner with different configurations, larger samples or a different machine should rerun the checks and timings.

336. What did we prove, and what did we not prove?

We implemented and executed the complete selected path. The formal evidence contains 43 full-model runs: 12 paired versions, four CPU controls, three changed scenarios and 24 batch scenarios. The repository suite passes 565 tests, including deliberate mutation, cache identity, restoration and invalid batch-evidence cases. The full suite was rerun at delivery after the timed series. Source fingerprints tie measured runs to the delivered production code; input and report fingerprints help detect changes to the evidence.

We did not prove that GPUs win every calculation. The unchanged strong CPU primitive comparisons still include CPU wins. We did not achieve an entirely GPU-resident model, a general persistent ActivitySim service, arbitrary network support, cold-install timings or a full 2.875-million-household run. The live independent CPU scheduling reference remains; Phase 62 reduces preparation around it rather than replacing it with previously saved answers.

We also did not add broader shared-input consumers just to increase the count. The selected store still serves two columns in 19 segments. Its smaller footprint is useful, but has a narrow measured boundary. GPU skim retention was an investigated opportunity that did not earn a place in the selected runtime. Reporting a rejected optimization is part of the result.

Finally, exactly reproducing a travel model is not the same as proving its predictions are true. Both CPU and GPU can reproduce an assumption that turns out to be inaccurate in the real world. Survey quality, behavioral calibration and validation against observed travel remain separate planning obligations.

337. What is the next meaningful opportunity?

The next performance work should target complete services that still take seconds, not simply add GPU kernels. Trip destination still takes about 7.55 seconds and mandatory scheduling about 6.90. A supported, reusable compiled reference plan could further reduce repeated Sharrow preparation while preserving independent live arithmetic. This phase only removed part of that cost; it did not deliver a general upstream Sharrow GPU backend.

The persistent worker needs a longer, memory-instrumented sequence and additional changed configurations before it can be called a durable scenario engine. That should include changing coefficients, samples and supported network inputs, checking output equivalence, peak memory and reset behavior. The first scenario and all setup must remain charged. Independent replication on another prepared machine is also needed before generalizing the speed ratios.

For a planner, the larger achievement is already meaningful: complete public model runs take about one third of the regular CPU time under the tested conditions, and repeated scenarios can save additional setup work. For a researcher, the remaining task is to turn that local, auditable result into broader evidence and a maintainable integration, without weakening the replication contract to make the numbers look better.

338. Why did Phase 63 work on preparation instead of adding another GPU kernel?

A faster calculator is not much help while someone is still finding the worksheet, translating the instructions and organizing the numbers. By this point, several GPU calculations were fast, but complete model runs still spent seconds preparing those calculations and their independent CPU checks.

Phase 63 therefore has two linked aims: reduce preparation that can safely be reused, and show that repeated scenarios do not accumulate old state or memory. It also turns the measured setup into a more reproducible package: pinned software, downloadable public inputs and a command that runs the comparison. The fresh-process goal remains below 70 seconds, with 65 seconds as a stretch. A goal is not a result; the measured results below must decide whether it was met.

This does not replace Phases 1-62. Those phases supplied the kernels, arithmetic rules, live CPU comparisons and complete-model integration. Phase 63 improves how that working system gets ready and runs repeatedly. It is still a CPU/GPU hybrid, not a new general-purpose Sharrow GPU backend.

339. What can we reuse without remembering yesterday's answers?

Think of a recipe and a cooked meal. Reusing the recipe is sensible. Serving yesterday's meal as if it were made from today's ingredients is not.

Sharrow turns model expressions into executable instructions. Phase 63 saves some of the preparation metadata used to construct those instructions: generated code text, names, positions and other structural information. It does not save a household's chosen destination, a timetable, a probability array or an old model-state object. The ordinary CPU instructions still run on live input data. Sharrow still prepares the current data relationships before the reusable part is considered.

The reuse key includes expression definitions, compiler options, relevant Sharrow source files and supported extra-function code and defaults. The cache also checks the generated program's actual source bytes. A digest is a short fingerprint computed from a file's contents: changing the contents should change the fingerprint. This helps detect stale or damaged entries; it is not a proof that the mathematical model is correct.

The program also avoids repeatedly parsing unchanged specification and coefficient files. A coefficient is a number controlling how strongly a factor affects a score, such as the tendency of a full-time worker to do a mandatory activity. Each request resolves the current file and fingerprints its contents. Changing a coefficient must therefore be noticed even if the file has the same size and timestamp. Returned tables, including their row and column labels, are private copies so later changes cannot alter the cached original.

Unsupported expression cases use the ordinary preparation path. Invalid cache entries rebuild rather than being accepted as valid. These checks and fallbacks are important parts of the optimization, not optional work to remove from the correctness experiment.

340. Why did changing memory measurements help, and what did we give up?

The model measures memory as it works. On this machine, one detailed measurement is expensive when repeated often. Phase 63 offers a less expensive in-step measurement, but only when the model is not using memory measurements to decide how to divide work into chunks.

RSS, or resident set size, measures memory pages currently present in physical RAM for a process, including pages it may share. USS, or unique set size, measures pages private to that process. Finding USS can require more detailed inspection. Neither number alone describes every kind of memory used by the whole system.

The selected option keeps RSS measurements inside model steps and still measures USS at scenario boundaries. A separate sampler checks RSS approximately every tenth of a second. Its highest reading is an **observed** peak: a shorter spike between readings could be missed. An in-step USS value that was not measured is reported as unavailable, never as zero actual memory.

The option refuses to run if adaptive chunking or chunk training is enabled. That guard matters for larger workloads, where changing the memory information could change how the model divides its work. Both CPU and hybrid controls get the same option in the repeated-scenario comparison. Savings from lighter diagnostics are not GPU-kernel arithmetic speedups.

341. What memory problem did the longer experiment find?

The first ten-scenario development run produced correct outputs but retained nearly 8.00 billion bytes of active GPU array allocations after its last reset. Correct answers alone were not enough: repeating the worker longer would risk running out of GPU memory. Host private-memory growth also exceeded the predeclared 512 MiB tolerance. MiB means 1,048,576 bytes.

The problem was like removing a book from a catalog while someone still held the book. Removing Python module names did not destroy old objects still referenced by compiled or library objects. Retained destination and scheduling services owned large random-state and calculation buffers.

The fix explicitly releases those per-scenario service owners, live Sharrow objects and named input/device caches before discarding the application modules. Only eligible source-keyed compiled programs remain reusable. Every new scenario gets fresh mutable application state and current input data.

In the corrected ten-scenario development run, all output audits passed. Active CuPy array-pool use returned to zero before every scenario and after the final reset. Host private-memory growth between the fifth and tenth scenarios was 185,577,472 bytes, below the 512 MiB limit. These observations motivated the formal repetition; a single successful development run was not the final proof.

Zero active pool bytes does **not** mean zero physical GPU use. The CUDA context, compiled programs and unused memory reserved by the allocator can remain. The corrected development worker still reserved about 2.24 billion GPU pool bytes after its final reset. The report distinguishes active allocations from reserved space rather than calling both a leak or calling both zero.

342. How does the new experiment make the comparison fairer?

There are four scenarios. A uses 50,000 households and the default random seed. B changes the seed to 17. C uses 10,000 households and seed 991. D changes a full-time-worker activity coefficient from 1.378734579 to 1.5. The CPU reference generator checks that D actually changes activity decisions, rather than merely changing an unused setting. It changed 6,373 person-activity choices in the newly generated CPU reference, while the sampled person IDs stayed the same.

Each strategy runs A, B, C, D, A, C, B, A, D, A. Returning to A is especially useful: its answers should not change because a different scenario ran earlier. There are four strategies: fresh hybrid processes, one persistent hybrid worker, fresh CPU processes and one persistent CPU worker. The entire four-strategy experiment is repeated with the strategy order reversed. That makes 80 complete scenario models, compared against four newly generated independent CPU references.

Both engines receive the same new preparation and memory-measurement options, 48-thread capacity, passive waiting policy and private copies of reusable raw input tables. Every scenario checks its public-input digests. Timing includes worker startup, the first scenario and per-scenario input checking. Independent output audits occur after the timed work and are excluded for both engines.

The memory contract was fixed before formal measurements: compare host USS and active GPU memory between the fifth and tenth scenarios, both A. Allowed growth is at most 512 MiB host USS and 128 MiB active device memory. The hybrid must also release all active default CuPy pool arrays before each scenario and after final reset. These are finite workload checks, not a mathematical guarantee that an unlimited service with arbitrary inputs can never leak.

Separately, six reversed-order Phase 62/63 pairs measure fresh complete processes, and two ordinary CPU48 runs provide a familiar ActivitySim baseline. That is 94 measured models in total, plus CPU-reference generation and raw-image preparation. The default-A cases also provide eight fresh processes per engine with matched new preparation. Their clocks include input checking and reset telemetry, so they are labeled separately from the fresh-pair harness.

A ten-scenario total mixes eight 50,000-household and two 10,000-household runs. Dividing it by ten would not give the time for a normal 50,000-household run. Equal preparation also does not isolate GPU hardware: the hybrid retains earlier CPU-side, data-layout and output-writing improvements. It is a useful comparison of complete working systems, not proof that every improvement comes from GPU arithmetic or that no better CPU implementation is possible.

343. What does rebuilding from public inputs establish?

We staged the source in a new directory without copying the original Python environment, runtime caches, raw data or model answers. A bootstrap command installed 98 exact library versions, obtained the pinned ActivitySim revision and applied the reviewed integration patch. It downloaded the public data archive and verified all five selected input files against their expected digests. The configuration snapshot includes 158 configuration/license files so a moving online example cannot silently change the experiment.

This exercise caught details that an existing working installation can hide. ActivitySim's generated version label depended on the clone's Git tag history. Patch application also changed newline bytes. The bootstrap now pins the version label and restores reviewed source bytes, refusing unexpected model-code edits. Git checkout rules preserve the relevant local-source and metadata bytes too.

The first clean test attempt also found missing benchmark source scripts in the staging recipe. We repaired the recipe and copied those scripts, not old results. The subsequent clean-directory suite passed 601 tests with one intentional skip: an old public checkpoint was not staged. A prior 602-pass run launched from the original directory could see that old checkpoint, so it is not presented as fully isolated evidence. Final test counts after additional reporting checks are recorded with the completed qualification below.

The reconstructed environment uses its own CuPy and Numba cache directories. Tests and preparation warm these caches. The first full CPU reference must still compile many Sharrow model expressions; its preparation cost is kept separate from the warmed comparison. A fresh **process** is not a fresh **installation**. Here, the first full CPU reference took 1,880.84 seconds, about 31 minutes. The later changed-seed, smaller-sample and changed-coefficient references took 201.26, 111.48 and 205.11 seconds. These are reference-generation costs, not the final matched performance comparison. Reporting only the later timings without disclosing preparation would hide a real cost of getting started.

This is stronger than merely saying that the code runs on the author's existing environment. It is not independent hardware replication: the rebuild still runs on the same Windows workstation, with the same processor and RTX A4000. Library version pins are not cryptographic guarantees of every installed binary. Another machine must execute the numerical gates and measure its own timings. The supported compiled destination path is specialized for this public specification and NVIDIA sm86; arbitrary GPU architectures are not silently certified.

The first comparison attempt also found a Windows file-path limit. One checkpoint filename reached 260 characters because the rebuilt checkout was nested inside the original project. We shortened the output location and added a check before models start, rather than changing the system's registry. The failed attempt is retained but excluded from the complete comparison. The new series uses the same independently generated references and unchanged production code. Its raw-image preparation run took 272.80 seconds and passed the output checks; that setup cost is reported separately too.

344. What are the final Phase 63 speed and accuracy results?

Think of elapsed time as a stopwatch from launching the model until it exits. Charged service time is a narrower ledger of model, validation and prewarming work. We publish both but judge the fresh-run goal using the full stopwatch. Auditing saved outputs happens afterwards and is outside both engines' timed model work. Installing software and compiling everything for the first time are separate preparation costs described above.

Complete fresh model	Median elapsed seconds	Median charged seconds
Ordinary CPU ActivitySim, 48-thread capacity	202.82	196.95
Previous hybrid, Phase 62 measured alongside Phase 63	82.19	75.25
Latest hybrid, Phase 63	73.94	67.15

Six paired old/new trials alternate which version goes first. The new version wins every pair on both clocks, with a 10.04% reduction in median elapsed time. The ordinary CPU comparison is 2.74 times faster, or 63.54% less waiting. There are two ordinary CPU controls. The six latest-hybrid times range from 73.17 to 80.00 seconds; we did not discard a slow successful run. The old Phase 62 76.81-second result elsewhere in this guide belongs to an earlier session, not this paired experiment.

The under-70-second median target is NOT met. Neither is the 65-second stretch. The 67.15-second charged number cannot be used to say otherwise: it is a different stopwatch. Likewise, a repeated-scenario average cannot stand in for a fresh 50,000-household run.

Would CPU catch up if it received the new preparation improvements? We also measured that. Eight default-A fresh worker processes per engine, drawn from the balanced scenario campaign, give 193.18 seconds CPU versus 76.42 hybrid: 2.53 times faster. Both get the same new recipe/file reuse and memory-tracing options. This worker clock includes input hashing and reset telemetry, unlike the fresh-pair harness, so only compare numbers within its own table/boundary.

Accuracy is checked at several levels. All 94 measured models run all 34 model steps. Every checked decision matches its independent CPU reference; logical values in all 115 travel matrices and all 24 summary reports match under the existing key-format contract. Some diagnostic scores allow previously declared tiny numerical differences. This does not mean every file is byte-identical or that all conceivable inputs have been proven equivalent. It means the declared contracts passed the entire published experiment, including genuinely changed choices under the coefficient scenario.

The [complete component comparison](#) lists all 34 steps. For orientation, ordinary CPU versus latest-hybrid medians are 23.75 versus 6.05 seconds for mandatory scheduling, 39.65 versus 6.85 for trip destination, and 9.25 versus 4.05 for trip mode choice. Auto ownership is slower in the hybrid (1.00 versus 0.80 seconds). These are rounded component logs, not isolated kernel benchmarks; their medians need not sum to the median whole-model time. Equal preparation still does not isolate GPU hardware, because earlier CPU, data-layout and output-writing improvements also remain.

345. What happens when a planner runs many scenarios?

A persistent worker keeps the program open between scenarios, like leaving a kitchen ready for the next recipe. Ingredients and decisions must still belong to the new order. Each ten-scenario sequence contains eight 50,000-household models and two 10,000-household models. Two trials reverse the strategy order.

Ten-scenario strategy	Median total seconds	Approximate minutes
CPU, new process per scenario	1753.83	29.23
CPU, one persistent worker	1466.35	24.44
Hybrid, new process per scenario	714.47	11.91
Hybrid, one persistent worker	590.70	9.85

Persistence saves 17.32% of hybrid time and 16.39% of CPU time. The persistent hybrid is 2.48 times faster than the persistent CPU. These totals include the first scenario and setup; none are steady-state-only averages. They exclude external output audits equally. In practical terms, this particular workload falls from roughly 24 minutes to under 10 when comparing the two persistent engines. Whether that helps a real planning team depends on its actual models, hardware and how much time it spends outside model execution.

The memory checks also pass, with limits chosen before the formal experiment. Between the fifth and tenth scenarios, hybrid private host-memory growth is 189,591,552 and 189,046,784 bytes in the two trials (about 181 and 180 MiB). CPU growth is 357,072,896 and 128,790,528 bytes. All are below the 512 MiB limit, but none should be described as zero host growth. Hybrid active GPU-pool growth is zero, and every hybrid scenario begins and final reset ends with zero active default CuPy pool array bytes. Reserved GPU space and physical device usage are different measurements and can remain nonzero.

Main tests finish with 619 passes. The independently rebuilt directory finishes with 618 passes and one expected skip because an old checkpoint was deliberately not copied. Both produce 91 dependency/deprecation warnings, not zero warnings. These tests supplement the 94 full model measurements; they do not replace them. Publication checks re-read evidence digests, reproduce the 80-model qualification and matched comparison, and verify the exact tested production source. Git checkout tests cover both common newline-conversion settings. PDF pages are rendered and visually checked, not approved from extracted text alone.

346. What remains unproven, and what would move the project forward?

This phase improves a complete public benchmark and makes its repeated use safer. It does not replace the earlier phases: those supplied the kernels, strict arithmetic, live CPU safety checks and integration that this runtime reuses. Phase 63 mostly removes repeated preparation and retained state around that work. Faster housekeeping is useful even when it is not new GPU math.

The evidence is limited to the supported public model, sample sizes, scenarios, Windows environment and this workstation. We did not run a complete calibrated 2.875-million-household model, certify another GPU architecture, or establish an indefinitely running multi-user server. Hooks are for sequential execution, not simultaneous callers. Parsed configuration reuse is limited to 256 entries; individual JSON plan files are limited to 16 MiB, but there is not yet a global disk-cache eviction policy. Boundary memory checks cannot prove an infinite service leak-free, and pinned versions alone do not verify every binary on someone else's machine. None of the timing ratios promise a universal speedup.

Three useful next advances are distinct. First, run the pinned recipe on a second compatible workstation, including newly generated CPU references and all numerical gates. Second, profile the remaining complete-model cost and test larger reductions in preparation, scheduling and destination work against strong matched CPU implementations; keep the under-70 goal on the full clock. Third, prepare a small upstream Sharrow proposal with the data-only reuse contract, mutation tests and independent arithmetic evidence so others can review the integration. These are proposed next steps, not completed results.

The bottom line is practical, not magical: the supported complete model is about 2.5-2.7 times faster than the measured CPU alternatives, and a tested ten-scenario service finishes in about ten minutes while preserving its checked answers. The reproducibility package exposes preparation costs, failed attempts and limits so a reviewer can challenge and rerun the claim.